

FROM DATA TO THESIS

Using R



**Modern Data Analysis in the
Age of AI**

by Polla Fattah
and Collection of LLMs



FROM DATA TO THESIS

Using R for Non-Technical

Polla Fattah and Collection of LLMs

Contents

Welcome	17
How to use this book	17
About the data	17
Dedication	19
Foreword	21
Modern Research in the Age of AI	21
How to Journey Through This Book	21
The Interactive Playground	22

I Foundations of R Programming and Data Handling

1	Getting Started with R	25
1.1	Research written in code	26
1.2	Installing R and RStudio	27
1.3	The RStudio window	27
1.4	Calculations	28
1.5	Storing values in objects	28
1.6	Types of data	29
1.6.1	Missing values	30
1.7	Functions	30
1.7.1	Getting help	31
1.8	Scripts	32
1.9	RStudio Projects	32
1.10	Packages	33
1.11	A first look at the data	33
1.12	Chapter review	36
1.12.1	Summary	36
1.12.2	Key terms	37
1.13	Exercises	37
1.14	Further reading	37
2	Data Structures in R	39
2.1	Data as measurement	39
2.2	Variables and their values	40
2.2.1	Operations on a whole vector	41
2.2.2	Selecting values	41
2.2.3	Changing values	42
2.3	Categories	42

2.3.1	Kinds of measurement and R's types	43
2.4	A table of cases	43
2.4.1	Selecting columns and rows	44
2.4.2	Adding columns	45
2.4.3	The student wellbeing data	45
2.5	Matrices and lists	46
2.5.1	Matrices	46
2.5.2	Lists	47
2.6	Importing data	47
2.6.1	CSV files	47
2.6.2	Excel files	48
2.6.3	SPSS files	49
2.7	A codebook	50
2.8	Exporting data	51
2.9	SPSS and Excel equivalents	51
2.10	Error messages	52
2.11	A first summary	52
2.12	Chapter review	54
2.12.1	Summary	54
2.12.2	Key terms	54
2.13	Exercises	54
2.14	Further reading	55
3	Data Manipulation	57
3.1	What clean data means	57
3.2	Tidy data	59
3.3	The tidyverse	59
3.3.1	Chaining steps with the pipe	60
3.4	Working with cases and variables	60
3.4.1	Choosing the cases you need	61
3.4.2	Choosing variables	62
3.4.3	Sorting cases	62
3.4.4	Creating new variables	62
3.4.5	Summarising, overall and by group	63
3.4.6	Combining the steps	64
3.5	Reshaping between wide and long	65
3.6	Joining tables	66
3.7	Cleaning the survey export	68
3.7.1	Inspecting the export	68
3.7.2	Test responses and duplicates	69
3.7.3	Usable variable names	69
3.7.4	Consistent categories	70
3.7.5	Missing codes and numbers stored as text	71
3.7.6	Impossible values	72
3.7.7	Scale scores	72
3.7.8	From wide to long	73
3.7.9	Checking the result	74
3.8	Chapter review	77
3.8.1	Summary	77
3.8.2	Key terms	77

3.9	Exercises	77
3.10	Further reading	78
4	Data Visualization	79
4.1	What a graph is for	80
4.2	How people read graphs	81
4.3	Choosing a graph	82
4.4	The grammar of graphics	83
4.5	The basic graphs	84
4.5.1	Histograms	84
4.5.2	Bar charts	85
4.5.3	Box plots	86
4.5.4	Scatter plots	86
4.5.5	Line charts	87
4.6	Mapping and setting	88
4.7	Preparing graphs for publication	89
4.7.1	Labels	89
4.7.2	Scales and colours	90
4.7.3	Themes	91
4.7.4	Annotations	91
4.8	Honest graphs	92
4.8.1	A graph improved	93
4.9	Further kinds of graph	94
4.9.1	Small multiples	94
4.9.2	Density plots	95
4.9.3	Heat maps	96
4.10	Exporting figures	97
4.11	Common misconceptions	98
4.12	Chapter review	99
4.12.1	Summary	99
4.12.2	Key terms	99
4.13	Exercises	99
4.14	Further reading	100

II

Statistical Analysis for Research

5	From Research Question to Data	103
5.1	What data analysis is for	104
5.2	From a topic to a research question	105
5.2.1	Three kinds of question	106
5.3	From a question to a hypothesis	107
5.3.1	Falsifiability	107
5.3.2	Directional and non-directional hypotheses	108
5.3.3	The null hypothesis	108
5.3.4	Hypotheses for the whole study	110
5.4	Variables and measurement	112
5.4.1	Cases and the unit of analysis	112
5.4.2	From constructs to variables	112

5.4.3	Levels of measurement	114
5.4.4	The roles variables play	116
5.5	Reliability and validity	118
5.5.1	Kinds of reliability	119
5.5.2	Kinds of validity	120
5.6	Study designs	121
5.6.1	Why random assignment works	121
5.6.2	Cross-sectional and longitudinal designs	123
5.7	Samples and populations	123
5.7.1	Ways to draw a sample	124
5.7.2	Random error and bias	124
5.7.3	Non-response and attrition	126
5.7.4	Generalising from the sample	126
5.8	Planning the analysis before the data	126
5.8.1	An analysis plan	127
5.8.2	Sample size	127
5.8.3	Preregistration and honest analysis	128
5.8.4	Ethics	128
5.9	Common misconceptions	130
5.10	Chapter review	131
5.10.1	Summary	131
5.10.2	Key terms	131
5.11	Exercises	132
5.12	Further reading	132
6	Descriptive Statistics and Exploratory Data Analysis	133
6.1	The place of exploration	134
6.2	The centre of a distribution	135
6.2.1	Choosing a summary	137
6.3	The spread of a distribution	138
6.3.1	The standard deviation	139
6.3.2	The interquartile range	140
6.4	The shape of a distribution	141
6.4.1	The normal distribution	142
6.5	Unusual values	143
6.6	Missing data	145
6.7	Relationships between variables	146
6.7.1	Correlation is not causation	148
6.8	Describing your sample in a thesis	148
6.9	Common misconceptions	150
6.10	Chapter review	151
6.10.1	Summary	151
6.10.2	Key terms	151
6.11	Exercises	151
6.12	Further reading	152
7	Hypothesis Testing and Statistical Inference	153
7.1	Samples and populations	154
7.1.1	Sampling distributions	154

7.1.2	The standard error	155
7.1.3	Confidence intervals	156
7.1.4	Bootstrap confidence intervals	157
7.2	The logic of hypothesis testing	157
7.2.1	A permutation test	158
7.2.2	The p-value	160
7.2.3	Type I and Type II errors	160
7.2.4	Statistical power	161
7.2.5	Multiple testing	163
7.2.6	One-tailed and two-tailed tests	163
7.2.7	Statistical and practical significance	163
7.3	The one-sample t-test	163
7.4	The two-sample t-test	164
7.4.1	A small example	165
7.4.2	Testing the workshop	165
7.4.3	Effect size	166
7.5	The paired t-test	167
7.6	Assumptions of the t-test	168
7.7	Non-parametric tests	168
7.8	Chi-square tests	169
7.8.1	Test of expected proportions	170
7.8.2	Test of independence	170
7.9	Choosing a test	171
7.10	Reporting the results	174
7.11	Common misconceptions	175
7.12	Chapter review	175
7.12.1	Summary	175
7.12.2	Key terms	176
7.13	Exercises	176
7.14	Further reading	176
8	ANOVA and Regression	177
8.1	One-way ANOVA	178
8.1.1	Partitioning variation	178
8.1.2	The student wellbeing data	180
8.1.3	Post-hoc tests	181
8.1.4	Assumptions of ANOVA	182
8.2	Two-way ANOVA	183
8.3	Linear regression	184
8.3.1	The regression line	184
8.3.2	The student wellbeing data	187
8.3.3	Multiple regression	187
8.3.4	Categorical predictors	189
8.3.5	Separating an effect from a confounder	189
8.3.6	Curved relationships	191
8.3.7	Interactions	192
8.3.8	Model diagnostics	193
8.3.9	Reporting a regression	195
8.4	Logistic regression	195
8.4.1	Odds and odds ratios	196

8.4.2	A model for dropout	196
8.4.3	Predicted probabilities	197
8.5	Common misconceptions	198
8.6	Chapter review	199
8.6.1	Summary	199
8.6.2	Key terms	199
8.7	Exercises	200
8.8	Further reading	200
9	Multivariate Statistical Methods	201
9.1	Constructs and latent variables	202
9.2	Many variables at once	202
9.3	Principal component analysis	204
9.3.1	Running PCA	204
9.3.2	The number of components	205
9.3.3	PCA and factor analysis compared	206
9.4	Factor analysis	206
9.4.1	Reliability and Cronbach's alpha	208
9.5	Cluster analysis	209
9.5.1	k-means clustering	210
9.5.2	The number of clusters	210
9.5.3	Describing the clusters	211
9.5.4	Hierarchical clustering	212
9.5.5	Clusters in random data	213
9.6	Common misconceptions	215
9.7	Chapter review	215
9.7.1	Summary	215
9.7.2	Key terms	216
9.8	Exercises	216
9.9	Further reading	216
10	Mixed-Effects Models	217
10.1	Data with a structure	217
10.1.1	Why independence matters	218
10.2	A study of sleep deprivation	219
10.2.1	Random intercepts	220
10.2.2	The intraclass correlation	221
10.2.3	Random slopes	221
10.3	Change in wellbeing over two years	222
10.3.1	Why ordinary regression misleads	224
10.3.2	Confidence intervals and p-values	224
10.3.3	Persistence of the workshop effect	225
10.3.4	Change in GPA	227
10.4	Students within supervisors	228
10.5	Counts and yes-or-no outcomes	228
10.6	Reporting a mixed-effects model	229
10.7	Common misconceptions	230
10.8	Chapter review	231
10.8.1	Summary	231

10.8.2	Key terms	231
10.9	Exercises	231
10.10	Further reading	232

III

Machine Learning with R

11	Introduction to Machine Learning in R	235
11.1	Machine learning and prediction	236
11.2	Generalisation and overfitting	236
11.3	Designing a fair test	239
11.4	Data for prediction	242
11.5	Splitting the data	243
11.6	Preparing the data	243
11.7	Specifying and fitting a model	245
11.8	Evaluating a classifier	245
11.9	Overfitting in practice	246
11.10	Cross-validation	247
11.11	Tuning a model	248
11.12	The final test	249
11.13	Common misconceptions	251
11.14	Chapter review	252
11.14.1	Summary	252
11.14.2	Key terms	252
11.15	Exercises	252
11.16	Further reading	253
12	Classification Models	255
12.1	Data and resampling	256
12.2	Decision trees	257
12.3	Random forests	259
12.4	k-nearest neighbours	260
12.5	Support vector machines	262
12.6	Comparing models fairly	263
12.7	The confusion matrix	264
12.8	Choosing the threshold	265
12.8.1	Costs of the two errors	268
12.9	Imbalanced outcomes	270
12.10	Chapter review	272
12.10.1	Summary	272
12.10.2	Key terms	273
12.11	Exercises	273
12.12	Further reading	274

13	Predictive Regression	275
13.1	Data for predicting final GPA	276
13.2	Measuring prediction error	277
13.3	Ordinary regression with many predictors	279
13.4	Regularised regression	280
13.4.1	Choosing the penalty	282
13.4.2	Predictors kept by the lasso	283
13.5	Boosting	285
13.6	Comparing the models	286
13.7	The final test	287
13.8	Common misconceptions	290
13.9	Chapter review	290
13.9.1	Summary	290
13.9.2	Key terms	290
13.10	Exercises	290
13.11	Further reading	291
14	Advanced Clustering	293
14.1	The profile data	294
14.2	What counts as a group	295
14.3	Gaussian mixture models	297
14.3.1	A one-variable example	297
14.3.2	Choosing the number of clusters with BIC	298
14.3.3	Describing the profiles	300
14.3.4	Certainty of assignment	301
14.4	Clusters as dense regions	302
14.4.1	A small example	302
14.4.2	Choosing eps	302
14.4.3	Students who fit no profile	303
14.5	Evaluating a clustering	305
14.6	Common misconceptions	308
14.7	Chapter review	308
14.7.1	Summary	308
14.7.2	Key terms	308
14.8	Exercises	309
14.9	Further reading	309
15	Neural Networks	311
15.1	A single neuron	311
15.2	From neurons to networks	312
15.2.1	How a network learns	314
15.3	A neural network for the dropout question	316
15.3.1	A network that memorises	317
15.3.2	Tuning size and weight decay	318
15.3.3	A fair comparison	319
15.4	When neural networks are worth using	320
15.5	Deep learning	320

15.6	Common misconceptions	323
15.7	Chapter review	323
15.7.1	Summary	323
15.7.2	Key terms	323
15.8	Exercises	324
15.9	Further reading	324
16	Time Series Forecasting	325
16.1	Time series data	326
16.1.1	A small example	326
16.1.2	The parts of a pattern	327
16.2	The counselling service data	328
16.2.1	Seasonal patterns	330
16.2.2	Autocorrelation	330
16.3	Decomposition	332
16.4	Forecasting	333
16.4.1	Simple benchmarks	333
16.4.2	Two families of models	334
16.4.3	Training and test periods	334
16.4.4	Forecasting next year	336
16.5	Common misconceptions	339
16.6	Chapter review	340
16.6.1	Summary	340
16.6.2	Key terms	340
16.7	Exercises	340
16.8	Further reading	341

IV

Reproducible Research and Applications

17	Reproducible Research	345
17.1	Reproducibility and credibility	346
17.2	Documents that contain their own analysis	346
17.2.1	A small example	347
17.2.2	What happens when you render	347
17.3	A results chapter in Quarto	348
17.3.1	References and citations	350
17.3.2	Equations	350
17.3.3	Output formats	351
17.3.4	Parameterised reports	351
17.4	Keeping the same package versions	352
17.5	Interactive dashboards with Shiny	352
17.5.1	A wellbeing dashboard	353
17.6	Open science	355
17.6.1	Sharing data and code	355
17.6.2	Protecting participants	355
17.6.3	Preregistration	356
17.6.4	Citing software	357
17.7	Keeping a history with git (optional)	358

17.8	Common misconceptions	359
17.9	Chapter review	359
17.9.1	Summary	359
17.9.2	Key terms	359
17.10	Exercises	360
17.11	Further reading	360
18	Using AI with R	361
18.1	The reliability of coding	362
18.2	How AI assistants work	362
18.3	AI as a coding assistant	363
18.3.1	Asking good questions	363
18.3.2	Always check the answer	363
18.3.3	AI inside the editor	364
18.4	Coding the open-ended answers	364
18.4.1	Coding by keywords	365
18.4.2	Calling a language model from R	367
18.4.3	Structured answers for many texts	367
18.4.4	Agreement with the hand coding	369
18.4.5	What the students said	371
18.5	Using AI responsibly	372
18.5.1	Check, and stay responsible	372
18.5.2	Protect your data	372
18.5.3	Keep it reproducible	373
18.5.4	Disclose it	373
18.5.5	Be aware of bias	373
18.6	Common misconceptions	374
18.7	Chapter review	374
18.7.1	Summary	374
18.7.2	Key terms	375
18.8	Exercises	375
18.9	Further reading	376
19	Putting It All Together	377
19.1	The shape of a research project	378
19.1.1	Organising the project	380
19.2	From raw export to clean data	380
19.3	Describing the sample	383
19.4	Answering the research questions	384
19.4.1	The workshop and its persistence (RQ3)	384
19.4.2	Explaining GPA (RQ5)	385
19.4.3	Considering dropout (RQ9)	385
19.5	One figure for the thesis	386
19.6	The chain of reasoning	387
19.7	The results chapter	390
19.8	Choosing a method	390
19.9	Challenges every researcher meets	392
19.10	Where to go next	393

19.11	Common misconceptions	395
19.12	Chapter review	396
19.12.1	Summary	396
19.12.2	Key terms	396
19.13	Exercises	396
19.14	Further reading	397
	References	399
	Appendices	401
A	Appendix A: Installation and Setup	403
A.1	What you need	403
A.2	Installing R	404
A.2.1	Windows	404
A.2.2	macOS	404
A.2.3	Linux	404
A.2.4	Checking the installation	404
A.3	Installing RStudio	405
A.3.1	Positron, an alternative	405
A.3.2	No installation at all	405
A.4	Recommended RStudio settings	405
A.5	Installing the book's packages	406
A.5.1	The student wellbeing data: the data2thesis package	407
A.5.2	The versions used in this book	407
A.6	Common problems and solutions	407
A.7	Keeping R up to date	408
B	Appendix B: Glossary	409
B.1	Key terms	409
	A	409
	B	410
	C	411
	D	413
	E	415
	F	416
	G	416
	H	417
	I	417
	J	418
	K	418
	L	419
	M	420
	N	422
	O	423
	P	423
	Q	425
	R	425
	S	428
	T	430
	U	431

V	432
W	432
X	433
Y	433
Z	433
B.2 R functions	433
C Appendix C: R Packages	435
C.1 The packages in this book	436
C.2 Finding and choosing packages	438
C.3 Citing packages	438
D Appendix D: Common Errors and How to Fix Them	441
D.1 Reading a message	441
D.2 Starting out	441
D.2.1 object ‘...’ not found	441
D.2.2 could not find function “...”	442
D.2.3 there is no package called ‘...’	442
D.2.4 unexpected symbol, unexpected ‘)’, unexpected string constant	442
D.2.5 The Console shows + and nothing happens	442
D.2.6 argument “...” is missing, with no default	443
D.2.7 A misspelled argument, with no error at all	443
D.3 Files and folders	443
D.3.1 cannot open file ...: No such file or directory	443
D.3.2 cannot change working directory	443
D.4 Data	444
D.4.1 non-numeric argument to binary operator	444
D.4.2 argument is not numeric or logical: returning NA	444
D.4.3 NAs introduced by coercion	444
D.4.4 The result is NA	444
D.4.5 object ‘...’ not found, inside dplyr	445
D.4.6 arguments imply differing number of rows; replacement has ... rows	445
D.4.7 \$ operator is invalid for atomic vectors	445
D.4.8 subscript out of bounds	445
D.4.9 invalid factor level, NA generated	446
D.4.10 Detected an unexpected many-to-many relationship	446
D.4.11 Can’t combine ... and ...	446
D.5 Plots	447
D.5.1 Cannot use + with a single argument	447
D.5.2 mapping must be created by aes() ... Did you use %>% or > instead of +?	447
D.5.3 object ‘...’ not found, in a plot	447
D.5.4 stat_count() must only have an x or y aesthetic	448
D.6 Statistical tests and models	448
D.6.1 grouping factor must have exactly 2 levels	448
D.6.2 not enough ‘x’ observations	448
D.6.3 contrasts can be applied only to factors with 2 or more levels	448
D.6.4 Chi-squared approximation may be incorrect	449
D.6.5 glm.fit: fitted probabilities numerically 0 or 1 occurred; algorithm did not converge	449
D.6.6 boundary (singular) fit: see help(‘isSingular’)	449
D.6.7 Observations deleted due to missingness	450
D.7 Machine learning	450

D.7.1	For a classification model, the outcome should be a	450
D.7.2	Can't select columns that don't exist ... <code>.pred_Yes</code>	451
D.8	When you are stuck	451
E	Appendix E: Further Learning Resources	453
E.1	Books	453
E.1.1	Getting started and data skills	453
E.1.2	Statistics	453
E.1.3	Mixed models, machine learning, and forecasting	453
E.1.4	Reporting and dashboards	454
E.2	Courses and practice	454
E.3	Communities	454
	Bibliography	457

Welcome

Complete Course & Companion Slides

All 19 chapters of *From Data to Thesis* are available, accompanied by interactive [Playground](#) exercises and [Lecture Slide Decks](#) for teaching and review.

Prefer reading everything in one continuous page?

Open the  [Complete Book Online Edition \(Single-Page Reader\)](#) or  [Download Companion PDF \(14.8 MB\)](#).

This book follows Elaf, a Master's student in Educational Psychology, through her thesis. She studies 600 graduate students over four semesters to learn how sleep, stress, and supervisor support relate to their wellbeing, grades, and thoughts of dropping out, and whether a short wellbeing workshop helps.

Every chapter answers one of her research questions. You learn each method the way you will use it: first on a handful of numbers, then on the study data, and then in your own field.

How to use this book

- Read the chapters in order if you are new to R. Each one builds on the last.
- Run the code yourself. Everything you need is in the `data2thesis` package:

```
install.packages("https://polla-fattah.github.io/data2thesis_r/downloads/data2thesis_1.1.0.tar.gz",  
                 repos = NULL, type = "source")
```

- Practise in the [playground](#) after each chapter.
- Present or review with the interactive [lecture slides](#).

About the data

Elaf, her study, and her data are fictional. The data was generated by a computer program to look like a realistic survey of graduate students, so that every method in this book has something to find. It describes no real people, and its patterns (for example, that the workshop raised wellbeing, or that students who sleep more have higher grades) were built into the program by the author, not discovered. Nothing in this book is evidence about the wellbeing of real students, and the data and results must not be cited or used as findings about students, univer-

sities, or any real situation. The same applies to the counselling service's records and to the students' written answers. The data was designed so that every method in the book finds something worth interpreting, including some honest "no effect" results. The script that generates it is published with the book, so anyone can check it or generate it again.

Dedication

Dedicated to Tano

Foreword

Writing a thesis is one of the most intellectually demanding undertakings in an academic life. Yet, for many graduate researchers, the most exhausting hurdle is not formulating hypotheses or collecting observations; it is the daunting chasm between a messy spreadsheet of raw data and a defended, publication-ready results chapter.

For decades, higher education has pushed researchers toward point-and-click statistical packages. While these tools offer an easy first step, they often leave the underlying analysis opaque, create error-prone manual copy-paste rituals, and lock research inside proprietary formats. When an examiner asks for an adjustment or an assumption check, rebuilding the analysis from scratch often becomes a nightmare.

This book was written to offer a different path.

R has evolved into the lingua franca of modern statistical science and reproducible research. Paired with the tidyverse and modern reporting frameworks like Quarto, R gives researchers complete command over their analytical journey—from the first row of survey data to polished, APA-style tables and high-resolution figures.

Modern Research in the Age of AI

We now live and research in the age of artificial intelligence. Large language models and AI coding assistants have permanently transformed how we write code, brainstorm models, and explore data. Throughout this book, and particularly in Chapter 18, we treat AI not as a replacement for critical thought, but as a tireless research collaborator. You will learn to use AI assistants to draft and debug R scripts, translate complex errors, and assist with qualitative coding, while maintaining the rigorous verification standards required for academic defense.

How to Journey Through This Book

Rather than presenting mathematical abstractions in isolation, every statistical and machine learning method in this volume is anchored in a single, realistic running case study following 600 graduate students over four semesters. Every chapter answers a real research question.

For each concept, you will follow a disciplined, five-step roadmap: 1. **A Concrete Research Question** 2. **A Minimal Example** to build immediate intuition 3. **The Case-Study Analysis** with real data 4. **In Your Field** applications (health, agriculture, business, and social sciences) 5. **Thesis Reporting & Interpretation** showing exactly how to write up your findings

The Interactive Playground

To ensure that no software installation stands between you and your first analysis, all companion exercises are hosted in an interactive browser environment. You can test code, experiment with parameters, and verify your answers directly on the book's companion website:

Web Playground: https://polla-fattah.github.io/data2thesis_r/playground/

Whether you are embarking on your Master's dissertation, finishing your doctoral dissertation, or publishing your first peer-reviewed paper, may this book serve as your steady desk companion from your first command to your final defense.

Dr. Polla Fattah

Salahaddin University-Erbil

1	Getting Started with R	25
1.1	Research written in code	26
1.2	Installing R and RStudio	27
1.3	The RStudio window	27
1.4	Calculations	28
1.5	Storing values in objects	28
1.6	Types of data	29
1.7	Functions	30
1.8	Scripts	32
1.9	RStudio Projects	32
1.10	Packages	33
1.11	A first look at the data	33
1.12	Chapter review	36
1.13	Exercises	37
1.14	Further reading	37
2	Data Structures in R	39
2.1	Data as measurement	39
2.2	Variables and their values	40
2.3	Categories	42
2.4	A table of cases	43
2.5	Matrices and lists	46
2.6	Importing data	47
2.7	A codebook	50
2.8	Exporting data	51
2.9	SPSS and Excel equivalents	51
2.10	Error messages	52
2.11	A first summary	52
2.12	Chapter review	54
2.13	Exercises	54
2.14	Further reading	55
3	Data Manipulation	57
3.1	What clean data means	57
3.2	Tidy data	59
3.3	The tidyverse	59
3.4	Working with cases and variables	60
3.5	Reshaping between wide and long	65
3.6	Joining tables	66
3.7	Cleaning the survey export	68
3.8	Chapter review	77
3.9	Exercises	77
3.10	Further reading	78
4	Data Visualization	79
4.1	What a graph is for	80
4.2	How people read graphs	81
4.3	Choosing a graph	82
4.4	The grammar of graphics	83
4.5	The basic graphs	84
4.6	Mapping and setting	88
4.7	Preparing graphs for publication	89
4.8	Honest graphs	92
4.9	Further kinds of graph	94
4.10	Exporting figures	97
4.11	Common misconceptions	98
4.12	Chapter review	99
4.13	Exercises	99
4.14	Further reading	100

1. Getting Started with R

Every data analysis is a long chain of decisions. Which cases are kept and which are excluded, how a messy answer is corrected, which variables are combined into a score, which test is used, and how its result is rounded: each choice shapes the final numbers, and each should be open to inspection. When an analysis is done by clicking through menus, most of these decisions leave no trace. Weeks later, not even the researcher can say exactly how a number in the thesis was produced, and an examiner who asks cannot be given a complete answer.

Writing the analysis as code changes this. A script records every step in order, in a form that a supervisor can read, an examiner can check, and the researcher can run again after correcting a mistake or receiving new data. This is the main reason research is increasingly done in a programming language, and it is the reason this book teaches R. The chapter starts from nothing: installing R, finding your way around the program used to write it, and learning the handful of ideas on which everything else rests.

The book follows one research project from its first day to its last. Elaf has just begun a Master's degree in Educational Psychology. During her first year of graduate study she noticed how tired many of her fellow students were: they slept little, lived on coffee, worried about their supervisors and their money, and a few spoke quietly about leaving. With her supervisor, she turned that observation into a thesis. It asks how sleep, stress, and supervisor support relate to graduate students' wellbeing, their grades, and their thoughts of dropping out, and whether a short wellbeing workshop can help. Her university agreed to support a two-year study of 600 graduate students from five faculties: a questionnaire at the start, a record of each student's sleep, study, and wellbeing at the end of every semester, and a six-week workshop to which half of the students would be invited at random.

At their first meeting, her supervisor set out what the thesis would demand. The research questions had to become precise hypotheses before any results were seen. The raw survey export, full of test entries, typing errors, and impossible values, had to be cleaned without losing track of a single change. The sample had to be described honestly, including the students who stopped taking part. Every analysis had to be chosen for a reason and its assumptions checked, and every result reported with its size and its uncertainty, not only with a p-value. At the end, examiners would read the results chapter and could ask how any number in it had been produced. "You should be able to show them," the supervisor said, "step by step."

Elaf had used Excel for coursework and had once clicked through an SPSS menu in a statistics class, but she

had never written a line of code. Her supervisor suggested R, for exactly the reason given above: an analysis written as code can be shown step by step. The first data has now arrived, and this chapter is Elaf's first day with R, and yours. Each later chapter takes her one step further, from cleaning the data to the finished thesis. By the end of this one, you will have R running on your computer, you will know your way around the program used to write it, and you will have opened her data and answered a first, simple question about it.

💡 By the end of this chapter you will be able to

- Explain why researchers use R, and install R and RStudio.
- Find your way around RStudio and run code in the Console and in a script.
- Store values in objects and recognise the basic types of data.
- Use functions and their arguments, and get help when you are stuck.
- Organise your work in an RStudio Project and install packages.
- Open the student wellbeing data and answer a first question about it.

Note

Teaching & Practice: [Interactive Lecture Slides](#) · [Browser Playground Exercises](#)

1.1 - Research written in code

R is a free program for working with data: cleaning it, analysing it, and turning it into tables and graphs. It began in the early 1990s at the University of Auckland, where two statisticians, Ross Ihaka and Robert Gentleman, wrote it for teaching [1]. Its first stable version appeared in 2000. Today it is used in universities, hospitals, governments, and companies around the world.

Four properties make R worth learning for a researcher. It is made for data: statistical tests, models, and graphs are part of the language rather than add-ons. It is free and open source, so you, your students, and anyone who wants to check your work can use it without buying a licence. Its instructions form a written record of the analysis, which anyone can run again to obtain exactly the same results; research that can be repeated in this way is called **reproducible**, and Chapter 17 returns to the idea. Finally, it is shared by a large community. Thousands of researchers publish free **packages**, add-ons that give R new abilities, so whatever method your field uses, someone has probably written a package for it.

The price is that R asks you to type instead of click. That feels slow at first. It becomes fast surprisingly quickly, and it pays you back every time you need to repeat, correct, or extend an analysis.

i For users of SPSS or Excel

In SPSS or Excel, you click, and the program changes your data or produces output. In R, you write an instruction, and R carries it out. The instructions are saved in a file, so the next time you need the same analysis, you run the file instead of clicking through the menus again. Chapter 2 shows how familiar SPSS and Excel tasks map to R.

1.2 - Installing R and RStudio

Working with R involves two programs. **R** is the engine that does the work. **RStudio** is the program in which you drive it: you write R code there, run it, and see the results. You will almost never open R itself; you open RStudio, and RStudio uses R for you.

R must be installed first. Download it from cran.r-project.org, choosing the version for your operating system (Windows, macOS, or Linux), and install it like any other program. Then download the free RStudio Desktop from posit.co and install it. Appendix A walks through both installations step by step, with solutions to common problems.

i Other editors

RStudio is the most widely used editor for R, and this book uses it throughout. **Positron**, a newer editor from the same company, works with R and Python in one window and is a good alternative if you plan to use both languages. Everything in this book works in either.

💡 Working in the browser

The exercises for this chapter run in your web browser. Try them in the [playground](#) first, and install R when you are ready.

1.3 - The RStudio window

When you open RStudio, the window is divided into panes. With a script open (you will open one shortly), there are four, described in Table 1.1.

Table 1.1: The four panes of RStudio

Where	Pane	What it is for
Top left	Source	Where you write and save your code, in files called scripts.
Bottom left	Console	Where R runs code and shows the results.
Top right	Environment	The objects you have created, such as your data.
Bottom right	Files, Plots, Packages, Help	Your files, your graphs, your installed packages, and R's help pages.

You can type code directly into the Console. Click in it, type `2 + 2`, and press Enter. R answers straight away. The `[1]` at the start of the answer simply means “this is the first value of the result”; you can ignore it for now.

1.4 - Calculations

The simplest use of R is as a calculator. Suppose you slept 6.5, 7, and 5.5 hours on three nights last week. Your average is:

```
(6.5 + 7 + 5.5) / 3
```

```
[1] 6.333333
```

R follows the usual order of operations: multiplication and division before addition and subtraction, and brackets first of all. Without the brackets, R would divide only the last number by 3:

```
6.5 + 7 + 5.5 / 3
```

```
[1] 15.33333
```

The usual arithmetic symbols work as you would expect: `+`, `-`, `*` (multiply), `/` (divide), and `^` (power).

1.5 - Storing values in objects

Typing the same numbers again and again is tiresome and invites mistakes. Instead, you can store a value under a name. The stored value is called an **object**, and you create it with the **assignment arrow** `<-`, which you can read as “gets”:

```
nights <- 3
```

Nothing is printed, but R now remembers that `nights` is 3. The object appears in the Environment pane, and you can use it by name:

```
nights
```

```
[1] 3
```

```
nights * 7
```

```
[1] 21
```

To store several values in one object, combine them with `c()`, which stands for *combine*:

```
sleep <- c(6.5, 7, 5.5)  
sleep
```

```
[1] 6.5 7.0 5.5
```

Chapter 2 explains these collections, called **vectors**, in detail. For now, notice how much clearer the code becomes when the numbers have a name:

```
sum(sleep) / nights
```

```
[1] 6.333333
```

Names follow a few rules. A name starts with a letter and can contain letters, numbers, dots, and underscores, such as `sleep`, `sleep_week1`, or `avg.sleep`. It cannot contain spaces, so an underscore takes their place: `sleep_hours`, not `sleep hours`. R is case-sensitive, which makes `Sleep` and `sleep` two different objects. Beyond these rules, the best names say what the object holds: `sleep_hours` is far better than `x`, and your future self will be grateful for it.

If you assign a new value to an existing name, the old value is replaced without warning:

```
nights <- 4  
nights
```

```
[1] 4
```

1.6 - Types of data

Research data is not all numbers. A survey records numbers (hours of sleep), words (the student's faculty), and yes-or-no answers (whether the student was invited to a workshop). R keeps track of the **type** of every value, because the type decides what can be done with it: hours of sleep can be averaged, but faculty names cannot. Table 1.2 lists the types you will meet most often.

Table 1.2: The most common types of data in R

Type	Holds	Example
numeric	numbers, with or without decimals	6.5, 600
character	text, written inside quotes	"Education"
logical	TRUE or FALSE	TRUE

The function `class()` tells you the type of an object:

```
hours <- 6.5
faculty <- "Education"
invited <- TRUE

class(hours)
```

```
[1] "numeric"
```

```
class(faculty)
```

```
[1] "character"
```

```
class(invited)
```

```
[1] "logical"
```

Text must be inside quotes. Without them, R thinks you mean an object with that name:

```
faculty <- Education
```

```
Error: object 'Education' not found
```

This is one of the most common errors for beginners. The message says exactly what went wrong: R looked for an object called `Education` and did not find one.

1.6.1 - Missing values

Real data always has gaps. A student skips a question, or a value is lost. R marks a missing value with `NA`, short for *not available*. It is neither zero nor empty text; it means “we do not know”:

```
sleep_with_gap <- c(6.5, NA, 5.5)
mean(sleep_with_gap)
```

```
[1] NA
```

If one value is unknown, the average is unknown too, so R returns `NA` rather than guessing. You will see shortly how to tell R to leave missing values out.

1.7 - Functions

You have already used several **functions**: `c()`, `sum()`, `mean()`, and `class()`. A function takes some input, does

something with it, and returns a result. You call a function by writing its name followed by brackets, with the input inside:

```
mean(sleep)
```

```
[1] 6.333333
```

```
max(sleep)
```

```
[1] 7
```

```
length(sleep)
```

```
[1] 3
```

Many functions accept more than one input. The inputs are called **arguments**, and they are separated by commas. The function `round()`, for example, takes a number and the number of decimal places to keep:

```
round(6.333333, digits = 1)
```

```
[1] 6.3
```

Arguments have names. You can leave the names out if you give the arguments in the expected order, so `round(6.333333, 1)` gives the same result, but writing the names makes your code easier to read.

The missing value from before can now be handled. The function `mean()` has an argument called `na.rm`, short for “NA remove”. Set it to `TRUE`, and R calculates the average of the values it does know:

```
mean(sleep_with_gap, na.rm = TRUE)
```

```
[1] 6
```

You can also put one function inside another. R works from the inside out:

```
round(mean(sleep), digits = 1)
```

```
[1] 6.3
```

1.7.1 - Getting help

Every function has a help page. Type a question mark before its name in the Console:

```
?mean
```

The page opens in the Help pane. Help pages are written for experienced users, so they can look dense at first. Start with three parts: **Usage** (how to call the function), **Arguments** (what each input means), and **Examples** at the bottom, which you can copy and run.

When a search engine or an AI assistant suggests code, treat it like advice from a knowledgeable stranger: often

right, sometimes wrong, and always worth checking. Chapter 18 shows how to use AI tools well.

1.8 - Scripts

Code typed in the Console is gone once you close RStudio. For real work, write your code in a **script**: a plain text file, ending in `.R`, that holds your instructions in order. A new script is created with **File > New File > R Script**, and it opens in the Source pane. You type one instruction per line, and run the current line, or the lines you have selected, with **Ctrl+Enter** (**Cmd+Enter** on a Mac): the code is sent to the Console, and the result appears there. **Ctrl+S** (**Cmd+S**) saves the script.

Lines that start with `#` are **comments**. R ignores them; they are notes for people. Use them to explain *why* you did something:

```
# Sleep last week, in hours per night
sleep <- c(6.5, 7, 5.5)

# Average, rounded for reporting
round(mean(sleep), digits = 1)
```

```
[1] 6.3
```

A script is the analysis written down. Months later, when a supervisor asks how a number was obtained, the script is the answer.

1.9 - RStudio Projects

Research involves many files: data, scripts, graphs, and drafts. An **RStudio Project** keeps everything for one piece of work together in one folder, and makes sure R looks for files in that folder.

Create one with **File > New Project > New Directory > New Project**, give it a name (for example, `wellbeing-thesis`), and choose where to put it. RStudio creates the folder, with a file ending in `.Rproj` inside. From then on, double-click that file to open the project, and RStudio starts in the right place with the right files.

The folder R looks in for files is called the **working directory**. In a project, it is the project folder, so a file stored there can be opened by its name alone:

```
students <- read.csv("students.csv")
```

As a project grows, it needs subfolders, such as `data/` for data and `figures/` for graphs. The **here** package builds file paths that start from the project folder, so the same code works on any computer:

```
library(here)
students <- read.csv(here("data",
"students.csv"))
```

 Avoid `setwd()`

Older tutorials start scripts with `setwd("C:/Users/Elaf/Documents/thesis")`. That line works only on the computer where it was written. Use a project, and your code works for your supervisor too.

1.10 - Packages

R comes with a lot built in, but much of its power comes from **packages**. Using a package takes two steps. It is installed once on your computer with `install.packages()`, which downloads it from CRAN, R's official collection of packages:

```
install.packages("here")
```

It is then loaded, with `library()`, in every session in which you want to use it:

```
library(here)
```

A useful comparison is a library book: installing a package is like buying a book and putting it on your shelf, and loading it is taking it off the shelf to read. You buy it once, but you take it down every time you need it.

The data for this book comes as a package too, called **data2thesis**. It is not on CRAN, so you install it from this book's website:

```
install.packages("https://polla-fattah.github.io/data2thesis_r/downloads/data2thesis_1.1.0.tar.gz",  
                 repos = NULL, type = "source")
```

1.11 - A first look at the data

With R, RStudio, and the package installed, the data can be loaded:

```
library(data2thesis)
```

 About Elaf's data

Elaf, her study, and her data are fictional. The data was generated by a computer program to look like a realistic survey of graduate students, so that every method in this book has something to find. It describes no real people, and its patterns (for example, that the workshop raised wellbeing, or that students who sleep more have higher grades) were built into the program by the author, not discovered. Nothing in this book is evidence about the wellbeing of real students, and the data and results must not be cited or used as findings about students, universities, or any real situation. The same applies to the counselling service's records and to the students' written answers.

The package contains several **data frames**: tables with one row per case and one column per variable, like a spreadsheet. The main one is `students`, with one row per student. Two functions give its size:

```
nrow(students)
```

```
[1] 600
```

```
ncol(students)
```

```
[1] 14
```

There are 600 students and 14 variables. The function `head()` shows the first six rows:

```
head(students)
```

```

      student_id supervisor_id age gender
faculty programme study_mode
1      S0001      SUP088  31 Female Health
Sciences Master's Full-time
2      S0002      SUP116  28 Female
Education PhD Full-time
3      S0003      SUP040  33 Male
Education PhD Full-time
4      S0004      SUP112  34 Female Health
Sciences Master's Part-time
5      S0005      SUP118  35 Female Social
Sciences Master's Part-time
6      S0006      SUP028  36 Male
Humanities PhD Full-time
      employment has_children lives_away
financial_worry workshop
1      None      No      Yes
2 Not invited
2 Part-time job      No      Yes
2      Invited
3      None      No      No
4      Invited
4 Full-time job      No      No
2      Invited
5      None      No      No
2      Invited
6 Part-time job      No      Yes
5 Not invited
      workshop_sessions considering_dropout
1      0      No
2      0      No
3      5      No
4      6      No
5      2      No
6      0      No

```

The names of the variables are listed by `names()`, and the help page, `?students`, describes each one.

```
names(students)
```

```

[1] "student_id"      "supervisor_id"
"age"

```

```
[4] "gender"                "faculty"
"programme"
[7] "study_mode"           "employment"
"has_children"
[10] "lives_away"           "financial_worry"
"workshop"
[13]
"workshop_sessions"    "considering_dropout"
```

To pick out one variable, write the data frame's name, a dollar sign, and the variable's name: `students$age` is the age of every student. A natural first question concerns the age of the students in the study:

```
mean(students$age)
```

```
[1] NA
```

The answer is `NA`. One student's age is missing, and R will not guess, but the solution is already familiar:

```
mean(students$age, na.rm = TRUE)
```

```
[1] 29.72454
```

The students are 29.7 years old on average. The number of missing values can be counted too: `is.na()` marks each missing value as `TRUE`, and `sum()` counts them, because R counts every `TRUE` as 1:

```
sum(is.na(students$age))
```

```
[1] 1
```

Only 1 age is missing. It is worth knowing: Chapter 3 shows why it is missing, and what to do about it.

Finally, `table()` counts how many times each value appears. It shows how many students come from each faculty, and how many were invited to the wellbeing workshop:

```
table(students$faculty)
```

```
      Education Health Sciences      Humanities
Natural Sciences      148      154
95      87
Social Sciences
      116
```

```
table(students$workshop)
```

```
Invited Not invited
300      300
```

Exactly 300 students, half of the study, were invited. That is not a coincidence: they were chosen at random, and Chapter 5 explains why random assignment matters so much for the conclusions a study can draw.

In your field: agriculture

R comes with dozens of real datasets for practice. `PlantGrowth` records the dried weight of plants grown under a control condition and two different treatments, from a classic experiment comparing crop yields. The same functions work on it:

```
head(PlantGrowth)
```

```
  weight group
1   4.17  ctrl
2   5.58  ctrl
3   5.18  ctrl
4   6.11  ctrl
5   4.50  ctrl
6   4.61  ctrl
```

```
table(PlantGrowth$group)
```

```
ctrl trt1 trt2
  10   10   10
```

```
mean(PlantGrowth$weight)
```

```
[1] 5.073
```

Type `data()` in the Console to see the full list of built-in datasets, and `?PlantGrowth` to read about this one.

1.12 - Chapter review

1.12.1 - Summary

- Every analysis is a chain of decisions. Writing it as code records each one, so the analysis can be checked and repeated: it becomes reproducible.
- R is a free language made for working with data. You write it in RStudio: the Source pane holds your scripts, the Console runs the code, the Environment pane lists your objects, and the fourth pane shows files, plots, packages, and help.
- `<-` stores a value in an object. `c()` combines several values.
- Every value has a type. The most common are numeric, character (in quotes), and logical (`TRUE` or `FALSE`). `NA` marks a missing value.
- Functions take arguments inside brackets. `?` opens a function's help page.
- Write your code in scripts, keep each piece of work in an RStudio Project, install packages once with `install.packages()`, and load them each session with `library()`.

1.12.2 - Key terms

R, RStudio, Console, script, comment, object, assignment, vector, data type, missing value (`NA`), function, argument, package, data frame, working directory, RStudio Project, reproducible research.

1.13 - Exercises

These short exercises check your understanding. The [playground](#) has more, with hints and solutions, which you can run in your browser or download as an RStudio project.

1. A supervisor sleeps 8, 7.5, 6, and 7 hours on four nights. Store the numbers in an object called `supervisor_sleep` and find the average, rounded to one decimal place.
2. Check the types of `"600"` and `600` with `class()`, and explain the difference.
3. Explain in your own words why `mean(c(4, NA, 6))` returns `NA`, and show how to get the average of the two known values.
4. Using `students`, find the youngest and the oldest student. (Hint: `min()` and `max()` also have an `na.rm` argument.)
5. Using `table()`, find out how many students are studying part-time (the variable is `study_mode`).
6. Think of an analysis you have done, or read about, with a point-and-click program. List three decisions in it that a reader could not check without a written record.

1.14 - Further reading

- *R for Data Science* [2], chapters “Workflow: basics”, “Workflow: scripts and projects”, and “Workflow: getting help”, covers the same ground with a different dataset.
- The [RStudio User Guide](#) describes every part of RStudio.

2. Data Structures in R

Data is the written result of measurement. Each value records one observation about one case: a student's age, the faculty she belongs to, how many hours she slept in a semester, her answer to a questionnaire item. Before any analysis, a researcher has to know what each value stands for, what kind of measurement produced it, and how the values are arranged. The same questions arise whether the data comes from a questionnaire, a laboratory instrument, or an administrative register, and they decide which summaries and tests make sense later.

Research data rarely arrives in one neat file. In the well-being study, the survey tool produced an Excel export, the university's records office sent the semester results as CSV files, and a colleague who helped with an earlier pilot study works in SPSS, so part of the data exists as an SPSS file too. Before any of these files can be opened, it helps to know how R holds data once it is inside. R has a small number of ways to hold data, called **data structures**. This chapter introduces them as the tools for recording measurements: vectors for the values of one variable, factors for categories, and data frames for whole tables of cases. It then shows how to bring files of every common format into R, and how to keep track of what each variable means.

💡 By the end of this chapter you will be able to

- Describe data as a table of cases and variables, and identify the unit of analysis.
- Create vectors and factors, and pick out and change their values.
- Explain what a data frame is, and select rows and columns from one.
- Recognise matrices and lists when you meet them.
- Import data from CSV, Excel, and SPSS files, keeping SPSS labels, and export data again.
- Build a codebook that links each variable to the question behind it.
- Map familiar SPSS and Excel tasks to R, and read R's most common error messages.

Note

Teaching & Practice: [Interactive Lecture Slides](#) · [Browser Playground Exercises](#)

2.1 - Data as measurement

Almost all research data can be arranged as a table in which each row is a **case** and each column a **variable**. A case is the thing being measured: a student, a patient,

a plot of land, a school. A variable is one characteristic measured on every case, such as age or faculty. Each cell then holds one value: the measurement of one variable on one case. This arrangement, sometimes called the **data matrix**, is the starting point of nearly every method in this book.

What counts as a case depends on the study, and the choice is called the **unit of analysis**. The wellbeing study has more than one. In the `students` table, each row is a student. In the `semesters` table, each row is one student in one semester, so each student appears up to four times:

```
head(semesters, 5)
```

	student_id	semester	gpa	sleep_hours
1	S0001	1	3.02	7.2
5				130
2	S0001	2	3.27	8.4
3				80
3	S0001	3	3.17	7.8
7				105
4	S0001	4	3.00	7.2
4				145
5	S0002	1	3.25	6.4
2				210
	supervisor_meetings	wellbeing		
1	0	74		
2	1	69		
3	3	71		
4	2	74		
5	3	60		

The first four rows all belong to student `S0001`, one for each semester. Confusing the two units is a common source of error. A question about students should count each student once; a question about change over time needs all the semester rows, together with methods that know that rows from the same student belong together (Chapter 10). Chapter 5 returns to the unit of analysis when research questions are turned into data.

2.2 - Variables and their values

In R, the values of one variable are held in a **vector**. You met vectors in Chapter 1: a set of values combined with `c()`. They are the basic building block of R, and almost everything else is built from them.

```
sleep <- c(6.5, 7, 5.5, 8, 6)
sleep
```

```
[1] 6.5 7.0 5.5 8.0 6.0
```

A vector holds values of **one type** only. If you mix types, R quietly converts everything to the most flexible type, which is usually text:

```
mixed <- c(6.5, "seven", 5.5)
mixed
```

```
[1] "6.5" "seven" "5.5"
```

```
class(mixed)
```

```
[1] "character"
```

The numbers are now text, in quotes, and they can no longer be averaged. This matters more than it seems. If one person in a survey types “seven” instead of 7, the whole column arrives in R as text. The Excel file later in this chapter shows exactly this.

2.2.1 - Operations on a whole vector

Most things you do to a vector happen to every value at once. Converting the sleep values from hours into minutes needs only one multiplication:

```
sleep * 60
```

```
[1] 390 420 330 480 360
```

Comparisons work the same way, and give one `TRUE` or `FALSE` for each value. The nights shorter than the recommended 7 hours are marked by:

```
sleep < 7
```

```
[1] TRUE FALSE TRUE FALSE TRUE
```

Because R counts `TRUE` as 1 and `FALSE` as 0, `sum()` counts the short nights, and `mean()` gives their share:

```
sum(sleep < 7)
```

```
[1] 3
```

```
mean(sleep < 7)
```

```
[1] 0.6
```

Three of the five nights, or 60%, were short. The same line of code would work just as well on 2,000 nights.

2.2.2 - Selecting values

Square brackets pick values from a vector by their position. Positions start at 1:

```
sleep[1] # the first night
```

```
[1] 6.5
```

```
sleep[c(1, 3)] # the first and third nights
```

```
[1] 6.5 5.5
```

```
sleep[-2] # every night except the second
```

```
[1] 6.5 5.5 8.0 6.0
```

Values can also be picked with a condition, by putting a TRUE/FALSE vector inside the brackets. R keeps the values where the condition is TRUE:

```
sleep[sleep < 7]
```

```
[1] 6.5 5.5 6.0
```

The line reads aloud as “sleep, where sleep is less than 7”. This way of selecting, called **logical indexing**, is one of the most useful ideas in R, and you will use it constantly.

2.2.3 - Changing values

Brackets on the left of the arrow change values. Suppose the second night turns out to have been 7.5 hours, not 7:

```
sleep[2] <- 7.5
sleep
```

```
[1] 6.5 7.5 5.5 8.0 6.0
```

2.3 - Categories

Research data is full of categories: faculty, gender, treatment group, agreement on a scale. R stores categories as **factors**. A factor looks like text but knows the complete set of possible values, called its **levels**:

```
faculty <- factor(c("Education",
  "Humanities",      "Education",      "Health
  Sciences"))
faculty
```

```
[1] Education      Humanities      Education
Health Sciences
Levels: Education Health Sciences Humanities
```

```
levels(faculty)
```

```
[1] "Education"      "Health Sciences"
"Humanities"
```

```
table(faculty)
```

```
faculty
  Education Health Sciences      Humanities
           2             1             1
```

By default, levels are in alphabetical order. When the categories have a natural order, set it yourself with the `levels` argument, so that tables and graphs show them in that order:

```
employment <- factor(
  c("None", "Full-time job", "Part-time job",
    "None"),
```

```
levels = c("None", "Part-time job", "Full-
time job")
)
table(employment)
```

```
employment
      None Part-time job Full-time job
      2         1         1
```

Factors matter most in statistics and graphs. When groups are compared in Chapter 7, or plotted in Chapter 4, R uses the factor's levels to decide which groups exist and in which order to show them.

2.3.1 - Kinds of measurement and R's types

The choice between a factor and a number is not a matter of convenience; it follows from the kind of measurement that produced the values. Categories without an order, such as faculty, are **nominal**. Categories with an order but uneven or unknown distances between them, such as none, part-time, and full-time employment, are **ordinal**. Measurements on a scale with equal distances, such as hours of sleep or a wellbeing index, are numeric. Table 2.1 shows how each kind is stored in R. Chapter 5 explains these levels of measurement in full, and why they decide which summaries are meaningful.

Table 2.1: Kinds of measurement and how R stores them

Kind of measurement	Example	Stored in R as
Categories, no order (nominal)	faculty, gender	factor
Ordered categories (ordinal)	employment, agreement on a scale	factor with levels in order, or an ordered factor
Quantities with equal distances	sleep hours, age, wellbeing	numeric
Yes or no	invited to the workshop	logical, or a factor with two levels

2.4 - A table of cases

Most research data is a table of this kind: one row for each case, one column for each variable. In R, such a table is a **data frame**. Each column is a vector, so each column has one type, but different columns can have different types.

A small data frame can be built with `data.frame()`:

```
pilot <- data.frame(
  student = c("S1", "S2", "S3", "S4"),
  faculty = c("Education", "Humanities",
```

```
"Education", "Health Sciences"),
  sleep = c(6.5, 7.5, 5.5, 8),
  invited = c(TRUE, FALSE, TRUE, FALSE)
)
pilot
```

	student	faculty	sleep	invited
1	S1	Education	6.5	TRUE
2	S2	Humanities	7.5	FALSE
3	S3	Education	5.5	TRUE
4	S4	Health Sciences	8.0	FALSE

The quickest way to understand a data frame is `str()`, short for *structure*. It lists every column with its type and first few values:

```
str(pilot)
```

```
'data.frame': 4 obs. of 4 variables:
 $ student: chr "S1" "S2" "S3" "S4"
 $ faculty: chr "Education" "Humanities"
"Education" "Health Sciences"
 $ sleep : num 6.5 7.5 5.5 8
 $ invited: logi TRUE FALSE TRUE FALSE
```

2.4.1 - Selecting columns and rows

A single column is picked with `$`, as in Chapter 1:

```
pilot$sleep
```

```
[1] 6.5 7.5 5.5 8.0
```

Square brackets work on data frames too, with two positions separated by a comma: rows first, then columns. An empty position means “all”:

```
pilot[2, ] # the second
row, all columns
```

	student	faculty	sleep	invited
2	S2	Humanities	7.5	FALSE

```
pilot[, c("student", "sleep")] # all rows,
two columns
```

	student	sleep
1	S1	6.5
2	S2	7.5
3	S3	5.5
4	S4	8.0

```
pilot[pilot$sleep < 7, ] # rows where
sleep is under 7
```

	student	faculty	sleep	invited
1	S1	Education	6.5	TRUE
3	S3	Education	5.5	TRUE

The last line combines a data frame with logical indexing: “pilot, the rows where sleep is under 7, all columns”. It is the R version of SPSS’s *Select Cases*.

2.4.2 - Adding columns

Assigning to a new column name adds a column:

```
pilot$short_sleep <- pilot$sleep < 7
pilot
```

	student		faculty	sleep	invited
	short_sleep				
1	S1	Education	6.5	TRUE	TRUE
2	S2	Humanities	7.5	FALSE	FALSE
3	S3	Education	5.5	TRUE	TRUE
4	S4	Health Sciences	8.0	FALSE	FALSE

2.4.3 - The student wellbeing data

The students data frame from the data2thesis package has the same structure as pilot, only larger:

```
str(students)
```

```
'data.frame': 600 obs. of 14 variables:
 $ student_id      : chr  "S0001" "S0002"
 "S0003" "S0004" ...
 $ supervisor_id   : chr  "SUP088" "SUP116"
 "SUP040" "SUP112" ...
 $ age             : int   31 28 33 34 35
 36 NA 25 29 26 ...
 $ gender          : chr   "Female" "Female"
 "Male" "Female" ...
 $ faculty         : chr   "Health Sciences"
 "Education" "Education" "Health Sciences" ...
 $ programme       : chr   "Master's" "PhD"
 "PhD" "Master's" ...
 $ study_mode      : chr   "Full-time" "Full-
 time" "Full-time" "Part-time" ...
 $ employment      : chr   "None" "Part-time
 job" "None" "Full-time job" ...
 $ has_children    : chr   "No" "No" "No"
 "No" ...
 $ lives_away      : chr   "Yes" "Yes"
 "No" "No" ...
 $ financial_worry : int    2 2 4 2 2 5 3 NA
 2 2 ...
 $ workshop        : chr   "Not invited"
 "Invited" "Invited" "Invited" ...
 $ workshop_sessions : int    0 0 5 6 2 0 4 0
 0 0 ...
 $ considering_dropout: chr   "No" "No" "No"
 "No" ...
```

In this listing, `chr` means character, `int` whole numbers, and `num` numbers with decimals. The categories, such as `faculty`, arrive as text, and they are converted into factors when they are needed as categories:

```
students$faculty <- factor(students$faculty)
levels(students$faculty)
```

```
[1] "Education"          "Health Sciences"
[2] "Humanities"         "Natural Sciences"
[3] "Social Sciences"
```

Selection works as before. The number of PhD students in Health Sciences, for example, is found with two conditions:

```
hs_phd <- students[students$faculty == "Health
Sciences" & students$programme == "PhD", ]
nrow(hs_phd)
```

```
[1] 44
```

Two symbols are new here. The double equals sign, `==`, means “is equal to”; a single `=` is used for arguments, so comparisons need two. The ampersand, `&`, means “and”: both conditions must be `TRUE`. Its partner `|` means “or”.

Looking at data like a spreadsheet

In RStudio, `View(students)` opens the data in a spreadsheet-style viewer, where you can scroll, sort, and filter. It is only for looking: changes you want to keep should be made with code, so they are recorded.

2.5 - Matrices and lists

Two more structures appear regularly, even though they are rarely built by hand.

2.5.1 - Matrices

A **matrix** is a grid of values of one type, usually numbers. Matrices appear most often as results. The function `cor()`, for example, calculates the correlations between several variables and returns them as a matrix. The correlations below use the first semester only, so that each student counts once:

```
first_semester <-
semesters[semesters$semester == 1, ]
vars <- first_semester[, c("gpa",
"sleep_hours", "study_hours", "wellbeing")]
correlations <- cor(vars, use =
"complete.obs")
round(correlations, 2)
```

	gpa	sleep_hours	study_hours	wellbeing
gpa	1.00	0.25	0.07	0.31
sleep_hours	0.25	1.00	-0.56	0.46
study_hours	0.07	-0.56	1.00	-0.24
wellbeing	0.31	0.46	-0.24	1.00

Every variable correlates perfectly with itself, hence the 1s on the diagonal. Chapter 6 explains how to read cor-

relations. For now, notice that values are picked from a matrix exactly as from a data frame, with `[row, column]`:

```
correlations["sleep_hours", "wellbeing"]
```

```
[1] 0.4602433
```

2.5.2 - Lists

A **list** is a container that can hold anything: vectors of different lengths, data frames, even other lists. Its parts are usually named:

```
student <- list(
  id       = "S0001",
  programme = "Master's",
  sleep    = c(7.2, 8.4, 7.8, 7.2)
)
student$sleep
```

```
[1] 7.2 8.4 7.8 7.2
```

Lists are rarely built by hand, but R's statistical functions return their results as lists. A preview of a test from Chapter 7, which examines whether students sleep 7 hours on average in their first semester, shows this:

```
result <- t.test(first_semester$sleep_hours,
  mu = 7)
names(result)
```

```
[1] "statistic"      "parameter"      "p.value"
"conf.int"       "estimate"
[6] "null.value"     "stderr"         "alternative"
"method"         "data.name"
```

```
result$estimate
```

```
mean of x
6.483502
```

The function `names()` lists everything the test calculated, and `$` picks out one part, here the students' average sleep. This is how a number from a statistical test is retrieved for a report.

2.6 - Importing data

Each file format has its own function for reading it, and every one of them gives back a data frame. The files used below come with the `data2thesis` package, and `data2thesis_example()` finds them on your computer. For your own data, you would put the file in your RStudio Project and give its name instead (see Chapter 1).

2.6.1 - CSV files

A CSV file (*comma-separated values*) is a plain text table, where commas separate the columns. It is the most

common format for sharing data, and R reads it with `read.csv()`:

```
path <- data2thesis_example("semesters.csv")
semester_file <- read.csv(path)
head(semester_file, 3)
```

```
      student_id semester   gpa sleep_hours
study_hours exercise_days caffeine_mg
1      S0001      1 3.02      7.2         26
5      130
2      S0001      2 3.27      8.4         17
3      80
3      S0001      3 3.17      7.8         25
7      105
      supervisor_meetings wellbeing
1              0          74
2              1          69
3              3          71
```

In your own project, this would simply be `read.csv("semesters.csv")`, or `read.csv(here::here("data", "semesters.csv"))` with a data folder.

2.6.2 - Excel files

The **readxl** package reads Excel files. Install it once with `install.packages("readxl")`:

```
library(readxl)
raw <- read_excel(data2thesis_example("wellbeing_raw.xlsx"))
dim(raw)
```

```
[1] 615  64
```

This is the export from the study's survey tool, and it shows what real survey exports look like. The function `dim()` gives its size: 615 rows (more than the 600 students) and 64 columns. The first rows are:

```
head(raw[, 1:6])
```

```
# A tibble: 6 × 6
  `Response ID` `Q1_Student ID` Q2_Age
Q3_Gender Q4_Faculty      Q5_Programme
<chr>      <chr>          <chr> <chr>
<chr>      <chr>
1 R0001     TEST          99    Female
<NA>       <NA>
2 R0002     test           <NA> <NA>
<NA>       <NA>
3 R0003     TEST2          30    Male
<NA>       <NA>
4 R0004     S0191          27    male
Humanities Masters
5 R0005     S0015          30    Male
Health Sciences Master's
6 R0006     S0474          38    F
health sciences Masters
```

The printout looks a little different from before, because `read_excel()` returns a **tibble**: a modern kind of data

frame that prints more compactly and shows each column's type under its name. It can be used exactly like a data frame.

Several problems are already visible. There are test responses (TEST), odd codes such as 99 for missing age, and the same answer spelled several ways (Male, male). Every column shown is <chr>, text, even the ages, because a few entries are not plain numbers. Chapter 3 is devoted to cleaning this file; for now, it is enough to be able to open it.

If a workbook has several sheets, choose one with the sheet argument, for example `read_excel("file.xlsx", sheet = "Semester 2")`.

2.6.3 - SPSS files

The **haven** package reads SPSS files (.sav), and also Stata and SAS files:

```
library(haven)
spss <- read_sav(data2thesis_example("wellbeing.sav"))
spss[1:3, c("student_id", "gender",
"faculty", "stress_1")]
```

```
# A tibble: 3 × 4
  student_id gender faculty
stress_1
  <chr>          <dbl+lbl> <dbl+lbl>
<dbl+lbl>
1 S0001      1 [Female] 2 [Health Sciences]
3 [Neutral]
2 S0002      1 [Female] 1 [Education]      4
[Agree]
3 S0003      2 [Male]   1 [Education]      4
[Agree]
```

SPSS files carry **labels**, and haven keeps them. Each value is stored as a number, with its label shown in square brackets: 2 [Health Sciences]. The question behind each variable is kept too:

```
attr(spss$stress_1, "label")
```

```
[1] "I feel unable to control important things
in my studies."
```

To analyse a labelled variable as categories, turn it into a factor with `as_factor()`, which uses the labels as levels:

```
table(as_factor(spss$faculty))
```

```
      Education Health Sciences      Humanities
Natural Sciences      148              154
95              87
Social Sciences      116
```

This is one of the great conveniences for anyone moving from SPSS: the work put into labelling the data is not lost.

2.7 - A codebook

A value such as 4 in the column `stress_1` means nothing on its own. It becomes a measurement only when the reader knows the question behind it (“I feel unable to control important things in my studies”), the scale of the answers (1 = strongly disagree, 5 = strongly agree), and how missing answers are recorded. A **codebook** is the document that records this for every variable: its name, its meaning, its type, its possible values, and how it was measured. It is the bridge between the questionnaire that the participants saw and the table that the researcher analyses, and most theses include one in an appendix.

The labels in the SPSS file already contain much of a codebook, and R can collect them into a table. The code below takes the label of every variable, and the type in which it is stored:

```
codebook <- data.frame(
  variable = names(spss),
  label    = sapply(spss, function(x) attr(x,
"label")),
  type     = sapply(spss, function(x) class(x)
[1])),
  row.names = NULL
)
head(codebook, 12)
```

label	type	variable
1 student_id	character	Student identifier
2 supervisor_id	character	Supervisor identifier
3 age	numeric	Age in years at baseline
4		gender
Gender haven_labelled		faculty
5		faculty
Faculty haven_labelled		Degree
6 programme		Degree
programme haven_labelled		Study
7 study_mode		Study
mode haven_labelled		Paid work
8 employment		Paid work
alongside study haven_labelled		
9 has_children		Has children
children haven_labelled		
10 lives_away		Moved away from family to study
11 financial_worry		How worried are you about money?
haven_labelled		
12 workshop		Randomly invited to the wellbeing workshop
haven_labelled		

The function `sapply()` applies a function to every column in turn and collects the results. A labelled categorical variable also records which number stands for which category, its **value labels**:

```
attr(spss$faculty, "labels")
```

	Education	Health Sciences	Humanities
	Natural Sciences		
		1	2
3		4	
	Social Sciences		
		5	

A codebook built in this way can be written to a file with the functions in the next section and completed by hand, adding the answer scales and any notes on how variables were coded. Keeping it next to the data means that anyone who opens the data, including the researcher a year later, can tell what every value means.

2.8 - Exporting data

A data frame is saved with the matching `write` function. Each of these creates a file in the project folder:

```
write.csv(students, "students_clean.csv",
row.names = FALSE) # CSV
writexl::write_xlsx(students,
"students_clean.xlsx") # Excel
(writexl package)
haven::write_sav(spss, "students_clean.sav")
# SPSS
```

The argument `row.names = FALSE` stops R from adding an extra column of row numbers to the CSV file.

Original data files should stay untouched. Cleaned or changed versions are saved under new names, and the script records how one became the other.

2.9 - SPSS and Excel equivalents

Most tasks done in SPSS or Excel have a direct equivalent in R. Table 2.2 collects the ones from this chapter and the last.

Table 2.2: SPSS and Excel tasks and their R equivalents

In SPSS or Excel	In R
Open a data file	<code>read_sav()</code> , <code>read_excel()</code> , <code>read.csv()</code>
Variable View: names, types, labels	<code>str(); attr(x, "label")</code> for a variable label
Data View	<code>View()</code>
Select Cases	<code>data[condition,]</code>
Compute Variable	<code>data\$new <- ...</code>
Frequencies	<code>table()</code>
Descriptives	<code>mean()</code> , <code>summary()</code>
Save As	<code>write_sav()</code> , <code>write_xlsx()</code> , <code>write.csv()</code>

The biggest change is not a particular command. In SPSS or Excel, you change the data and the change is saved; the steps you took are forgotten. In R, the data file stays as it was, and the steps are saved in your script. Run the script again, and you get the same result again.

2.10 - Error messages

Everyone who writes R sees error messages every day. They are not a sign of failure; they are R reporting precisely what it could not do. The message usually says what went wrong, and often where, so it should always be read. Four messages account for most of the errors a beginner meets.

Object not found. A name is misspelled, or the object was never created (perhaps the line that creates it was not run):

```
mean(slep)
```

```
Error: object 'slep' not found
```

Could not find function. A function name is misspelled, or its package has not been loaded with `library()`:

```
reed_csv("students.csv")
```

```
Error in reed_csv("students.csv"): could not  
find function "reed_csv"
```

Non-numeric argument. A calculation was attempted on text, often because a column arrived as text:

```
mixed * 2
```

```
Error in mixed * 2: non-numeric argument to  
binary operator
```

Undefined columns selected. A column name in brackets does not exist, often because of a typo or wrong capitals:

```
students[, "Faculty"]
```

```
Error in `[.data.frame`(students, ,  
"Faculty"): undefined columns selected
```

When the problem is not obvious, check spelling and capitals first, then check that every line above was run, then read the help page of the function. Copying the error message into a search engine works surprisingly often, because someone has almost always met it before.

2.11 - A first summary

With the data open, `summary()` gives a quick overview of every column: averages and ranges for numbers, counts for factors, and the number of missing values. Here it is for the first semester, one row per student:

```
summary(first_semester[, c("gpa",
"sleep_hours", "study_hours", "wellbeing")])
```

```
      gpa      sleep_hours      study_hours
wellbeing
Min.   :2.170   Min.   : 3.500   Min.   :
2.00   Min.   :24.00
1st Qu.:2.888   1st Qu.: 5.800   1st Qu.:18.00
1st Qu.:52.00
Median :3.100   Median : 6.500   Median :26.00
Median :60.00
Mean   :3.102   Mean    : 6.484   Mean    :27.64
Mean   :60.46
3rd Qu.:3.340   3rd Qu.: 7.200   3rd Qu.:37.00
3rd Qu.:69.00
Max.   :4.000   Max.    :10.000   Max.    :67.00
Max.   :93.00
                                     NA's    :6
NA's    :8
```

Already, the summary shows that the average student sleeps less than 7 hours, and that a few values are missing (NA's). Chapter 6 turns this quick look into a full description of the data.

In your field: health research

`ToothGrowth` is a real dataset that comes with R, from an experiment on the effect of vitamin C on tooth growth in 60 guinea pigs. Its structure is typical of experiments in health research: an outcome, a treatment group, and a dose.

```
str(ToothGrowth)
```

```
'data.frame':  60 obs. of  3 variables:
 $ len : num  4.2 11.5 7.3 5.8 6.4 10 11.2
11.2 5.2 7 ...
 $ supp: Factor w/ 2 levels "OJ","VC": 2
2 2 2 2 2 2 2 2 2 ...
 $ dose: num  0.5 0.5 0.5 0.5 0.5 0.5 0.5
0.5 0.5 0.5 ...
```

```
table(ToothGrowth$supp, ToothGrowth$dose)
```

```
      0.5  1  2
OJ   10 10 10
VC   10 10 10
```

The variable `supp` is already a factor with two levels, the two ways the vitamin was given (orange juice, OJ, or ascorbic acid, VC), and the table shows 10 animals for every combination of method and dose. The unit of analysis is the animal: each row is one guinea pig.

2.12 - Chapter review

2.12.1 - Summary

- Data is the record of measurements: a table of cases (rows) and variables (columns). The unit of analysis is what one row represents, and the same study can have more than one.
- A **vector** holds the values of one variable, all of one type. Mixing types turns everything into text.
- Most operations work on every value of a vector at once. `[ ]` picks values by position or by a condition (logical indexing).
- A **factor** holds categories with a fixed set of **levels**, whose order you can set. The kind of measurement decides whether a variable should be a factor or a number.
- A **data frame** is a table: one row per case, one column (a vector) per variable. Select with `$` or `[rows, columns]`, and use `str()` to see its structure.
- A **matrix** is a grid of one type, often a result such as a correlation matrix. A **list** can hold anything; statistical results are lists.
- `read.csv()`, `read_excel()` (readxl), and `read_sav()` (haven) import data. haven keeps SPSS labels, and `as_factor()` turns them into factors. `write.csv()`, `write_xlsx()`, and `write_sav()` export data.
- A **codebook** records what every variable means and how it was measured; SPSS labels provide much of it.
- Error messages say what R could not do. Read them, and check spelling and capitals first.

2.12.2 - Key terms

Case, variable, data matrix, unit of analysis, data structure, vector, type conversion, logical indexing, factor, level, data frame, tibble, matrix, list, CSV file, variable label, value label, codebook, error message.

2.13 - Exercises

The [playground](#) has these and more, with hints and solutions, in your browser or as an RStudio project to download.

1. Create a vector with the ages 24, 31, 28, 45, and 26. Use logical indexing to show only the ages over 30, and count how many there are.
2. Create a factor from `c("Agree", "Disagree", "Neutral", "Agree")` whose levels run from "Disagree" through "Neutral" to "Agree". Check the order with `table()`.
3. Using `students`, count the students who are part-time and have children. (Use `&`.)
4. Add a column to `students` called `over_30` that is `TRUE` for students older than 30. Count those students, and explain why the count might need `na.rm = TRUE`.
5. Read the SPSS file with `read_sav()` and find the question behind `support_3`.
6. State the unit of analysis of each table in the package: `students`, `semesters`, `questionnaire`, and

`supervisors`. For each, name one research question it could answer on its own.

2.14 - Further reading

- *R for Data Science* [2]: chapters “Data import”, “Spreadsheets”, and “Factors” go further with importing and with categories, and “A field guide to base R” covers the square-bracket style used in this chapter.
- The documentation of the `readxl` and `haven` packages.

3. Data Manipulation

Raw data is never ready for analysis. Survey tools record test responses alongside real ones, participants submit the same form twice, one person types “female” and another “F”, missing answers are stored as 99, and a slip of the finger turns an age of 25 into 250. None of this is unusual, and none of it is harmless: a single 99 among answers from 1 to 5 can shift an average noticeably, and a duplicated row counts one participant twice. Preparing data for analysis, usually called **data cleaning**, often takes longer than the analysis itself, and its decisions shape every result that follows.

Cleaning is therefore part of the analysis, not a chore before it, and it deserves the same care. This chapter first sets out what clean data means and the principles that guide cleaning decisions. It then introduces the tools for working with tables in R: choosing cases and variables, creating new variables, summarising by group, reshaping, and combining tables. Finally, it applies them, step by step, to the raw export of the wellbeing survey that Chapter 2 opened: the file with test responses, duplicated rows, the same answer spelled four ways, ages stored as text, and missing values coded as 99.

💡 By the end of this chapter you will be able to

- Explain what clean data means, and why every cleaning decision must be recorded.
- Explain what tidy data is, and recognise data that is not tidy.
- Chain steps together with the pipe `|>`.
- Choose cases and variables, sort, create new variables, and summarise, overall and by group.
- Reshape data between wide and long formats, and join tables that share an identifier.
- Clean a real survey export: remove test and duplicate rows, rename variables, fix inconsistent categories, convert text to numbers, and handle missing codes and impossible values.
- Compute questionnaire scale scores, including reversed items, and report the cleaning in a thesis.

Note

Teaching & Practice: [Interactive Lecture Slides](#) · [Browser Playground Exercises](#)

3.1 - What clean data means

Data is clean when it can be trusted to represent what was measured. That broad idea can be broken into five qualities, each of which the survey export fails in some

way. Clean data is **valid**: every value is possible for its variable, so there are no ages of 250 or 26 hours of sleep a night. It is **accurate**: values record what the participant actually meant, so a missing-answer code is not mistaken for a real answer. It is **complete** as far as possible, and where values are missing, they are marked as missing rather than hidden behind codes. It is **consistent**: the same answer is always recorded in the same way, so “F”, “female”, and “Female” become one category. And it is **unique**: each case appears exactly once, with no test entries and no duplicates.

Each failure distorts results in a different way. A missing-answer code that stays in the data is the easiest to see. Suppose five students answered a question on a 1-to-5 scale, and one of them skipped it, which the survey tool recorded as 99:

```
library(dplyr)
```

```
Attaching package: 'dplyr'
```

```
The following objects are masked from
'package:stats':
```

```
filter, lag
```

```
The following objects are masked from
'package:base':
```

```
intersect, setdiff, setequal, union
```

```
answers <- c(4, 3, 5, 99, 2)
mean(answers)
```

```
[1] 22.6
```

```
mean(na_if(answers, 99), na.rm = TRUE)
```

```
[1] 3.5
```

With the code left in, the average is 22.6, a value that is impossible on a 1-to-5 scale. Marked as missing with `na_if()`, the code drops out, and the average of the four real answers is 3.5. In a real dataset the distortion is usually smaller and therefore harder to notice, which is exactly what makes it dangerous. Duplicates work in a similar way, but on the sample size and the balance of the sample: a participant who appears twice counts twice.

Three principles follow for every cleaning decision. The **raw file is never edited**; it stays exactly as it was received, and all changes are made by code that starts from it. **Every decision is recorded**, in the script and, where it affects the results, in the thesis: how many cases were removed and why, which values were set to missing, how scale scores were built. And **the result is checked after every step**, because many cleaning mistakes produce no error message at all. The final section of this chapter shows how to report the cleaning in a thesis.

3.2 - Tidy data

Clean data should also be arranged in a way that makes analysis straightforward. The arrangement used throughout this book is called **tidy data** [3]. In tidy data, each variable forms a column, each observation forms a row, and each value sits in its own cell. The idea is easiest to see in a small example. Table 3.1 stores two semesters of wellbeing for three students side by side, as a spreadsheet often would.

Table 3.1: Wellbeing in wide format: one variable, wellbeing, is spread over two columns

student	wellbeing_s1	wellbeing_s2
S1	58	63
S2	71	70
S3	64	69

The table looks neat, but one variable, wellbeing, is spread over two columns, and a second variable, the semester, is hidden inside the column names. Table 3.2 holds the same six measurements in tidy form.

Table 3.2: The same data in tidy (long) format: one column per variable, one row per observation

student	semester	wellbeing
S1	1	58
S1	2	63
S2	1	71
S2	2	70
S3	1	64
S3	2	69

Now each of the three variables has its own column, and each row is one observation: one student in one semester. Questions such as “average wellbeing in each semester” become a simple grouped summary. The `students` table in the package is tidy, with one row per student and one column per variable. The survey export is not: it squeezes four semesters of measurements side by side into one row per student. Much of cleaning consists of making data tidy, because once it is tidy, every tool in this chapter works on it.

3.3 - The tidyverse

The tools in this chapter come from the **tidyverse**, a collection of packages that share one way of working and are designed for tidy data. This chapter uses four of them: **dplyr** for working with rows and columns, **tidyr**

for reshaping data, **stringr** for working with text, and **readr** for turning text into numbers (and for reading and writing files). The whole collection can be installed at once with `install.packages("tidyverse")` and loaded with `library(tidyverse)`, but this book loads only the packages each chapter needs:

```
library(dplyr)
library(tidyr)
library(stringr)
library(readr)
```

3.3.1 - Chaining steps with the pipe

Analysis is a series of steps: take the data, keep some rows, calculate something, round the result. In Chapter 1, such steps were written by putting functions inside each other:

```
sleep <- c(6.5, 7, 5.5, 8, 6)
round(mean(sleep), 1)
```

```
[1] 6.6
```

This reads from the inside out, which gets hard to follow as the steps multiply. The **pipe**, written `|>`, writes the steps in the order they happen. It takes the result on its left and passes it as the first argument to the function on its right:

```
sleep |> mean() |> round(1)
```

```
[1] 6.6
```

The pipe reads as “and then”: take `sleep`, *and then* take the mean, *and then* round it to one decimal place. In RStudio, **Ctrl+Shift+M** (**Cmd+Shift+M** on a Mac) types the pipe for you.

The older pipe

Before R had its own pipe, the tidyverse used `%>%`, from the `magrittr` package. You will see it in many tutorials and older scripts. For everything in this book, it does the same job as `|>`.

3.4 - Working with cases and variables

The `dplyr` package provides a small set of **verbs**, each doing one job. Every verb takes a data frame first and returns a data frame, so verbs can be chained with the pipe. The verbs are easiest to understand on a tiny table. The function `tibble()` builds one; a **tibble** is the tidyverse’s version of a data frame, which prints more neatly:

```
pilot <- tibble(
  student = c("S1", "S2", "S3", "S4", "S5"),
  faculty   = c("Education",
```

```
"Humanities", "Education", "Health Sciences",
"Humanities"),
  sleep = c(6.5, 7.5, 5.5, 8, 6),
  stress = c(3.2, 2.1, 4.5, 1.8, 3.9)
)
pilot
```

```
# A tibble: 5 × 4
  student faculty    sleep stress
  <chr>   <chr>      <dbl> <dbl>
1 S1     Education    6.5    3.2
2 S2     Humanities    7.5    2.1
3 S3     Education    5.5    4.5
4 S4     Health Sciences  8      1.8
5 S5     Humanities    6      3.9
```

3.4.1 - Choosing the cases you need

Most analyses use only some of the cases: the students of one faculty, the first semester, the participants who completed the study. The verb `filter()` keeps the rows that meet a condition:

```
pilot |> filter(sleep < 7)
```

```
# A tibble: 3 × 4
  student faculty    sleep stress
  <chr>   <chr>      <dbl> <dbl>
1 S1     Education    6.5    3.2
2 S3     Education    5.5    4.5
3 S5     Humanities    6      3.9
```

Several conditions separated by commas must all be true:

```
pilot |> filter(sleep < 7, faculty ==
"Education")
```

```
# A tibble: 2 × 4
  student faculty    sleep stress
  <chr>   <chr>      <dbl> <dbl>
1 S1     Education    6.5    3.2
2 S3     Education    5.5    4.5
```

To keep rows matching any of several values, use `%in%`:

```
pilot |> filter(faculty %in% c("Education",
"Humanities"))
```

```
# A tibble: 4 × 4
  student faculty    sleep stress
  <chr>   <chr>      <dbl> <dbl>
1 S1     Education    6.5    3.2
2 S2     Humanities    7.5    2.1
3 S3     Education    5.5    4.5
4 S5     Humanities    6      3.9
```

Every filter is a decision about the sample, and should be reported: an analysis of “students who completed all four semesters” answers a different question from an analysis of all students.

3.4.2 - Choosing variables

Large datasets have far more variables than any one analysis needs. The verb `select()` keeps the columns you name, in the order you name them, and a minus sign drops a column instead:

```
pilot |> select(student, sleep)
```

```
# A tibble: 5 × 2
  student sleep
  <chr>   <dbl>
1 S1      6.5
2 S2      7.5
3 S3      5.5
4 S4       8
5 S5       6
```

```
pilot |> select(-faculty)
```

```
# A tibble: 5 × 3
  student sleep stress
  <chr>   <dbl> <dbl>
1 S1      6.5   3.2
2 S2      7.5   2.1
3 S3      5.5   4.5
4 S4       8   1.8
5 S5       6   3.9
```

3.4.3 - Sorting cases

Sorting helps when looking at data, for example to see the most extreme values first. The verb `arrange()` sorts rows by one or more columns, and `desc()` sorts from largest to smallest:

```
pilot |> arrange(desc(stress))
```

```
# A tibble: 5 × 4
  student faculty      sleep stress
  <chr>   <chr>     <dbl> <dbl>
1 S3      Education    5.5   4.5
2 S5      Humanities    6     3.9
3 S1      Education    6.5   3.2
4 S2      Humanities    7.5   2.1
5 S4      Health Sciences  8     1.8
```

3.4.4 - Creating new variables

Many variables used in analysis are calculated from others: a total score, a conversion to other units, a category derived from a number. The verb `mutate()` adds new columns calculated from existing ones, or changes existing ones:

```
pilot |> mutate(
  sleep_minutes = sleep * 60,
  short_sleep   = sleep < 7
)
```



```
# A tibble: 5 × 6
  student faculty      sleep stress
sleep_minutes short_sleep
  <chr>      <chr>      <dbl>  <dbl>
<dbl> <lg1>
1 S1      Education      6.5    3.2
390 TRUE
2 S2      Humanities      7.5    2.1
450 FALSE
3 S3      Education      5.5    4.5
330 TRUE
4 S4      Health Sciences    8      1.8
480 FALSE
5 S5      Humanities      6      3.9
360 TRUE
```

To sort values into categories, `case_when()` checks conditions in order and uses the first one that is true, and `.default` catches everything else:

```
pilot |> mutate(
  stress_level = case_when(
    stress >= 4 ~ "High",
    stress >= 2.5 ~ "Medium",
    .default = "Low"
  )
)
```

```
# A tibble: 5 × 5
  student faculty      sleep stress
stress_level
  <chr>      <chr>      <dbl>  <dbl> <chr>
1 S1      Education      6.5    3.2 Medium
2 S2      Humanities      7.5    2.1 Low
3 S3      Education      5.5    4.5 High
4 S4      Health Sciences    8      1.8 Low
5 S5      Humanities      6      3.9 Medium
```

Turning a number into categories in this way loses information, since a stress score of 3.9 and one of 2.6 become the same “Medium”. It is useful for description, but analyses should normally keep the original number.

3.4.5 - Summarising, overall and by group

Summaries reduce many values to a few numbers. The verb `summarise()` reduces a table to one row of summaries, and `n()` counts the rows:

```
pilot |> summarise(
  mean_sleep = mean(sleep),
  students   = n()
)
```

```
# A tibble: 1 × 2
  mean_sleep students
  <dbl>      <int>
1      6.7         5
```

The real power comes with groups. The `.by` argument calculates the summary separately for each group:

```
pilot |> summarise(
  mean_sleep = mean(sleep),
  students   = n(),
  .by = faculty
)
```

```
# A tibble: 3 × 3
  faculty      mean_sleep students
  <chr>          <dbl>     <int>
1 Education         6         2
2 Humanities       6.75         2
3 Health Sciences   8         1
```

For simple counts, `count()` is a shortcut:

```
pilot |> count(faculty)
```

```
# A tibble: 3 × 2
  faculty      n
  <chr>    <int>
1 Education     2
2 Health Sciences 1
3 Humanities     2
```

i Grouping in older code

In many tutorials you will see `group_by(faculty) |> summarise(...)` instead of the `.by` argument. Both give the same result. The `.by` argument is newer and keeps each step self-contained, so this book uses it.

3.4.6 - Combining the steps

Each verb is simple, but chained together they answer real questions. One such question concerns how average wellbeing changed over the four semesters, separately for students who were and were not invited to the workshop. The semester records are in `semesters`, and the workshop group is in `students`, so the workshop group is first added to the semester records. (The `left_join()` in the first line is explained in Section 3.6.)

```
semesters |>
  left_join(students |> select(student_id,
    workshop), join_by(student_id)) |>
  summarise(
    mean_wellbeing = mean(wellbeing),
    students       = n(),
    .by = c(semester, workshop)
  ) |>
  arrange(semester, workshop)
```

	semester	workshop	mean_wellbeing	students
1	1	Invited	60.64667	300
2	1	Not invited	60.26667	300
3	2	Invited	65.47000	300
4	2	Not invited	60.22000	300
5	3	Invited	63.22261	283
6	3	Not invited	59.22857	280

7	4	Invited	60.93286	283
8	4	Not invited	58.42500	280

Read aloud, the code says: take the semester records, *and then* add each student's workshop group, *and then* calculate the average wellbeing and the number of students for each semester and workshop group, *and then* sort the result. In semester 1, before the workshop, the two groups are almost level. From semester 2, the invited students are ahead. Chapter 7 tests whether that difference is larger than chance.

The number of students also drops in semesters 3 and 4. Some students left the study, which is worth remembering: it returns later in this chapter and in Chapter 6.

3.5 - Reshaping between wide and long

The same data can be laid out in two ways, as the section on tidy data showed. In **wide** format, repeated measurements sit side by side, one column per occasion. In **long** format, each measurement has its own row. Here is the wide table from Table 3.1 in R:

```
wide <- tibble(
  student = c("S1", "S2", "S3"),
  wellbeing_1 = c(58, 71, 64),
  wellbeing_2 = c(63, 70, 69)
)
wide
```

```
# A tibble: 3 × 3
  student wellbeing_1 wellbeing_2
  <chr>         <dbl>         <dbl>
1 S1             58             63
2 S2             71             70
3 S3             64             69
```

Spreadsheets and SPSS often use wide format. For analysis in R, long format is usually easier, because each variable (student, semester, wellbeing) is then one column. The function `pivot_longer()` turns wide into long:

```
long <- wide |>
  pivot_longer(
    cols = c(wellbeing_1, wellbeing_2),
    names_to = "semester",
    names_prefix = "wellbeing_",
    values_to = "wellbeing"
  )
long
```

```
# A tibble: 6 × 3
  student semester wellbeing
  <chr>    <chr>         <dbl>
1 S1      1             58
2 S1      2             63
3 S2      1             71
4 S2      2             70
```

5	S3	1	64
6	S3	2	69

The arguments say which columns to stack (`cols`), where the old column names should go (`names_to`), which part of the names to drop (`names_prefix`), and where the values should go (`values_to`). The result is Table 3.2.

The opposite function, `pivot_wider()`, is especially useful for presenting results, because a wide table is easier to read:

```
long |> pivot_wider(names_from = semester,
  values_from = wellbeing)
```

```
# A tibble: 3 × 3
  student `1` `2`
  <chr>   <dbl> <dbl>
1 S1      58    63
2 S2      71    70
3 S3      64    69
```

The average wellbeing table from the previous section, for example, reads more easily with one column per workshop group:

```
semesters |>
  left_join(students |> select(student_id,
    workshop), join_by(student_id)) |>
  summarise(mean_wellbeing =
    round(mean(wellbeing), 1), .by = c(semester,
    workshop)) |>
  pivot_wider(names_from = workshop,
    values_from = mean_wellbeing)
```

```
# A tibble: 4 × 3
  semester `Not invited` Invited
  <int>         <dbl>   <dbl>
1     1         60.3    60.6
2     2         60.2    65.5
3     3         59.2    63.2
4     4         58.4    60.9
```

3.6 - Joining tables

Research data is often spread over several tables that share an identifier. The wellbeing study keeps students in one table, their semester records in another, and their supervisors in a third. **Joins** combine tables by matching the identifier. Two small tables show how:

```
people <- tibble(
  student = c("S1", "S2", "S3"),
  faculty = c("Education", "Humanities",
    "Education")
)
scores <- tibble(
  student = c("S1", "S1", "S2", "S4"),
  wellbeing = c(58, 63, 71, 66)
)
```

The function `left_join()` keeps every row of the first table and adds the matching information from the second:

```
scores |> left_join(people, join_by(student))
```

```
# A tibble: 4 × 3
  student wellbeing faculty
  <chr>      <dbl> <chr>
1 S1          58 Education
2 S1          63 Education
3 S2          71 Humanities
4 S4          66 <NA>
```

The argument `join_by(student)` names the column that links the two tables. Student S4 has a score but no entry in `people`, so their faculty is NA.

The function `anti_join()` does the opposite: it keeps the rows of the first table that have *no* match in the second. It is the quickest way to find what is missing, here the students without any scores:

```
people |> anti_join(scores, join_by(student))
```

```
# A tibble: 1 × 2
  student faculty
  <chr>    <chr>
1 S3      Education
```

On the wellbeing data, joins answer questions that no single table can. Each student's supervisor has an academic rank, stored in `supervisors`, and a join brings it next to each student's answer about dropping out:

```
by_rank <- students |>
  left_join(supervisors |
    > select(supervisor_id, rank),
    join_by(supervisor_id)) |>
  summarise(
    considering_dropout =
  mean(considering_dropout == "Yes"),
    students = n(),
    .by = rank
  )
by_rank
```

	rank	considering_dropout	students
1	Lecturer	0.1367521	234
2	Professor	0.1217391	115
3	Assistant Professor	0.1752988	251

The shares differ a little, from about 12% to about 18%. Differences this small could easily be chance; Chapter 7 shows how to test that. A second join finds the students who have no records for semester 3:

```
left <- students |>
  anti_join(semesters |> filter(semester ==
3), join_by(student_id))
nrow(left)
```

```
[1] 37
```

```
left |> count(considering_dropout)
```

```
  considering_dropout  n
1                No  15
2                Yes  22
```

In all, 37 students have no semester 3 record: they left the study after the first year, and 22 of them had said they were considering dropping out. The students who left are not a random selection, which matters for any analysis of the later semesters. Chapter 6 returns to this.

3.7 - Cleaning the survey export

The survey export can now be cleaned. The steps below are the ones needed for almost any survey data, and each is preceded by the problem it solves. Each step is short. What matters is doing them in a sensible order, and checking the result after each one.

3.7.1 - Inspecting the export

Cleaning starts with looking, because problems that have not been seen cannot be fixed:

```
library(readxl)
raw <- read_excel(data2thesis_example("wellbeing_raw.xlsx"))
dim(raw)
```

```
[1] 615  64
```

The file has 615 rows, although there are only 600 students. A few lines of code show what needs fixing:

```
raw |> count(Q3_Gender)
```

```
# A tibble: 7 × 2
  Q3_Gender      n
  <chr>      <int>
1 F          28
2 Female    233
3 M          46
4 Male     173
5 female     56
6 male       78
7 <NA>        1
```

```
raw
|> filter(str_detect(str_to_upper(`Q1_Student
ID`), "TEST")) |> select(1:4)
```

```
# A tibble: 3 × 4
  `Response ID` `Q1_Student ID` Q2_Age
Q3_Gender
  <chr>      <chr>      <chr> <chr>
1 R0001     TEST       99   Female
2 R0002     test       <NA> <NA>
3 R0003     TEST2      30   Male
```

Gender is spelled six ways (F, female, Female, and so on), plus a blank from a test response, and there are test responses from when the survey was set up. (The raw file also contains "Female " with a trailing space; `read_excel()` removes such spaces automatically, but other import functions do not, which is why the cleaning below still uses `str_trim()`.) Column names such as `Q1_Student ID` contain spaces, so they must be written inside backticks (```) in code. And as Chapter 2 showed, many numeric columns arrived as text.

3.7.2 - Test responses and duplicates

Test responses are not data about students, and duplicated submissions would count some students twice, so both must go before anything else. Test responses have a student ID starting with "TEST", in upper or lower case. The function `str_to_upper()` makes the case irrelevant, and `str_detect()` checks for the pattern. The function `distinct()` then removes rows that are exact copies of another row. The response ID is different for each submission, even for duplicates, so it is dropped first:

```
responses <- raw |>
  filter(!str_detect(str_to_upper(`Q1_Student
ID`), "^TEST")) |>
  select(-`Response ID`) |>
  distinct()
nrow(responses)
```

```
[1] 600
```

In the pattern, `!` means "not" and `^` means "at the start". Exactly 600 rows are left: one per student.

3.7.3 - Usable variable names

Names such as `Q10_How worried are you about money? (1-5)` are awkward to type and easy to get wrong. Short, consistent names make the code readable and reduce errors. The function `rename()` changes names, new name first, and the 22 questionnaire items, `Q11_1` to `Q11_22`, are renamed in one go with `rename_with()`, using the scale each item belongs to:

```
item_names <- c(paste0("stress_", 1:6),
  paste0("burnout_", 1:6),
  paste0("support_", 1:6),
  paste0("satisfaction_", 1:4))

responses <- responses |>
  rename(
    student_id = `Q1_Student ID`,
    age        = Q2_Age,
    gender     = Q3_Gender,
    faculty    = Q4_Faculty,
    programme  = Q5_Programme,
    study_mode = `Q6_Study mode`,
    employment = `Q7_Paid work`,
    has_children = Q8_Children,
    lives_away  = `Q9_Moved away`
```

```

from family`,
  financial_worry    = `Q10_How worried are
you about money? (1-5)`,
  workshop           = `Workshop group`,
  workshop_sessions  = `Workshop sessions
attended`,
  considering_dropout = `Y1_Considered
leaving?`
) |>
  rename_with(~ item_names, .cols =
Q11_1:Q11_22)

```

The call `paste0("stress_", 1:6)` builds the names `stress_1` to `stress_6`, which saves typing them all. The questionnaire's wording is not lost: it belongs in the codebook (Chapter 2).

3.7.4 - Consistent categories

If “F”, “female”, and “Female” stay separate, every table and every test treats them as three different groups. Categories are fixed by looking for a pattern rather than listing every spelling, after `str_to_lower()` and `str_trim()` have removed differences in case and stray spaces:

```

responses <- responses |>
  mutate(
    gender = if_else(str_starts(str_to_lower(str_trim(gender))),
" f"), "Female", "Male"),
    faculty = case_when(
      str_detect(str_to_lower(faculty), "educ")
~ "Education",
      str_detect(str_to_lower(faculty),
"health") ~ "Health Sciences",
      str_detect(str_to_lower(faculty),
"humanit") ~ "Humanities",
      str_detect(str_to_lower(faculty),
"social") ~ "Social Sciences",
      str_detect(str_to_lower(faculty),
"natural|^sciences$") ~ "Natural Sciences"
    ),
    programme =
if_else(str_detect(str_to_lower(programme),
"ph"), "PhD", "Master's"),
    study_mode =
if_else(str_detect(str_to_lower(study_mode),
"part"), "Part-time", "Full-time"),
    employment =
if_else(str_to_lower(employment) %in%
c("none", "no job"), "None", employment),
      across(c(has_children, lives_away,
considering_dropout),
~
if_else(str_starts(str_to_lower(.x), "y"),
"Yes", "No"))
    )
responses |> count(gender)

```

```

# A tibble: 2 × 2
  gender      n
  <chr>   <int>

```



```
1 Female    312
2 Male      288
```

```
responses |> count(faculty)
```

```
# A tibble: 5 × 2
  faculty      n
  <chr>    <int>
1 Education    148
2 Health Sciences 154
3 Humanities    95
4 Natural Sciences 87
5 Social Sciences 116
```

Two details deserve attention. The order of conditions in `case_when()` matters: “social sciences” also contains “sciences”, so Social Sciences is checked before Natural Sciences. And `across()` applies the same change to several columns at once, here turning `Yes`, `yes`, and `Y` into `Yes`.

After each fix, the categories should be counted again, as above. If a spelling was missed, it shows up immediately.

3.7.5 - Missing codes and numbers stored as text

Survey tools often mark missing answers with codes such as `99` or `-9`. As the opening example showed, these codes must become `NA` before any calculation, or R will treat them as real values. The function `na_if()` turns a code into `NA`, and `as.integer()` then turns the text into whole numbers:

```
responses <- responses |>
  mutate(
    financial_worry = financial_worry |>
      na_if("99") |> na_if("-9") |> as.integer(),
    across(all_of(item_names), ~
      as.integer(na_if(.x, "99"))),
    workshop_sessions =
      as.integer(workshop_sessions),
    age = parse_number(age)
  )
```

⚠ Do not replace 99 everywhere

It is tempting to turn every `99` in the file into `NA` in one go. But a `99` is not always a code. In a wellbeing score from 0 to 100, it is a perfectly real value. Replace missing codes only in the columns where you know they are codes.

The function `parse_number()` from `readr` reads the number out of a text value, ignoring extra text around it. It is useful for the sleep columns, where some students typed “7 hrs”. A decimal comma, however, defeats it:

```
parse_number(c("7 hrs", "6.5", "6,5"))
```

```
[1] 7.0 6.5 65.0
```

The value "6,5" becomes 65, not 6.5, because in English a comma separates thousands. The comma has to be changed to a point first:

```
parse_number(str_replace(c("7 hrs", "6.5", "6,5"), ",", "."))
```

```
[1] 7.0 6.5 6.5
```

Mistakes like this produce no error message. The only protection is to check the data after each step.

3.7.6 - Impossible values

Some values cannot be true, such as an age of 250 or 26 hours of sleep a night. Left in the data, a single one can dominate an average or a regression. They are typing mistakes, and since the true value cannot be known, the honest fix is to set them to `NA` rather than to guess (an age of 250 might have been 25, or 50):

```
responses |> filter(age > 100) |>
select(student_id, age)
```

```
# A tibble: 1 × 2
  student_id age
  <chr>      <dbl>
1 S0007      250
```

```
responses <- responses |> mutate(age =
  if_else(age > 100, NA, age))
```

This is where the missing age in Chapter 1 came from. Unusual but possible values, such as a student who sleeps 4 hours a night, are a different matter: they are real data and are kept. Chapter 6 discusses how to handle them.

3.7.7 - Scale scores

A questionnaire scale is scored as the average of its items, and the items must all point in the same direction before they are averaged. One stress item, `stress_4` ("I feel confident handling problems in my studies"), is worded in the opposite direction to the others: agreeing with it means *less* stress. Averaged as it stands, it would cancel part of what the other five items measure. It must first be **reversed**, so that 1 becomes 5, 2 becomes 4, and so on. On a 1-to-5 scale, subtracting from 6 does exactly that:

```
questionnaire_clean <- responses |>
  select(student_id, all_of(item_names))

scores <- questionnaire_clean |>
  mutate(
    stress_4 = 6 - stress_4,
    stress_score =
  rowMeans(pick(stress_1:stress_6), na.rm =
    TRUE),
    burnout_score =
```

```

rowMeans(pick(burnout_1:burnout_6), na.rm =
TRUE),
      support_score
      =
rowMeans(pick(support_1:support_6), na.rm =
TRUE),
      satisfaction_score = rowMeans(pick(satisfaction_1:satisfaction_4),
na.rm = TRUE)
) |>
  select(student_id, ends_with("_score"))
head(scores)

```

```

# A tibble: 6 × 5
  student_id stress_score burnout_score
support_score satisfaction_score
  <chr>          <dbl>          <dbl>
<dbl>          <dbl>          <dbl>
1 S0191          2.83            1.5
3.83            3.33
2 S0015          4.5            3.83
1.67            3
3 S0474          3.5            4.5
1.4            2.75
4 S0418          3.17            3
1.67            3.75
5 S0049          3            3.5
3.17            3.33
6 S0537          2.83            2.5
3.2            4.75

```

The function `pick()` selects the columns of one scale, and `rowMeans()` averages each student's answers across them. With `na.rm = TRUE`, a student who skipped one item still gets a score from the items they answered. Whether that is acceptable, and how many skipped items is too many, is a decision to report in the thesis. Chapter 9 checks that the items really do measure four separate scales.

The original items should stay in the data, with the scores stored separately as here, because the items are needed again for checking the scales.

3.7.8 - From wide to long

The semester measurements sit side by side: `GPA_S1`, `Sleep_S1`, ..., `Wellbeing_S4`, 28 columns in all. The data is not tidy, since each row holds four observations, and analyses of change over time need one row per student per semester. This is a `pivot_longer()` with one extra idea: each column name holds two pieces of information, the variable (GPA) and the semester (1), separated by `_S`. The special name `.value` tells R to use the first piece as a column name:

```

semesters_clean <- responses |>
  select(student_id, matches("_S[1-4]$")) |>
  pivot_longer(
    cols      = -student_id,
    names_to  = c(".value", "semester"),
    names_sep = "_S"
  )
head(semesters_clean)

```

```
# A tibble: 6 × 9
  student_id semester GPA      Sleep Study
Exercise Caffeine Meetings Wellbeing
  <chr>      <chr>    <chr> <chr> <chr> <chr>
<chr>    <chr>    <chr>
1 S0191      1      3.07  5.6   31    3
140        2      64
2 S0191      2      3.29  5,7   30    4
140        7      65
3 S0191      3      3.35  5.5   32    3
130        5      60
4 S0191      4      3.58  6.6   29    2
155        7      69
5 S0015      1      2.8   6.5   15    2
210        2      34
6 S0015      2      3.2   6.5   18    0
275        5      38
```

The selection `matches("_S[1-4]$")` picks the columns whose names end in `_S1` to `_S4`. The rest is cleaning already familiar from the steps above: tidy names, text to numbers (with the decimal-comma fix for sleep), impossible values to `NA`, and finally removing the empty rows of students who had left the study:

```
semesters_clean <- semesters_clean |>
  rename(gpa = GPA, sleep_hours = Sleep,
         study_hours = Study,
         exercise_days = Exercise, caffeine_mg
         = Caffeine,
         supervisor_meetings = Meetings,
         wellbeing = Wellbeing) |>
  mutate(
    semester = as.integer(semester),
    sleep_hours =
      parse_number(str_replace(sleep_hours, ",",
                              ".")),
    across(c(gpa, study_hours, exercise_days,
             caffeine_mg, supervisor_meetings, wellbeing),
           parse_number),
    sleep_hours = if_else(sleep_hours > 24,
                          NA, sleep_hours),
    study_hours = if_else(study_hours > 168,
                          NA, study_hours),
    gpa = if_else(gpa > 4, NA, gpa)
  ) |>
  filter(!if_all(gpa:wellbeing, is.na))
nrow(semesters_clean)
```

```
[1] 2326
```

The condition `if_all(gpa:wellbeing, is.na)` is `TRUE` for rows where every measurement is missing, and `!` keeps the others.

3.7.9 - Checking the result

The last step checks what the cleaning has produced. The amount of missing data comes first, and `across()` inside `summarise()` counts the `NA`s in every column at once:

```
semesters_clean
|> summarise(across(everything(), ~
sum(is.na(.x))))

# A tibble: 1 × 9
  student_id semester    gpa sleep_hours
study_hours exercise_days caffeine_mg
    <int>      <int>    <int>      <int>
<int>      <int>      <int>
1         0         0      1         21
23        20        25
# i 2 more variables: supervisor_meetings
<int>, wellbeing <int>
```

Up to 25 values per measurement are still missing, about 1% of the records: answers students really did not give, and the impossible values set to NA. These are genuine gaps, and Chapter 6 discusses how to handle them.

The most satisfying check of all is a comparison with a trusted result. The clean tables in the `data2thesis` package were produced from this same export, so the cleaned data should match them exactly:

```
same <- function(mine, theirs) {
  mine <- as.data.frame(mine)[, names(mine)]
  theirs <- as.data.frame(theirs)[,
names(mine)]
  isTRUE(all.equal(mine, theirs,
check.attributes = FALSE))
}
same(questionnaire_clean |
> arrange(student_id),
questionnaire |> arrange(student_id))
```

```
[1] TRUE
```

```
same(semesters_clean |> arrange(student_id,
semester),
semesters |> arrange(student_id,
semester))
```

```
[1] TRUE
```

Both are TRUE. The cleaning is complete, and from here on the book uses the clean tables from the package.

💡 Writing it up

A thesis reports the cleaning briefly, usually in the methods chapter, so that a reader can judge what was done to the data. For the survey export, with the numbers filled in by the code:

The survey export contained 615 responses. Test responses (3) and duplicate submissions (12) were removed, leaving 600 students. Inconsistent spellings of categories were standardised, and missing-answer codes (99 and -9) were recoded as missing. Impossible values (4 in all), such as an age of 250 years or 26 hours of sleep a night, were set to missing, because the true values could not be known. The reverse-worded item `stress_4` was recoded before scale scores were computed as the mean of each scale's items, using all items a student had answered.

💡 Save your cleaning as a script

Put all of these steps in one script, for example `01-clean-data.R`, that starts from the raw file and ends by saving the clean tables:

```
write_csv(semesters_clean,
  here::here("data",
    "semesters_clean.csv"))
```

Never edit the raw file by hand. If you find a new problem next month, you fix it in the script and run it again, and every step stays on record.

In your field: environmental and public health

`airquality` is a real dataset that comes with R: daily air measurements in New York from May to September 1973. Like most real data, it has gaps. The same verbs summarise it by month:

```
airquality |>
  summarise(
    mean_ozone = mean(Ozone, na.rm =
      TRUE),
    missing_days = sum(is.na(Ozone)),
    .by = Month
  )
```

	Month	mean_ozone	missing_days
1	5	23.61538	5
2	6	29.44444	21
3	7	59.11538	5
4	8	59.96154	5
5	9	31.44828	1

Ozone levels peak in July and August, and June has the most missing days, a pattern worth knowing before any analysis: if the missing days were mostly hot ones, the June average would be too low.

3.8 - Chapter review

3.8.1 - Summary

- Clean data is valid, accurate, complete as far as possible, consistent, and unique. Cleaning decisions shape the results, so they are made in code from an untouched raw file, recorded, reported, and checked after every step.
- Tidy data has one variable per column, one observation per row, and one value per cell.
- The pipe `|>` passes a result to the next function: read it as “and then”.
- dplyr verbs: `filter()` chooses cases, `select()` chooses variables, `arrange()` sorts, `mutate()` creates variables, `summarise()` with `.by` summarises by group, and `count()` counts. `case_when()` sorts values into categories.
- `pivot_longer()` and `pivot_wider()` reshape data between wide and long formats.
- `left_join()` combines tables by an identifier; `anti_join()` finds rows without a match.
- Cleaning survey data: remove test and duplicate rows, rename variables, fix categories by pattern, turn missing codes into `NA` column by column, convert text to numbers (watch decimal commas), set impossible values to `NA`, reverse reversed items before computing scale scores, and check the result.

3.8.2 - Key terms

Data cleaning, tidyverse, tidy data, tibble, pipe, verb, grouped summary, wide format, long format, join, identifier, missing code, reversed item, scale score.

3.9 - Exercises

The [playground](#) has these and more, with hints and solutions.

1. Using `students`, find the average age of students in each faculty, sorted from oldest to youngest. (Remember `na.rm = TRUE`.)
2. Add a column to `semesters` that is "Short" when `sleep_hours` is under 6, "Recommended" when it is 7 or more, and "Moderate" otherwise. Count the semester records in each category.
3. Using `semesters`, make a table of average GPA with one row per student and one column per semester, using `pivot_wider()`. Show the first rows.
4. Join `semesters` to `students` and compare the average semester GPA of full-time and part-time students.
5. Explain why the `stress_4` item must be reversed before the stress score is calculated, and describe what would happen to the scores if it were not.
6. For each of the five qualities of clean data, give one example of a problem from your own field, and the cleaning decision it would require.

3.10 - Further reading

- *R for Data Science* [2]: chapters “Data transformation”, “Data tidying”, “Joins”, “Strings”, and “Missing values”.
- “Tidy data” [3], the paper that introduced the idea, with many examples of untidy data and how to fix it.
- The [dplyr](#) and [tidyr](#) websites, with a reference page for every function.

4. Data Visualization

A table of 600 numbers says almost nothing at a glance; a well-made graph of the same numbers can show in seconds where most values lie, which values are unusual, how groups differ, and how one variable moves with another. Graphs are therefore not decoration added to an analysis at the end. They are one of its main instruments: the first thing a careful researcher does with new data is look at it, and many errors in data, and in reasoning about data, are found only by looking.

A graph is also an argument. The choices behind it, which graph, which scale, which colours, what to leave out, decide what a reader sees first and what they conclude. Those choices can make a pattern clear or hide it, and they can make a small difference look dramatic or a large one look trivial. This chapter therefore begins with the ideas behind good graphs: what graphs are for, how people read them, how to choose a graph for a question, and what makes a graph honest. It then teaches **ggplot2**, the most widely used package for graphics in R, as a direct expression of those ideas, and uses it to answer the study's first research question, what graduate student life looks like (RQ1).

💡 By the end of this chapter you will be able to

- Explain why graphs are needed alongside numerical summaries, and the difference between graphs for exploring and graphs for explaining.
- Explain how people read graphs, and why position and length are read more accurately than angle, area, or colour.
- Choose a graph from the research question and the types of the variables involved.
- Build graphs with **ggplot2** from data, aesthetic mappings, and geometric layers.
- Draw histograms, bar charts, box plots, scatter plots, line charts, facets, density plots, and heat maps.
- Judge whether a graph is honest, and improve a misleading one.
- Prepare and export figures at the size and resolution a journal or thesis requires.

Note

Teaching & Practice: [Interactive Lecture Slides](#) · [Browser Playground Exercises](#)

4.1 - What a graph is for

Numerical summaries compress data into a few numbers, and that compression can hide almost anything. The statistician Francis Anscombe made the point with four small datasets that he constructed in 1973 [4]. R includes them as `anscombe`. Each has 11 pairs of values, and the usual summaries are identical for all four:

```
library(ggplot2)
library(dplyr)
```

```
Attaching package: 'dplyr'
```

```
The following objects are masked from
'package:stats':
```

```
filter, lag
```

```
The following objects are masked from
'package:base':
```

```
intersect, setdiff, setequal, union
```

```
anscombe_long <- tibble(
  set = rep(paste("Set", 1:4), each = 11),
  x    = c(anscombe$x1, anscombe$x2,
  anscombe$x3, anscombe$x4),
  y    = c(anscombe$y1, anscombe$y2,
  anscombe$y3, anscombe$y4)
)

anscombe_long |>
  summarise(mean_x = mean(x), mean_y = mean(y),
  sd_y = sd(y), correlation = cor(x, y),
  .by = set) |>
  mutate(across(where(is.numeric), \ (v)
  round(v, 2)))
```

```
# A tibble: 4 × 5
  set   mean_x mean_y sd_y correlation
<chr> <dbl>   <dbl> <dbl>         <dbl>
1 Set 1     9     7.5  2.03         0.82
2 Set 2     9     7.5  2.03         0.82
3 Set 3     9     7.5  2.03         0.82
4 Set 4     9     7.5  2.03         0.82
```

The same means, the same spread, the same correlation of 0.82: judged by the numbers, the four datasets are the same. Figure 4.1 shows that they are not.

```
ggplot(anscombe_long, aes(x = x, y = y)) +
  geom_point(size = 2) +
  geom_smooth(method = "lm", se = FALSE, colour
  = "grey50") +
  facet_wrap(~ set) +
  theme_minimal(base_size = 12)
```

```
`geom_smooth()` using formula = 'y ~ x'
```




Figure 4.1: Anscombe's four datasets: identical means, standard deviations, and correlations, but four entirely different relationships.

Only the first set is the kind of data the correlation describes well: a straight-line relationship with random scatter. The second is a smooth curve, the third a perfect line spoiled by one outlier, and in the fourth a single unusual point creates the entire relationship. Anyone who analysed these data without looking would draw the same, wrong, conclusion from all four. Looking first is not optional.

Graphs serve two different purposes, and it helps to know which one a graph is for. **Exploratory graphs** are made for the researcher, quickly and in large numbers, to understand the data: to check distributions, find unusual values, and notice patterns worth testing. They need no polish. **Explanatory graphs** are made for readers, to show a finding that the researcher already understands, and they need care: a clear message, accurate labels, and nothing that distracts. A thesis contains a few explanatory graphs, chosen from the many exploratory ones made along the way.

4.2 - How people read graphs

A graph works by turning numbers into visual properties: positions, lengths, angles, areas, colours. People do not judge all of these equally well. In a series of experiments, the statisticians William Cleveland and Robert McGill asked people to compare values shown in different ways, and found a clear ranking [5]. Positions along a common scale are judged most accurately, followed by lengths, then angles and slopes, then areas, and finally shades of colour. A good graph therefore puts its most important comparison into position or length.

The difference is easy to experience. The code below shows the same five percentages, which differ only slightly, as a pie chart and as a bar chart:

```
library(patchwork)

shares <- tibble(group = LETTERS[1:5], percent
= c(23, 21, 20, 19, 17))
```

```
pie_chart <- ggplot(shares, aes(x = "", y =
percent, fill = group)) +
  geom_col(width = 1, colour = "white") +
  coord_polar(theta = "y") +
  scale_fill_viridis_d(end = 0.9) +
  theme_void()

bar_chart <- ggplot(shares, aes(x = group, y
= percent)) +
  geom_col(fill = "grey50") +
  labs(x = NULL, y = "Percent") +
  theme_minimal(base_size = 12)

pie_chart + bar_chart
```




Figure 4.2: The same five percentages as a pie chart (angles and areas) and as a bar chart (lengths on a common scale). The order of the groups is much easier to see in the bars.

In the pie chart, the slices look almost identical, and ranking them requires reading a legend. In the bar chart, the steady decline from A to E is obvious at once, because the comparison is made by length along a shared axis. For this reason, researchers generally prefer bar charts and dot plots to pie charts, and scatter plots to bubble charts. The `patchwork` package, loaded above, places `ggplot2` graphs side by side with a simple `+`.

The same research gives two further rules of thumb. Comparisons are easiest when the values to be compared sit next to each other on the same scale, so the most important comparison should be placed within one panel, not across panels. And colour is best used to distinguish a few categories, or to show one quantity that does not need to be read precisely, rather than to carry the main message.

4.3 - Choosing a graph

The right graph follows from the research question and from the kinds of variables involved (Chapter 2). A question about the distribution of one numeric variable needs a different graph from a question about the relationship between two. Table 4.1 is a starting point for the graphs in this chapter.

Table 4.1: Choosing a graph from the question and the variables

Question	Variables	Graph	geom
How are the values distributed?	One numeric	Histogram density plot	<code>geom_histogram()</code> , <code>geom_density()</code>
How many cases fall in each category?	One categorical	Bar chart	<code>geom_bar()</code>
Do groups differ?	Numeric by categorical	Box plot points with averages	<code>geom_boxplot()</code> , <code>geom_jitter()</code>
How are two variables related?	Two numeric	Scatter plot	<code>geom_point()</code>
How does something change over time?	Numeric over time	Line chart	<code>geom_line()</code>
How do many variables relate?	Several numeric	Heat map of correlations	<code>geom_tile()</code>

4.4 - The grammar of graphics

The `ggplot2` package is built on an idea called the **grammar of graphics** [6]. It follows directly from the previous sections: a graph maps variables to visual properties. Every `ggplot2` graph therefore combines three parts. The **data** is a data frame. The **aesthetic mappings**, written inside `aes()`, state which variable goes to which visual property: the x axis, the y axis, the colour, and so on. The **geometric layers**, called *geoms*, state which shape represents the data, such as points, bars, or lines. The parts are added together with `+`. Five students from a pilot study show how:

```
pilot <- tibble(
  student = c("S1", "S2", "S3", "S4", "S5"),
  sleep   = c(6.5, 7.5, 5.5, 8, 6),
  stress  = c(3.2, 2.1, 4.5, 1.8, 3.9),
  invited = c("Yes", "No", "Yes", "No", "Yes")
)
```

The function `ggplot()`, given the data and the mappings, sets up an empty canvas with the axes:

```
ggplot(pilot, aes(x = sleep, y = stress))
```




Figure 4.3: Data and mappings alone give an empty canvas with the axes.

Adding a geom draws the data. The geom `geom_point()` draws one point per row:

```
ggplot(pilot, aes(x = sleep, y = stress)) +  
  geom_point(size = 3)
```




Figure 4.4: Adding a layer of points: students who sleep less report more stress.

Every further detail is another `+`: another layer, a label, a colour scale, a theme. That is the whole idea, and every graph in this chapter, however complex it looks, is more of the same.

For the wellbeing data, one row per student makes most graphs simplest. The code below combines each student's first-semester record with their background information:

```
first_sem <- semesters |>  
  filter(semester == 1) |>  
  left_join(students, join_by(student_id))
```

4.5 - The basic graphs

4.5.1 - Histograms

The first question about any numeric variable concerns its distribution: where most values lie, how widely they

spread, and whether some are unusual. A **histogram** answers it by dividing the values into ranges, called *bins*, and showing how many cases fall into each. Here is the distribution of sleep:

```
ggplot(first_sem, aes(x = sleep_hours)) +  
  geom_histogram(binwidth = 0.5)
```

```
Warning: Removed 6 rows containing non-finite  
outside the scale range  
(`stat_bin()`).
```




Figure 4.5: Hours of sleep per night in the first semester.

Figure 4.5 shows a roughly symmetrical, bell-shaped distribution, centred a little above 6 hours. Most students sleep less than the recommended 7 hours: 65% of them in the first semester. A few sleep very little, under 4.5 hours; Chapter 6 looks at such unusual values.

The argument `binwidth` sets the width of each bin, here half an hour. The choice matters: bins that are too wide hide the shape, and bins that are too narrow make it noisy, so it is worth trying a few widths.

4.5.2 - Bar charts

A **bar chart** shows how many cases fall into each category. The geom `geom_bar()` does the counting:

```
ggplot(students, aes(x = faculty)) +  
  geom_bar()
```




Figure 4.6: Number of students in each faculty.

When the heights have already been calculated, for example averages, `geom_col()` is used instead, with the calculated value mapped to `y`. Because a bar represents its value by its length, the axis of a bar chart must start at zero; the section on honest graphs returns to this.

4.5.3 - Box plots

A **box plot** summarises a numeric variable for each group, which makes it the standard graph for comparing groups. The box covers the middle half of the values, the line inside is the median, the whiskers reach out to the typical range, and points beyond the whiskers mark unusual values. The comparison below concerns wellbeing in full-time and part-time students:

```
ggplot(first_sem, aes(x = study_mode, y = wellbeing)) +
  geom_boxplot()
```




Figure 4.7: Wellbeing in the first semester, by study mode.

Part-time students' wellbeing is a little lower, but the boxes overlap considerably: many part-time students are doing better than many full-time students. Whether the difference is larger than chance is a question for Chapter 8.

4.5.4 - Scatter plots

The relationship between two numeric variables, such as sleep and grades, is shown with a **scatter plot**, which puts one variable on each axis. With 600 students, points pile on top of each other, a problem called **overplotting**, so `alpha` makes them partly transparent, and darker areas show where many students are:

```
ggplot(first_sem, aes(x = sleep_hours, y = gpa)) +
  geom_point(alpha = 0.4) +
  geom_smooth(method = "lm")
```

```
`geom_smooth()` using formula = 'y ~ x'
```



```
Warning: Removed 6 rows containing non-finite
outside the scale range
(`stat_smooth()`).
```

```
Warning: Removed 6 rows containing missing
values or values outside the scale range
(`geom_point()`).
```




Figure 4.8: Sleep and GPA in the first semester, with a straight trend line.

The layer `geom_smooth(method = "lm")` adds a straight trend line, with a grey band showing its uncertainty. The line rises: students who sleep more tend to have slightly higher GPAs. The points, however, scatter widely around it. Sleep is related to GPA, but it is far from the whole story, and Chapter 8 measures the relationship.

4.5.5 - Line charts

The study followed its students for four semesters, and a **line chart** shows how the average changed. The averages are calculated first (Chapter 3) and then plotted, with one line per workshop group:

```
wellbeing_by_sem <- semesters |>
  left_join(students, join_by(student_id)) |>
  summarise(mean_wellbeing =
    mean(wellbeing), .by = c(semester, workshop))

ggplot(wellbeing_by_sem, aes(x = semester, y =
  mean_wellbeing, colour = workshop)) +
  geom_line(linewidth = 1) +
  geom_point(size = 2.5)
```




Figure 4.9: Average wellbeing over four semesters, by workshop group.

Figure 4.9 tells the story of the workshop at a glance: the groups start level, the invited students pull ahead in semester 2, and the gap narrows afterwards. Chapters 7 and 10 test whether this pattern is real.

4.6 - Mapping and setting

In Figure 4.9, `colour = workshop` sat inside `aes()`. That **maps** colour to a variable: each workshop group gets its own colour, and ggplot2 adds a legend to explain them. To make everything one fixed colour, the colour is **set** outside `aes()` instead:

```
ggplot(first_sem, aes(x = sleep_hours)) +
  geom_histogram(binwidth = 0.5, fill =
    "steelblue", colour = "white")
```

```
Warning: Removed 6 rows containing non-finite
outside the scale range
(`stat_bin()`).
```




Figure 4.10: Setting a fixed colour outside `aes()`: all bars are the same colour, and there is no legend.

A very common mistake is to put a fixed colour inside `aes()`. ggplot2 then treats "steelblue" as a variable with one value, colours it with its first default colour (a salmon pink), and adds a pointless legend:

```
ggplot(first_sem, aes(x = sleep_hours, fill =
"steelblue")) +
  geom_histogram(binwidth = 0.5)
```

```
Warning: Removed 6 rows containing non-finite
outside the scale range
(`stat_bin()`).
```




Figure 4.11: The mistake: a fixed colour inside `aes()` is treated as data.

The rule follows from the grammar: *inside* `aes()` for anything that should depend on the data, *outside* for anything that should be the same everywhere. Note also the difference between `fill` and `colour`: `fill` is the inside of a shape such as a bar or box, and `colour` is its outline, or the colour of points and lines.

4.7 - Preparing graphs for publication

The graphs so far are exploratory: good enough for the researcher. An explanatory graph for a thesis or paper needs more care, and every addition below is a design decision with a reason behind it.

4.7.1 - Labels

A reader cannot interpret a graph whose axes say `mean_wellbeing` and `semester`. Axis labels should state what was measured, with units, and the title or caption should state the message. The function `labs()` sets the title, subtitle, axis labels, legend title, and caption:

```
p <- ggplot(wellbeing_by_sem, aes(x =
semester, y = mean_wellbeing, colour =
workshop)) +
  geom_line(linewidth = 1) +
  geom_point(size = 2.5) +
  labs(
    title = "Wellbeing over the two years",
    subtitle = "Average score per semester,
by workshop group",
    x = "Semester",
    y = "Average wellbeing (0 to 100)",
    colour = "Workshop"
```




Figure 4.12: The line chart with clear labels.

Storing the graph in an object, here `p`, makes it possible to add to it step by step without repeating the code.

4.7.2 - Scales and colours

Scales control how data values become positions, colours, and sizes, and each has a `scale_` function. Two decisions are needed here. The x axis should show only whole semesters, since there is no semester 2.5. And the colours should come from a palette that people with colour vision deficiency, about one man in twelve, can tell apart. The **viridis** palettes, built into `ggplot2`, are designed for exactly that, and they also print well in greyscale:

```
p <- p +
  scale_x_continuous(breaks = 1:4) +
  scale_colour_viridis_d(end = 0.8)
p
```




Figure 4.13: Whole-number semesters on the x axis and a colour-blind-friendly palette.

Colours can also be chosen by hand, with `scale_colour_manual(values = c("Invited" = "#1b9e77", "Not invited" = "#7570b3"))`. Whatever

the choice, colour should never be the *only* way to tell groups apart in a printed figure: different shapes or line types, or labels placed directly on the lines, keep the graph readable in black and white.

4.7.3 - Themes

A **theme** controls everything that is not data: the background, grid lines, fonts, and the position of the legend. The default grey background of `ggplot2` is useful on screen, but most journals prefer a plain look. The complete themes `theme_minimal()` and `theme_bw()` provide it, `base_size` sets the font size, and `theme()` adjusts individual details:

```
p <- p +
  theme_minimal(base_size = 13) +
  theme(legend.position = "bottom")
p
```



Figure 4.14: A clean theme, larger text, and the legend below the graph.

Compared with Figure 4.9, Figure 4.14 shows the same data, but it is now ready for a thesis.

4.7.4 - Annotations

A reference value often helps a reader interpret a graph: a recommended amount, a threshold, a mean. The functions `geom_vline()` and `geom_hline()` draw reference lines, and `annotate()` places text:

```
ggplot(first_sem, aes(x = sleep_hours)) +
  geom_histogram(binwidth = 0.5, fill =
    "steelblue", colour = "white") +
  geom_vline(xintercept = 7, linetype =
    "dashed") +
  annotate("text", x = 7.1, y = Inf, label =
    "Recommended: 7 hours", hjust = 0, vjust =
    1.5) +
  labs(x = "Hours of sleep per night", y =
    "Number of students") +
  theme_minimal(base_size = 13)
```

```
Warning: Removed 6 rows containing non-finite
outside the scale range
(`stat_bin()`).
```




Figure 4.15: Sleep in the first semester, with the recommended 7 hours marked.

In the annotation, `y = Inf` places the text at the top of the plot, whatever the height of the bars; `hjust = 0` makes it start just right of the line, and `vjust = 1.5` moves it slightly down from the edge. The line makes the graph's message, that most students sleep less than recommended, visible without any further explanation.

4.8 - Honest graphs

A graph can be accurate in every detail and still mislead. Most misleading graphs are not deliberate; they come from defaults that nobody questioned. Four choices deserve particular attention.

The first is the **axis range**. A bar represents its value by its length, so a bar chart whose axis does not start at zero misrepresents every value: a bar twice as long no longer means twice as much. Points and lines represent values by position, so their axes may be zoomed in to show change clearly, as long as the reader is told. The y axis of Figure 4.14, for example, runs only from about 58 to 66, because `ggplot2` zooms in on the data, which makes a 5-point gap look large. Either the caption should say so, or the full range of the scale can be shown with `coord_cartesian(ylim = c(0, 100))`, letting readers judge the size of the difference themselves.

The second is the **aspect ratio**, the shape of the graph. The same line looks steep in a tall, narrow graph and flat in a wide, short one. A good default is a graph somewhat wider than it is tall, and the same shape for graphs that will be compared.

The third is **clutter**. Everything in a graph that does not carry information, such as heavy grid lines, a legend that repeats the axis labels, or three-dimensional effects, competes with the data for the reader's attention. Three-dimensional bars and pies are worst of all, since the perspective distorts exactly the lengths and angles that carry the values.

The fourth is **colour**. Colours for categories should be clearly different from each other but none should stand

out, since no category is more important than another; a qualitative palette such as viridis does this. Colours for quantities should run in one direction, from light to dark (a sequential palette), or away from a meaningful mid-point such as zero in two directions (a diverging palette, as in the heat map later in this chapter). A rainbow palette is poor for both, because its bright bands suggest boundaries that are not in the data.

4.8.1 - A graph improved

The principles are clearest when applied to one graph. Suppose the question is whether wellbeing differs between faculties. A first attempt, using defaults and a zoomed axis, might look like Figure 4.16:

```
faculty_means <- first_sem |>
  summarise(wellbeing = mean(wellbeing), .by = faculty)

ggplot(faculty_means, aes(x = faculty, y = wellbeing, fill = faculty)) +
  geom_col() +
  coord_cartesian(ylim = c(58, 62.5))
```




Figure 4.16: A misleading graph of average wellbeing by faculty: the axis starts at 58, the colours repeat the axis labels, and the spread of the data is invisible.

The graph suggests large differences: the Education bar looks several times taller than the Social Sciences bar. But the axis starts at 58, so the lengths of the bars mean nothing; the averages actually differ by about 3 points on a 100-point scale. The colours add nothing that the axis labels do not already say, the legend repeats them a second time, the faculties appear in alphabetical rather than meaningful order, and the axis titles are variable names. Most seriously, the graph shows only averages, and hides how much students within each faculty differ.

Figure 4.17 shows the same data, redesigned. Every student appears as a faint point, so the spread within each faculty is visible, and the average of each faculty is marked in red. The faculties are ordered by their average, the labels say what was measured, and the legend is gone:

```
ggplot(first_sem, aes(x = wellbeing, y = reorder(faculty, wellbeing))) +
  geom_jitter(height = 0.15, alpha = 0.25,
```

```
colour = "grey45") +
  stat_summary(fun = mean, geom = "point",
    size = 3.5, colour = "#b2182b") +
  labs(x = "Wellbeing in semester 1 (0 to 100)", y = NULL) +
  theme_minimal(base_size = 12)
```




Figure 4.17: The same data, redesigned: every student as a grey point, each faculty's average in red, faculties ordered by their average. The differences between faculties are small compared with the differences within them.

The function `reorder()` orders the faculties by their average wellbeing, `geom_jitter()` spreads the points a little vertically so that they do not hide each other, and `stat_summary()` calculates and draws each faculty's mean. The redesigned graph tells the truth that the first one hid: the differences between faculties are small compared with the differences between students within each faculty. Chapter 8 tests whether they are larger than chance.

4.9 - Further kinds of graph

4.9.1 - Small multiples

Facets, also called small multiples, split one graph into a panel for each group, all with the same axes, so the groups are easy to compare. The function `facet_wrap()` takes the variable to split by:

```
ggplot(first_sem, aes(x = sleep_hours, y = gpa)) +
  geom_point(alpha = 0.4) +
  geom_smooth(method = "lm") +
  facet_wrap(~ programme) +
  labs(x = "Hours of sleep per night", y = "GPA") +
  theme_minimal(base_size = 12)
```

```
`geom_smooth()` using formula = 'y ~ x'
```

```
Warning: Removed 6 rows containing non-finite
outside the scale range
(`stat_smooth()`).
```


Warning: Removed 6 rows containing missing values or values outside the scale range (``geom_point()``).




Figure 4.18: Sleep and GPA in the first semester, one panel per programme.

The relationship looks similar in both programmes. For two grouping variables, `facet_grid(rows ~ columns)` makes a grid of panels. Because the panels share their axes, facets respect the rule from the section on reading graphs: comparisons are made by position on a common scale.

4.9.2 - Density plots

A **density plot** is a smoothed histogram. Its advantage is that several distributions can be drawn on top of each other and compared. Here, `fill` is mapped to the workshop group and `alpha` lets the two curves show through each other:

```
semesters |>
  filter(semester == 2) |>
  left_join(students, join_by(student_id)) |>
  ggplot(aes(x = wellbeing, fill = workshop)) +
  geom_density(alpha = 0.5) +
  scale_fill_viridis_d(end = 0.8) +
  labs(x = "Wellbeing (0 to 100)", y =
    "Density", fill = "Workshop") +
  theme_minimal(base_size = 13)
```




Figure 4.19: Distribution of wellbeing in the second semester, by workshop group.

In this code, the data flows straight into `ggplot()` through the pipe. The two curves have the same shape, but the invited group's is shifted to the right: the whole distribution moved, not just a few students.

4.9.3 - Heat maps

A **heat map** shows a table of numbers as coloured tiles, and it suits correlation matrices well. Chapter 2 calculated the correlations between four first-semester measurements. To plot them, the matrix is turned into a long table with one row per pair (Chapter 3), and then `geom_tile()` draws a tile for each pair and `geom_text()` writes the value on it:

```
library(tidyr)

correlations <- first_sem |>
  select(gpa, sleep_hours, study_hours,
         wellbeing) |>
  cor(use = "complete.obs")

correlations |>
  as.data.frame() |>
  mutate(var1 = rownames(correlations)) |>
  pivot_longer(-var1, names_to = "var2",
               values_to = "r") |>
  ggplot(aes(x = var1, y = var2, fill = r)) +
  geom_tile() +
  geom_text(aes(label = round(r, 2))) +
  scale_fill_gradient2(limits = c(-1, 1)) +
  labs(x = NULL, y = NULL, fill = "Correlation") +
  theme_minimal(base_size = 12)
```




Figure 4.20: Correlations between four first-semester measurements.

The scale `scale_fill_gradient2()` is a diverging palette: two colours for negative and positive values, with white at zero, so the direction and strength of each correlation are visible at once. Students who study more hours sleep less (a negative correlation), and students who sleep

more report higher wellbeing (positive). The numbers written on the tiles matter, because colour alone is read imprecisely. Chapter 6 explains how to read correlations.

4.10 - Exporting figures

The function `ggsave()` saves the most recent graph, or one that is named, to a file. The file type follows from the file name:

```
ggsave("figures/wellbeing-lines.png", plot =  
p, width = 16, height = 10, units = "cm", dpi  
= 300)  
ggsave("figures/wellbeing-lines.pdf", plot =  
p, width = 16, height = 10, units = "cm")
```

A figure should be saved at the size at which it will be printed, in the units of the page: a figure for a single column of a journal is often about 8 cm wide, and a full page about 16 cm. Setting the size when saving, rather than stretching the image later, keeps the text readable and the proportions right. Images such as PNG files need a resolution of at least 300 dpi (dots per inch), which is what most journals require, while vector formats such as PDF and SVG stay sharp at any size and are preferable whenever the journal accepts them. Finally, the text should be checked at the final size, since fonts that look fine on screen are often too small in print; `base_size` in the theme fixes that.

The code that makes each figure belongs in the analysis script. When a supervisor asks for a larger font or a different colour, one line changes and the figure is saved again.

In your field: economics and business

ggplot2 includes `economics`, real monthly data on the US economy from 1967 to 2015. The same line chart shows the rise and fall of unemployment, with each recession visible as a peak:

```
ggplot(economics, aes(x = date, y =
unemploy)) +
  geom_line() +
  labs(x = NULL, y = "Unemployed
(thousands)") +
  theme_minimal(base_size = 12)
```



Figure 4.21: US unemployment, 1967 to 2015 (ggplot2's `economics` data).

The graph shows the number of unemployed people, not the unemployment rate, and the population grew greatly over these 48 years. A reader comparing 1970 with 2010 should know that, which is why the axis label names exactly what is shown.

4.11 - Common misconceptions

Several beliefs about graphs are widespread, and each leads to graphs that mislead or fail to inform.

- **“The summary statistics tell the whole story.”** Anscombe's datasets share every summary and differ completely. Always look.
- **“A graph with more elements is more informative.”** Decoration, three-dimensional effects, and repeated labels compete with the data.
- **“Pie charts are good for proportions.”** Angles and areas are judged poorly; a bar chart or dot plot of the same proportions is almost always clearer.
- **“Zooming in on the axis is always wrong.”** It is wrong for bar charts, whose lengths encode values, but acceptable for points and lines, if the reader is told.
- **“Colour makes a graph clearer.”** Colour helps to separate a few categories; used for everything, it becomes noise, and it disappears in black-and-white printing.

4.12 - Chapter review

4.12.1 - Summary

- Numerical summaries can hide very different patterns, as Anscombe's datasets show, so data should always be graphed. Exploratory graphs help the researcher understand the data; explanatory graphs show readers a finding.
- People judge positions and lengths most accurately, then angles, areas, and colours. The most important comparison belongs in position or length, within one panel.
- The graph follows from the question and the kinds of variables: histograms for distributions, bar charts for counts, box plots or points for group comparisons, scatter plots for relationships, line charts for change over time.
- Every `ggplot2` graph combines data, aesthetic mappings in `aes()`, and geometric layers, joined with `+`. *Map* a colour inside `aes()` when it should depend on the data; *set* it outside `aes()` when it should be fixed.
- `labs()` adds labels, `scale_` functions control axes and colours (viridis palettes are colour-blind-friendly), and themes control the look.
- Honest graphs start bar axes at zero, state when an axis is zoomed, avoid clutter, and use colour appropriately: qualitative palettes for categories, sequential or diverging palettes for quantities.
- Facets split a graph into comparable panels; density plots compare distributions; heat maps show tables of numbers.
- `ggsave()` exports figures. Set the size in page units, use at least 300 dpi for images, and prefer PDF or SVG when accepted.

4.12.2 - Key terms

Exploratory graph, explanatory graph, graphical perception, grammar of graphics, aesthetic mapping, geom, layer, histogram, bin, bar chart, box plot, median, scatter plot, overplotting, trend line, line chart, mapping, setting, scale, theme, annotation, aspect ratio, qualitative palette, sequential palette, diverging palette, facet, density plot, heat map, dpi, vector format.

4.13 - Exercises

The [playground](#) has these and more, with hints and solutions.

1. Draw a histogram of `study_hours` in the first semester. Try binwidths of 1, 5, and 10, and explain which shows the shape best.
2. Draw a bar chart of how many students have each level of `employment`. Make the bars a single colour of your choice, and order the levels from no job to a full-time job.
3. Draw box plots of first-semester GPA for each faculty, with clear axis labels.

4. Draw a scatter plot of caffeine against sleep in the first semester, with a trend line, and describe what it suggests. (Chapter 8 returns to this relationship.)
5. Take any graph from this chapter and save it as a PNG file, 16 cm wide and 10 cm high, at 300 dpi.
6. Redraw Figure 4.16 as a bar chart whose axis starts at zero. Compare it with Figure 4.16 and Figure 4.17, and explain which of the three you would put in a thesis, and why.

4.14 - Further reading

- *R for Data Science* [2]: chapters “Data visualization”, “Layers”, and “Communication”.
- *Data Visualization: A Practical Introduction* [7], free online, explains the principles of good graphs and builds them with ggplot2.
- *ggplot2: Elegant Graphics for Data Analysis* [6], by the author of ggplot2, explains the grammar of graphics in depth.
- “Graphical perception” [5], the experiments behind the ranking of visual properties.

	Exploratory Data Analysis	133
6.1	The place of exploration	134
6.2	The centre of a distribution	135
6.3	The spread of a distribution	138
6.4	The shape of a distribution	141
6.5	Unusual values	143
6.6	Missing data	145
6.7	Relationships between variables	146
6.8	Describing your sample in a thesis	148
6.9	Common misconceptions	149
6.10	Chapter review	151
6.11	Exercises	151
6.12	Further reading	152
7	Hypothesis Testing and Statistical Inference	153
7.1	Samples and populations	154
7.2	The logic of hypothesis testing	157
7.3	The one-sample t-test	163
7.4	The two-sample t-test	164
7.5	The paired t-test	167
7.6	Assumptions of the t-test	168
7.7	Non-parametric tests	168
7.8	Chi-square tests	169
7.9	Choosing a test	171
7.10	Reporting the results	174
7.11	Common misconceptions	175
7.12	Chapter review	175
7.13	Exercises	176
7.14	Further reading	176
8	ANOVA and Regression	177
8.1	One-way ANOVA	178
8.2	Two-way ANOVA	183
8.3	Linear regression	184
8.4	Logistic regression	195
8.5	Common misconceptions	198
8.6	Chapter review	199
8.7	Exercises	200
8.8	Further reading	200
9	Multivariate Statistical Methods	201
9.1	Constructs and latent variables	202
9.2	Many variables at once	202
9.3	Principal component analysis	204
9.4	Factor analysis	206
9.5	Cluster analysis	209
9.6	Common misconceptions	215
9.7	Chapter review	215
9.8	Exercises	216
9.9	Further reading	216
10	Mixed-Effects Models	217
10.1	Data with a structure	217
10.2	A study of sleep deprivation	219
10.3	Change in wellbeing over two years	222
10.4	Students within supervisors	228
10.5	Counts and yes-or-no outcomes	228
10.6	Reporting a mixed-effects model	229
10.7	Common misconceptions	230
10.8	Chapter review	231
10.9	Exercises	231
10.10	Further reading	232

5. From Research Question to Data

Part 1 gave you the tools: you can import data, clean it, and draw it. Part 2 uses those tools to answer research questions. Before any statistics, though, comes a step that no software can do for you: deciding *what exactly you are asking*, *what result would answer it*, and *whether your data can give that answer at all*.

Most problems that examiners find in a thesis start here, not in the analysis: a question too vague to answer, a hypothesis that no result could contradict, a questionnaire that measures something other than what the title promises, or a sample that cannot support the conclusion. No statistical test can fix those problems afterwards. This chapter follows the path from a research topic to data that can answer a question, using the student wellbeing study as the example throughout. It is more conceptual than the other chapters, but it still uses R wherever R helps you see an idea: simulations, checks of how variables are measured, and a first calculation of sample size.

💡 By the end of this chapter you will be able to

- Explain what data analysis is for in research: describing, explaining, and predicting.
- Turn a research topic into a focused research question, and tell a good question from a weak one.
- Write a hypothesis that is testable and falsifiable, and state its null hypothesis.
- Explain how an abstract idea, such as stress, becomes a measured variable, and judge its level of measurement.
- Name the roles variables play: outcome, predictor, confounder, moderator, and mediator.
- Distinguish reliability from validity.
- Explain what experimental, observational, cross-sectional, and longitudinal designs can and cannot show.
- Explain why a biased sample cannot be fixed by making it larger.
- Plan an analysis before collecting data, including a first estimate of the sample size, and use R to check levels of measurement, balance between groups, and sampling by simulation.

Note

Teaching & Practice: [Interactive Lecture Slides](#) · [Browser Playground Exercises](#)

5.1 - What data analysis is for

In one sentence: data analysis tests claims against evidence.

Research makes claims about the world: graduate students sleep too little; a workshop improves wellbeing; supervisor support protects against dropout. A claim is only as strong as the evidence behind it, and data analysis is the part of research that confronts claims with evidence. It does three kinds of work. **Describing** establishes what the world looks like: how long graduate students sleep, or how many consider dropping out. **Explaining** asks why it looks like that: whether the workshop *causes* higher wellbeing, or which factors go together with a higher GPA. **Predicting** concerns what will happen in new cases, such as which of next year's students are likely to consider leaving.

Each kind of work asks a different kind of question, is judged by a different standard, and leads to different methods, as Table 5.1 shows.

Table 5.1: Three aims of data analysis

Aim	Ex-ample ques-tion	Judged by	Meth-ods	Chap-ters
De-scribe	How long do graduate students sleep?	Accu-racy of the summary; who it applies to	Sum-maries, graphs, confidence intervals	4, 6, 7
Explain	Does the workshop improve wellbeing?	Whether rival explanations are ruled out	Tests, regres-sion, mixed models	7 to 10
Predict	Who will consider dropping out?	Accu-racy on new cases	Ma-chine learning	11 to 16

Keep the aim in mind, because the same numbers can serve different aims. A regression model can *explain* GPA (which predictors matter, and how much?) or *predict* it (how close are the predictions for new students?). The model may be identical; the question, and how the answer is judged, are not.

5.2 - From a topic to a research question

In one sentence: a research question is a topic narrowed until data can answer it.

Nobody starts with a research question. They start with a **topic**, something that interests or worries them, such as “the wellbeing of graduate students”. A topic is too broad to study: it contains hundreds of possible questions. The work of the first months of a thesis is narrowing it:




Figure 5.1: Narrowing a topic into a research question. Each step makes the question more specific, until it is clear what data would answer it.

The **problem** says why the topic matters: something is wrong, unknown, or disputed. The **research question** says precisely what the study will find out. A good research question is **specific**: it names who (graduate students at one university), what (wellbeing, measured with an index), and when (the end of semester 2). It is **answerable with data**, in the sense that the researcher can say what data would answer it and can collect that data. It is **not already answered**, because the literature leaves a gap or the setting is new. It is **feasible** within the time, money, skills, and access available. And it is **worth answering**: someone would act differently depending on the answer. Table 5.2 shows weak questions and how they can be improved.

Table 5.2: Weak research questions, and how to improve them

Weak question	Problem	Better question
Are students stressed?	Stressed compared with what? Which students?	Do part-time graduate students report higher stress than full-time students?
What affects students' success?	Too broad; "success" is undefined	Is average sleep in a semester associated with semester GPA, allowing for study hours?
Is the workshop good?	"Good" cannot be measured	Does the workshop raise well-being scores at the end of the following semester?
Why do students drop out?	Needs students who left, and years of follow-up	Which baseline characteristics are associated with <i>considering</i> dropout in the first year?

5.2.1 - Three kinds of question

Research questions come in three kinds, and each kind needs a different design and analysis. **Descriptive questions** ask *what is*, such as how many hours graduate students sleep; they need a sample that represents the population well. **Relational questions** ask *what goes with what*, such as whether students who sleep more have higher GPAs; they need measurements of both variables, and they establish association, not cause. **Causal questions** ask *what leads to what*, such as whether the workshop improves wellbeing; they need a design that rules out other explanations, ideally an experiment.

The kind of question decides the kind of claim you can make. Much confusion in published research comes from asking a relational question and answering it in causal words (“sleep *improves* grades”). The section on study designs below returns to this.

5.3 - From a question to a hypothesis

In one sentence: a hypothesis is a prediction precise enough to be wrong.

A research question asks; a **hypothesis** answers in advance. It is the researcher’s best prediction, based on theory and earlier studies, of what the data will show. The workshop question becomes:

Hypothesis: Students invited to the workshop will have higher wellbeing at the end of semester 2 than students not invited.

The hypothesis is useful because the data can disagree with it.

5.3.1 - Falsifiability

The philosopher Karl Popper argued that what separates a scientific claim from other claims is not that it can be proved, but that it can be **refuted**: it forbids certain results [8]. A claim that is compatible with every possible result tells us nothing. Compare:

Table 5.3: Unfalsifiable and falsifiable hypotheses

Hypothesis	Could any result contradict it?
“The workshop affects students in some way.”	No. Whatever happens, some effect on someone can be found. Unfalsifiable.
“The workshop helps students who are ready for it.”	No, unless “ready” is defined before the study; otherwise any student who did not improve was “not ready”. Unfalsifiable.
“Invited students will have higher wellbeing at the end of semester 2 than students not invited.”	Yes: equal or lower wellbeing in the invited group would contradict it. Falsifiable.

A good test of your own hypothesis is to write down, before seeing any data, **what result would count against it**. If you cannot, the hypothesis is not yet precise enough. If you can, you have also protected yourself against a common temptation: finding, after the fact, a reason why the disappointing result supports your idea after all.

A good hypothesis therefore has four properties. It is about a **stated population and stated variables**: graduate students at this university, and wellbeing measured with the index. It is **testable with the data available**, because the variables are measured and there are enough cases. It is **falsifiable**: it names the result that would contradict it. And it is **stated before the data is analysed**, since a hypothesis written after seeing the data is a description of the data, not a test of it.

5.3.2 - Directional and non-directional hypotheses

A **directional** hypothesis predicts the direction of an effect (“invited students will have *higher* wellbeing”). A **non-directional** hypothesis predicts only that there is a difference (“wellbeing will *differ* between faculties”). Use a directional hypothesis when theory gives you a clear prediction; use a non-directional one when it does not. Note that the direction of the hypothesis and the choice of a one-tailed or two-tailed test are separate decisions: many researchers state a directional hypothesis but still use a two-tailed test, the cautious choice explained in Chapter 7.

5.3.3 - The null hypothesis

Statistical tests do not test the researcher’s hypothesis directly. They test its opposite, the **null hypothesis** (H_0): the claim that nothing is going on, no difference and no relationship. The researcher’s hypothesis becomes the **alternative hypothesis** (H_1). For the workshop, H_1 states that invited and not-invited students differ in average wellbeing at the end of semester 2, and H_0 that they have the same average wellbeing.

Testing the opposite of what one believes may seem strange, but the null hypothesis has one great advantage: it is precise enough to calculate with. “The workshop helps” does not say how much, but “the workshop makes no difference” makes a sharp prediction: any difference between the groups is due to chance alone, to which students happened to be invited. If we know how large chance differences usually are, we can ask whether the difference in the data is larger than chance would easily produce.

You can see this world of chance by simulation. Suppose the workshop truly does nothing. Then being invited is just a label on students whose wellbeing is unaffected. The code below creates 300 imaginary students with wellbeing scores similar to those in the study (average 60, standard deviation about 11), splits them into two random groups of 150, and records the difference between the group averages. It then repeats this 5,000 times:

```
set.seed(2026)
chance_differences <- replicate(5000, {
  wellbeing <- rnorm(300, mean = 60, sd = 11)
  group      <- sample(rep(c("Invited", "Not
invited"), 150))
  mean(wellbeing[group == "Invited"]) -
```

```
mean(wellbeing[group == "Not invited"])
})
summary(chance_differences)
```

	Min.	1st Qu.	Median	Mean	3rd Qu.
	-5.188460	-0.871667	0.001230	-0.008296	0.838404
		4.688888			

The function `replicate()` runs the code in curly brackets many times and collects the results, `rnorm()` draws random numbers from a normal distribution, and `sample()` shuffles the labels. Figure 5.2 shows the 5,000 differences:

```
library(ggplot2)
ggplot(data.frame(difference =
  chance_differences), aes(x = difference)) +
  geom_histogram(binwidth = 0.5, fill =
    "grey70", colour = "white") +
  geom_vline(xintercept =
    quantile(chance_differences,
    c(0.025,
    0.975)),
    linetype = "dashed") +
  labs(x = "Difference in average wellbeing
    (invited minus not invited)",
    y = "Number of simulated studies")
```




Figure 5.2: The null hypothesis world: differences between two random groups of 150 students when the workshop does nothing. Most are within about 2.5 points of zero.

Even when the workshop does nothing, the two groups almost never have exactly the same average: chance alone produces differences, usually small ones. In 95% of the simulated studies, the difference lies between -2.5 and 2.6 points (the dashed lines). A real study that found a difference of 5 points would therefore be hard to explain by chance, and would count as evidence against H_0 . A difference of 1 point would not: it is exactly what chance produces. Chapter 7 turns this idea into the p-value, using the real data.

Two things follow from this picture. First, **rejecting H_0 is not proving H_1** : it says only that “nothing is going on” is a poor explanation of the data. Second, **not rejecting H_0 is not proving it**: a small study can fail to

detect a real effect, just as a blurred photograph can fail to show a real face.

5.3.4 - Hypotheses for the whole study

Table 5.4 writes out the study's research questions in this form. Some questions are **confirmatory**: they test a hypothesis stated in advance. Others are **exploratory**: they look for patterns without a prior prediction, and their results suggest hypotheses for future studies rather than confirming them. Both are legitimate, as long as they are labelled honestly.

Table 5.4: The study's research questions as hypotheses. Descriptive, exploratory, and predictive questions have no null hypothesis; they are judged in other ways.

Re- search ques- tion	Hy- pothesis (H_1)	Null hy- pothesis (H_0)	What would count against H_1	Chap- ters
RQ1 What does gradu- ate life look like?	De- scrip- tive: no hypothesis	none	none	4, 6
RQ2 Do stu- dents sleep less than 7 hours?	Average sleep is below 7 hours	Average sleep is 7 hours	An av- erage of 7 hours or more, or a confi- dence interval that in- cludes 7	7
RQ3 Does the work- shop improve wellbe- ing?	Invited stu- dents have higher wellbe- ing in semes- ter 2	The groups have the same average wellbe- ing	A dif- ference near zero or neg- ative, with a narrow confi- dence interval	7, 10
RQ4 Do facul- ties and study modes differ?	Stress dif- fers be- tween faculties (non- direc- tional)	All fac- ulties have the same average stress	Simi- lar av- erages in every faculty	7, 8
RQ5 What explains GPA?	More sleep goes with a higher GPA, allow- ing for study hours, stress, and support	The sleep coeffi- cient is zero	A coef- ficient near zero or nega- tive	8
RQ6 Do the ques- tions measure	The 22 items form a reliable scale	none (checked by the model fitting process)	Items that do not load on the scale	9
RQ7 Do the ques- tions measure	Exploratory: The items form a reliable scale	Predictive: The items form a reliable scale	Items that do not load on the scale	9, 10

The later chapters return to this table: each restates its hypothesis before the analysis, and ends by saying whether the data contradicts the null hypothesis.

i Think before you analyse

From this chapter on, every main analysis in the book starts with the same four questions. Answer them in writing, before you run any code:

1. **What is the question**, and is it descriptive, relational, or causal?
2. **What is the unit of analysis?** Students, semesters, supervisors?
3. **Which variables**, of what level of measurement, play which role?
4. **What result would support the hypothesis, and what would count against it?**

5.4 - Variables and measurement

In one sentence: a variable is a decision about how to turn an idea into a number.

5.4.1 - Cases and the unit of analysis

As Chapter 2 showed, data is a table of **cases** (rows) and **variables** (columns). The first question about any dataset is what one row represents: the **unit of analysis**. The wellbeing study has three, in different tables:

```
nrow(students)      # one row per student
```

```
[1] 600
```

```
nrow(semesters)     # one row per student
per semester
```

```
[1] 2326
```

```
nrow(supervisors)   # one row per supervisor
```

```
[1] 120
```

The unit of analysis decides what a question means. “Do students who sleep more have higher GPAs?” is a question about students, so each student should count once, for example with their first-semester values or their average over the semesters. Treating the 2326 semester rows as 2326 independent cases would count each student up to four times, and make the evidence look stronger than it is. Chapter 10 shows how to use all the rows correctly.

5.4.2 - From constructs to variables

Many things researchers care about cannot be observed directly: stress, wellbeing, motivation, intelligence, job satisfaction, quality of life. Such ideas are called **constructs**. To study a construct, you must decide how to

measure it, a step called **operationalisation**. Every operationalisation is a choice, and a different choice could give a different answer.

Table 5.5 shows how the wellbeing study operationalises some of its constructs.

Table 5.5: How some of the study's constructs are measured

Construct	Operationalisation	Variable(s)
Sleep	Self-reported average hours per night in the semester	<code>sleep_hours</code>
Stress	The average of six questionnaire items on a 1 to 5 scale	<code>stress_1</code> to <code>stress_6</code>
Wellbeing	A 0 to 100 wellbeing index, from a validated instrument	<code>wellbeing</code>
Academic performance	Semester GPA from university records	<code>gpa</code>
Thoughts of dropping out	"Have you seriously considered leaving your programme?"	<code>considering_dropout</code>

Stress is a good example. No single question captures it, so the questionnaire asks six, each about one aspect, and averages the answers. Look at the items, and at how the answers of the first few students differ from item to item:

```
library(dplyr)
```

```
Attaching package: 'dplyr'
```

```
The following objects are masked from
'package:stats':
```

```
filter, lag
```

```
The following objects are masked from
'package:base':
```

```
intersect, setdiff, setequal, union
```

```
questionnaire |>
  select(student_id, stress_1:stress_6) |>
  head(5)
```

	student_id	stress_1	stress_2	stress_3	stress_4	stress_5	stress_6
1	S0001	3	2	4	4		
3	5						
2	S0002	4	3	NA	2		
5	3						
3	S0003	4	4	2	1		
4	3						
4	S0004	4	5	5	1		
4	5						
5	S0005	4	4	4	3		
4	2						

One item, `stress_4` (“I feel confident handling problems in my studies”), is worded in the opposite direction: agreeing with it means *less* stress. Before averaging, it must be reversed, so that 5 becomes 1 and 1 becomes 5. You did this in Chapter 3:

```
stress <- questionnaire |>
  mutate(stress_4 = 6 - stress_4,
         stress =
rowMeans(pick(stress_1:stress_6), na.rm =
TRUE)) |>
  select(student_id, stress)
summary(stress$stress)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
1.333	2.667	3.167	3.216	3.667	5.000

The resulting score is the study’s operational definition of stress. Anyone who reads the thesis should be able to see exactly how it was built, which is why the methods chapter of a thesis describes every measure, and why the questionnaire items go in an appendix.

i Self-report is a measurement choice

Sleep, study hours, and caffeine in the wellbeing study are *self-reported*: students estimate them. People tend to overestimate their sleep and underestimate their caffeine, and they may give the answers they think are expected. Records (such as GPA from the university) or devices (such as a sleep tracker) avoid some of these problems, but are more expensive or intrusive. Every measure has weaknesses; say what they are.

5.4.3 - Levels of measurement

Not all numbers are numbers in the same sense. The psychologist S. S. Stevens proposed four **levels of measurement**, each allowing more than the one before [9]:

Table 5.6: Levels of measurement

Level	What the values mean	Meaningful summaries	Examples in the study	In R
Nominal	Categories with no order	Counts, percentages, mode	<code>faculty</code> , <code>gender</code> , <code>workshop</code>	<code>factor</code>
Ordinal	Ordered categories; distances between them unknown	Also median and percentiles	<code>employment</code> , <code>financial_worry</code> , <code>single_questionnaire</code> items	ordered <code>factor</code> , or numbers used with care
Interval	Equal distances; no true zero	Also mean, standard deviation, differences	<code>wellbeing</code> (0 to 100 index), scale scores (arbitrarily)	numeric
Ratio	Equal distances and a true zero	Also ratios (“twice as much”)	<code>sleep_hours</code> , <code>study_hours</code> , <code>caffeine_mg</code> , <code>age</code>	numeric

The level decides which summaries and tests make sense. The average faculty is meaningless. “Twice as much caffeine” is meaningful (caffeine has a true zero); “twice as much wellbeing” is not (a score of 0 does not mean no wellbeing at all).

R does not know the level of measurement; you have to tell it. Look at how R stores `employment`:

```
table(students$employment)
```

```
Full-time job      None Part-time job
          97          307          196
```

R lists the categories in alphabetical order, which puts “Full-time job” before “None”. For a nominal variable, that would not matter. But `employment` is **ordinal**: none, part-time, full-time is a real order, and tables and graphs should show it. Tell R with an **ordered factor**:

```
students <- students |>
  mutate(employment = factor(employment,
                             levels = c("None",
```

```
"Part-time job", "Full-time job"),
      ordered = TRUE))
table(students$employment)
```

```
None Part-time job Full-time job
307          196          97
```

Now tables and graphs follow the natural order, and comparisons such as `employment > "None"` work.

The opposite mistake is more common: an ordinal variable stored as numbers, such as `financial_worry` (1 = not at all worried, 5 = extremely worried). R will happily calculate its mean, but the distance from “not at all” to “a little” need not equal the distance from “very” to “extremely”. For a single ordinal item, report the distribution or the median:

```
table(students$financial_worry, useNA = "ifany")
```

```
1    2    3    4    5 <NA>
84  139  172  111  56   38
```

```
median(students$financial_worry, na.rm = TRUE)
```

```
[1] 3
```

Scale scores, the average of several items, are usually treated as interval data, and the book does so. Averaging several items smooths out the unequal steps of each one. This is a convention with good support, not a law; if an examiner asks, say that you followed it, and why.

5.4.4 - The roles variables play

In an analysis, variables also have **roles**. The **outcome**, or dependent variable, is what the researcher wants to explain or predict, such as wellbeing, GPA, or considering dropout. A **predictor**, also called an independent or explanatory variable, is what is used to explain or predict it, such as the workshop, sleep, or stress. Three further roles concern a third variable that affects the relationship between the two. A **confounder** is related to both the predictor and the outcome, and can create a misleading association between them. A **moderator** changes the strength or direction of a relationship, so that the effect of the predictor depends on it. A **mediator** lies on the path between predictor and outcome: it is *how* the predictor has its effect.

The roles are not properties of the variables; they come from the research question. Stress is an outcome in one question (whether the workshop reduces stress) and a predictor in another (whether stress predicts dropout).

Figure 5.3 draws the three less familiar roles as diagrams, using examples from the study: arrows show which variable is thought to influence which.




Figure 5.3: Three roles a third variable can play. Confounder: sleep affects both caffeine intake and GPA. Moderator: the effect of support on GPA depends on the programme. Mediator: the workshop may raise wellbeing by lowering stress.

The confounder is the most important of the three, because it can produce an association that is not a cause. In the wellbeing data, students who take more caffeine have lower GPAs. But students who take more caffeine also sleep less:

```
first_sem <- semesters |> filter(semester == 1)
first_sem |>
  select(caffeine_mg, sleep_hours, gpa) |>
  cor(use = "complete.obs") |>
  round(2)
```

	caffeine_mg	sleep_hours	gpa
caffeine_mg	1.00	-0.64	-0.18
sleep_hours	-0.64	1.00	0.25
gpa	-0.18	0.25	1.00

Caffeine and GPA are negatively correlated (-0.19), but caffeine and sleep are strongly negatively correlated (-0.64). Caffeine might be harming grades, or short-sleeping students might both drink more coffee and get lower grades; the correlation alone cannot tell which. Chapter 8 answers the question with a regression that compares students with the same amount of sleep. The general lesson is that before you interpret any relationship, you should ask what else could produce it, and measure it if you can.

5.5 - Reliability and validity

In one sentence: a reliable measure gives consistent results; a valid measure measures the right thing.

Two properties decide whether a measure can be trusted. **Reliability** is consistency: a student should get a similar stress score if they filled in the questionnaire again next week, with nothing changed, and the six stress items should agree with each other. **Validity** is whether the measure captures what it claims to measure: the stress score should reflect stress, and not something else, such as general negativity or tiredness on the day of the survey.

The classic picture is a target. Each shot is a measurement, and the centre is the true value. The code below simulates four measures, each shooting 30 times, and draws them; it is intuition code, written to show an idea rather than to analyse data:

```
set.seed(1)
shots <- function(label, centre_x, centre_y,
  spread) {
  data.frame(measure = label,
    x = rnorm(30, centre_x, spread),
    y = rnorm(30, centre_y, spread))
}
target_data <- rbind(
  shots("Reliable and valid", 0,
0, 0.15),
  shots("Reliable but not valid", 0.9,
0.8, 0.15),
  shots("Valid on average, not reliable", 0,
0, 0.7),
```



```
shots("Neither reliable nor valid", 0.8,
      -0.6, 0.7)
)
rings <- expand.grid(angle = seq(0, 2 * pi,
length.out = 100), radius = 1:3 / 2)
rings$x <- rings$radius * cos(rings$angle)
rings$y <- rings$radius * sin(rings$angle)

ggplot(target_data, aes(x, y)) +
  geom_path(data = rings, aes(group = radius),
colour = "grey70") +
  geom_point(colour = "#b2182b", alpha = 0.8) +
  annotate("point", x = 0, y = 0, shape = 3,
size = 4) +
  facet_wrap(~ measure) +
  coord_equal(xlim = c(-1.7, 1.7), ylim =
c(-1.7, 1.7)) +
  theme_void() +
  theme(strip.text = element_text(size = 11,
margin = margin(b = 4)))
```




Figure 5.4: Reliability and validity as shots at a target. Reliable measures are tightly grouped; valid measures are centred on the true value. A measure can be reliable without being valid, but not valid without being reasonably reliable.

A measure that is reliable but not valid is precise about the wrong thing: a bathroom scale that always shows two kilograms too much. A measure that is unreliable cannot be very valid, because its random errors hide whatever it measures. Reliability is therefore necessary for validity, but not enough.

5.5.1 - Kinds of reliability

Reliability can be assessed in three main ways. **Internal consistency** is the extent to which the items of a scale

agree with each other; it is measured with Cronbach's alpha (Chapter 9). **Test-retest reliability** is the extent to which the same person gets a similar score on two occasions when nothing has changed. **Inter-rater reliability** is the extent to which two people who code the same material agree; it is measured with Cohen's kappa (Chapter 18, where a language model is one of the coders).

5.5.2 - Kinds of validity

Validity has several meanings. For a measure, **content validity** concerns whether the items cover the whole construct: a stress scale about deadlines only would miss stress about money or family. **Construct validity** concerns whether the score behaves as the construct should. It should correlate with measures of related constructs (**convergent validity**) and correlate less with unrelated ones (**discriminant validity**) [10].

The second kind can be checked in the wellbeing data. If the stress score measures stress, it should correlate strongly with burnout (a closely related construct), less strongly and negatively with supervisor support and satisfaction:

```
scale_scores <- questionnaire |>
  mutate(
    stress_4      = 6 - stress_4,
    stress        = rowMeans(pick(stress_1:stress_6), na.rm = TRUE),
    burnout       = rowMeans(pick(burnout_1:burnout_6), na.rm = TRUE),
    support       = rowMeans(pick(support_1:support_6), na.rm = TRUE),
    satisfaction   = rowMeans(pick(satisfaction_1:satisfaction_4), na.rm = TRUE)
  ) |>
  select(stress, burnout, support, satisfaction)
round(cor(scale_scores), 2)
```

	stress	burnout	support	satisfaction
stress	1.00	0.65	-0.33	-0.41
burnout	0.65	1.00	-0.24	-0.38
support	-0.33	-0.24	1.00	0.41
satisfaction	-0.41	-0.38	0.41	1.00

Stress and burnout correlate at 0.65: related, as expected, but far from identical, so the two scales are not measuring the same thing twice. The negative correlations with support and satisfaction fit the idea that stressed students feel less supported and less satisfied. This is not proof of validity, which is built from many pieces of evidence, but it is the kind of evidence a thesis should report. Chapter 9 tests the structure of the whole questionnaire with factor analysis.

For a study as a whole, two further kinds of validity matter. **Internal validity** is the extent to which a study can rule out other explanations for its results; it is highest in randomised experiments. **External validity** is the

extent to which the results apply beyond the study, to other people, places, and times; it depends on the sample. The next two sections are about these.

5.6 - Study designs

In one sentence: the design, not the statistics, decides whether a study can show cause.

A **study design** is the plan for who is measured, on what, when, and under which conditions. Table 5.7 summarises the main designs.

Table 5.7: The main study designs

Design	What happens	Can it show cause?	In the wellbeing study
Randomised experiment	The researcher assigns the treatment by chance	Yes, for the treatment	The workshop invitation
Quasi-experiment	Groups receive different treatments, but not by chance	Only with care; other differences must be ruled out	Comparing students who <i>chose</i> to attend many sessions
Observational, cross-sectional	Variables measured once, at one time	No; association only	The baseline questionnaire
Observational, longitudinal	The same people measured repeatedly over time	Stronger evidence of order, but still association	The four semesters

5.6.1 - Why random assignment works

Suppose the workshop had been open to anyone who wanted to come. The students who came would probably have had more energy, more time, and less stress to start with. If the attenders later had higher wellbeing, the cause could be the workshop or who they already were, and the two explanations could not be separated. This is **self-selection**, a form of confounding.

A simulation makes the problem visible. Take the real stress scores of the 600 students, and imagine that less stressed students are more likely to volunteer. Then compare the baseline stress of volunteers and non-volunteers, and do the same for a random invitation:

```
set.seed(11)
stress_scores <- scale_scores$stress
```

```
# Volunteering: the chance of volunteering
falls as stress rises
chance_to_volunteer <- plogis(-2 *
(stress_scores - mean(stress_scores)))
volunteered <- runif(600) <
chance_to_volunteer

# Random invitation: a coin toss for every
student
invited <- sample(rep(c(TRUE, FALSE), 300))

data.frame(
  design = c("Volunteers", "Random
invitation"),
  stress_in =
c(mean(stress_scores[volunteered]),
mean(stress_scores[invited])),
  stress_out = c(mean(stress_scores[!
volunteered]), mean(stress_scores[!invited]))
) |>
  mutate(difference = stress_in - stress_out)
|>
  mutate(across(where(is.numeric), \ (x)
round(x, 2)))
```

	design	stress_in	stress_out	difference
1	Volunteers	2.83	3.63	-0.80
2	Random invitation	3.25	3.19	0.06

In this code, `plogis()` turns any number into a probability between 0 and 1, and `runif()` draws a random number between 0 and 1 for each student; a student volunteers when their random number is below their probability. The volunteers start the study 0.8 points less stressed than the others, *before any workshop*. Any later difference in wellbeing would mix the workshop's effect with this head start. With random invitation, the two groups start almost the same, because chance does not favour any kind of student.

Randomisation balances not only the variables you measured, but also the ones you did not: motivation, family support, health, and everything else. That is why a randomised experiment can support a causal claim. The real invitation in the wellbeing study, which was random, can be checked in the same way:

```
students |>
  left_join(stress, join_by(student_id)) |>
  group_by(workshop) |>
  summarise(students = n(),
            average_age = mean(age, na.rm
= TRUE),
            percent_female = 100 * mean(gender
== "Female"),
            percent_phd = 100 * mean(programme
== "PhD"),
            financial_worry =
mean(financial_worry, na.rm = TRUE),
            stress = mean(stress,
na.rm = TRUE)) |>
```

```
mutate(across(where(is.double), \(\x)
round(x, 1)))
```

```
# A tibble: 2 × 7
  workshop      students average_age
percent_female percent_phd financial_worry
  <chr>      <int>      <dbl>      <dbl>
<dbl>      <dbl>
1 Invited      300      29.6      50.3
28.7          2.8
2 Not invited   300      29.8      53.7
29.3          2.9
# i 1 more variable: stress <dbl>
```

The two groups look very similar at baseline. A balance table like this belongs in any thesis with an experiment: it shows the reader that the randomisation worked.

Randomisation is about the treatment only

The random invitation allows causal claims about **the workshop invitation**, and nothing else. The study's other questions, about sleep, stress, caffeine, and GPA, are observational: nobody assigned students their sleep. For those, the thesis can report associations, allow for the confounders that were measured, and discuss the ones that were not. Note, too, that what was randomised is the *invitation*, not attendance: many invited students attended only some sessions. Comparing students by the number of sessions they attended is a quasi-experiment again, because students chose how many sessions to attend.

5.6.2 - Cross-sectional and longitudinal designs

A **cross-sectional** design measures everything at one time. It is quick and cheap, but it cannot show which came first: stress might reduce sleep, or short sleep might produce stress. A **longitudinal** design measures the same people repeatedly, so it can show change and the order of events. The four semesters make the wellbeing study longitudinal, which Chapter 10 uses to model how wellbeing changes. The price is **attrition**: some people leave the study, and those who leave are rarely a random selection, as the next section shows.

5.7 - Samples and populations

In one sentence: a large sample reduces random error, but not bias.

The **population** is everyone the conclusions are meant to apply to: graduate students, perhaps at one university, perhaps everywhere. The **sample** is the people actually studied. Since the results come from the sample and the conclusions are about the population, the central question is always how well the sample represents the population.

5.7.1 - Ways to draw a sample

In **simple random sampling**, every member of the population has the same chance of being chosen. It is the ideal, and rarely possible, because it needs a complete list of the population, called a **sampling frame**. In **stratified sampling**, the population is divided into groups (strata), such as faculties, and a random sample is drawn from each, to make sure every group is represented. In **cluster sampling**, whole groups are sampled, such as all the students of randomly chosen supervisors; it is cheaper, but students in the same cluster are alike, which the analysis must allow for (Chapter 10). **Convenience sampling** takes whoever is easy to reach, such as the people who answer an online survey shared on social media, or the students in the researcher's own classes. It is very common in theses, and it is the weakest basis for generalising.

5.7.2 - Random error and bias

A sample estimate can be wrong in two ways. **Random error** is the difference caused by which people happened to be selected: a different random sample would give a slightly different answer. **Bias** is a systematic difference: the sampling method tends to select certain kinds of people, so the estimate is off in one direction, every time.

The difference matters because they respond differently to sample size. A simulation shows it. Treat the study's 600 students as a complete population, whose true average first-semester wellbeing is 60.5. Draw many samples in two ways: at random, and as a convenience sample where students with higher wellbeing are more likely to answer (a common pattern: people who are struggling are less likely to fill in surveys). Try samples of 50 and 200:

```
population <- first_sem$wellbeing
true_mean <- mean(population)

# A student's chance of answering a voluntary
# survey rises with their wellbeing
chance_to_answer <- plogis((population -
  true_mean) / 10)

draw_samples <- function(n, method) {
  replicate(2000, {
    if (method == "Random") {
      chosen <- sample(length(population), n)
    } else {
      chosen <- sample(length(population), n,
        prob = chance_to_answer)
    }
    mean(population[chosen])
  })
}

set.seed(5)
sampling_results <- expand_grid(n = c(50,
  200), method = c("Random", "Voluntary")) |>
  rowwise() |>
```

```
mutate(estimate = list(draw_samples(n,
method))) |>
tidyr::unnest(estimate) |>
mutate(sample_size = paste(n, "students"))

sampling_results |>
group_by(method, sample_size) |>
summarise(average_estimate =
round(mean(estimate), 1),
spread = round(sd(estimate),
2),
.groups = "drop")
```

```
# A tibble: 4 × 4
  method      sample_size average_estimate
spread
<fct>      <chr>          <dbl> <dbl>
1 Random    200 students        60.4  0.73
2 Random    50 students         60.5  1.63
3 Voluntary 200 students        65.3  0.6
4 Voluntary 50 students         65.9  1.43
```

In this simulation, `sample()` with `prob` chooses students with unequal chances, like a voluntary survey; `rowwise()` and `unnest()` run the simulation once for each combination of sample size and method and collect the results in one table. Figure 5.5 draws them:

```
ggplot(sampling_results, aes(x = estimate,
fill = method)) +
  geom_histogram(binwidth = 0.4, alpha = 0.7,
position = "identity") +
  geom_vline(xintercept = true_mean, linetype =
"dashed") +
  facet_wrap(~ sample_size, ncol = 1) +
  scale_fill_manual(values = c(Random =
"grey50", Voluntary = "#d6604d")) +
  labs(x = "Estimated average wellbeing", y =
"Number of samples", fill = "Sample")
```




Figure 5.5: Estimates of average wellbeing from 2,000 samples, drawn at random or by voluntary response. The dashed line is the true population average. A larger sample narrows the spread (random error) but does not move a biased estimate back to the truth.

Random samples are centred on the true value; larger random samples are simply more precise. Voluntary samples miss the true value, and the larger voluntary sample misses it just as much, only more confidently. **A biased sample cannot be fixed by making it bigger.** The only remedies are a better sampling method, or knowing enough about the bias to correct for it.

5.7.3 - Non-response and attrition

The same problem appears in the real data. Not every student answered the final open-ended question, and some left the study after the first year. The comparison below shows whether the students who remained are like those who did not:

```
leavers <- setdiff(students$student_id,
  semesters$student_id[semesters$semester ==
  4])

students |>
  mutate(stayed = if_else(student_id %in%
  leavers, "Left after year 1", "Stayed")) |>
  group_by(stayed) |>
  summarise(students = n(),
    percent_considered_dropout
  = round(100 * mean(considering_dropout ==
  "Yes")))
```

```
# A tibble: 2 × 3
  stayed          students
percent_considered_dropout
<chr>          <int>      <dbl>
1 Left after year 1      37        59
2 Stayed              563        12
```

Of the 37 students who left, 59% had considered dropping out, compared with 12% of those who stayed. Any analysis of semesters 3 and 4 is therefore based on a group that is less likely to be struggling than the students who started. Chapter 6 describes this missing data in detail, and Chapter 10 uses a model that makes the best use of the incomplete records.

5.7.4 - Generalising from the sample

The wellbeing study's sample is all graduate students at one university who agreed to take part. Its results apply most directly to that university, and to others like it, only by argument: similar students, similar programmes, similar pressures. A thesis should say this plainly in its limitations section. Claiming less than the data allows is rarely criticised; claiming more is.

5.8 - Planning the analysis before the data

In one sentence: decide how you will judge your hypotheses before the data can influence you.

5.8.1 - An analysis plan

For every hypothesis, the plan names the variables, the analysis, and what would count as support. Writing it before collecting data forces you to check that you will collect everything you need, in a form you can analyse:

Table 5.8: Part of the study's analysis plan

Hy- pothesis	Out- come	Predic- tor(s)	Analy- sis	Support if
Stu- dents sleep less than 7 hours	<code>sleep_hours</code> (semes- ter 1)	none	One- sample t-test against 7	The 95% confi- dence interval lies be- low 7
The work- shop raises wellbe- ing	<code>wellbeing</code> (semes- ter 2)	workshop	Two- sample t-test; mixed model over se- mesters	Invited stu- dents higher; inter- val ex- cludes 0
Stress raises the odds of con- sidering dropout	<code>considering_dropout</code>	stress score, with back- ground vari- ables	Logistic regres- sion	Odds ratio above 1; inter- val ex- cludes 1

5.8.2 - Sample size

The simulation of the null world showed that chance differences are smaller with larger groups. So the size of the sample decides how small an effect a study can detect. The **power** of a study is the probability that it detects an effect of a given size, if the effect is real. A common target is 80%.

Power depends on three things: the sample size, the size of the effect, and the significance level. Effect sizes are often expressed as Cohen's d : the difference between groups divided by the standard deviation (Chapter 7). Before her study, Elaf's reading suggested that brief well-being programmes improve wellbeing by roughly 0.4 to 0.5 standard deviations. Base R's `power.t.test()` calculates how many students each group needs to detect $d = 0.45$ with 80% power:

```
power.t.test(delta = 0.45, sd = 1, sig.level
= 0.05, power = 0.80)
```

Two-sample t test power calculation

n = 78.49181

```
delta = 0.45
sd = 1
sig.level = 0.05
power = 0.8
alternative = two.sided
```

NOTE: *n* is number in *each* group

With `sd = 1`, `delta` is the effect in standard deviations, which is Cohen's *d*. The answer, about 79 students per group, is well below the study's 300 per group. Smaller effects need far larger samples:

```
effect_sizes <- c(small = 0.2, medium = 0.5,
large = 0.8)
sapply(effect_sizes,
\ (d) ceiling(power.t.test(delta = d, power =
0.80)$n))
```

small	medium	large
394	64	26

Halving the effect size roughly quadruples the sample needed. If the effect you expect is small, and your sample is small, a non-significant result tells you very little: the study could not have detected the effect even if it were there. Plan the sample size *before* collecting data, and report the calculation in the methods chapter. Chapter 7 returns to power, with a simulation that shows what it means. The `pwr` package covers many more designs than `power.t.test()`.

5.8.3 - Preregistration and honest analysis

Every analysis involves many small decisions: which students to exclude, which variables to control for, how to handle outliers, which test to use. Made after seeing the data, each decision can be nudged, often unconsciously, towards the result the researcher hoped for. With enough such choices, “significant” findings appear from pure noise.

The protection is to decide in advance. **Preregistration** means writing the hypotheses and analysis plan down, with a date, before analysing the data, for example on the Open Science Framework [11]. Afterwards, analyses that follow the plan are reported as confirmatory, and anything else is reported honestly as exploratory. Preregistration does not forbid exploring; it only makes clear which results were predicted and which were found. Chapter 17 returns to this as part of reproducible research.

5.8.4 - Ethics

Research with people requires ethical approval before any data is collected. The committee will ask the questions in this chapter, too: what the question is, why the data is needed, how participants are recruited and informed, and how their data is stored and anonymised. A clear research

question and analysis plan make the application much easier, and they justify collecting only the data you need.

💡 Writing it up

The methods chapter of a thesis reports the decisions of this chapter, usually under the headings Design, Participants, Measures, and Analysis. A short example for the wellbeing study, with the numbers filled in by R:

Design. A two-year longitudinal study with an embedded randomised experiment: at the end of semester 1, half the participants were randomly invited to a six-week wellbeing workshop.

Participants. 600 graduate students from 5 faculties took part, supervised by 120 supervisors; 37 (6%) left the study after the first year.

Measures. Stress, burnout, supervisor support, and academic satisfaction were measured at baseline with a 22-item questionnaire (1 = strongly disagree, 5 = strongly agree; items in Appendix A); scale scores are item averages, with one reverse-worded item reversed. Wellbeing (0 to 100), sleep, study hours, caffeine, and exercise were self-reported each semester; GPA was taken from university records.

Analysis. Hypotheses and the analysis plan were written before the data was analysed. With 300 students per group, the workshop comparison had 80% power to detect an effect of $d = 0.23$ at $\alpha = .05$.

The last sentence uses `power.t.test()` the other way round: given the sample size and the power, it finds the smallest effect the study could reliably detect.

In your field: a laboratory experiment

R's built-in `ToothGrowth` data comes from a classic experiment on guinea pigs [12]. Sixty animals were assigned to receive vitamin C at one of three doses (0.5, 1, or 2 mg a day), by one of two methods (orange juice or ascorbic acid), and the length of the cells responsible for tooth growth was measured.

```
table(ToothGrowth$supp, ToothGrowth$dose)
```

```
    0.5  1  2
OJ   10 10 10
VC   10 10 10
```

```
ToothGrowth |>
  group_by(supp, dose) |>
    summarise(average_length =
      round(mean(len), 1), .groups = "drop")
```

```
# A tibble: 6 × 3
  supp  dose average_length
<fct> <dbl>         <dbl>
1 OJ    0.5          13.2
2 OJ    1           22.7
3 OJ    2           26.1
4 VC    0.5           8
5 VC    1           16.8
6 VC    2           26.1
```

The same planning applies. The **design** is an experiment with two factors and ten animals per combination. The **outcome**, `len`, is a ratio measure; the method of delivery, `supp`, is nominal; the dose is ratio in principle, but with only three values it is often treated as an ordered factor. A falsifiable **hypothesis**: *at the same dose, orange juice produces longer tooth cells than ascorbic acid*; H_0 : *the method makes no difference*. It would count against the hypothesis if the ascorbic-acid groups were as long or longer. With only ten animals per group, the study has good power only for large effects, which is typical of laboratory work and a limitation worth stating.

5.9 - Common misconceptions

- **“A hypothesis is a question.”** A hypothesis is a prediction: it says what the answer will be, precisely enough to be wrong.
- **“Rejecting the null hypothesis proves my hypothesis.”** It shows only that “nothing is going on” explains the data poorly. Other explanations, including confounders, remain possible unless the design rules them out.
- **“A non-significant result proves there is no effect.”** It may only mean the study was too small to detect it.

- **“A big sample makes the results representative.”** Size reduces random error, not bias. A large convenience sample is still a convenience sample.
- **“If it is a number, I can average it.”** The level of measurement decides which summaries are meaningful.
- **“A reliable measure is a valid measure.”** A measure can be consistent and still measure the wrong thing.
- **“Correlation with a control variable proves cause.”** Only randomisation balances the confounders you did not measure.

5.10 - Chapter review

5.10.1 - Summary

- Data analysis tests claims against evidence. It describes, explains, or predicts, and each aim is judged differently.
- A research question narrows a topic until data can answer it. Good questions are specific, answerable, feasible, and worth answering; they are descriptive, relational, or causal.
- A hypothesis is a prediction stated in advance. It must be falsifiable: it names the result that would count against it. Directional hypotheses predict a direction; non-directional ones predict a difference.
- Tests compare the data with the null hypothesis, the world where nothing is going on. Simulation shows the differences chance alone produces.
- Constructs are operationalised as variables. The unit of analysis says what a row is. Levels of measurement (nominal, ordinal, interval, ratio) decide which summaries make sense, and R must be told about ordered categories.
- Variables play roles: outcome, predictor, confounder, moderator, mediator. Confounders can create associations that are not causes.
- Reliability is consistency; validity is measuring the right thing. Internal validity is about ruling out other explanations; external validity is about who the results apply to.
- Only randomised experiments support causal claims directly, because randomisation balances measured and unmeasured variables. Observational designs show associations; longitudinal designs show change and order.
- A larger sample reduces random error but not bias. Non-response and attrition can bias a sample.
- Plan the analysis before the data: hypotheses, variables, analyses, and sample size (power). Preregistration separates confirmatory from exploratory results.

5.10.2 - Key terms

Research question, descriptive question, relational question, causal question, hypothesis, falsifiability, directional hypothesis, non-directional hypothesis, null hypothesis, alternative hypothesis, confirmatory analysis, exploratory analysis, unit of analysis, construct, operationalisation, self-report, level of measurement, nominal, ordinal, interval, ratio, ordered factor, outcome, predictor, confounder, moderator, mediator, reliability, validity, internal consistency, test-retest reliability, inter-rater re-

liability, content validity, construct validity, internal validity, external validity, study design, randomised experiment, quasi-experiment, self-selection, cross-sectional design, longitudinal design, attrition, population, sample, sampling frame, simple random sampling, stratified sampling, cluster sampling, convenience sampling, random error, bias, power, preregistration.

5.11 - Exercises

The [playground](#) has these and more, with hints and solutions.

1. Improve these research questions so that data could answer them: (a) “Is exercise good for students?” (b) “What makes supervisors effective?” For each, say whether your version is descriptive, relational, or causal.
2. Write a falsifiable hypothesis, its null hypothesis, and the result that would count against it, for the question: “Do students with children study fewer hours per week?”
3. For each variable in `students`, give its level of measurement. Then convert `financial_worry` into an ordered factor with the labels “Not at all”, “A little”, “Moderately”, “Very”, and “Extremely”, and make a table of it.
4. Change the null-world simulation to groups of 30 students instead of 150. Between which values do 95% of the chance differences now lie? What does this mean for a small study of the workshop?
5. Check the balance of the random workshop invitation on three other baseline variables: `study_mode`, `has_children`, and `lives_away`.
6. Use `power.t.test()` to find how many students per group a study needs to detect a difference of 0.3 standard deviations with 80% power, and with 90% power.

5.12 - Further reading

- *Introduction to Modern Statistics* [13], chapters 1 and 2, introduces variables, study designs, and sampling with many short examples.
- *Experimental and Quasi-Experimental Designs for Generalized Causal Inference* [14] is the standard reference on designs and the kinds of validity.
- “The preregistration revolution” [11] explains why and how to preregister, in a few pages.
- *The Logic of Scientific Discovery* [8], for falsifiability in Popper’s own words.

6. Descriptive Statistics and Exploratory Data Analysis

Before a study can test anything, it has to say what it found in the plainest sense: who took part, what a typical value of each variable looks like, how much the participants differ, and whether anything in the data is unusual or missing. These are the questions of **descriptive statistics**, the numbers that summarise what data looks like. They are the first results in almost every thesis, and they are not a formality. The shape of a variable decides which later tests are appropriate, an unusual value can dominate an analysis, and missing data can quietly change who the results are about.

This chapter develops the ideas behind the main descriptive statistics: what “typical” means, what spread measures, why the shape of a distribution matters, how to think about unusual and missing values, and how two variables can be described together. It answers the study’s first research question, what graduate student life looks like (RQ1), with numbers, and checks the data for the problems that could mislead every later analysis. In the story, the data is now clean and has been graphed; Elaf is eager to test her hypotheses, and her supervisor asks her to describe the sample first.

💡 By the end of this chapter you will be able to

- Explain the place of exploratory analysis in the research workflow.
- Choose a summary that suits the level of measurement and the shape of a variable.
- Summarise the centre and spread of a variable, and explain what the standard deviation measures.
- Describe the shape of a distribution, and check it against the normal distribution.
- Find unusual values, and decide what to do with them.
- Describe how much data is missing, why it is missing, and why that matters.
- Measure and interpret correlations, and explain why correlation is not causation.
- Present a description of your sample in a thesis, with every number accompanied by its spread and its sample size.

Note

Teaching & Practice: [Interactive Lecture Slides](#) · [Browser Playground Exercises](#)

6.1 - The place of exploration

Chapter 5 distinguished three aims of data analysis: describing, explaining, and predicting. They correspond to three kinds of analysis. **Exploratory analysis** gets to know the data through summaries, graphs, unusual values, and missing values; it is open-ended, it proves nothing, but it often suggests questions worth testing, and this chapter is exploratory. **Inferential analysis** uses a sample to draw conclusions about a wider population and asks whether a result could be due to chance, such as whether the workshop improves wellbeing for graduate students in general, not just for these 600; Chapters 7 to 10 are inferential. **Predictive analysis** builds models that make accurate predictions for new cases, such as which students are likely to consider dropping out next year; Part 3 is predictive.

The three fit into one workflow, shown in Figure 6.1. Exploration comes first, and the researcher returns to it whenever a later result is surprising.




Figure 6.1: The research data analysis workflow. Exploration comes before testing and prediction, and you return to it when results surprise you.

6.2 - The centre of a distribution

The first thing to know about any variable is its typical value. Three measures describe it. The **mean** is the average: add up the values and divide by how many there are. In symbols, for values $x_1, x_2, \dots, x_n$:

$$\bar{x} = \frac{1}{n} \sum_{i=1}^n x_i \quad (6.1)$$

The **median** is the middle value when the values are sorted: half the values are below it and half above. The

mode is the most common value, which is mainly useful for categories.

The mean and median can tell very different stories. Here is the daily caffeine intake, in milligrams, of five students:

```
caffeine <- c(0, 120, 150, 180, 900)
mean(caffeine)
```

```
[1] 270
```

```
median(caffeine)
```

```
[1] 150
```

One student who takes 900 mg a day pulls the mean up to 270 mg, higher than four of the five students. The median, 150 mg, stays with the typical student. The mean is sensitive to extreme values; the median is not.

The same two measures for the study data use each student's first-semester record, together with their background information:

```
library(dplyr)
```

```
Attaching package: 'dplyr'
```

```
The following objects are masked from
'package:stats':
```

```
filter, lag
```

```
The following objects are masked from
'package:base':
```

```
intersect, setdiff, setequal, union
```

```
library(ggplot2)
```

```
first_sem <- semesters |>
  filter(semester == 1) |>
  left_join(students, join_by(student_id))
```

```
first_sem |>
  summarise(
    mean_sleep      = mean(sleep_hours, na.rm
= TRUE),
    median_sleep    = median(sleep_hours, na.rm
= TRUE),
    mean_caffeine   = mean(caffeine_mg, na.rm
= TRUE),
    median_caffeine = median(caffeine_mg, na.rm
= TRUE)
  )
```

	mean_sleep	median_sleep	mean_caffeine	median_caffeine
1	6.483502	6.5	189.4054	160

For sleep, the mean and median are almost identical. For caffeine, the mean is well above the median, just as in the five-student example: a minority of heavy caffeine users pull the mean up. Figure 6.2 shows why.

```
ggplot(first_sem, aes(x = caffeine_mg)) +
  geom_histogram(binwidth = 50, fill =
    "grey70", colour = "white") +
  geom_vline(xintercept =
    mean(first_sem$caffeine_mg, na.rm = TRUE),
    linewidth = 1) +
  geom_vline(xintercept =
    median(first_sem$caffeine_mg, na.rm = TRUE),
    linewidth = 1, linetype =
    "dashed") +
  labs(x = "Caffeine per day (mg)", y = "Number
    of students") +
  theme_minimal(base_size = 13)
```

```
Warning: Removed 3 rows containing non-finite
outside the scale range
(`stat_bin()`).
```




Figure 6.2: Daily caffeine intake in the first semester, with the mean (solid line) and median (dashed line).

For categories, the mode is simply the most common category, which `count()` shows at the top when sorted:

```
students |> count(faculty, sort = TRUE)
```

	faculty	n
1	Health Sciences	154
2	Education	148
3	Social Sciences	116
4	Humanities	95
5	Natural Sciences	87

6.2.1 - Choosing a summary

Which summary is meaningful depends first on the level of measurement of the variable (Chapter 5), and then on its shape. The average faculty does not exist, so a nominal variable is described by counts and percentages. An ordinal variable, such as the five answers to the question about money worries, has a meaningful middle but uneven steps, so the median and percentages suit it better

than the mean. Numeric variables can be described by the mean, but only when they are roughly symmetrical and without extreme values; otherwise the median describes the typical case more faithfully. Table 6.1 brings these rules together.

Table 6.1: Choosing a summary from the level of measurement and the shape of a variable

Variable	Centre	Spread	Exam- ples in the study
Nominal	Mode; per- centage in each cate- gory	none	faculty, gender
Ordinal	Median; percentage in each cat- egory	Interquar- tile range	financial worry, sin- gle ques- tionnaire items
Numeric, roughly symmetri- cal	Mean	Standard deviation	sleep, well- being
Numeric, skewed or with ex- treme val- ues	Median	Interquar- tile range	caffeine, age

6.3 - The spread of a distribution

Two groups can have the same average and still be very different. The sleep of two small groups of students shows how:

```
group_a <- c(6.0, 6.5, 6.5, 7.0, 6.5)
group_b <- c(4.0, 8.5, 5.0, 9.0, 6.0)
mean(group_a)
```

```
[1] 6.5
```

```
mean(group_b)
```

```
[1] 6.5
```

Both average 6.5 hours, but in group A everyone sleeps about the same, while group B ranges from 4 to 9 hours. Measures of **spread** capture this difference, and a mean reported without one hides half of the story.

The **range** is the difference between the largest and smallest values. It is simple, but it depends only on the two most extreme values:

```
range(group_b)
```

[1] 4 9

6.3.1 - The standard deviation

The **standard deviation** (SD) measures how far values typically lie from the mean. The idea is easiest to see in a picture. Figure 6.3 shows the five students of group B, the mean as a dashed line, and each student's distance from the mean as a red line. These distances are the **deviations**, $x_i - \bar{x}$.

```
deviations <- tibble(student = 1:5, sleep =
  group_b)

ggplot(deviations, aes(x = student, y =
  sleep)) +
  geom_hline(yintercept = mean(group_b),
    linetype = "dashed") +
  geom_segment(aes(xend = student, yend =
    mean(group_b)), colour = "#b2182b", linewidth
    = 1) +
  geom_point(size = 3) +
  labs(x = "Student", y = "Hours of sleep") +
  theme_minimal(base_size = 12)
```




Figure 6.3: The sleep of five students (points), their mean (dashed line), and each student's deviation from the mean (red lines). The standard deviation is roughly the typical length of the red lines.

The standard deviation summarises the red lines. Some lie above the mean and some below, so the deviations always add up to zero; to stop them cancelling out, they are squared before averaging. The average squared deviation is the **variance**, and its square root is the standard deviation:

$$s = \sqrt{\frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})^2} \quad (6.2)$$

Dividing by $n - 1$ rather than n gives a slightly better estimate when working with a sample. The square root brings the result back to the original units, so the standard deviation of sleep is measured in hours:

```
sd(group_a)
```

```
[1] 0.3535534
```

```
sd(group_b)
```

```
[1] 2.179449
```

```
mean(abs(group_b - mean(group_b))) # the
average length of the red lines
```

```
[1] 1.8
```

The standard deviation of group B, 2.18 hours, is close to the average length of the red lines, 1.8 hours; it is a little larger because squaring gives more weight to the longest lines. Group A's small standard deviation says that its students hardly differ at all.

6.3.2 - The interquartile range

The **interquartile range** (IQR) is the spread of the middle half of the values. The **quartiles** split the sorted values into four equal parts: a quarter of the values lie below the first quartile (Q_1), half below the median, and three quarters below the third quartile (Q_3). The IQR is $Q_3 - Q_1$. Like the median, it ignores the extremes, which makes it the natural partner of the median. A box plot's box is exactly the IQR.

```
quantile(first_sem$caffeine_mg, c(0.25, 0.5,
0.75), na.rm = TRUE)
```

```
25% 50% 75%
105 160 240
```

```
IQR(first_sem$caffeine_mg, na.rm = TRUE)
```

```
[1] 135
```

The same `summarise()` calculates spread for several variables at once:

```
first_sem |>
  summarise(
    sd_sleep = sd(sleep_hours, na.rm = TRUE),
    sd_gpa    = sd(gpa, na.rm = TRUE),
    iqr_caffeine = IQR(caffeine_mg, na.rm =
TRUE)
  )
```

```
sd_sleep sd_gpa iqr_caffeine
1 1.033795 0.329459 135
```

The standard deviation of sleep is about 1 hour: a typical student sleeps within about an hour of the average of 6.5 hours. The mean belongs with the SD, and the median with the IQR.

6.4 - The shape of a distribution

The centre and spread do not describe everything. The **shape** of a distribution matters too, because it decides which summaries and which tests are appropriate. A distribution is **symmetrical** if its two sides mirror each other, in which case the mean and median are about equal. It is **skewed to the right**, or positively skewed, if it has a long tail of high values, as caffeine does; the mean then lies above the median. It is **skewed to the left**, or negatively skewed, if it has a long tail of low values, and the mean lies below the median. A distribution with one peak is **unimodal**, and one with two peaks is **bimodal**, which often means that two different groups are mixed together. Figure 6.4 compares the shape of sleep and caffeine in the study.

```
library(tidyr)

first_sem |>
  select(sleep_hours, caffeine_mg) |>
  pivot_longer(everything(), names_to =
    "variable", values_to = "value") |>
  ggplot(aes(x = value)) +
  geom_density(fill = "grey80") +
  facet_wrap(~ variable, scales = "free") +
  labs(x = NULL, y = "Density") +
  theme_minimal(base_size = 12)
```

```
Warning: Removed 9 rows containing non-finite
outside the scale range
(`stat_density()`).
```




Figure 6.4: Sleep is roughly symmetrical; caffeine is skewed to the right.

The argument `scales = "free"` lets each panel have its own axes, since hours and milligrams are on very different scales.

The **skewness** statistic puts a number on the asymmetry: 0 for perfect symmetry, positive for a tail to the right, negative for a tail to the left. As a rough guide, values between -1 and $+1$ are mild. The **psych** package calculates it (install it once with `install.packages("psych")`):

```
library(psych)
```

```
Attaching package: 'psych'
```

The following objects are masked from
'package:ggplot2':

```
%+%, alpha
```

```
skew(first_sem$sleep_hours, na.rm = TRUE)
```

```
[1] -0.09604954
```

```
skew(first_sem$caffeine_mg, na.rm = TRUE)
```

```
[1] 1.739451
```

```
skew(students$age, na.rm = TRUE)
```

```
[1] 1.340088
```

Sleep is almost perfectly symmetrical, while caffeine and age are clearly skewed to the right: most graduate students are in their twenties, and fewer are older.

6.4.1 - The normal distribution

Many variables in nature and in research have a similar symmetrical, bell-shaped distribution, called the **normal distribution**. It matters because many statistical tests assume that data, or at least averages of data, follow it approximately (Chapter 7 explains why).

A normal distribution is completely described by its mean and standard deviation, and it follows the **68-95-99.7 rule**: about 68% of values lie within one standard deviation of the mean, 95% within two, and 99.7% within three. The sleep data can be checked against the rule:

```
sleep      <-      first_sem$sleep_hours[!
is.na(first_sem$sleep_hours)]
z <- (sleep - mean(sleep)) / sd(sleep)
c(within_1_sd = mean(abs(z) < 1),
  within_2_sd = mean(abs(z) < 2),
  within_3_sd = mean(abs(z) < 3))
```

```
within_1_sd within_2_sd within_3_sd
0.7070707   0.9579125   0.9949495
```

The proportions are very close to the rule. The values *z* are **z-scores**: how many standard deviations each value is from the mean. They put any variable on the same scale, which makes them useful for spotting unusual values too.

A graph gives a better check than any single number. A **Q-Q plot** (quantile-quantile plot) compares the data with what a normal distribution would give. If the data is normal, the points fall along a straight line:

```
first_sem |>
  select(sleep_hours, caffeine_mg) |>
  pivot_longer(everything(), names_to =
"variable", values_to = "value") |>
  ggplot(aes(sample = value)) +
```



```
stat_qq(alpha = 0.4) +
stat_qq_line() +
facet_wrap(~ variable, scales = "free") +
  labs(x = "Expected if normal", y =
"Observed") +
theme_minimal(base_size = 12)
```

```
Warning: Removed 9 rows containing non-finite
outside the scale range
(`stat_qq()`).
```

```
Warning: Removed 9 rows containing non-finite
outside the scale range
(`stat_qq_line()`).
```




Figure 6.5: Q-Q plots: sleep follows the straight line closely; caffeine curves away from it.

Sleep sits on the line. Caffeine bends away at the upper end, because its largest values are much larger than a normal distribution would produce: the long right tail again.

6.5 - Unusual values

An **outlier** is a value far from the rest. Chapter 3 set values that were *impossible*, such as an age of 250, to missing. Outliers are different: they are possible, just unusual, and they are information before they are a problem. An unusual value may be an error that cleaning missed, a participant who misunderstood a question, or a real case that the research should pay attention to. The first task is therefore to find unusual values and look at them, not to remove them.

Two common rules flag them. The **box plot rule** flags values more than $1.5 \times \text{IQR}$ below the first quartile or above the third quartile; these are the points a box plot draws separately. The **z-score rule** flags values more than 3 standard deviations from the mean, and suits roughly normal variables. For caffeine, which is skewed, the box plot rule is the better choice:

```
q1 <- quantile(first_sem$caffeine_mg, 0.25,
na.rm = TRUE)
q3 <- quantile(first_sem$caffeine_mg, 0.75,
na.rm = TRUE)
```

```
upper_fence <- q3 + 1.5 * (q3 - q1)
upper_fence
```

```
75%
442.5
```

```
sum(first_sem$caffeine_mg > upper_fence,
na.rm = TRUE)
```

```
[1] 35
```

In all, 35 students take more caffeine than the upper fence of 442 mg. For sleep, which is roughly normal, the z-score rule flags 3 students. The most extreme cases combine both:

```
first_sem |>
  filter(sleep_hours <= 4.5, caffeine_mg >=
600) |>
  select(student_id, sleep_hours, caffeine_mg,
study_hours, wellbeing)
```

	student_id	sleep_hours	caffeine_mg
study_hours	wellbeing		
1	S0105	3.5	740
58	52		
2	S0124	4.0	845
46	58		
3	S0195	3.6	840
59	58		
4	S0310	3.5	885
57	38		
5	S0312	3.9	785
67	51		
6	S0371	4.1	650
36	53		
7	S0530	4.3	720
65	44		
8	S0554	3.5	900
26	56		
9	S0599	3.6	610
50	47		

These students sleep very little, take a lot of caffeine, and study long hours. Nothing about these values is impossible, and they describe a real and worrying group of students: exactly the kind of case the thesis is about.

! Never delete a value just because it is unusual

First check whether it is an error (Chapter 3). If it is a real value, keep it: it is part of what you are studying. Report it, and choose summaries and tests that are not thrown off by it, such as the median, or check whether your conclusions change when the unusual cases are left out (a *sensitivity analysis*). Removing inconvenient values to get a cleaner result is a form of research misconduct.

6.6 - Missing data

Almost every dataset has gaps, and the study's data has three kinds, each with its own lesson. The first step is to count them in every variable:

```
semesters |> summarise(across(everything(), ~
  sum(is.na(.x))))
```

```
      student_id semester gpa sleep_hours
study_hours exercise_days caffeine_mg
1             0         0  1         21      23
20            25
supervisor_meetings wellbeing
1             0         0
```

```
students |> summarise(across(everything(), ~
  sum(is.na(.x))))
```

```
      student_id supervisor_id age gender faculty
programme study_mode employment
1             0         0  1         0         0
0             0         0
has_children lives_away financial_worry
workshop workshop_sessions
1             0         0         38
0             0
considering_dropout
1             0
```

A few semester measurements are missing (about 1% each), and 38 students did not answer the question about money worries.

What matters most is not *how much* is missing, but *why*. Statisticians distinguish three situations. Data is **missing completely at random** when the gaps have nothing to do with anything, as when a student skips a question by accident; analyses of the remaining data are then unbiased, just a little less precise. Data is **missing at random** when the gaps depend on something that was measured, such as study mode; analyses can be corrected for it if they include what it depends on. Data is **missing not at random** when the gaps depend on the missing value itself. Students with the most serious money worries, for example, might be the least willing to answer the question about money. This is the hardest case, because the students who answered are no longer typical.

The most important gap in the study's data is not a skipped question: it is the students who left the study after the first year. Chapter 3 found 37 of them. Their first-semester records show whether they resembled everyone else:

```
left_ids <- setdiff(students$student_id,
  semesters$student_id[semesters$semester ==
  3])

first_sem |>
  mutate(left_study = student_id %in%
    left_ids) |>
```

```

summarise(
  students      = n(),
  wellbeing     = mean(wellbeing),
  gpa           = mean(gpa, na.rm
= TRUE),
  considering_dropout =
mean(considering_dropout == "Yes"),
  .by = left_study
)

```

	left_study	students	wellbeing	gpa
	considering_dropout			
1	FALSE	563	60.66075	3.107620
			0.1207815	
2	TRUE	37	57.35135	3.012973
			0.5945946	

They did not. The students who left already had lower wellbeing in their first semester, and most of them had said they were considering dropping out. The students still in the study in semesters 3 and 4 are therefore, on average, the ones who were doing better, and a simple average of the later semesters would overestimate how well graduate students were doing.

At a minimum, a thesis should report how many participants left and how they differed, as above. Some methods cope with this kind of gap better than others: the mixed-effects models of Chapter 10 use every observation each student provided, including the first year of those who left. More advanced techniques, such as *multiple imputation*, are beyond this book, but worth knowing about when a large share of the data is missing.

6.7 - Relationships between variables

Descriptive statistics also describe how variables go together. The **correlation coefficient**, written r , measures how closely two numeric variables follow a straight-line relationship. It runs from -1 to $+1$. A value near $+1$ means that as one variable rises, the other rises too; a value near -1 means that as one rises, the other falls; and a value near 0 means that there is no straight-line relationship, although, as Anscombe's datasets in Chapter 4 showed, there may still be a curved one. In the social and behavioural sciences, correlations of about 0.1 , 0.3 , and 0.5 (positive or negative) are often described as small, medium, and large [15]. These are rough guides, not rules.

The questionnaire scores are needed here. They are calculated as in Chapter 3, reversing `stress_4` first:

```

scores <- questionnaire |>
mutate(
  stress_4      = 6 - stress_4,
  stress        =
rowMeans(pick(stress_1:stress_6), na.rm
= TRUE),
  burnout       =
rowMeans(pick(burnout_1:burnout_6), na.rm
= TRUE),

```

```

      support =
rowMeans(pick(support_1:support_6), na.rm =
TRUE),
      satisfaction = rowMeans(pick(satisfaction_1:satisfaction_4),
na.rm = TRUE)
) |>
  select(student_id, stress, burnout, support,
satisfaction)

first_sem <- first_sem |> left_join(scores,
join_by(student_id))

first_sem |>
  select(stress, burnout, support, wellbeing,
gpa) |>
  cor(use = "pairwise.complete.obs") |>
  round(2)

```

	stress	burnout	support	wellbeing	gpa
stress	1.00	0.65	-0.33	-0.53	-0.31
burnout	0.65	1.00	-0.24	-0.51	-0.27
support	-0.33	-0.24	1.00	0.36	0.36
wellbeing	-0.53	-0.51	0.36	1.00	0.32
gpa	-0.31	-0.27	0.36	0.32	1.00

The option `use = "pairwise.complete.obs"` uses every student who has both values for each pair. The table is rich. Stress and burnout go together strongly; students who feel supported by their supervisor report less stress; stress and burnout go with lower wellbeing; support goes with higher GPA.

A **scatter plot matrix** shows the same relationships as graphs, one small scatter plot for each pair of variables:

```

pairs(first_sem[, c("stress", "support",
"wellbeing", "gpa")],
      pch = 16, col = rgb(0, 0, 0, 0.25))

```




Figure 6.6: Scatter plot matrix of stress, support, wellbeing, and GPA in the first semester.

The function `pairs()` comes with R; `pch = 16` draws filled dots, and the `rgb()` colour makes them transparent.

6.7.1 - Correlation is not causation

A correlation says that two variables go together. It does not say that one causes the other. Chapter 5 introduced the example of caffeine and GPA:

```
first_sem |>
  select(caffeine_mg, sleep_hours, gpa) |>
  cor(use = "pairwise.complete.obs") |>
  round(2)
```

	caffeine_mg	sleep_hours	gpa
caffeine_mg	1.00	-0.64	-0.19
sleep_hours	-0.64	1.00	0.25
gpa	-0.19	0.25	1.00

Students who take more caffeine have lower GPAs, but that does not show that caffeine harms grades. Caffeine also goes strongly with *less sleep*, and sleep goes with *higher* grades, so caffeine may only look harmful because heavy caffeine users sleep less. A third variable that produces a misleading correlation like this is a **confounder**. Chapter 8 shows how regression can separate the two explanations. For now, the lesson is to describe correlations as associations, not as causes.

6.8 - Describing your sample in a thesis

The methods chapter of almost every thesis includes a description of the sample: who took part, and what the main variables look like. It is usually a table, which dplyr builds directly:

```
first_sem |>
  summarise(
    students      = n(),
    female_pct    = 100 * mean(gender ==
"Female"),
    phd_pct       = 100 * mean(programme
== "PhD"),
    part_time_pct = 100 * mean(study_mode ==
"Part-time"),
    age_mean      = mean(age, na.rm = TRUE),
    age_sd        = sd(age, na.rm = TRUE)
  ) |>
  round(1)
```

	students	female_pct	phd_pct	part_time_pct
age_mean				
age_sd				
1	600	52	29	29.5
29.7	4.8			

A summary is complete only with its spread and the number of cases it is based on. A mean of 6.4 hours says little until the reader knows whether students differ by minutes or by hours, and whether the mean comes from 20 students or 600. When values are missing, the number

of cases also differs from variable to variable, so it must be reported for each. The second table therefore gives, for each main variable, the number of values, the mean and SD for symmetrical variables, the median and IQR for skewed ones, and the number missing:

```
describe_var <- function(x) {
  c(n = sum(!is.na(x)),
    mean = mean(x, na.rm = TRUE), sd = sd(x,
na.rm = TRUE),
    median = median(x, na.rm = TRUE), iqr =
IQR(x, na.rm = TRUE),
    missing = sum(is.na(x)))
}
apply(first_sem[, c("sleep_hours",
"study_hours", "caffeine_mg", "gpa",
"wellbeing")],
      describe_var) |>
  t() |>
  round(2)
```

	n	mean	sd	median	iqr	missing
sleep_hours	594	6.48	1.03	6.5	1.40	6
study_hours	592	27.64	13.06	26.0	19.00	8
caffeine_mg	597	189.41	143.62	160.0	135.00	3
gpa	600	3.10	0.33	3.1	0.45	0
wellbeing	600	60.46	12.01	60.0	17.00	0

The small function `describe_var()`, written for this chapter, calculates the six summaries for one variable, and `apply()` applies it to each column; `t()` turns the result so that each variable is a row. Writing your own functions is a skill for later; for now, you can reuse this one.

In the text, the sample is described in a sentence or two, for example:

The sample consisted of 600 graduate students (52% female; 29% PhD students), aged 23 to 52 ($M = 29.7$, $SD = 4.8$). Students slept 6.5 hours per night on average ($SD = 1.0$), and 65% slept less than the recommended 7 hours.

Every number in that sentence comes from the code, so it cannot drift out of step with the data.

💡 An exploration checklist

Before any test, check for each variable you will use:

1. How many values are missing, and why?
2. What is the typical value, and how spread out are the values?
3. What is the shape: symmetrical or skewed, one peak or two?
4. Are there unusual values, and are they errors or real?
5. How does it relate to the other variables?

In your field: earth sciences

R's `faithful` dataset records 272 eruptions of the Old Faithful geyser in Yellowstone National Park. The duration of the eruptions is a classic **bimodal** distribution: most eruptions last either about 2 minutes or about 4.5 minutes, and few last in between.

```
ggplot(faithful, aes(x = eruptions)) +
  geom_histogram(binwidth = 0.2, fill =
    "grey70", colour = "white") +
  labs(x = "Eruption duration (minutes)",
    y = "Number of eruptions") +
  theme_minimal(base_size = 12)
```




Figure 6.7: Old Faithful eruption durations: a bimodal distribution.

The mean of these durations, about 3.5 minutes, describes almost no actual eruption. For a bimodal variable, report the two groups separately, and look for what separates them.

6.9 - Common misconceptions

Descriptive statistics look simple, and that is exactly why their mistakes are common.

- **“The mean is the typical value.”** Only for roughly symmetrical variables. For skewed ones, such as caffeine or income, the median is.
- **“A mean on its own is a result.”** Without its spread and the number of cases, a reader cannot judge it.
- **“Outliers should be removed.”** Real unusual values are part of the data; they are reported and handled with suitable methods, not deleted.
- **“A few missing values do not matter.”** Why data is missing matters more than how much. A small but systematic gap can bias the results.
- **“A correlation of zero means no relationship.”** It means no straight-line relationship. A strong curved relationship can have a correlation near zero.

6.10 - Chapter review

6.10.1 - Summary

- Exploratory analysis describes the data; inferential analysis draws conclusions about a population; predictive analysis predicts new cases. Exploration comes first.
- The level of measurement and the shape of a variable decide which summary is meaningful: percentages for categories, the median for ordinal and skewed variables, the mean for symmetrical ones.
- The mean is sensitive to extreme values; the median is not. The standard deviation is roughly the typical distance of the values from the mean. Report the mean with the SD, and the median with the IQR.
- The shape of a distribution can be symmetrical or skewed, and unimodal or bimodal. Skewness measures asymmetry; histograms, density plots, and Q-Q plots show it.
- The normal distribution follows the 68-95-99.7 rule. z-scores give each value's distance from the mean in standard deviations.
- Outliers are information before they are a problem. Flag them with the box plot rule or z-scores, check them, and never delete real values just because they are unusual.
- Why data is missing matters more than how much. Students who leave a study are rarely a random selection.
- The correlation r measures a straight-line relationship, from -1 to $+1$. Correlation is not causation: a confounder can create a misleading association.
- Every summary in a thesis is reported with its spread and its number of cases.

6.10.2 - Key terms

Descriptive statistics, exploratory analysis, inferential analysis, predictive analysis, mean, median, mode, range, deviation, variance, standard deviation, quartile, interquartile range, skewness, symmetrical, unimodal, bimodal, normal distribution, z-score, Q-Q plot, outlier, missing completely at random, missing at random, missing not at random, attrition, correlation coefficient, confounder.

6.11 - Exercises

The [playground](#) has these and more, with hints and solutions.

1. Calculate the mean, median, SD, and IQR of first-semester [study_hours](#). Decide whether the variable is closer to symmetrical or skewed, and which pair of summaries you would report.
2. Draw a Q-Q plot of first-semester [wellbeing](#), and judge whether it looks normal.
3. Using the box plot rule, count the students with unusually high [study_hours](#) in the first semester, and look at their other values.

4. Compare the first-semester `sleep_hours` of students who later left the study with those who stayed, and say whether they differ in sleep as they did in wellbeing.
5. Find the correlation between `support` and `satisfaction`, classify it with Cohen's guidelines, and explain why it does not show that support causes satisfaction.
6. For each variable in `students`, choose the summary you would report in a thesis, using Table 6.1, and justify each choice in one sentence.

6.12 - Further reading

- *R for Data Science* [2]: chapters “Exploratory data analysis” and “Missing values”.
- *Statistical Power Analysis for the Behavioral Sciences* [15] is the source of the small, medium, and large guidelines for effect sizes, including correlations.

7. Hypothesis Testing and Statistical Inference

Research is almost always about more people than it can measure. A study of 600 graduate students is interesting because of what it says about graduate students in general, and a clinical trial of 200 patients matters because of the patients who will be treated after it. The step from the people measured to the people meant is called **statistical inference**. It rests on one uncomfortable fact: a different sample would have given different numbers. Inference is the discipline of saying how different, and of deciding when a pattern in a sample is strong enough to be believed.

This chapter builds that discipline from the ground up. It starts with what happens when samples are drawn again and again, which leads to standard errors and confidence intervals. It then develops the logic of hypothesis testing by simulation, shuffling group labels until the meaning of a p-value can be seen, before turning to the formal tests that appear in almost every thesis. The hypotheses come from Chapter 5: whether graduate students sleep less than the recommended 7 hours (RQ2), and whether the wellbeing workshop improved wellbeing (RQ3).

💡 By the end of this chapter you will be able to

- Explain sampling distributions, the standard error, and the central limit theorem.
- Calculate and interpret confidence intervals, including bootstrap intervals.
- Explain the logic of a hypothesis test through a permutation test, and interpret a p-value correctly.
- Explain the two kinds of error, statistical power, and the problem of multiple testing, and demonstrate each by simulation.
- Carry out one-sample, two-sample, and paired t-tests, and measure the size of an effect.
- Judge the assumptions of a test, and choose a non-parametric test when they fail.
- Test relationships between categorical variables with chi-square tests.
- Choose an appropriate test for a question, and report the result in a thesis.

Note

Teaching & Practice: [Interactive Lecture Slides](#) · [Browser Playground Exercises](#)

7.1 - Samples and populations

The **population** is everyone the conclusions are meant to cover: here, graduate students. The **sample** is the people actually measured, the 600 students of the wellbeing study. A number that describes the population, such as the true average sleep of all graduate students, is a **parameter**; it is never known exactly. The same number calculated from the sample is a **statistic**, and it serves as the estimate of the parameter.

A second sample would give a slightly different estimate, and a third another. The size of this variation cannot be seen in a single study, but it can be seen in a simulation. For a moment, treat the 600 students as if they were the whole population, draw many small samples from them, and watch how the estimates behave.

7.1.1 - Sampling distributions

Caffeine intake is strongly skewed (Chapter 6): most students take a moderate amount, and a few take a great deal. The simulation draws a sample of 5 students and records their average caffeine, repeats this a thousand times, and then does the same with samples of 30. In the code, `sample()` draws a random sample and `replicate()` repeats the whole step:

```
library(dplyr)
```

```
Attaching package: 'dplyr'
```

```
The following objects are masked from
'package:stats':
```

```
filter, lag
```

```
The following objects are masked from
'package:base':
```

```
intersect, setdiff, setequal, union
```

```
library(ggplot2)
```

```
first_sem <- semesters |>
  filter(semester == 1) |>
  left_join(students, join_by(student_id))

caffeine <- first_sem$caffeine_mg[!
is.na(first_sem$caffeine_mg)]
```

```
set.seed(2026)
means_5 <- replicate(1000,
  mean(sample(caffeine, 5)))
means_30 <- replicate(1000,
  mean(sample(caffeine, 30)))
```

The line `set.seed(2026)` fixes R's random number generator, so that the "random" samples are the same every time the code runs and the results match this book. Any

analysis that involves randomness should fix the seed in this way.

The thousand averages form a **sampling distribution**: the distribution of a statistic over many possible samples. Figure 7.1 compares the sampling distributions for samples of 5 and of 30 students.

```
tibble(
  mean_caffeine = c(means_5, means_30),
  sample_size   = rep(c("Samples of 5",
    "Samples of 30"), each = 1000)
) |>
  mutate(sample_size = factor(sample_size,
    levels = c("Samples of 5", "Samples of 30")))
|>
  ggplot(aes(x = mean_caffeine)) +
  geom_histogram(bins = 40, fill = "grey70",
    colour = "white") +
  facet_wrap(~ sample_size) +
  labs(x = "Average caffeine in the sample
    (mg)", y = "Number of samples") +
  theme_minimal(base_size = 12)
```




Figure 7.1: The sampling distribution of average caffeine intake, for samples of 5 and of 30 students drawn from the student wellbeing data.

The figure carries the three ideas on which the rest of the chapter depends. Both distributions are centred on the average of all 600 students, 189 mg, which means that a sample average neither overestimates nor underestimates on average: it is an **unbiased** estimate. The averages of 30 students are packed far more tightly than the averages of 5, so larger samples give more precise estimates. Finally, although caffeine itself is strongly skewed, and so are the averages of 5 students, the averages of 30 students are almost symmetrical and bell-shaped.

This last observation is the **central limit theorem**: whatever the shape of the data, the averages of reasonably large samples follow an approximately normal distribution. It explains why methods based on the normal distribution work for so many kinds of data. As a rule of thumb, samples of 30 or more are large enough unless the data is extremely skewed.

7.1.2 - The standard error

The spread of a sampling distribution has its own name, the **standard error** (SE). It measures how much an estimate would vary from sample to sample, which is pre-

cisely the uncertainty a researcher needs to know. A real study has only one sample and cannot draw a thousand, but for an average the standard error can be calculated from that single sample:

$$SE = \frac{s}{\sqrt{n}} \quad (7.1)$$

where s is the standard deviation and n the sample size. The formula and the simulation can be compared directly:

```
sd(means_30)                                # spread
of the 1,000 simulated averages
```

```
[1] 25.69209
```

```
sd(caffeine) / sqrt(30)                     # the
formula, from the data
```

```
[1] 26.22184
```

The two agree closely. The square root in the formula has a practical consequence that every researcher planning a study should remember: halving the standard error requires four times as many participants.

7.1.3 - Confidence intervals

An estimate on its own says nothing about its precision. A **confidence interval** (CI) adds that information: a range of values that are plausible for the population parameter, given the data. A 95% confidence interval for a mean is approximately

$$\bar{x} \pm 2 \times SE \quad (7.2)$$

where, more precisely, the 2 is a value from the t -distribution that R works out. The function `t.test()` calculates the interval; here it is for the average first-semester sleep:

```
t.test(first_sem$sleep_hours)$conf.int
```

```
[1] 6.400196 6.566808
attr(,"conf.level")
[1] 0.95
```

The average sleep of graduate students is plausibly between 6.40 and 6.57 hours. The interval is narrow because the sample is large.

The “95%” describes the method rather than this particular interval. If the study were repeated many times, and a 95% confidence interval calculated each time, about 95% of those intervals would contain the true population value. Any single interval either contains it or does not; the confidence lies in the procedure that produced it.

 Mean $\pm$ SE is not a confidence interval

Theses often report “mean $\pm$ SE” and treat it as if it were a confidence interval. It is not: the SE is only half the width of an approximate 95% CI. Report the confidence interval itself, and say what it is.

7.1.4 - Bootstrap confidence intervals

The formula above works for averages. For other statistics, such as the median, there is no simple formula. The **bootstrap** solves this with an idea of remarkable simplicity: treat the sample as if it were the population, and draw many new samples from it *with replacement*, so that each new sample contains some people twice and leaves others out. The spread of the statistic across these resamples estimates its sampling distribution.

Suppose the survey had reached only 40 students, and a confidence interval were needed for their median caffeine intake:

```
set.seed(1)
small_sample <- sample(caffeine, 40)
median(small_sample)
```

```
[1] 150
```

```
boot_medians <- replicate(2000,
  median(sample(small_sample, replace = TRUE)))
quantile(boot_medians, c(0.025, 0.975))
```

```
2.5% 97.5%
110.0 177.5
```

Sampling with replacement, the argument `replace = TRUE`, is what makes this a bootstrap. The middle 95% of the 2,000 bootstrap medians, from the 2.5th to the 97.5th percentile, forms the 95% bootstrap confidence interval. The median of all 600 students, 160 mg, lies inside it. The same few lines work for almost any statistic.

7.2 - The logic of hypothesis testing

A confidence interval describes which values are plausible. A **hypothesis test** addresses a sharper question: whether the data is consistent with a particular claim. The claim that is tested is the **null hypothesis**, H_0 , the statement that nothing is going on: no difference, no relationship. The researcher’s own expectation is the **alternative hypothesis**, H_1 . For the workshop, H_0 states that the workshop has no effect on wellbeing, and H_1 that it does.

Chapter 5 explained why tests are built around the null hypothesis: it is precise enough to calculate with. If the workshop had no effect, any difference between the invited and the not-invited students would be due only to which students happened to receive an invitation. The

test asks whether the observed difference is larger than such chance differences usually are. That question can be answered directly, by simulation, before any formula is introduced.

i Think before you analyse

Question: does the workshop improve wellbeing (RQ3)? It is a causal question, and the random invitation allows a causal answer. **Unit of analysis:** the student, with one wellbeing score each, at the end of semester 2. **Variables:** wellbeing (0 to 100, interval) is the outcome; the invitation (nominal, two groups) is the predictor. **Evidence against H_1 :** a difference near zero, or in favour of the students who were not invited.

7.2.1 - A permutation test

The data needed is each student's workshop group together with their wellbeing at the end of semester 2:

```
semester2 <- semesters |>
  filter(semester == 2) |>
  left_join(students, join_by(student_id))
nrow(semester2)
```

```
[1] 600
```

```
observed_gap <-
mean(semester2$wellbeing[semester2$workshop
== "Invited"]) -
  mean(semester2$wellbeing[semester2$workshop
== "Not invited"])
observed_gap
```

```
[1] 5.25
```

Invited students scored 5.2 points higher. If the null hypothesis were true, the labels “Invited” and “Not invited” would carry no information about wellbeing: each student would have had the same score whichever label they had received. Under that assumption the labels can be shuffled at random among the students, and the difference recalculated. Each shuffle shows one difference that chance alone could have produced. Repeating the shuffle thousands of times builds the **null distribution**, the range of differences to expect when the workshop does nothing:

```
set.seed(2026)
shuffled_gaps <- replicate(5000, {
  shuffled <- sample(semester2$workshop)
  mean(semester2$wellbeing[shuffled ==
"Invited"]) -
  mean(semester2$wellbeing[shuffled == "Not
invited"])
})
summary(shuffled_gaps)
```


Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
-3.6033	-0.6367	0.0100	0.0130	0.6633	3.2700

Here `sample()` without a size simply shuffles the whole column. Figure 7.2 places the observed difference against the 5,000 shuffled ones.

```
ggplot(data.frame(gap = shuffled_gaps), aes(x = gap)) +
  geom_histogram(binwidth = 0.25, fill = "grey70", colour = "white") +
  geom_vline(xintercept = observed_gap, colour = "#b2182b", linewidth = 1) +
  labs(x = "Difference in average wellbeing (invited minus not invited)",
       y = "Number of shuffles") +
  theme_minimal(base_size = 12)
```




Figure 7.2: The null distribution from 5,000 shuffles of the workshop labels, and the observed difference (red line). No shuffle came close to the observed difference.

Shuffled differences cluster around zero and rarely exceed 3 points in either direction. The observed difference lies far beyond all of them. The proportion of shuffles that produce a difference at least as large as the observed one, in either direction, is the **p-value**:

```
mean(abs(shuffled_gaps) >= abs(observed_gap))
```

```
[1] 0
```

Not one of the 5,000 shuffles matched the observed difference. The p-value is not literally zero: a finite number of shuffles can only show that it is smaller than about 1 in 5,000, so it is reported as $p < .001$. Chance alone is a very poor explanation of the data, and the null hypothesis can be rejected.

A test does not always end this way, and the contrast is instructive. Chapter 5 listed GPA by gender as a comparison where no difference was expected. The same shuffling procedure gives a very different picture:

```
gpa_gap
<- mean(first_sem$gpa[first_sem$gender == "Female"]) -
  mean(first_sem$gpa[first_sem$gender == "Male"])
```

```
set.seed(2026)
shuffled_gpa <- replicate(5000, {
  shuffled <- sample(first_sem$gender)
  mean(first_sem$gpa[shuffled == "Female"]) -
  mean(first_sem$gpa[shuffled == "Male"])
})
c(observed = gpa_gap, p_value =
  mean(abs(shuffled_gpa) >= abs(gpa_gap)))
```

```
      observed      p_value
-0.01286325    0.63240000
```

The average GPA of women and men differs by only 0.01 grade points, and more than half of the shuffles produce a difference at least that large. Such a difference is exactly what chance produces, so the data gives no reason to reject the null hypothesis.

This procedure is a **permutation test**. It contains the whole logic of hypothesis testing, and every test in the rest of the chapter follows it: assume the null hypothesis, work out which results it makes likely, and see where the observed result falls. The formal tests differ mainly in using mathematics instead of shuffling to find the null distribution, which is faster and, when their assumptions hold, gives almost the same answer.

7.2.2 - The p-value

The p-value is the probability of a result at least as extreme as the one observed, *if the null hypothesis were true*. A small p-value means the data would be surprising in a world where nothing is going on, and the null hypothesis is **rejected**. By convention, “small” means below 0.05, a threshold called the **significance level**, α , and such a result is called **statistically significant**. A large p-value means the data is consistent with the null hypothesis, which is then **not rejected**. This is not proof that the null hypothesis is true; it means only that the data does not provide strong evidence against it.

! What a p-value is not

A p-value is **not** the probability that the null hypothesis is true, and **not** the probability that the result happened by chance. It is the probability of data like yours, *assuming* the null hypothesis is true. A p-value of 0.03 means: if the workshop had no effect, a difference at least this large would appear in only 3% of studies like this one.

7.2.3 - Type I and Type II errors

Because a test works with probabilities, its decision can be wrong in two ways, shown in Table 7.1.

Table 7.1: The two kinds of error in hypothesis testing

	H_0 is really true	H_0 is really false
Reject H_0	Type I error (a false alarm)	Correct
Do not reject H_0	Correct	Type II error (a missed effect)

The significance level controls Type I errors: with $\alpha = 0.05$, about 5% of tests of a true null hypothesis will be “significant” by chance alone. This can be demonstrated. The simulation below splits the students into two groups completely at random, so that no real difference can exist, tests whether their wellbeing differs, and repeats the whole procedure a thousand times:

```
set.seed(7)
p_values <- replicate(1000, {
  random_group <- sample(rep(c("A", "B"),
    300))
  t.test(first_sem$wellbeing ~
    random_group)$p.value
})
mean(p_values < 0.05)
```

```
[1] 0.065
```

About 6% of the tests are “significant”, even though every difference is pure chance, close to the 5% that the significance level predicts. A single significant result is therefore never the final word.

7.2.4 - Statistical power

A Type II error, missing an effect that is really there, depends mainly on two things: the size of the effect and the size of the sample. A stricter significance level also makes effects harder to detect, which is one reason not to lower it without cause. The probability that a study detects an effect of a given size, when the effect is real, is its **power** (Chapter 5). Power is easiest to understand by simulation. The function below imagines a study of the workshop in which the true effect is 5 points (with a standard deviation of 11, as in the real data), runs a t-test, and reports whether the result was significant. Running it 2,000 times for each group size shows how often studies of that size succeed:

```
simulate_study <- function(n_per_group,
  effect = 5, sd = 11) {
  invited <- rnorm(n_per_group, mean = 60
    + effect, sd = sd)
  not_invited <- rnorm(n_per_group, mean = 60,
    sd = sd)
  t.test(invited, not_invited)$p.value < 0.05
}
```

```
set.seed(3)
group_sizes <- c(10, 20, 40, 80, 150)
simulated_power <- sapply(group_sizes, \(n)
  mean(replicate(2000, simulate_study(n))))
data.frame(students_per_group = group_sizes,
  power = round(simulated_power, 2))
```

	students_per_group	power
1	10	0.17
2	20	0.29
3	40	0.53
4	80	0.82
5	150	0.97

```
power_curve <- data.frame(n = 5:160)
power_curve$power <- power.t.test(n =
  power_curve$n, delta = 5, sd = 11)$power

ggplot(power_curve, aes(x = n, y = power)) +
  geom_line(colour = "grey50") +
  geom_point(data = data.frame(n =
    group_sizes, power = simulated_power), size =
    2.5) +
  geom_hline(yintercept = 0.8, linetype =
    "dashed") +
  labs(x = "Students per group", y = "Power") +
  theme_minimal(base_size = 12)
```




Figure 7.3: Power of a study to detect a 5-point workshop effect, by the number of students in each group: simulated (points) and calculated with `power.t.test()` (line).

With 10 students per group, a real 5-point effect is detected only about 17% of the time; such a study would usually end with “no significant difference”, even though the workshop works. Around 80 students per group are needed to reach the conventional 80% power, marked by the dashed line. The study’s 300 per group gives power close to 100%. The line, calculated with `power.t.test()`, agrees with the simulation, which shows what that function computes. A non-significant result from a small study should therefore be read with great caution: it may say more about the size of the study than about the effect.

7.2.5 - Multiple testing

The 5% false alarm rate applies to *each* test. A study that runs many tests will almost certainly produce some false alarms. With 20 independent tests of true null hypotheses, the chance of at least one “significant” result is

```
1 - 0.95^20
```

```
[1] 0.6415141
```

about 64%. A simulation confirms it. Each simulated study below runs 20 t-tests on pure noise and records whether any of them came out significant:

```
set.seed(8)
any_false_alarm <- replicate(1000, {
  p <- replicate(20, t.test(rnorm(30),
    rnorm(30))$p.value)
  any(p < 0.05)
})
mean(any_false_alarm)
```

```
[1] 0.653
```

Roughly two studies in three find “something”, although there is nothing to find. This is why the main hypotheses of a thesis should be stated in advance (Chapter 5), why a study that reports only its significant results misleads, and why the tests in Chapter 8 correct for multiple comparisons.

7.2.6 - One-tailed and two-tailed tests

A **two-tailed** test looks for a difference in either direction: the workshop might raise *or* lower wellbeing. A **one-tailed** test looks in one direction only. One-tailed tests are easier to pass, which makes them tempting, but they are justified only when the opposite direction is truly impossible or irrelevant, and when the direction was decided before seeing the data. Two-tailed tests are the standard choice, and R’s tests are two-tailed by default. The permutation test above was two-tailed as well: it counted shuffles that were extreme in either direction.

7.2.7 - Statistical and practical significance

A statistically significant result is not necessarily an important one. With a large enough sample, even a tiny, meaningless difference becomes significant. Every result should therefore be reported with *its size*, not only with its significance, and the rest of this chapter does so for each test.

7.3 - The one-sample t-test

The second research question asks whether graduate students sleep less than the recommended 7 hours (RQ2). Chapter 5 stated the hypothesis with a direction: average sleep is *below* 7 hours. The test is nevertheless two-tailed,

with the null hypothesis that average sleep is 7 hours and the alternative that it is not, the cautious choice explained in the section on one-tailed and two-tailed tests. The direction of the result is then read from the estimate. A **one-sample t-test** compares a sample average with a fixed value, given by the argument `mu`:

```
sleep_test <- t.test(first_sem$sleep_hours,
  mu = 7)
sleep_test
```

One Sample t-test

```
data: first_sem$sleep_hours
t = -12.177, df = 593, p-value < 2.2e-16
alternative hypothesis: true mean is not equal
to 7
95 percent confidence interval:
 6.400196 6.566808
sample estimates:
mean of x
 6.483502
```

The output contains three results. The estimate is the sample average, 6.48 hours. The confidence interval runs from 6.40 to 6.57 hours and does not include 7. The test itself reports the t statistic, which measures how far the average lies from 7 in units of standard errors (here 12.2 standard errors below), and a p-value so small that R prints it as “< 2.2e-16”, meaning less than 0.0000000000000002.

Graduate students sleep about half an hour less than recommended, and the difference is far too large to be chance. The confidence interval also answers the practical question: the shortfall is somewhere between about 26 and 36 minutes a night.

i One value per student

The test uses first-semester records only, one value per student. A t-test assumes that the observations are **independent**. Using all four semesters would count each student up to four times, as if they were four different people, and make the result look more certain than it is. Chapter 10 shows how to analyse repeated measurements properly.

7.4 - The two-sample t-test

The workshop question (RQ3) compares two independent groups: the students invited at random to the six-week workshop after their first semester, and those who were not invited. The permutation test has already answered it; the two-sample t-test reaches the same answer with a formula.

7.4.1 - A small example

A small invented example shows the mechanics. Suppose there were only five students in each group, with these wellbeing scores out of 100:

```
invited      <- c(68, 72, 65, 75, 70)
not_invited  <- c(62, 66, 60, 69, 64)

mean(invited)
```

```
[1] 70
```

```
mean(not_invited)
```

```
[1] 64.2
```

The invited group is 5.8 points higher on average. A **two-sample t-test** assesses whether a difference between the averages of two independent groups is larger than chance would produce:

```
small_test <- t.test(invited, not_invited)
small_test
```

```
Welch Two Sample t-test
```

```
data: invited and not_invited
t = 2.5099, df = 7.9411, p-value = 0.03658
alternative hypothesis: true difference in
means is not equal to 0
95 percent confidence interval:
 0.4642934 11.1357066
sample estimates:
mean of x mean of y
 70.0      64.2
```

The p-value is 0.037: if the workshop had no effect, a difference this large would appear in about 4 studies in 100. By the usual 0.05 threshold the result is statistically significant, although with only five students per group it is fragile evidence.

7.4.2 - Testing the workshop

The real study uses the `semester2` table built for the permutation test. Before any test, the data should be inspected, and Figure 7.4 shows wellbeing in each group.

```
ggplot(semester2, aes(x = workshop, y =
wellbeing, fill = workshop)) +
  geom_boxplot(show.legend = FALSE, width =
0.5) +
  labs(x = NULL, y = "Wellbeing (0 to 100)") +
  theme_minimal(base_size = 13)
```




Figure 7.4: Wellbeing at the end of semester 2, by workshop group. Each box shows the middle half of the students; the line inside is the median.

The invited group sits a little higher, but the groups overlap considerably. In the test, the formula `wellbeing ~ workshop` reads “wellbeing by workshop group”:

```
result <- t.test(wellbeing ~ workshop, data
= semester2)
result
```

Welch Two Sample t-test

```
data: wellbeing by workshop
t = 5.469, df = 597.52, p-value = 6.66e-08
alternative hypothesis: true difference in
means between group Invited and group Not
invited is not equal to 0
95 percent confidence interval:
 3.364699 7.135301
sample estimates:
 mean in group Invited mean in group Not
invited
                        65.47                      60.22
```

Invited students scored 5.2 points higher on average (65.5 against 60.2). The 95% confidence interval for the difference runs from 3.4 to 7.1 points, and the p-value is below 0.001, in agreement with the permutation test. Because students were invited **at random**, the study can conclude that the workshop *caused* the improvement, not merely that it accompanied it: randomisation makes the two groups alike in everything except the invitation (Chapter 5).

The heading of the output, “Welch Two Sample t-test”, names the version R uses by default. It does not assume that the two groups have equal spread, which makes it the safer choice and the appropriate default.

7.4.3 - Effect size

A difference of five points on a 100-point scale is hard to judge on its own. A common standard is **Cohen’s d**: the difference between the means divided by the standard

deviation. It expresses the difference in standard deviations, so it can be compared across studies and scales:

```
inv <- semester2$wellbeing[semester2$workshop
== "Invited"]
not <- semester2$wellbeing[semester2$workshop
== "Not invited"]
d <- (mean(inv) - mean(not)) / sqrt((var(inv)
+ var(not)) / 2)
d
```

```
[1] 0.4465413
```

Cohen's guidelines call 0.2 a small effect, 0.5 medium, and 0.8 large [15]. At about 0.45, the workshop's effect is small to medium: a real improvement, but not a transformation. For a six-week workshop, that is a worthwhile result, and exactly the kind of judgement a thesis discussion should make.

7.5 - The paired t-test

The two-sample test compared two different groups of students. A different question concerns change within the same people: whether the invited students' own wellbeing rose from semester 1 to semester 2. When the same students are measured twice, the two sets of scores are **paired**. Each student is compared with themselves, which removes the large differences *between* students and makes the test more sensitive.

The data needs one row per student, with the two semesters side by side (Chapter 3):

```
library(tidyr)

invited_wide <- semesters |>
  filter(semester %in% c(1, 2)) |>
  left_join(students, join_by(student_id)) |>
  filter(workshop == "Invited") |>
  select(student_id, semester, wellbeing) |>
  pivot_wider(names_from = semester,
              values_from = wellbeing,
              names_prefix = "semester_")

t.test(invited_wide$semester_2,
invited_wide$semester_1, paired = TRUE)
```

Paired t-test

```
data:      invited_wide$semester_2      and
invited_wide$semester_1
t = 10.936, df = 299, p-value < 2.2e-16
alternative hypothesis: true mean difference
is not equal to 0
95 percent confidence interval:
 3.955360 5.691307
sample estimates:
mean difference
 4.823333
```

The invited students' wellbeing rose by 4.8 points on average, and the change is clearly significant.

This test cannot say *why* wellbeing rose. The rise could come from the workshop, or semester 2 could simply be less stressful than semester 1 for everyone. The paired test has no comparison group; it is the two-sample test, comparing invited with not-invited students, that answers the causal question. Choosing the design that actually answers the question matters as much as running the test correctly.

7.6 - Assumptions of the t-test

Every test rests on assumptions, and a result is only as trustworthy as the assumptions behind it. The first assumption of the t-test is **independence**: each observation comes from a different person, or, for the paired test, each pair does. Independence is a property of the study design rather than of the numbers, so it cannot be tested; it must be secured when the data is collected, which is why the tests in this chapter use one value per student. The second is **approximate normality** of the data in each group. Thanks to the central limit theorem, this matters little in large samples: with 30 or more students per group, moderate skewness is not a problem. The third, **similar spread** in the two groups, is required only by the classic version of the two-sample test; Welch's version, R's default, does not need it.

Normality is best judged with graphs: a histogram and a Q-Q plot (Chapter 6). Formal tests of normality exist, such as the Shapiro-Wilk test, but they contain a trap:

```
shapiro.test(first_sem$sleep_hours)
```

```
Shapiro-Wilk normality test
```

```
data: first_sem$sleep_hours  
W = 0.99291, p-value = 0.006536
```

The test declares that sleep is *not* normal ($p < 0.05$), yet in Chapter 6 the histogram and Q-Q plot of sleep looked almost perfectly normal. Both are right. With 600 students, the test is sensitive enough to detect tiny departures from normality that make no practical difference. In large samples, normality tests reject almost everything; in small samples, where normality matters most, they can miss real problems. The graphs, together with the sample size, are the better guide.

7.7 - Non-parametric tests

When data is clearly not normal and the sample is small, or when the data is ordinal (ranks, or answers on a short scale), **non-parametric tests** offer an alternative. Instead of the values themselves, they compare the **ranks** of the values, so extreme values and skewness do not distort them. The price is some loss of power when the data is in fact close to normal.

The **Mann-Whitney U test**, called the Wilcoxon rank-sum test in R, replaces the two-sample t-test. Caffeine is strongly skewed, which makes the comparison of caffeine intake between women and men a suitable case:

```
tapply(first_sem$caffeine_mg,
first_sem$gender, median, na.rm = TRUE)
```

```
Female   Male
 152.5   170.0
```

```
wilcox.test(caffeine_mg ~ gender, data =
first_sem)
```

```
Wilcoxon rank sum test with continuity
correction
```

```
data: caffeine_mg by gender
W = 40551, p-value = 0.06319
alternative hypothesis: true location shift is
not equal to 0
```

The medians differ a little, but the p-value is above 0.05, so there is no clear evidence of a difference. Medians, rather than means, should be reported alongside a non-parametric test.

The **Wilcoxon signed-rank test** replaces the paired t-test. Applied to the invited students' wellbeing in semesters 1 and 2, it gives:

```
wilcox.test(invited_wide$semester_2,
invited_wide$semester_1, paired = TRUE)
```

```
Wilcoxon signed rank test with continuity
correction
```

```
data: invited_wide$semester_2 and
invited_wide$semester_1
V = 33026, p-value < 2.2e-16
alternative hypothesis: true location shift is
not equal to 0
```

The result agrees with the paired t-test. When the parametric and non-parametric tests agree, either can be reported with confidence; when they disagree, the data deserves a second look, and the test whose assumptions fit should be trusted.

7.8 - Chi-square tests

The t-test compares averages of a numeric variable. For categorical variables, the **chi-square test** (χ^2) compares counts: how many cases fall into each category, against how many would be expected if the null hypothesis were true.

7.8.1 - Test of expected proportions

The first form of the test, usually called the **goodness-of-fit test**, compares the counts of one categorical variable with proportions stated in advance. Whether the sample is evenly split between women and men is a simple example:

```
table(students$gender)
```

```
Female  Male
   312   288
```

```
chisq.test(table(students$gender), p = c(0.5,
0.5))
```

Chi-squared test for given probabilities

```
data:  table(students$gender)
X-squared = 0.96, df = 1, p-value = 0.3272
```

With 312 women and 288 men, the split is a little uneven, but the p-value is large: a difference of this size is quite likely in a random sample from an evenly split population. There is no evidence against a 50/50 split, and a non-significant result of this kind is still a result.

7.8.2 - Test of independence

The **test of independence** examines whether two categorical variables are related, such as having a job and considering dropping out:

```
jobs <- table(students$employment,
students$considering_dropout)
jobs
```

```
           No Yes
Full-time job  72 25
None          265 42
Part-time job 173 23
```

```
prop.table(jobs, margin = 1) |> round(2)
```

```
           No  Yes
Full-time job 0.74 0.26
None          0.86 0.14
Part-time job 0.88 0.12
```

Dividing by the row totals, with `prop.table(..., margin = 1)`, turns the counts into proportions within each row. About a quarter of students with full-time jobs have considered dropping out, compared with about one in eight of the others. The test shows whether this difference exceeds what chance would produce:

```
job_test <- chisq.test(jobs)
job_test
```

Pearson's Chi-squared test

```
data: jobs
X-squared = 10.888, df = 2, p-value = 0.004322
```

With a p-value of 0.004, the difference is statistically significant: students with full-time jobs are more likely to consider dropping out. The result is an association, not proof that jobs cause thoughts of dropping out, since students who work full-time may differ in other ways too, such as financial pressure.

The chi-square test is reliable when the **expected** counts, the counts that would appear if the variables were unrelated, are at least 5 in every cell. The test stores them:

```
round(job_test$expected, 1)
```

	No	Yes
Full-time job	82.4	14.6
None	261.0	46.0
Part-time job	166.6	29.4

All are well above 5. When some are not, Fisher's exact test, `fisher.test()`, is the appropriate alternative for small counts.

7.9 - Choosing a test

Table 7.2 and Figure 7.5 summarise the tests of this chapter and the next.

Table 7.2: Choosing a test

Question	Data	Test	In R
Is a mean equal to a value?	Numeric, one group	One-sample t-test	<code>t.test(x, mu =)</code>
Do two groups differ?	Numeric, two independent groups	Two-sample t-test	<code>t.test(y ~ group)</code>
Do the same people differ at two times?	Numeric, paired	Paired t-test	<code>t.test(x2, x1, paired = TRUE)</code>
Do three or more groups differ?	Numeric, 3+ groups	ANOVA (Chapter 8)	<code>aov()</code>
Do two groups differ (skewed or ordinal)?	Ranks, two groups	Mann-Whitney U	<code>wilcox.test(y ~ group)</code>
Paired, skewed or ordinal?	Ranks, paired	Wilcoxon signed-rank	<code>wilcox.test(x2, x1, paired = TRUE)</code>
Do counts match expected proportions?	One categorical variable	Chi-square goodness of fit	<code>chisq.test(table, p =)</code>
Are two categorical variables related?	Two categorical variables	Chi-square test of independence	<code>chisq.test(table)</code>




Figure 7.5: Choosing a test for comparing groups.

7.10 - Reporting the results

A test is reported with the statistic, its degrees of freedom (df, printed by R), the p-value, and, most importantly, the size of the effect with a confidence interval. Exact p-values are given to two or three decimal places, and very small ones as “ $p < .001$ ”. The report should also return to the hypothesis stated in advance and say whether the data contradicts the null hypothesis. The examples below follow these conventions.

💡 Writing it up

One-sample t-test. Students slept less than the recommended 7 hours ($M = 6.48$, 95% CI [6.40, 6.57]), $t(593) = -12.18$, $p < .001$.

Two-sample t-test. Students invited to the well-being workshop reported higher wellbeing at the end of semester 2 ($M = 65.5$) than those not invited ($M = 60.2$), a difference of 5.2 points, 95% CI [3.4, 7.1], $t(598) = 5.47$, $p < .001$, $d = 0.45$. A permutation test with 5,000 shuffles gave the same conclusion. The null hypothesis of no effect was therefore rejected.

Chi-square test. Considering dropout was related to employment, $\chi^2(2) = 10.89$, $p = .004$: 26% of students with full-time jobs had considered dropping out, compared with 14% of those without a job.

All the numbers in these examples are filled in by the code, so they cannot drift out of step with the analysis.

In your field: health research

The t-test was invented to analyse small medical experiments [16]. R still includes one of the original datasets, `sleep`: the extra hours of sleep that 10 patients gained with each of two sleeping drugs. Because the *same* patients took both drugs, this calls for a **paired** t-test:

```
drug1 <- sleep$extra[sleep$group == 1]
drug2 <- sleep$extra[sleep$group == 2]
t.test(drug2, drug1, paired = TRUE)
```

Paired t-test

```
data: drug2 and drug1
t = 4.0621, df = 9, p-value = 0.002833
alternative hypothesis: true mean
difference is not equal to 0
95 percent confidence interval:
 0.7001142 2.4598858
sample estimates:
mean difference
      1.58
```

Patients slept on average 1.6 hours longer with the second drug. With only 10 patients, the confidence interval is wide, from about 0.7 to 2.5 hours: a clear effect, but an imprecise estimate of its size.

7.11 - Common misconceptions

Several misunderstandings about hypothesis tests are common enough in published research, and in theses, to deserve a list of their own. Each of them has been addressed in this chapter, and each is worth checking against before a results chapter is submitted.

- **“The p-value is the probability that the null hypothesis is true.”** It is the probability of data at least this extreme *if* the null hypothesis is true.
- **“A non-significant result shows that there is no effect.”** A small study can miss a real effect; its power may have been low.
- **“A significant result is an important result.”** Significance says that an effect is unlikely to be zero, not that it is large. The effect size answers the second question.
- **“ $p = 0.049$ and $p = 0.051$ are different findings.”** They are almost identical evidence. The 0.05 threshold is a convention, not a boundary in nature.
- **“Normality tests decide whether a t-test may be used.”** In large samples they reject trivial departures; graphs and sample size are the better guide.
- **“Running many tests is harmless if each one uses $\alpha = 0.05$.”** Each test carries its own 5% risk, and the risks accumulate.

7.12 - Chapter review

7.12.1 - Summary

- A statistic estimates a population parameter. Its sampling distribution shows how it would vary across samples; its spread is the standard error, $s/\sqrt{n}$.
- The central limit theorem: averages of reasonably large samples are approximately normal, whatever the shape of the data.
- A 95% confidence interval gives the plausible values of a parameter. The bootstrap gives confidence intervals for almost any statistic.
- A hypothesis test asks how surprising the data would be if the null hypothesis were true. A permutation test answers this directly by shuffling group labels; formal tests use mathematics to reach the same answer.
- The p-value is the probability of a result at least as extreme as the observed one under the null hypothesis; below 0.05 is conventionally called significant. Significance is not importance: always report the size of the effect.
- Type I errors are false alarms (about 5% of tests of true null hypotheses); Type II errors are missed effects. Power, the chance of detecting a real effect, grows with the sample size. Many tests produce false alarms unless the hypotheses are fixed in advance.
- t-tests compare a mean with a value, two independent groups, or paired measurements. Random assignment allows causal conclusions.
- Assumptions are best judged with graphs; in large samples, formal normality tests reject trivial departures.

- Chi-square tests compare counts of categorical variables; Mann-Whitney and Wilcoxon tests are non-parametric alternatives to t-tests.

7.12.2 - Key terms

Population, sample, parameter, statistic, sampling distribution, central limit theorem, standard error, confidence interval, bootstrap, null hypothesis, alternative hypothesis, permutation test, null distribution, p-value, significance level, statistically significant, Type I error, Type II error, power, multiple testing, one-tailed test, two-tailed test, one-sample t-test, two-sample t-test, Welch test, paired t-test, effect size, Cohen's *d*, independence, chi-square test, expected count, non-parametric test, Mann-Whitney U test, Wilcoxon signed-rank test.

7.13 - Exercises

The [playground](#) has these and more, with hints and solutions.

1. Draw 1,000 samples of 10 students' first-semester `study_hours` and plot the sample averages. Then try samples of 50. Describe how the two sampling distributions differ.
2. Calculate a 95% confidence interval for the average first-semester GPA, and explain in one sentence what it tells you.
3. Test whether the average first-semester wellbeing differs from 60. State the hypotheses, run the test, and write a one-sentence conclusion.
4. Compare the first-semester GPA of full-time and part-time students, first with a permutation test of 5,000 shuffles, then with a two-sample t-test. Calculate Cohen's *d* and compare the two p-values.
5. Use the `simulate_study()` function to find the power of a study with 40 students per group when the true effect is 3 points instead of 5. Explain what the result means for a researcher planning such a study.
6. Test whether living away from family (`lives_away`) is related to considering dropout. Run a chi-square test and check the expected counts.

7.14 - Further reading

- *Introduction to Modern Statistics* [13], a free online textbook, explains sampling distributions, bootstrapping, permutation tests, and hypothesis tests with many more examples, using simulation throughout.
- *Statistical Power Analysis for the Behavioral Sciences* [15] on effect sizes and power.

8. ANOVA and Regression

Most research questions involve more than two groups or more than one explanation. A researcher may want to compare five faculties rather than two, to know which of several factors go together with better grades, or to find what makes a student more likely to consider leaving a programme. Each of these questions asks, in one way or another, how much of the variation in an outcome can be traced to other variables. Students differ in stress, grades, and wellbeing; the task is to find which part of those differences is connected with faculty, sleep, or support, and which part remains unexplained.

This chapter introduces the two families of methods built on that idea, which together are probably the most widely used tools in quantitative research. **Analysis of variance** (ANOVA) compares the averages of three or more groups by splitting variation into the part between groups and the part within them. **Regression** models how an outcome depends on one or more predictors: linear regression for numeric outcomes such as GPA, and logistic regression for yes-or-no outcomes such as considering dropout. In the study, they answer three research questions: whether stress differs between faculties and study modes (RQ4), what explains students' grades (RQ5), and who considers dropping out (RQ9).

💡 By the end of this chapter you will be able to

- Explain how ANOVA partitions variation into between-group and within-group parts, and what the F statistic compares.
- Compare three or more groups with one-way ANOVA, and find which groups differ with post-hoc tests.
- Analyse two grouping factors and their interaction with two-way ANOVA.
- Explain regression as a model of the average outcome at each value of the predictors, and fit and interpret simple and multiple linear regression models, including categorical predictors.
- Use regression to separate a real effect from a confounder, and model curves and interactions.
- Recognise violated assumptions in diagnostic plots.
- Fit and interpret a logistic regression model in terms of odds ratios and probabilities.
- Report ANOVA and regression results in a thesis.

Note

Teaching & Practice: [Interactive Lecture Slides](#) · [Browser Playground Exercises](#)

The chapter uses each student's first-semester records, their background information, and their questionnaire scores, calculated as in Chapter 3:

```
library(dplyr)
```

```
Attaching package: 'dplyr'
```

```
The following objects are masked from
'package:stats':
```

```
filter, lag
```

```
The following objects are masked from
'package:base':
```

```
intersect, setdiff, setequal, union
```

```
library(ggplot2)

scores <- questionnaire |>
  mutate(
    stress_4 = 6 - stress_4,
    stress_5 = stress_5,
    rowMeans(pick(stress_1:stress_6), na.rm = TRUE),
    support_4 = support_4,
    support_5 = support_5,
    rowMeans(pick(support_1:support_6), na.rm = TRUE)
  ) |>
  select(student_id, stress, support)

study <- students |>
  left_join(scores, join_by(student_id)) |>
  left_join(semesters |> filter(semester == 1), join_by(student_id))
```

8.1 - One-way ANOVA

The first question is whether stress differs between the five faculties. One approach would be a t-test for every pair of faculties, but that is 10 tests. Chapter 7 showed that each test has a 5% chance of a false alarm, and that the risks accumulate: across 10 independent tests, the chance of at least one false alarm is about 40%. **Analysis of variance** (ANOVA) avoids the problem by asking one question first: whether the faculty averages differ *at all*.

8.1.1 - Partitioning variation

ANOVA compares two kinds of variation. Stress scores for three small groups show the idea:

```
group_scores <- tibble(
  group = rep(c("A", "B", "C"), each = 4),
  stress = c(2.5, 3.0, 2.8, 3.1, 3.4, 3.8,
            3.5, 3.9, 2.9, 3.3, 3.0, 3.4)
)
group_scores |> summarise(mean =
  mean(stress), .by = group)
```

```
# A tibble: 3 × 2
  group mean
  <chr> <dbl>
1 A     2.85
2 B     3.65
3 C     3.15
```

Figure 8.1 draws the twelve scores, the average of each group, and the overall average. The group averages differ from the overall average: that is variation **between** groups. The scores also differ from their own group's average: that is variation **within** groups.

```
group_means <- group_scores |> summarise(mean
  = mean(stress), .by = group)

ggplot(group_scores, aes(x = group, y =
  stress)) +
  geom_hline(yintercept =
    mean(group_scores$stress), linetype =
    "dashed") +
  geom_point(size = 2.5, position =
    position_nudge(x = rep(c(-0.09, -0.03, 0.03,
    0.09), 3))) +
  geom_errorbar(data = group_means, aes(y =
    mean, ymin = mean, ymax = mean), width = 0.4,
    linewidth = 1) +
  labs(x = "Group", y = "Stress score") +
  theme_minimal(base_size = 12)
```




Figure 8.1: Variation between and within groups. Points are individual scores, short black lines are group averages, and the dashed line is the overall average. The gaps between the black lines and the dashed line are between-group variation; the spread of the points around their black line is within-group variation.

If the groups really come from populations with the same average, the between-group variation should be no larger than the within-group variation would lead one to expect: group averages always differ a little by chance. ANOVA's **F statistic** is the ratio of the two:

$$F = \frac{\text{variation between groups}}{\text{variation within groups}} \quad (8.1)$$

An F close to 1 means that the group averages differ no more than chance would produce; a large F means that they differ more. The function `aov()` fits the model, and `summary()` shows the test:

```
summary(aov(stress ~ group, data =
group_scores))
```

```
          Df Sum Sq Mean Sq F value Pr(>F)
group      2  1.307   0.6533   10.69 0.00419 **
Residuals  9   0.550   0.0611
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*'
0.05 '.' 0.1 ' ' 1
```

The p-value is small: these three groups differ.

The ratio matters more than the size of the differences between averages. The next code keeps the three group averages exactly as they are, but places the scores either close to their averages or far from them, and calculates F each time:

```
offsets <- rep(c(-0.3, -0.1, 0.1, 0.3), 3)
group_avg <- rep(group_means$mean, each = 4)

tight <- group_scores |> mutate(stress =
group_avg + 0.5 * offsets)
wide <- group_scores |> mutate(stress =
group_avg + 4 * offsets)

c(tight = summary(aov(stress ~ group, data =
tight))[[1]][["F value"]][1],
  wide = summary(aov(stress ~ group, data =
wide))[[1]][["F value"]][1])
```

```
      tight    wide
39.2000  0.6125
```

The group averages are identical in both versions, yet F is large when the scores cluster tightly around their averages and small when they spread widely. Differences of about half a point between groups are convincing when students within a group hardly differ, and unremarkable when they differ by more than two points. This is the whole logic of ANOVA, and it is why group averages should never be compared without looking at the spread within groups (Chapter 4).

8.1.2 - The student wellbeing data

The data comes first. Figure 8.2 shows stress in each faculty.

```
ggplot(study, aes(x = faculty, y = stress)) +
  geom_boxplot() +
  labs(x = NULL, y = "Stress score (1 to 5)") +
  theme_minimal(base_size = 12)
```




Figure 8.2: Stress scores by faculty.

The faculties overlap a lot, but Health Sciences sits a little higher and Humanities a little lower. The ANOVA tests whether these differences exceed what chance would produce:

```
stress_anova <- aov(stress ~ faculty, data =
study)
summary(stress_anova)
```

```
          Df Sum Sq Mean Sq F value Pr(>F)
faculty    4   8.72   2.1788   4.152 0.00252 **
Residuals 595 312.19   0.5247
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*'
0.05 '.' 0.1 ' ' 1
```

The p-value is 0.003, so the faculties do differ in stress. The size of the difference is measured by **eta squared** (η^2), the share of all the variation in stress that lies between faculties:

```
ss <- summary(stress_anova)[[1]][["Sum Sq"]]
ss[1] / sum(ss)
```

```
[1] 0.02715776
```

Only about 3% of the differences in stress between students have to do with their faculty; the rest is about the individual students. By common guidelines, an η^2 of 0.01 is small, 0.06 medium, and 0.14 large [15], so this is a real but small effect. For the research question, the answer is that faculty matters, but knowing a student's faculty says very little about how stressed that student is.

8.1.3 - Post-hoc tests

ANOVA says that the faculties differ, but not which ones. **Post-hoc tests** compare every pair while keeping the overall chance of a false alarm at 5%. Tukey's test is the most common:

```
TukeyHSD(stress_anova)
```

```
Tukey multiple comparisons of means
95% family-wise confidence level

Fit: aov(formula = stress ~ faculty, data =
study)

$faculty
      lwr      upr      p adj      diff
Health Sciences-Education      0.19429917
-0.033842346  0.42244068 0.1367660
Humanities-Education      -0.14304647
-0.403603347  0.11751041 0.5614602
Natural Sciences-Education      -0.05877343
-0.326527149  0.20898029 0.9749738
Social Sciences-Education      0.12216527
-0.123607736  0.36793827 0.6535324
Humanities-Health Sciences      -0.33734564
-0.595910542 -0.07878073 0.0035331
Natural Sciences-Health Sciences -0.25307260
-0.518888282  0.01274309 0.0707118
Social Sciences-Health Sciences -0.07213390
-0.315794100  0.17152630 0.9275445
Natural Sciences-Humanities      0.08427304
-0.209834619  0.37838070 0.9352210
Social Sciences-Humanities      0.26521174
-0.009035651  0.53945912 0.0635851
Social Sciences-Natural Sciences 0.18093870
-0.100155231  0.46203263 0.3973784
```

Each row compares two faculties: the difference in average stress, its confidence interval, and an adjusted p-value (`p adj`). Only one comparison is significant: Health Sciences students are more stressed than Humanities students, by about 0.34 points on the 1-to-5 scale. The other differences are small enough to be chance.

8.1.4 - Assumptions of ANOVA

ANOVA makes the same assumptions as the t-test: independent observations, roughly normal data within each group (or large groups), and similar spread in each group. The spread can be checked with **Levene's test**, from the `car` package:

```
library(car)
```

```
Loading required package: carData
```

```
Attaching package: 'car'
```

```
The following object is masked from
'package:dplyr':
```

```
recode
```



```
leveneTest(stress ~ factor(faculty), data =
study)
```

```
Levene's Test for Homogeneity of Variance
(center = median)
      Df F value Pr(>F)
group  4  0.9429 0.4386
      595
```

The p-value is large, so there is no evidence that the spread differs between faculties. When the assumptions fail, the **Kruskal-Wallis test** is the non-parametric alternative, comparing ranks instead of averages:

```
kruskal.test(stress ~ faculty, data = study)
```

```
Kruskal-Wallis rank sum test

data: stress by faculty
Kruskal-Wallis chi-squared = 15.46, df = 4, p-
value = 0.003836
```

It agrees with the ANOVA.

8.2 - Two-way ANOVA

Wellbeing might depend on the programme (Master's or PhD), on study mode (full-time or part-time), or on a combination of both. A **two-way ANOVA** examines all three possibilities at once. The **main effect** of programme is the difference between Master's and PhD students on average, and the main effect of study mode is the difference between full-time and part-time students. The **interaction** asks whether the effect of study mode *depends on* the programme: part-time study might, for example, be harder on PhD students than on Master's students.

An **interaction plot** shows the group averages as lines. If the lines are parallel, there is no interaction: the effect of study mode is the same in both programmes.

```
study |>
  summarise(wellbeing = mean(wellbeing), .by
= c(programme, study_mode)) |>
  ggplot(aes(x = programme, y = wellbeing,
colour = study_mode, group = study_mode)) +
  geom_line(linewidth = 1) +
  geom_point(size = 3) +
  scale_colour_viridis_d(end = 0.8) +
  labs(x = NULL, y = "Average wellbeing",
colour = "Study mode") +
  theme_minimal(base_size = 13)
```




Figure 8.3: Average first-semester wellbeing by programme and study mode.

Part-time students have lower wellbeing in both programmes, and the lines are close to parallel. In the formula, `programme * study_mode` asks for both main effects and their interaction:

```
summary(aov(wellbeing ~ programme *
study_mode, data = study))
```

```

              Df Sum Sq Mean Sq F
value  Pr(>F)
programme              1      316    315.6
2.230 0.135913
study_mode              1    1640    1640.2  11.587
0.000709 ***
programme:study_mode    1      119     118.6   0.838
0.360341
Residuals              596    84364
141.6
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*'
0.05 '.' 0.1 ' ' 1
```

Study mode has a clear effect. Programme does not, and neither does the interaction: part-time study lowers wellbeing by about the same amount whether a student is doing a Master's or a PhD. When an interaction is significant, it should be interpreted before the main effects, because it means that no single “effect of study mode” describes everyone.

8.3 - Linear regression

ANOVA compares groups. **Regression** goes further: it models how an outcome changes with one or more **predictors**, which can be numeric or categorical. It is the workhorse of quantitative research, and ANOVA is in fact a special case of it.

8.3.1 - The regression line

At its heart, regression describes the **average outcome for given values of the predictors**. For sleep and GPA, the question is what the average GPA is among students who sleep 5 hours, among those who sleep 6 hours, and so

on. These averages can be calculated directly, by grouping students into half-hour bands of sleep:

```
sleep_bands <- study |>
  filter(!is.na(sleep_hours)) |>
  mutate(sleep_band = round(sleep_hours * 2) /
2) |>
  summarise(mean_gpa = mean(gpa), students =
n(), .by = sleep_band) |>
  arrange(sleep_band)
sleep_bands
```

	sleep_band	mean_gpa	students
1	3.5	2.882000	5
2	4.0	2.901250	8
3	4.5	2.939231	13
4	5.0	3.022340	47
5	5.5	3.021385	65
6	6.0	3.054951	103
7	6.5	3.094519	104
8	7.0	3.120000	102
9	7.5	3.175957	94
10	8.0	3.320606	33
11	8.5	3.206429	14
12	9.0	3.720000	2
13	9.5	3.160000	1
14	10.0	3.146667	3

Figure 8.4 places these averages over the individual students, together with a straight regression line.

```
ggplot(study, aes(x = sleep_hours, y = gpa)) +
  geom_point(alpha = 0.2, colour = "grey40") +
  geom_smooth(method = "lm", se = FALSE) +
  geom_point(data = sleep_bands, aes(x =
sleep_band, y = mean_gpa, size = students),
  colour = "#b2182b") +
  labs(x = "Hours of sleep per night", y =
"GPA", size = "Students") +
  theme_minimal(base_size = 12)
```

```
`geom_smooth()` using formula = 'y ~ x'
```

```
Warning: Removed 6 rows containing non-finite
outside the scale range
(`stat_smooth()`).
```

```
Warning: Removed 6 rows containing missing
values or values outside the scale range
(`geom_point()`).
```




Figure 8.4: GPA and sleep. Grey points are students; red points are the average GPA in each half-hour band of sleep, sized by the number of students; the line is the linear regression. The line is a smooth summary of the red averages.

The band averages rise gently with sleep, and the line passes close to them, especially where there are many students. The regression line is a smooth summary of these conditional averages: it assumes that the average GPA changes by the same amount for each extra hour of sleep, and estimates that amount from all the students at once. The spread of the grey points around the line is what the model does not explain.

A small example shows the calculation. Here are five students' sleep and GPA:

```
five <- tibble(
  sleep = c(5, 6, 6.5, 7, 8),
  gpa    = c(2.8, 3.0, 3.2, 3.1, 3.4)
)
```

Simple linear regression fits the straight line that best describes how GPA changes with sleep:

$$\text{GPA} = b_0 + b_1 \times \text{sleep} + \text{error} \quad (8.2)$$

The coefficient b_0 is the **intercept**, the predicted GPA for a student who sleeps zero hours, and b_1 is the **slope**, the change in predicted GPA for each extra hour of sleep. “Best” means the line that makes the squared vertical distances between the points and the line, the **residuals**, as small as possible, a method called **least squares**. The function `lm()`, for *linear model*, finds it:

```
lm(gpa ~ sleep, data = five)
```

```
Call:
lm(formula = gpa ~ sleep, data = five)

Coefficients:
(Intercept)      sleep
      1.865         0.190
```

Each extra hour of sleep goes with a GPA that is about 0.19 higher. The intercept is where the line would cross

zero hours of sleep, which no one does; it is needed to place the line, but it rarely means anything on its own.

8.3.2 - The student wellbeing data

For all 600 students, `summary()` gives the full results:

```
gpa_simple <- lm(gpa ~ sleep_hours, data = study)
summary(gpa_simple)
```

```
Call:
lm(formula = gpa ~ sleep_hours, data = study)

Residuals:
    Min       1Q   Median       3Q      Max
-0.97283 -0.20475 -0.00188  0.21898  0.94606

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  2.57713     0.08316   30.99  < 2e-16 ***
sleep_hours   0.08081     0.01267    6.38 3.57e-10 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.3189 on 592 degrees of freedom
(6 observations deleted due to missingness)
Multiple R-squared:  0.06433,    Adjusted R-squared:  0.06275
F-statistic: 40.7 on 1 and 592 DF,  p-value: 3.573e-10
```

The output contains three results that matter for the research question. The **coefficients** table gives the estimate of each coefficient, its standard error, a t statistic, and a p-value testing whether it is zero. The slope for `sleep_hours` is about 0.08: each extra hour of sleep goes with a GPA about 0.08 higher, and the p-value shows that this is not chance. **R-squared** is the share of the variation in GPA that the model explains; here it is only 0.06, so sleep explains about 6% of the differences in GPA. Sleep matters, but it is far from the whole story, as the scatter plot suggested. Finally, the **residual standard error** is the typical distance between a student's actual GPA and the line, about 0.32 grade points.

8.3.3 - Multiple regression

Many things affect grades at once. **Multiple regression** includes several predictors in one model:

```
gpa_model <- lm(gpa ~ sleep_hours + study_hours + stress + support, data = study)
summary(gpa_model)
```

```

Call:
lm(formula = gpa ~ sleep_hours + study_hours
+ stress + support,
    data = study)

Residuals:
    Min       1Q   Median       3Q      Max
-0.86953 -0.19437  0.00526  0.18957  0.76429

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  2.137403    0.145279  14.712  <
2e-16 ***
sleep_hours   0.106041    0.014163   7.487
2.62e-13 ***
study_hours   0.006928    0.001105   6.270
7.06e-10 ***
stress        -0.083526    0.017827  -4.685
3.48e-06 ***
support       0.110733    0.016696   6.632
7.57e-11 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*'
0.05 '.' 0.1 ' ' 1

Residual standard error: 0.2857 on 582 degrees
of freedom
(13 observations deleted due to missingness)
Multiple R-squared:  0.2525,    Adjusted R-
squared:  0.2473
F-statistic: 49.14 on 4 and 582 DF,  p-value:
< 2.2e-16

```

Each coefficient now means the change in predicted GPA for a one-unit increase in that predictor, **holding the other predictors constant**. Among students with the same study hours, stress, and support, each extra hour of sleep goes with a GPA about 0.11 higher. Each extra hour of study per week adds about 0.007, each point of stress lowers GPA by about 0.08, and each point of supervisor support raises it by about 0.11. All four are significant.

Together, the four predictors explain about 25% of the variation in GPA. That may sound modest, but grades depend on many things no survey measures, from ability to luck on exam day. In social and educational research, models that explain a quarter of the variation are common and useful. The **adjusted R-squared** is slightly lower, because it corrects for the number of predictors: adding any predictor, even a useless one, raises R-squared a little, but only a useful one raises the adjusted version.

The output also contains the line (13 observations deleted due to missingness): `lm()` leaves out students with a missing value in any variable of the model. The number of students each model is based on should always be reported.

8.3.4 - Categorical predictors

Categorical variables can be predictors too. Gender, for example, can be added to the model to see whether it matters once the other predictors are taken into account:

```
gpa_gender <- lm(gpa ~ sleep_hours +
  study_hours + stress + support + gender, data
  = study)
summary(gpa_gender)$coefficients
```

	Estimate	Std. Error	t value
Pr(> t)			
(Intercept)	2.123737610	0.147260520	
	14.4216360	1.592391e-40	
sleep_hours	0.106600187	0.014203710	
	7.5050947	2.318148e-13	
study_hours	0.006995075	0.001111694	
	6.2922647	6.163042e-10	
stress	-0.082830266	0.017877519	-4.6332081
	4.447816e-06		
support	0.110525008	0.016709773	6.6143931
	8.474218e-11		
genderMale	0.013815196	0.023831664	
	0.5796992	5.623422e-01	

R turns a categorical predictor into a comparison with a **reference category**, by default the first in alphabetical order: here *Female*. The coefficient `genderMale` is the difference between men and women, holding everything else constant. It is tiny and far from significant: there is no evidence that gender affects GPA. That is a finding too, and worth reporting.

For a variable with more categories, such as faculty, R compares each category with the reference, giving one coefficient for each of the others. A different reference is chosen with `relevel(factor(faculty), ref = "Humanities")`.

8.3.5 - Separating an effect from a confounder

Chapters 5 and 6 found that students who take more caffeine have lower GPAs, but also sleep less. Two explanations are possible, and Figure 8.5 draws them. In the first, caffeine itself harms grades. In the second, sleep is a common cause: students who sleep less both drink more coffee and get lower grades, and caffeine has no effect of its own.





Figure 8.5: Two explanations for the link between caffeine and GPA. Left: caffeine affects GPA. Right: sleep affects both caffeine and GPA, which creates a link between them even if caffeine has no effect.

Regression can tell the two apart, by comparing students who sleep the same amount. Caffeine alone comes first:

```
summary(lm(gpa ~ caffeine_mg, data = study))$coefficients
```


	Estimate	Std. Error	t
value			
Pr(> t)			
(Intercept)	3.1849088129	2.195535e-02	
145.063026	0.000000e+00		
caffeine_mg	-0.0004331688	9.239351e-05	
-4.688303	3.418851e-06		

Caffeine has a significant negative association with GPA. Sleep is then added to the model:

```
summary(lm(gpa ~ caffeine_mg + sleep_hours,
data = study))$coefficients
```

	Estimate	Std. Error	t
value			
Pr(> t)			
(Intercept)	2.6380932284	0.1238298172	
21.3041841	4.534485e-75		
caffeine_mg	-0.0000780726	0.0001194238	
-0.6537443	5.135321e-01		
sleep_hours	0.0738597698	0.0165663746	
4.4584148	9.891313e-06		

Once sleep is held constant, caffeine's coefficient shrinks to almost nothing and is no longer significant, while sleep stays significant. Among students who sleep the same amount, caffeine makes no difference to grades. Caffeine looked harmful only because heavy caffeine users sleep less: sleep was a **confounder**, and the data supports the second explanation.

The general rule is that a confounder is a common cause of the predictor and the outcome. Adjusting for it in a regression compares like with like, students who are the same on the confounder, and removes the misleading part of the association.

! Regression adjusts only for what you include

Regression can remove the effect of a confounder only if you measured it and put it in the model. It cannot adjust for anything you did not measure. That is why, outside a randomised experiment like the workshop, regression shows associations *adjusted for* the variables in the model, not proof of cause.

8.3.6 - Curved relationships

A straight line assumes that every extra hour of study adds the same amount to GPA, which is hard to believe: the difference between 5 and 15 hours a week is probably larger than the difference between 45 and 55. Adding a squared term lets the line curve:

```
gpa_curve <- lm(gpa ~ study_hours +
I(study_hours^2), data = study)
summary(gpa_curve)$coefficients
```

	Estimate	Std. Error
t value		
Pr(> t)		
(Intercept)	2.8630064457	5.380474e-02
53.211044	3.718109e-227	

```
study_hours      0.0183480169 3.914705e-03
4.686947 3.448155e-06
I(study_hours^2) -0.0002874497 6.488183e-05
-4.430358 1.121887e-05
```

In the formula, `I()` tells R to calculate `study_hours^2` before fitting. The squared term is negative and significant: the line bends downwards. Figure 8.6 shows the shape.

```
ggplot(study, aes(x = study_hours, y = gpa)) +
  geom_point(alpha = 0.3) +
  geom_smooth(method = "lm", formula = y ~ x
+ I(x^2)) +
  labs(x = "Study hours per week", y = "GPA") +
  theme_minimal(base_size = 13)
```

```
Warning: Removed 8 rows containing non-finite
outside the scale range
(`stat_smooth()`).
```

```
Warning: Removed 8 rows containing missing
values or values outside the scale range
(`geom_point()`).
```




Figure 8.6: GPA and weekly study hours, with a curved (quadratic) trend line.

The first hours of study make a real difference; the curve then flattens, and reaches its highest point at about 32 hours a week, beyond which more study brings nothing extra. Economists call this **diminishing returns**, and it is worth knowing for any student tempted to study all night instead of sleeping.

8.3.7 - Interactions

The effect of supervisor support need not be the same for every student. An **interaction** term lets the effect of one predictor depend on another. In the formula, `support * programme` gives the effect of support, the effect of programme, and how the effect of support differs between programmes. To make the coefficients easier to read, support is first **centred**: the average is subtracted, so that zero means “average support”:

```
study <- study |> mutate(support_c = support
- mean(support))
gpa_interaction <- lm(gpa ~ support_c *
programme + sleep_hours + study_hours, data
= study)
summary(gpa_interaction)$coefficients
```

	t value	Pr(> t)	Estimate	Std. Error
(Intercept)	19.2065312	5.044803e-64	2.15459277	0.112180214
support_c	5.5524856	4.288901e-08	0.10210747	0.018389506
programmePhD	0.4488851	6.536819e-01	0.01173221	0.026136331
sleep_hours	8.3851901	3.842792e-16	0.11771209	0.014038094
study_hours	5.9142870	5.691946e-09	0.00656252	0.001109605
support_c:programmePhD	3.7429227	1.999717e-04	0.13416356	0.035844599

For Master's students (the reference category), each point of support goes with a GPA about 0.10 higher. The interaction coefficient, `support_c:programmePhD`, says that for PhD students the effect is larger, by about 0.13, making it about 0.24. Supervisor support matters about 2.3 times as much for PhD students, which makes sense: a PhD depends more heavily on the supervisor.

8.3.8 - Model diagnostics

Linear regression assumes that the relationship is linear (or modelled as curved, as above), that the residuals have similar spread everywhere, and that they are roughly normal. Violations show up as patterns in a plot of the residuals against the fitted values. It helps to know what those patterns look like before judging a real model, and simulated data can show them. The code below creates three datasets: one that meets the assumptions, one with a curved relationship fitted by a straight line, and one whose spread grows with the predictor. It fits a straight line to each and plots the residuals:

```
set.seed(10)
x <- runif(200, 0, 10)
simulated <- list(
  "Assumptions met" = 2 + 0.5 * x + rnorm(200),
  "Curved" = 2 + 0.25 * x^2
+ rnorm(200),
  "Spread increases" = 2 + 0.5 * x + rnorm(200,
sd = 0.2 + 0.3 * x)
)

residual_data
<- bind_rows(lapply(names(simulated),
function(name) {
  fit <- lm(simulated[[name]] ~ x)
  tibble(case = name, fitted = fitted(fit),
residual = resid(fit))
}))
```

```
residual_data$case
<- factor(residual_data$case, levels =
names(simulated))

ggplot(residual_data, aes(x = fitted, y =
residual)) +
  geom_hline(yintercept = 0, linetype =
"dashed") +
  geom_point(alpha = 0.5, size = 1) +
  facet_wrap(~ case, scales = "free") +
  labs(x = "Fitted values", y = "Residuals") +
  theme_minimal(base_size = 11)
```




Figure 8.7: Residuals against fitted values for three simulated models. Left: assumptions met, a shapeless band. Middle: a curve fitted with a straight line, a U-shaped pattern. Right: spread that grows with the fitted values, a funnel.

A U-shaped or curved pattern means that the model has missed a curve, and a squared term or a transformation may be needed. A funnel means that the spread is not constant, which makes standard errors, and therefore p-values and confidence intervals, unreliable. With these pictures in mind, the real model can be checked. Calling `plot()` on a model draws four diagnostic plots, of which the first two matter most:

```
par(mfrow = c(1, 2))
plot(gpa_model, which = 1:2)
par(mfrow = c(1, 1))
```




Figure 8.8: Diagnostic plots for the GPA model: residuals against fitted values (left), and a Q-Q plot of the residuals (right).

The left plot resembles the first simulated panel: a shapeless band around zero, with no curve and no funnel. In the right plot, the residuals follow the line, so they are close

to normal. Both look good for the GPA model. Points far off the line in either plot would point to unusual cases worth looking at.

8.3.9 - Reporting a regression

The **broom** package turns model results into tidy tables, ready for a thesis. The function `tidy()` gives the coefficients with confidence intervals, and `glance()` the overall fit:

```
library(broom)
tidy(gpa_model, conf.int = TRUE)
```

```
# A tibble: 5 × 7
  term          estimate std.error statistic
p.value conf.low conf.high
<chr>      <dbl>      <dbl>      <dbl>
<dbl>    <dbl>    <dbl>
1 (Intercept)  2.14        0.145      14.7
6.89e-42  1.85        2.42
2 sleep_hours  0.106       0.0142     7.49
2.62e-13  0.0782     0.134
3 study_hours  0.00693     0.00111    6.27
7.06e-10  0.00476    0.00910
4 stress      -0.0835     0.0178    -4.69
3.48e- 6 -0.119   -0.0485
5 support      0.111     0.0167     6.63
7.57e-11  0.0779     0.144
```

```
glance(gpa_model)
```

```
# A tibble: 1 × 12
  r.squared adj.r.squared sigma statistic
p.value    df logLik    AIC    BIC
<dbl>    <dbl>    <dbl> <dbl>    <dbl>
<dbl> <dbl> <dbl> <dbl> <dbl>
1      0.252      4 -95.0  202.  228.
# i 3 more variables: deviance <dbl>,
df.residual <int>, nobs <int>
```

💡 Writing it up

A multiple regression of first-semester GPA on sleep, study hours, stress, and supervisor support explained 25% of the variance in GPA, $F(4, 582) = 49.14$, $p < .001$, $N = 587$. Holding the other predictors constant, each additional hour of sleep was associated with a GPA 0.11 points higher (95% CI [0.08, 0.13]). The hypothesis that more sleep goes with a higher GPA (Chapter 5) was therefore supported: the null hypothesis of a zero sleep coefficient was rejected.

8.4 - Logistic regression

The third research question concerns a yes-or-no outcome: whether a student has considered dropping out. Linear regression is the wrong tool here, because it could

predict probabilities below 0 or above 1, which make no sense. **Logistic regression** models the probability of a “yes” in a way that always stays between 0 and 1.

8.4.1 - Odds and odds ratios

Logistic regression works with **odds** rather than probabilities. The odds of an event are the probability that it happens divided by the probability that it does not. The table from Chapter 7 shows how they are calculated:

```
jobs <- table(students$employment,
students$considering_dropout)
jobs
```

	No	Yes
Full-time job	72	25
None	265	42
Part-time job	173	23

```
odds_full <- jobs["Full-time job", "Yes"] /
jobs["Full-time job", "No"]
odds_none <- jobs["None", "Yes"] /
jobs["None", "No"]
c(full_time_job = odds_full, no_job =
odds_none, odds_ratio = odds_full / odds_none)
```

full_time_job	no_job	odds_ratio
0.3472222	0.1584906	2.1908069

Among students with a full-time job, 25 have considered dropping out and 72 have not, which gives odds of about 0.35. Among students without a job, the odds are 42 to 265, about 0.16. The **odds ratio** compares the two: students with a full-time job have about 2.2 times the odds of considering dropout. An odds ratio of 1 means no difference, above 1 means higher odds, and below 1 means lower odds.

8.4.2 - A model for dropout

The function `glm()`, for *generalised linear model*, fits logistic regression when told `family = binomial`. The outcome must be coded 0 and 1:

```
study <- study |> mutate(dropout =
as.integer(considering_dropout == "Yes"))

dropout_model <- glm(
  dropout ~ stress + support + financial_worry
+ employment + study_mode,
  data = study, family = binomial
)
summary(dropout_model)$coefficients
```

	z value	Pr(> z)	Estimate	Std. Error
(Intercept)			-4.0185658	1.2316928
	-3.262636	1.103811e-03		
stress			1.2822937	0.2398794

```

5.345576 9.012973e-08
support -1.1102419 0.2039547
-5.443572 5.222263e-08
financial_worry 0.4537396 0.1216263
3.730605 1.910209e-04
employmentNone -0.5892840 0.4378768
-1.345776 1.783748e-01
employmentPart-time job -0.5387011 0.4522237
-1.191227 2.335644e-01
study_modePart-time 0.2763818 0.3674127
0.752238 4.519079e-01

```

The coefficients are on the scale of **log-odds**, which is hard to interpret directly. The function `exp()` turns them into odds ratios, and `confint.default()` gives their confidence intervals:

```

exp(cbind(odds_ratio = coef(dropout_model),
  confint.default(dropout_model))) |>
  round(2)

```

	odds_ratio	2.5 %	97.5 %
(Intercept)	0.02	0.00	0.20
stress	3.60	2.25	5.77
support	0.33	0.22	0.49
financial_worry	1.57	1.24	2.00
employmentNone	0.55	0.24	1.31
employmentPart-time job	0.58	0.24	1.42
study_modePart-time	1.32	0.64	2.71

Holding the other predictors constant, each extra point of **stress** multiplies the odds of considering dropout by about 3.6, and each extra point of **financial worry** by about 1.6. Each extra point of **supervisor support** multiplies them by about 0.33, which cuts them by about 67%. The hypothesis stated in Chapter 5, that higher stress raises the odds of considering dropout, is supported: the confidence interval for stress lies entirely above 1.

Something interesting has happened to **employment**. In Chapter 7, and in the odds ratio above, students with full-time jobs had about twice the odds of considering dropping out. In this model, the confidence intervals for both employment categories include 1: once stress and financial worry are taken into account, having a job adds little on its own. The likely explanation is that a full-time job matters *because* of the stress and money pressure that come with it. This is the kind of insight multiple regression makes possible.

8.4.3 - Predicted probabilities

Odds ratios are hard for many readers; predicted probabilities are easier. The model can predict the probability of considering dropout for two example students who are identical except for their stress and support:

```

examples <- tibble(
  stress      = c(2.5, 4.0),
  support     = c(4.0, 2.0),
  financial_worry = 3,
  employment  = "None",

```

```
study_mode      = "Full-time"
)
predict(dropout_model, newdata = examples,
type = "response") |> round(2)
```

```
1 2
0.01 0.42
```

The argument `type = "response"` asks for probabilities rather than log-odds. A student with low stress and good support has about a 1% chance of considering dropout; a stressed student with little support, about 42%. For a thesis's recommendations, that contrast is more persuasive than any coefficient.

Whether the model can *predict* who will consider dropping out, and how that should be tested fairly, is a different question, about prediction rather than explanation. It is the subject of Chapters 11 and 12.

In your field: forestry and agriculture

R's `trees` dataset records the girth, height, and timber volume of 31 felled black cherry trees. Foresters need to estimate a standing tree's volume from measurements they can take without cutting it down, a classic job for multiple regression:

```
tree_model <- lm(Volume ~ Girth + Height,
data = trees)
summary(tree_model)$coefficients
```

	Estimate	Std. Error	t
value			
(Intercept)	-57.9876589	8.6382259	
Girth	4.7081605	0.2642646	17.816084
Height	0.3392512	0.1301512	2.606594

```
summary(tree_model)$r.squared
```

```
[1] 0.94795
```

Girth and height together explain about 95% of the variation in volume. Physical measurements are often far more predictable than human behaviour: compare this with the 25% of the GPA model.

8.5 - Common misconceptions

Regression output is easy to produce and easy to misread. Four misreadings are especially common.

- **“A significant coefficient shows a cause.”** Outside a randomised experiment, a coefficient is an association adjusted for the variables in the model, and only for those.
- **“A low R-squared means the model is useless.”** A predictor can have a real and important effect while

explaining little of the total variation, as sleep does for GPA.

- **“Each coefficient describes the predictor on its own.”** In multiple regression, each coefficient is the effect *holding the others constant*, and it can change when predictors are added or removed, as caffeine’s did.
- **“A significant ANOVA shows that all groups differ.”** It shows only that at least one group differs from the others; post-hoc tests say which.

8.6 - Chapter review

8.6.1 - Summary

- ANOVA partitions variation into the part between groups and the part within them. The F statistic is their ratio, so the same group averages are convincing when groups are tight and unremarkable when they are spread out. η^2 measures the size of the effect.
- One-way ANOVA tests whether three or more group averages differ with one test instead of many. Post-hoc tests, such as Tukey’s, find which groups differ while controlling false alarms. Levene’s test checks equal spread; Kruskal-Wallis is the non-parametric alternative.
- Two-way ANOVA tests two main effects and their interaction. An interaction means the effect of one factor depends on the other; parallel lines in an interaction plot mean no interaction.
- Regression models the average outcome at each value of the predictors. A slope is the change in the average outcome for a one-unit change in the predictor; in multiple regression, holding the other predictors constant. R-squared is the share of variation explained.
- Categorical predictors are compared with a reference category. A confounder is a common cause of predictor and outcome; adding it to the model can remove a misleading association, but only confounders that were measured.
- Squared terms model curves; interaction terms let one predictor’s effect depend on another. Residual plots reveal missed curves (a U shape) and unequal spread (a funnel).
- Logistic regression models the probability of a yes-or-no outcome. `exp()` of its coefficients gives odds ratios; `predict(..., type = "response")` gives probabilities.

8.6.2 - Key terms

ANOVA, F statistic, between-group variation, within-group variation, eta squared, post-hoc test, Tukey’s test, Levene’s test, Kruskal-Wallis test, two-way ANOVA, main effect, interaction, interaction plot, linear regression, predictor, outcome, intercept, slope, residual, least squares, R-squared, adjusted R-squared, multiple regression, reference category, confounder, quadratic term, diminishing returns, centring, diagnostic plot, logistic regression, odds, odds ratio, log-odds, predicted probability.

8.7 - Exercises

The [playground](#) has these and more, with hints and solutions.

1. Test whether first-semester wellbeing differs between faculties with a one-way ANOVA, calculate η^2 , and, if the ANOVA is significant, run Tukey's test.
2. Fit a simple regression of first-semester wellbeing on sleep hours, and interpret the slope and R-squared.
3. Add stress and support to the model from Exercise 2, and describe what happens to the slope of sleep and to R-squared.
4. Add `study_mode` to the model from Exercise 3. Name the reference category and explain what its coefficient means.
5. Fit a logistic regression of considering dropout on `burnout` score alone (calculate it from the questionnaire first), report its odds ratio, and explain what it means.
6. Draw a diagram, like Figure 8.5, of a confounder that might explain an association in your own field, and describe how a regression could separate the two explanations.

8.8 - Further reading

- *Introduction to Modern Statistics* [13]: chapters “Linear regression with a single predictor”, “Linear regression with multiple predictors”, “Logistic regression”, and “Inference for comparing many means”.
- *Regression and Other Stories* [17] is an excellent next step, with a strong focus on interpreting and checking regression models in real research.

9. Multivariate Statistical Methods

Many of the things researchers most want to study cannot be measured directly. Stress, burnout, satisfaction, intelligence, motivation, and quality of life are **constructs** (Chapter 5): ideas that exist only through their effects on what people say and do. A questionnaire approaches such a construct indirectly, through several items that are each expected to reflect it, and the researcher then has to show that the items really do measure what they are meant to. At the same time, a dataset with many variables often contains patterns that no single variable reveals, such as groups of people with a similar profile across all of them.

Multivariate methods analyse many variables at once, and this chapter introduces three of them. **Principal component analysis** summarises many variables with a few. **Factor analysis** checks whether questionnaire items measure the underlying traits they are meant to, and **Cronbach's alpha** checks whether they do so consistently. **Cluster analysis** sorts people into groups with similar profiles. In the study, they answer two research questions: whether the 22 questionnaire items measure stress, burnout, supervisor support, and satisfaction as intended (RQ6), a question every examiner of a questionnaire study will ask, and whether there are distinct profiles of students (RQ7).

💡 By the end of this chapter you will be able to

- Explain why a construct is measured with several items, and what a latent variable is.
- Explore the correlations among many variables at once.
- Run a principal component analysis, and decide how many components to keep.
- Explain the difference between principal component analysis and factor analysis.
- Run an exploratory factor analysis, and interpret factor loadings, cross-loadings, and reversed items.
- Check the reliability of a questionnaire scale with Cronbach's alpha.
- Group cases with k-means and hierarchical clustering, choose the number of clusters, and describe the clusters as summaries rather than natural groups.

Note

Teaching & Practice: [Interactive Lecture Slides](#) · [Browser Playground Exercises](#)

9.1 - Constructs and latent variables

A single questionnaire item is a poor measure of a construct. The answer to “I feel unable to control important things in my studies” depends on the student’s stress, but also on how they read the question that day, on their mood, and on how they use the answer scale. Each item therefore carries two things: a signal from the construct, and noise of its own. A **latent variable** is the construct thought to lie behind the items: it is not observed, but it is assumed to cause part of every answer.

Averaging several items keeps the signal, which all items share, and lets the noise, which differs from item to item, partly cancel out. A simulation shows how much this helps. It creates a “true” stress level for 600 imaginary students, then six items, each equal to the true level plus its own random noise, and compares how closely one item and the average of several items follow the true level:

```
set.seed(42)
true_stress <- rnorm(600)
simulated_items <- sapply(1:6, function(i)
  true_stress + rnorm(600, sd = 1))

sapply(1:6, function(k)
  cor(rowMeans(simulated_items[, 1:k, drop =
    FALSE]), true_stress)) |>
  setNames(paste(1:6, "items")) |>
  round(2)
```

```
1 items 2 items 3 items 4 items 5 items 6 items
0.68    0.81    0.86    0.89    0.91    0.92
```

In this simulation, one item correlates only about 0.68 with the true stress level, while the average of six items correlates about 0.92. This is why questionnaires use several items for each construct, and why the items are combined into a scale score (Chapter 3).

The approach rests on an assumption that must be checked: that the items of a scale really reflect one common construct, and not several. If some stress items actually measured tiredness, averaging them would mix two things. Chapter 5 called this **construct validity**. The methods in this chapter provide evidence for it: factor analysis shows whether the items group as intended, and Cronbach’s alpha shows whether the items of each group agree with each other, which Chapter 5 called **internal consistency**.

9.2 - Many variables at once

The 22 questionnaire items are in `questionnaire`:

```
library(dplyr)
```

```
Attaching package: 'dplyr'
```

```
The following objects are masked from
'package:stats':
```

```
filter, lag
```

```
The following objects are masked from
'package:base':
```

```
intersect, setdiff, setequal, union
```

```
library(ggplot2)
```

```
items <- questionnaire |> select(-student_id)
ncol(items)
```

```
[1] 22
```

With 22 items, there are 231 correlations between pairs of items, far too many to read one by one. A picture helps. The **corrplot** package draws a correlation matrix as a grid of coloured squares, and `order = "hclust"` sorts the items so that those that correlate strongly sit together:

```
library(corrplot)
```

```
corrplot 0.95 loaded
```

```
corrplot(cor(items,                                use
= "pairwise.complete.obs"),
          method = "color", order = "hclust",
          tl.col = "black", tl.cex = 0.7)
```




Figure 9.1: Correlations among the 22 questionnaire items, sorted so that similar items are next to each other.

Blocks of related items stand out along the diagonal, one for each scale, just as the idea of latent variables predicts:

items that share a construct correlate with each other. Blue squares show positive correlations and red negative. The stress and burnout blocks also correlate with each other, and one item, `stress_4`, is red against the other stress items, because it is worded the other way round. The methods in this chapter turn this picture into numbers.

9.3 - Principal component analysis

Principal component analysis (PCA) replaces many correlated variables with a few new ones, called **principal components**, that together keep most of the information. The first component is the combination of the variables that captures as much of their variation as possible. The second captures as much as possible of what is left, and so on.

Two items that correlate strongly, such as “I feel emotionally drained by my studies” and “I feel exhausted when I think about my thesis”, illustrate the idea. Most students who score high on one score high on the other, so one combined score, their average for example, keeps most of what the two items tell. PCA does this for all the variables at once, and finds the best combinations automatically.

9.3.1 - Running PCA

The function `prcomp()` runs PCA, and `scale. = TRUE` puts every variable on the same scale first, which is almost always wanted. PCA needs complete data, so `na.omit()` keeps only the students who answered every item:

```
items_complete <- na.omit(items)
nrow(items_complete)
```

```
[1] 396
```

```
pca <- prcomp(items_complete, scale. = TRUE)
summary(pca)$importance[, 1:6] |> round(3)
```

		PC1	PC2	PC3	PC4
PC5	PC6				
Standard deviation		2.580	1.694	1.250	1.110
		0.899	0.852		
Proportion of Variance		0.302	0.130	0.071	0.056
		0.037	0.033		
Cumulative Proportion		0.302	0.433	0.504	0.560
		0.597	0.630		

Only 396 of the 600 students answered all 22 items: a small share of missing answers per item adds up to many incomplete students. Factor analysis, below, can use the incomplete students too.

The table shows how much of the total variation each component captures. The first component alone captures 30%, and the first four together 56%.

9.3.2 - The number of components

No single rule decides how many components to keep, so researchers use several. The **Kaiser rule** keeps components whose **eigenvalue**, the amount of variation they capture, is greater than 1, that is, more than one original variable's worth; here, 4 components pass. The **scree plot** shows the eigenvalues in order, and the researcher looks for the "elbow" where the line flattens out. **Parallel analysis** compares the eigenvalues with those from random data of the same size, and keeps the components that beat random data; it is the most reliable of the three.

The **factoextra** package draws a scree plot directly:

```
library(factoextra)
```

```
Welcome! Want to learn more? See two
factoextra-related books at https://goo.gl/ve
3WBa
```

```
fviz_eig(pca, addlabels = TRUE, ncp = 10)
```

```
Warning in geom_bar(stat = "identity", fill =
barfill, color = barcolor, :
Ignoring empty aesthetic: `width`.
```




Figure 9.2: Scree plot: the share of variation captured by each principal component.

The line drops steeply for four components and flattens after that. Parallel analysis, from the **psych** package, agrees:

```
library(psych)
```

```
Attaching package: 'psych'
```

```
The following objects are masked from
'package:ggplot2':
```

```
%+%, alpha
```

```
set.seed(1)
parallel <- fa.parallel(items, fa = "fa", plot
= FALSE)
```

```
Parallel analysis suggests that the number of
factors = 4 and the number of components
= NA
```

```
parallel$nfact
```

```
[1] 4
```

All three methods point to four dimensions, matching the four scales the questionnaire was designed to measure.

9.3.3 - PCA and factor analysis compared

PCA and factor analysis are often confused, and many theses use one when they mean the other. The difference lies in what they assume. **PCA** simply summarises the variables: the components are combinations of the items, with no claim about why the items are related, and it is used to reduce many variables to a few, for example before a further analysis. **Factor analysis** assumes the model of the first section: that each item is *caused* by one or more latent variables, called **factors**, plus error of its own. A student answers the stress items the way they do *because* they are stressed. Factor analysis is therefore the method for checking whether a questionnaire measures the constructs it is meant to, and that is the study's question.

9.4 - Factor analysis

Exploratory factor analysis (EFA) estimates how strongly each item is related to each factor. These relationships are called **loadings**. Loadings run roughly from -1 to 1: close to 0 means that the item is unrelated to the factor, and a value above about 0.4 in size is usually considered a clear relationship.

The function `fa()` from the `psych` package runs the analysis. Its main arguments are the number of factors (four, from parallel analysis), the **rotation**, and the estimation method. Rotation turns the factors to make them easier to interpret; "oblimin" allows the factors to correlate with each other, which is realistic, since stress and burnout surely go together. The option `fm = "ml"` uses maximum likelihood, a common estimation method:

```
efa <- fa(items, nfactors = 4, rotate =
"oblimin", fm = "ml")
```

```
Loading required namespace: GPArotation
```

```
print(efa$loadings, cutoff = 0.3, sort = TRUE)
```



```

Loadings:
      ML1    ML3    ML2    ML4
support_1  0.749
support_2  0.718
support_3  0.669
support_4  0.638
support_5  0.726
support_6  0.619
stress_1           0.623
stress_3           0.627
stress_4        -0.503
stress_5           0.827
stress_6           0.547
burnout_1           0.647
burnout_2           0.719
burnout_4           0.622
burnout_5           0.544
burnout_6           0.675
satisfaction_1           0.659
satisfaction_2           0.684
satisfaction_3           0.585
satisfaction_4           0.673
stress_2           0.496
burnout_3           0.325  0.446

      ML1    ML3    ML2    ML4
SS loadings  2.87 2.41 2.384 1.735
Proportion Var 0.13 0.11 0.108 0.079
Cumulative Var 0.13 0.24 0.348 0.427

```

In the printout, `cutoff = 0.3` hides small loadings so that the pattern is easy to see, and `sort = TRUE` groups the items by the factor they load on most. Unlike PCA, `fa()` uses every student, including those who skipped an item.

The table shows four clear factors. Each collects the items of one scale: the six support items on one, the stress items on another, the burnout items on a third, and the four satisfaction items on the fourth. The names ML1 to ML4 are only labels; naming the factors is the researcher's job. The reversed item, `stress_4` ("I feel confident handling problems in my studies"), loads *negatively* on the stress factor, which is exactly right for a reversed item: agreeing with it means less stress, and it is why the item must be reversed before the stress score is calculated (Chapter 3). One item, `burnout_3` ("Deadlines make me feel overwhelmed"), loads on both the burnout and the stress factor. Reading the item, that makes sense: feeling overwhelmed by deadlines is both. Such a **cross-loading** item is worth discussing in a thesis, and sometimes worth rewording or removing in future studies.

Because the rotation allowed the factors to correlate, `fa()` also estimates how strongly they do:

```
round(efa$Phi, 2)
```

```

      ML1    ML3    ML2    ML4
ML1  1.00 -0.37 -0.23  0.48
ML3 -0.37  1.00  0.66 -0.44
ML2 -0.23  0.66  1.00 -0.39
ML4  0.48 -0.44 -0.39  1.00

```

The stress and burnout factors correlate strongly (about 0.66), but not so strongly that they are the same thing. The support and satisfaction factors correlate positively with each other and negatively with stress and burnout. These relationships make sense, which is itself evidence that the scales measure what they should.

9.4.1 - Reliability and Cronbach's alpha

Factor analysis shows that the items of each scale belong together. **Reliability** asks a related question: whether the items of a scale give consistent results (Chapter 5). The most common measure of this internal consistency is **Cronbach's alpha**, which ranges from 0 to 1. It rises when the items correlate strongly with each other, and also when there are more items, for the reason shown by the simulation at the start of the chapter. As a rough guide, 0.7 or above is acceptable and 0.8 or above good for research use. The function `alpha()` from the `psych` package calculates it, after reversed items have been reversed:

```
q <- questionnaire
q$stress_4 <- 6 - q$stress_4

psych::alpha(q[,
  paste0("stress_", 1:6)])$total$raw_alpha
```

```
[1] 0.826084
```

```
psych::alpha(q[,      paste0("burnout_",
  1:6)])$total$raw_alpha
```

```
[1] 0.8313142
```

```
psych::alpha(q[,      paste0("support_",
  1:6)])$total$raw_alpha
```

```
[1] 0.8454391
```

```
psych::alpha(q[,      paste0("satisfaction_",
  1:4)])$total$raw_alpha
```

```
[1] 0.7642397
```

All four scales are reliable, with alphas between 0.76 and 0.85. The full output of `alpha()` also shows what alpha would be if each item were dropped, which helps to spot a weak item.

The code writes `psych::alpha()` rather than `alpha()` because `ggplot2` also has a function called `alpha()`, for making colours transparent. Whichever package is loaded last wins, so after loading `ggplot2` (or a package that loads it, such as `factoextra`) a plain `alpha()` no longer calculates Cronbach's alpha. The `package::function()` form always calls the intended function, whether or not the package is loaded.

A high alpha shows consistency, not validity. A scale can be highly reliable and still measure the wrong thing, as the target picture in Chapter 5 showed; the factor

structure and the relationships with other constructs are the evidence for validity.

💡 Reporting the questionnaire in a thesis

A typical report: “An exploratory factor analysis (maximum likelihood, oblimin rotation) supported four factors, as indicated by parallel analysis. All items loaded on their intended factor (loadings 0.45 to 0.83 in absolute value), with one item (*burnout_3*) also loading on the stress factor. Internal consistency was acceptable to good (Cronbach’s $\alpha = 0.76$ to 0.85).”

A further step, **confirmatory factor analysis**, tests whether a structure decided in advance fits the data, rather than exploring it. It is done with the **lavaan** package and is beyond this book.

9.5 - Cluster analysis

Factor analysis groups *variables*. **Cluster analysis** groups *people*: it looks for groups of students whose profiles are similar to each other and different from other groups. It is a descriptive method. It does not test a hypothesis, and the groups it finds are summaries of the data, useful for describing and thinking about students, not discoveries of natural kinds of student. The study’s question about student profiles (RQ7) is accordingly exploratory (Chapter 5).

Each student is described by their average sleep, study hours, caffeine, and exercise across the semesters, and their stress, support, and satisfaction scores:

```
scores <- q |>
  mutate(
    stress =
      rowMeans(pick(stress_1:stress_6), na.rm =
        TRUE),
    support =
      rowMeans(pick(support_1:support_6), na.rm =
        TRUE),
    satisfaction = rowMeans(pick(satisfaction_1:satisfaction_4),
      na.rm = TRUE)
  ) |>
  select(student_id, stress, support,
    satisfaction)

profiles <- semesters |>
  summarise(across(c(sleep_hours,
    study_hours, caffeine_mg, exercise_days),
    ~ mean(.x, na.rm = TRUE)),
    .by = student_id) |>
  left_join(scores, join_by(student_id)) |>
  na.omit()

profile_data <- scale(profiles |> select(-
  student_id))
```

Clustering works with **distances** between students: how different two students’ profiles are. Caffeine is measured in hundreds of milligrams and sleep in single hours, so without adjustment caffeine would dominate every

distance. The function `scale()` turns every variable into z-scores (Chapter 6), so each counts equally.

9.5.1 - k-means clustering

k-means is the most widely used clustering method. The researcher chooses the number of clusters, k , and the algorithm places k starting points at random, assigns each student to the nearest point, moves each point to the centre (the mean) of its students, and repeats the last two steps until nothing changes. Because the result depends on the random start, `nstart = 25` runs the algorithm from 25 different starts and keeps the best result.

9.5.2 - The number of clusters

As with components, no single rule decides the number of clusters, and two guides are common. The **elbow method** uses the fact that the total distance of students from their cluster centres always falls as k grows, and looks for the point where adding clusters stops helping much. The **silhouette** measures, for each student, how much closer they are to their own cluster than to the next nearest one, from -1 to 1 ; a higher average silhouette means clearer clusters.

```
set.seed(123)
fviz_nbclust(profile_data, kmeans, method =
"wss", nstart = 25)
```

```
Warning: The `size` argument of
`element_line()` is deprecated as of ggplot2
3.4.0.
i Please use the `linewidth` argument instead.
i The deprecated feature was likely used in
the ggpubr package.
Please report the issue at <https://github.
com/kassambara/ggpubr/issues>.
```

```
Warning: The `size` argument of
`element_rect()` is deprecated as of ggplot2
3.4.0.
i Please use the `linewidth` argument instead.
i The deprecated feature was likely used in
the ggpubr package.
Please report the issue at <https://github.
com/kassambara/ggpubr/issues>.
```

```
Warning: Using `size` aesthetic for lines was
deprecated in ggplot2 3.4.0.
i Please use `linewidth` instead.
i The deprecated feature was likely used in
the ggpubr package.
Please report the issue at <https://github.
com/kassambara/ggpubr/issues>.
```

```
fviz_nbclust(profile_data, kmeans, method =
"silhouette", nstart = 25)
```




Figure 9.3: Choosing the number of clusters.

The honest conclusion is that the data does not point clearly to one number. The elbow is gentle, and the silhouette is highest for 2 clusters but only a little lower for 3 or 4. This is common with real data: groups of people overlap, and there are no sharp boundaries. The choice then rests on what is useful and interpretable, and on prior theory. The thesis, based on earlier research, expects around four student profiles, so $k = 4$ is tried:

```
set.seed(123)
clusters <- kmeans(profile_data, centers = 4,
  nstart = 25)
table(clusters$cluster)
```

```
 1  2  3  4
85 197 200 117
```

9.5.3 - Describing the clusters

A cluster is only useful if its meaning can be stated. Adding the cluster to the original, unscaled data and averaging each variable by cluster shows each group's profile:

```
profiles |>
  mutate(cluster = clusters$cluster) |>
  summarise(students = n(),
    across(sleep_hours:satisfaction, ~
      round(mean(.x), 1)),
    .by = cluster) |>
  arrange(cluster)
```

	cluster	students	sleep_hours	study_hours	caffeine_mg	exercise_days	stress
1	1	85	5.1	41.6	423.7	1.2	3.6
2	2	197	6.7	18.0	149.8	2.2	3.4
3	3	200	7.2	23.6	122.9	3.5	2.6
4	4	117	5.8	40.8	188.1	1.5	3.6
					support	satisfaction	
1			2.9	3.0			
2			2.7	2.5			
3			3.6	3.8			
4			3.5	3.3			

The averages describe the groups. One is clearly **balanced**: the most sleep, moderate study, the most exercise,

low stress, and good support. Two are **overloaded**: long study weeks, little sleep and exercise, high stress, and one of them (the smaller) with very high caffeine intake. The fourth combines low study hours with low support and low satisfaction: students who seem **disengaged or isolated**. Cluster numbers are arbitrary labels, and can change if the code is run with a different seed.

Figure 9.4 shows the clusters on the first two principal components of the profile data, a common way to see many variables in two dimensions:

```
fviz_cluster(clusters, data = profile_data,
  geom = "point", ellipse.type = "convex",
  ggtheme = theme_minimal())
```



Figure 9.4: The four k-means clusters, shown on the first two principal components of the profile data.

The clusters overlap. That is not a failure: real students do not fall into neat boxes. The profiles are useful summaries, not natural categories.

9.5.4 - Hierarchical clustering

Hierarchical clustering takes a different approach. It starts with every student as their own cluster and repeatedly merges the two most similar clusters, until everyone is in one. The result is a tree, called a **dendrogram**, which can be cut at any height to give any number of clusters, without choosing k in advance. Ward's method ("ward.D2") merges clusters so that they stay as compact as possible:

```
tree <- hclust(dist(profile_data), method =
  "ward.D2")
fviz_dend(tree, k = 4, show_labels = FALSE,
  rect = TRUE)
```

```
Warning: `aes_string()` was deprecated in
ggplot2 3.0.0.
i Please use tidy evaluation idioms with
`aes()`.
i See also `vignette("ggplot2-in-packages")`
for more information.
```

i The deprecated feature was likely used in the factoextra package.
Please report the issue at <<https://github.com/kassambara/factoextra/issues>>.

Warning: The `<scale>` argument of `guides()` cannot be `FALSE`. Use "none" instead as of ggplot2 3.3.4.
i The deprecated feature was likely used in the factoextra package.
Please report the issue at <<https://github.com/kassambara/factoextra/issues>>.




Figure 9.5: Dendrogram from hierarchical clustering (Ward's method), cut into four clusters.

The function `dist()` calculates the distances between all pairs of students, and `cutree()` cuts the tree into groups. A cross-table shows how far the two methods agree:

```
hier_clusters <- cutree(tree, k = 4)
table(hierarchical = hier_clusters, kmeans =
clusters$cluster)
```

	kmeans			
hierarchical	1	2	3	4
1	6	152	4	3
2	0	43	186	4
3	7	1	10	75
4	72	1	0	35

They agree on the broad picture, most clearly on the balanced students, but many students switch groups between the two methods. When two reasonable methods disagree about a student, that student is near a boundary. The clearer groups, found by both methods, are the ones to trust most.

9.5.5 - Clusters in random data

A clustering method will split any data into groups, even data that contains no groups at all. The final code makes the point with 300 points drawn entirely at random, with no structure of any kind, and asks k-means for four clusters:

```

set.seed(7)
random_points <- tibble(x = rnorm(300), y =
  rnorm(300))
random_points$cluster
<- factor(kmeans(random_points, centers = 4,
  nstart = 25)$cluster)

ggplot(random_points, aes(x, y, colour =
  cluster)) +
  geom_point() +
  scale_colour_viridis_d(end = 0.9) +
  coord_equal() +
  theme_minimal(base_size = 12)

```




Figure 9.6: k-means finds four ‘clusters’ in 300 random points that contain no groups at all.

The method obliges, and divides one round cloud into four neat regions. Nothing in its output says that the groups are artificial. The existence of clusters therefore proves nothing on its own. Before clusters are reported as meaningful, it should be checked that they make sense, that different methods broadly agree, and that the groups differ in ways that matter for the research. Chapter 14 introduces methods that allow for overlapping groups and for students who fit no group at all.

In your field: criminology and social policy

R's `USArrests` dataset records rates of assault, murder, and rape per 100,000 residents, and the percentage of people living in urban areas, for the 50 US states in 1973. PCA summarises the four variables:

```
arrests_pca <- prcomp(USArrests, scale. =
TRUE)
summary(arrests_pca)$importance[, 1:2] |
> round(2)
```

	PC1	PC2
Standard deviation	1.57	0.99
Proportion of Variance	0.62	0.25
Cumulative Proportion	0.62	0.87

The first component captures about 62% of the variation: it is essentially an overall violent crime rate. The second, about 25%, mostly reflects how urban a state is. Two numbers now describe what took four.

9.6 - Common misconceptions

Multivariate methods produce impressive output, which makes their limits easy to forget.

- **“PCA and factor analysis are the same.”** PCA summarises variables; factor analysis models latent variables that cause them. Only factor analysis addresses whether a questionnaire measures its constructs.
- **“A high Cronbach’s alpha proves a scale is valid.”** Alpha measures consistency. A scale can be consistent and still measure the wrong thing.
- **“Clusters are natural groups.”** Clustering finds groups in any data, including random data. Clusters are summaries, and their usefulness must be argued.
- **“The number of factors or clusters is decided by the software.”** The criteria often disagree, and the final choice is a judgement that should be justified in the thesis.

9.7 - Chapter review

9.7.1 - Summary

- A construct is measured with several items because each item carries noise of its own; averaging keeps the shared signal. A latent variable is the unobserved construct assumed to cause the items.
- A sorted correlation plot shows the structure among many variables.
- PCA replaces many correlated variables with a few components that keep most of the variation. Decide how many to keep with eigenvalues above 1, the scree plot, and, most reliably, parallel analysis.
- PCA summarises; factor analysis models latent variables that cause the items. Use factor analysis to check a questionnaire.

- In an EFA, loadings show how strongly items relate to factors. Check that items load on their intended factor, that reversed items load negatively, and discuss cross-loadings. Oblique rotation lets factors correlate.
- Cronbach's alpha measures a scale's internal consistency: 0.7 or above is acceptable, 0.8 or above good. Reverse reversed items first. Reliability is necessary for validity but not enough.
- Cluster analysis groups people with similar profiles. Scale the variables first. k-means needs the number of clusters in advance; hierarchical clustering produces a tree that can be cut anywhere. Choose the number with the elbow and silhouette methods, theory, and interpretability.
- Clustering finds groups even in random data. Clusters are useful summaries, not proof of natural groups; describe each cluster by its averages, and check that methods agree.

9.7.2 - Key terms

Multivariate analysis, construct, latent variable, correlation matrix, principal component analysis, principal component, eigenvalue, Kaiser rule, scree plot, parallel analysis, factor analysis, factor, loading, rotation, oblique rotation, cross-loading, reversed item, reliability, internal consistency, Cronbach's alpha, confirmatory factor analysis, cluster analysis, distance, scaling, k-means, elbow method, silhouette, hierarchical clustering, dendrogram, Ward's method.

9.8 - Exercises

The [playground](#) has these and more, with hints and solutions.

1. Run a PCA on the six support items only. Report how much variation the first component captures, and what that suggests about the scale.
2. Run an EFA with three factors instead of four. Identify the scales that end up sharing a factor, and suggest why.
3. Calculate Cronbach's alpha for the stress scale *without* reversing `stress_4` first, and explain what happens.
4. Run k-means with 3 clusters on the profile data, describe the clusters, and identify which profiles from the 4-cluster solution were merged.
5. Explain, in two or three sentences for a thesis methods section, why you chose the number of clusters you did.
6. Change the simulation at the start of the chapter so that each item has more noise (`sd = 2`). Find how many items are now needed for the average to correlate at least 0.8 with the true stress level.

9.9 - Further reading

- *An Introduction to Statistical Learning* [18], chapter "Unsupervised Learning", explains PCA and clustering clearly, with R examples.
- The [psych package's online guides](#) by William Revelle cover factor analysis and reliability in depth.

10. Mixed-Effects Models

Every test in the previous chapters assumed that the observations were independent: that knowing one observation tells nothing about another. Much research data breaks this rule by design. When the same people are measured several times, a person who scores high on one occasion is likely to score high on the next. When people belong to groups, such as students with the same supervisor, pupils in the same class, or patients in the same hospital, members of a group tend to resemble each other. Such data contains less independent information than its number of rows suggests, and methods that ignore this get the uncertainty wrong, sometimes badly.

Mixed-effects models, also called **multilevel models**, are designed for exactly this kind of data. They separate what is common to a group from what varies within it, and so can study change within people and differences between groups at the same time. The wellbeing study was built this way: its students were followed for two years, and they share supervisors. The chapter answers the question of how wellbeing and GPA change over the two years and how much supervisors matter (RQ8), and returns to the workshop to ask whether its effect faded.

💡 By the end of this chapter you will be able to

- Recognise repeated-measures and nested data, and explain, with a simulation, why ordinary regression gets their uncertainty wrong.
- Fit mixed-effects models with random intercepts and random slopes using lme4.
- Interpret fixed effects, random effects, and the intraclass correlation.
- Test whether an effect changes over time, and compare models.
- Model a count outcome with a generalised linear mixed model.
- Report a mixed-effects model in a thesis.

Note

Teaching & Practice: [Interactive Lecture Slides](#) · [Browser Playground Exercises](#)

10.1 - Data with a structure

Two kinds of structure are common in research. In **repeated measures** data, the same people are measured several times, as the students in the wellbeing study are over four semesters, so measurements are grouped within

people. In **nested data**, people belong to groups, such as students within supervisors, pupils within classes within schools, or patients within hospitals. In both cases, observations within a group are more alike than observations from different groups.

10.1.1 - Why independence matters

The consequence of ignoring such structure is easiest to see in a simulation. Imagine a study in which 20 supervisors each have 10 students, and half of the supervisors, chosen at random, attend a training course. Suppose the course has no effect at all. Students of the same supervisor are somewhat alike, because each supervisor has a style of their own. The code below creates such a study, analyses it in two ways, and repeats the whole study 500 times. The first analysis is an ordinary regression that treats the 200 students as independent; the second is a mixed-effects model that knows which students share a supervisor:

```
library(lme4)
```

```
Loading required package: Matrix
```

```
library(dplyr)
```

```
Attaching package: 'dplyr'
```

```
The following objects are masked from
'package:stats':
```

```
filter, lag
```

```
The following objects are masked from
'package:base':
```

```
intersect, setdiff, setequal, union
```

```
library(ggplot2)
```

```
simulate_clustered <- function() {
  supervisor <- rep(1:20, each = 10)
  trained    <- rep(rep(c(0, 1), 10), each =
10)    # half the supervisors, no real effect
  style      <- rnorm(20, sd = 5)[supervisor]
  # what students of one supervisor share
  wellbeing  <- 60 + style + rnorm(200, sd
= 10)

  p_ordinary <- summary(lm(wellbeing ~
trained))$coefficients["trained", 4]
  mixed      <- suppressMessages(lmer(wellbeing
~ trained + (1 | supervisor)))
  t_mixed    <- coef(summary(mixed))["trained",
"t value"]
  c(ordinary = p_ordinary < 0.05, mixed = 2 *
pnorm(-abs(t_mixed)) < 0.05)
}
```

```
set.seed(12)
false_alarms <- replicate(500,
  simulate_clustered())
rowMeans(false_alarms)
```

```
ordinary    mixed
0.268      0.074
```

The course has no effect, so a correct analysis should declare it significant in about 5% of the simulated studies. The ordinary regression does so in 27% of them. The mixed model, which allows for the supervisors, comes close to the intended rate, at 7%. It is slightly above 5% because the quick p-value used in the simulation ignores that there are only 20 supervisors; the `lmerTest` package, introduced later in the chapter, corrects for this.

The reason is that the 200 students are not 200 independent pieces of evidence about the course. The course was given to supervisors, and there are only 20 of them; students of the same supervisor largely repeat each other's information. The ordinary regression counts every student as new evidence, underestimates the standard error, and finds "effects" that are really differences between a few supervisors. The mixed model builds the grouping into the model and gets the uncertainty right. In other situations, as later sections show, ignoring the structure can also hide a real effect. Either way, ordinary regression gets the uncertainty wrong for grouped data.

10.2 - A study of sleep deprivation

A small real study shows how mixed models work. R's `sleepstudy` data, from the `lme4` package, records the reaction times of 18 people over 10 days of sleep restriction, one measurement per person per day [19]:

```
head(sleepstudy)
```

```
  Reaction Days Subject
1 249.5600    0     308
2 258.7047    1     308
3 250.8006    2     308
4 321.4398    3     308
5 356.8519    4     308
6 414.6901    5     308
```

Figure 10.1 shows each person's reaction times, with their own trend line.

```
ggplot(sleepstudy, aes(x = Days, y =
  Reaction)) +
  geom_point(size = 1) +
  geom_smooth(method = "lm", se = FALSE,
    linewidth = 0.7) +
  facet_wrap(~ Subject, ncol = 6) +
  labs(x = "Days of sleep restriction", y =
    "Reaction time (ms)") +
  theme_minimal(base_size = 10)
```

```
`geom_smooth()` using formula = 'y ~ x'
```




Figure 10.1: Reaction times over 10 days of sleep restriction, one panel per person, with each person's own trend line.

Two things are clear. Almost everyone gets slower as the days go by. And people differ: some start faster than others (different **intercepts**), and some slow down more than others (different **slopes**).

10.2.1 - Random intercepts

An ordinary regression would fit one line through all 180 points, as if they came from 180 different people. A **random intercept model** fits one average line, but lets each person have their own starting level. In the formula, `(1 | Subject)` means “a separate intercept for each subject”, and `lmer()` fits the model:

```
sleep_ri <- lmer(Reaction ~ Days + (1 | Subject), data = sleepstudy)
summary(sleep_ri)
```

```
Linear mixed model fit by REML ['lmerMod']
Formula: Reaction ~ Days + (1 | Subject)
Data: sleepstudy
```

```
REML criterion at convergence: 1786.5
```

```
Scaled residuals:
```

```
      Min       1Q   Median       3Q      Max
-3.2257 -0.5529  0.0109  0.5188  4.2506
```

```
Random effects:
```

```
Groups   Name             Variance Std.Dev.
Subject (Intercept) 1378.2    37.12
Residual              960.5    30.99
Number of obs: 180, groups: Subject, 18
```

```
Fixed effects:
```

```
              Estimate Std. Error t value
(Intercept) 251.4051    9.7467   25.79
Days         10.4673    0.8042   13.02
```

```
Correlation of Fixed Effects:
      (Intr)
Days -0.371
```

The output has two important parts. The **fixed effects** describe the average line, as in ordinary regression: on day 0, the average reaction time is about 251 ms, and each day of sleep restriction adds about 10.5 ms. The **random effects** describe how much people vary around the average line. The standard deviation of the intercepts, about 37 ms, says that people's starting levels typically differ from the average by that much; the residual standard deviation, about 31 ms, is the day-to-day variation within a person.

10.2.2 - The intraclass correlation

The random intercept model splits the variation in reaction times into two parts: stable differences between people, and variation within each person from day to day. The **intraclass correlation** (ICC) compares them:

$$\text{ICC} = \frac{\text{variance between groups}}{\text{variance between groups} + \text{variance within groups}} \quad (10.1)$$

The function `VarCorr()` extracts the variances:

```
vc <- as.data.frame(VarCorr(sleep_ri))
vc$vcov[1] / sum(vc$vcov)
```

```
[1] 0.5893089
```

About 59% of the variation in reaction times is due to stable differences between people. The ICC can be read as a comparison of how much people differ from each other with how much each person varies. An ICC of 0 would mean the grouping does not matter: two measurements of the same person are no more alike than measurements of two different people. The higher the ICC, the more the measurements of one person repeat each other, and the more wrong an ordinary regression would be.

10.2.3 - Random slopes

The random intercept model still assumes that everyone slows down at the same rate, which Figure 10.1 shows is not true. A **random slope model** lets each person have their own slope as well. In the formula, `(Days | Subject)` means “a separate intercept *and* slope of Days for each subject”:

```
sleep_rs <- lmer(Reaction ~ Days + (Days |
Subject), data = sleepstudy)
summary(sleep_rs)$coefficients
```

```
              Estimate Std. Error  t value
(Intercept) 251.40510    6.824597 36.838090
Days         10.46729    1.545790  6.771481
```

```
VarCorr(sleep_rs)
```

Groups	Name	Std.Dev.	Corr
Subject	(Intercept)	24.7407	
	Days	5.9221	0.066
Residual		25.5918	

The average effect of a day of sleep restriction is the same, about 10.5 ms, but its standard error has grown from 0.8 to 1.55. People's slopes vary with a standard deviation of about 5.9 ms per day, so for most people a day of sleep restriction adds somewhere between about 5 and 16 ms. Ignoring that variation made the random intercept model too confident about the average.

A **likelihood ratio test**, run with `anova()`, shows whether the random slope is worth adding:

```
anova(sleep_ri, sleep_rs)
```

```
refitting model(s) with ML (instead of REML)
```

```
Data: sleepstudy
Models:
sleep_ri: Reaction ~ Days + (1 | Subject)
sleep_rs: Reaction ~ Days + (Days | Subject)
      npar    AIC    BIC  logLik -2*log(L)
Chisq Df Pr(>Chisq)
sleep_ri      4  1802.1  1814.8  -897.04
1794.1
sleep_rs     6 1763.9 1783.1  -875.97   1751.9
42.139  2  7.072e-10 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*'
0.05 '.' 0.1 ' ' 1
```

The p-value is tiny: the random slopes improve the model clearly. The **AIC** column agrees. AIC balances fit against complexity, and lower is better.

10.3 - Change in wellbeing over two years

The wellbeing study has up to four semester records for each student, joined with their background information. The variable `time` counts semesters from the first, so that time 0 is semester 1:

```
panel <- semesters |>
  left_join(students, join_by(student_id)) |>
  mutate(time = semester - 1)
nrow(panel)
```

```
[1] 2326
```

Figure 10.2 shows the wellbeing of 30 randomly chosen students across the four semesters.

```
set.seed(9)
some_students <-
  sample(unique(panel$student_id), 30)
```



```

panel |>
  filter(student_id %in% some_students) |>
  ggplot(aes(x = semester, y = wellbeing, group
    = student_id)) +
  geom_line(alpha = 0.6) +
  labs(x = "Semester", y = "Wellbeing (0 to
    100)") +
  theme_minimal(base_size = 13)

```




Figure 10.2: Wellbeing over four semesters for 30 randomly chosen students.

Students differ enormously in their overall level, much more than any student changes from one semester to the next. A model with no predictors, only a random intercept for each student, puts a number on this through the ICC:

```

wb_null <- lmer(wellbeing ~ 1 + (1 |
  student_id), data = panel)
vc_null <- as.data.frame(VarCorr(wb_null))
vc_null$vcov[1] / sum(vc_null$vcov)

```

```
[1] 0.7809117
```

About 78% of all the variation in wellbeing is between students: students differ from each other far more than they change. The consequence for the amount of information in the data can be calculated. The **design effect**, $1 + (m - 1) \times \text{ICC}$, where m is the number of records per student, says how many times more records are needed to give the same information as independent observations:

```

icc <- vc_null$vcov[1] / sum(vc_null$vcov)
records_per_student <- nrow(panel) /
  n_distinct(panel$student_id)
design_effect <- 1 + (records_per_student - 1)
  * icc
c(records = nrow(panel), design_effect =
  design_effect,
  independent_equivalent = nrow(panel) /
  design_effect)

```

records	design_effect
independent_equivalent	

	2326.000000	3.246423
716.480955		

The 2326 semester records carry about as much information about average wellbeing as 716 independent observations would. Treating them as 2326 independent records would therefore be seriously wrong.

10.3.1 - Why ordinary regression misleads

The first question about change is whether wellbeing declines over time. The *wrong* way to answer it is an ordinary regression that ignores the structure:

```
summary(lm(wellbeing ~ time, data = panel))$coefficients
```

	Estimate	Std. Error	t value
Pr(> t)			
(Intercept)	61.6319636	0.4124650	149.423511
0.00000000			
time	-0.3794867	0.2235408	-1.697618
0.08971384			

The slope is slightly negative, but not significant. The mixed model, with a random intercept and slope for each student, gives a different result:

```
wb_growth <- lmer(wellbeing ~ time + (time | student_id), data = panel)
summary(wb_growth)$coefficients
```

	Estimate	Std. Error	t value
(Intercept)	61.6455545	0.4764479	129.385717
time	-0.4342392	0.1107653	-3.920355

The average slope is similar, about -0.43 points per semester, but its standard error has fallen from 0.22 to 0.11, and the decline is now clearly significant (a t value of about -3.9). The reason is that the ordinary regression lumps the large, stable differences *between* students into its error term, which drowns the small change *within* each student. The mixed model separates the two, and sees the change clearly.

In the simulation at the start of the chapter, ignoring the structure made a non-existent effect look significant; in the sleep study, ignoring the differences in slopes made the average look more certain than it was; here, ignoring the structure makes a real decline look uncertain. The direction of the error depends on the data, but the error is always there.

10.3.2 - Confidence intervals and p-values

The lme4 package deliberately prints no p-values for fixed effects, because calculating them exactly for mixed models is not straightforward. There are three common solutions. The first is to report **confidence intervals**, which are often more informative than p-values anyway.

The function `confint()` calculates them, and `method = "Wald"` is quick:

```
confint(wb_growth, parm = "beta_", method = "Wald")
```

```

              2.5 %      97.5 %
(Intercept) 60.7117337 62.5793752
time        -0.6513351 -0.2171432

```

The argument `parm = "beta_"` asks for the fixed effects only. The interval for `time` lies entirely below zero: wellbeing declines.

The second solution is a **likelihood ratio test**, comparing models with and without a term with `anova()`, as for the random slopes above. The third is the **lmerTest** package, which adds p-values to the `summary()` output and is used in most published research. Once it is loaded, `lmer()` models fitted afterwards include a p-value for each fixed effect:

```
library(lmerTest)
```

```
Attaching package: 'lmerTest'
```

```
The following object is masked from
'package:lme4':
```

```
lmer
```

```
The following object is masked from
'package:stats':
```

```
step
```

```
wb_growth_p <- lmer(wellbeing ~ time + (time
| student_id), data = panel)
summary(wb_growth_p)$coefficients
```

```

              Estimate Std. Error      df
t value      Pr(>|t|)
(Intercept) 61.6455545    0.4764479 598.9342
129.385717 0.000000e+00
time        -0.4342392    0.1107653 571.4997
-3.920355 9.914491e-05

```

The new columns are the degrees of freedom (estimated by Satterthwaite's method) and the p-value, `Pr(>|t|)`. The decline over time is clearly significant, in agreement with the confidence interval. Chapter 5's hypothesis for RQ8, that wellbeing declines over the two years, is supported.

10.3.3 - Persistence of the workshop effect

Chapter 7 showed that the workshop raised wellbeing in semester 2. A mixed model can compare the two groups in every semester at once, using every record, including the

first year of the students who later left. Treating `semester` as a factor lets each semester have its own average, and the interaction with `workshop` lets the gap between the groups differ from semester to semester:

```
wb_workshop <- lmer(wellbeing ~
  factor(semester) * workshop + (time |
  student_id),
  data = panel)
round(summary(wb_workshop)$coefficients, 2)
```

	Std. Error	df	t value	Estimate
(Intercept)				60.65
0.69	667.20	87.89		
factor(semester)2				4.82
0.42	1452.39	11.38		
factor(semester)3				2.49
0.45	1595.59	5.51		
factor(semester)4				0.21
0.48	667.71	0.43		
workshopNot invited				-0.38
0.98	667.20	-0.39		
factor(semester)2:workshopNot invited				-4.87
0.60	1452.39	-8.13		
factor(semester)3:workshopNot invited				-3.75
0.64	1597.01	-5.85		
factor(semester)4:workshopNot invited				-2.26
0.68	668.47	-3.30		

	Pr(> t)
(Intercept)	0.00
factor(semester)2	0.00
factor(semester)3	0.00
factor(semester)4	0.67
workshopNot invited	0.70
factor(semester)2:workshopNot invited	0.00
factor(semester)3:workshopNot invited	0.00
factor(semester)4:workshopNot invited	0.00

The coefficients are easier to understand as predicted averages. The function `predict()` with `re.form = NA` gives the predictions for the average student, ignoring the random effects:

```
predicted <- expand.grid(semester = 1:4,
  workshop = c("Invited", "Not invited")) |>
  mutate(time = semester - 1)
predicted$wellbeing <- predict(wb_workshop,
  newdata = predicted, re.form = NA)

ggplot(predicted, aes(x = semester, y =
  wellbeing, colour = workshop)) +
  geom_line(linewidth = 1) +
  geom_point(size = 2.5) +
  scale_colour_viridis_d(end = 0.8) +
  labs(x = "Semester", y = "Predicted
  wellbeing", colour = "Workshop") +
  theme_minimal(base_size = 13)
```




Figure 10.3: Predicted wellbeing in each semester by workshop group, from the mixed-effects model.

The gap between the groups is 0.4 points in semester 1, before the workshop, 5.2 in semester 2, 4.1 in semester 3, and 2.6 in semester 4. The workshop's effect is real, but it fades over the following year. For the thesis's recommendations, that matters: a short follow-up session each year might keep the benefit going.

The interaction coefficients test whether the gap in each semester differs from the gap in semester 1. All three do (the largest p-value is 0.001).

10.3.4 - Change in GPA

The same growth model can be fitted to GPA:

```
gpa_growth <- lmer(gpa ~ time + (time | student_id), data = panel)
```

```
boundary      (singular)      fit:      see
help('isSingular')
```

R warns that the fit is **singular**. This message is common, and worth understanding. It means the model is more complex than the data supports: here, the students' GPA slopes hardly vary at all, so the model cannot estimate their spread. The solution is to simplify, keeping only the random intercept:

```
gpa_growth <- lmer(gpa ~ time + (1 | student_id), data = panel)
summary(gpa_growth)$coefficients
```

	Estimate	Std. Error	df
t value			
Pr(> t)			
(Intercept)	3.0921628374	0.013015341	
806.5702	237.5783101	0.0000000	
time	-0.0004749256	0.003392086	1738.6474
-0.1400099	0.8886684		

The slope of `time` is almost exactly zero: grades stay stable over the two years, even though wellbeing falls. A non-change is a finding too, and here an interesting one.

10.4 - Students within supervisors

The students are also nested within supervisors, and a model can include random intercepts for both levels at once:

```
wb_levels <- lmer(wellbeing ~ time + (1 |
  supervisor_id) + (1 | student_id), data =
  panel)
VarCorr(wb_levels)
```

Groups	Name	Std.Dev.
student_id	(Intercept)	9.7748
supervisor_id	(Intercept)	4.4098
Residual		5.6189

Because every student ID is unique, R knows that students are nested within supervisors. Dividing each variance by the total shows where the variation lies:

```
vc_levels
<- as.data.frame(VarCorr(wb_levels))
data.frame(level = vc_levels$grp, share =
  round(vc_levels$vcov / sum(vc_levels$vcov),
  2))
```

	level	share
1	student_id	0.65
2	supervisor_id	0.13
3	Residual	0.22

About 13% of all the variation in wellbeing lies between supervisors: students of the same supervisor are noticeably more alike. Most of the variation, 65%, lies between students within the same supervisor, and the remaining 22% is change within students from semester to semester. For the university, the supervisor share is a practical finding: who supervises a student makes a measurable difference to their wellbeing. It also means that any comparison of supervisors, or of anything assigned to supervisors, must allow for this grouping, as the simulation at the start of the chapter showed.

10.5 - Counts and yes-or-no outcomes

Like ordinary regression, mixed models extend to other kinds of outcome. The function `glmer()` fits a **generalised linear mixed model**: `family = binomial` for yes-or-no outcomes, as in Chapter 8's logistic regression, and `family = poisson` for **counts**, such as the number of meetings a student had with their supervisor each semester.

The model below examines whether students who feel more supported meet their supervisor more often. Support is calculated from the questionnaire, as before:

```
support_scores <- questionnaire |>
  mutate(support =
  rowMeans(pick(support_1:support_6), na.rm =
  TRUE)) |>
```

```
select(student_id, support)

meetings_model <- glmer(
  supervisor_meetings ~ support + programme +
  (1 | supervisor_id) + (1 | student_id),
  data = panel |> left_join(support_scores,
  join_by(student_id)),
  family = poisson
)
summary(meetings_model)$coefficients
```

	Estimate	Std. Error	z value
Pr(> z)			
(Intercept)	0.06224346	0.07134373	0.8724446
support	0.39258255	0.01975713	19.8704237
programmePhD	0.17676057	0.02881998	6.1332638

Poisson coefficients are on a log scale, and `exp()` turns them into **rate ratios**:

```
exp(fixef(meetings_model))
```

(Intercept)	support	programmePhD
1.064221	1.480800	1.193345

Each extra point of support goes with about 48% more meetings per semester, and PhD students have about 19% more meetings than Master's students, holding support constant. As always with observational data, the direction is not certain: more meetings may also make students feel more supported.

i Missing data and mixed models

Chapter 6 found that the students who left after the first year were not a random group. Mixed models handle this better than most methods: they use every record each student provided, including the first year of those who left, instead of dropping those students entirely. The results remain trustworthy as long as leaving depends on things the model can see, such as the students' earlier wellbeing, which is the "missing at random" situation of Chapter 6.

10.6 - Reporting a mixed-effects model

A mixed-effects model is reported with both of its parts: the fixed effects, with confidence intervals, and the random effects, as standard deviations or variance shares. The report should also say how many observations and groups the model used, and which random effects it included.

💡 Writing it up

A linear mixed-effects model with random intercepts and slopes for students (2326 observations of 600 students) showed that wellbeing declined slightly over the four semesters, by 0.43 points per semester (95% CI [0.22, 0.65]). Students differed considerably in their overall wellbeing (*SD* of intercepts = 10.7); 78% of the variation in wellbeing was between students.

In your field: agriculture

R's `ChickWeight` data records the weight of 50 chicks, measured every few days from birth to 21 days old, on one of four diets. It has exactly the structure of the wellbeing data: repeated measurements of the same individuals, in groups. A random-slope model compares how fast chicks grow on each diet:

```
chick_model <- lmer(weight ~ Time * Diet
+ (Time | Chick), data = ChickWeight)
round(summary(chick_model)$coefficients,
2)
```

	Estimate	Std. Error	df	t
value Pr(> t)				
(Intercept)	33.66	2.92	47.54	
11.53	0.00			
Time	6.28	0.76	46.99	
8.24	0.00			
Diet2	-5.03	5.01	46.15	
-1.00	0.32			
Diet3	-15.41	5.01	46.15	
-3.08	0.00			
Diet4	-1.75	5.02	46.40	
-0.35	0.73			
Time:Diet2	2.33	1.30	45.76	
1.79	0.08			
Time:Diet3	5.15	1.30	45.76	
3.95	0.00			
Time:Diet4	3.25	1.31	45.86	
2.49	0.02			

Chicks on diet 1 gain about 6.3 grams a day. The interaction terms show how much faster chicks grow on the other diets: diet 3 adds about 5.1 grams a day, which is 82% more than the rate on diet 1.

10.7 - Common misconceptions

Mixed models are powerful, and several misunderstandings about grouped data survive even among experienced researchers.

- **“More rows means more evidence.”** Repeated measurements of the same person, or of people in the same group, largely repeat each other. The design effect shows how much less information they carry.

- **“Ignoring the grouping only makes results a little less precise.”** It can produce false alarms many times more often than intended, as the simulation showed, or hide real effects.
- **“Averaging each person’s records solves the problem.”** It removes the dependence, but throws away the information about change, and gives unequal weight to people with different numbers of records.
- **“A singular fit means the model is wrong.”** It means the model is more complex than the data can support; simplifying the random effects is usually enough.

10.8 - Chapter review

10.8.1 - Summary

- Repeated measures and nested data break the independence assumption of ordinary regression, which then gets the uncertainty wrong, in either direction. A simulation shows false alarms far above 5% when grouping is ignored.
- Mixed-effects models have fixed effects (the average relationships) and random effects (how groups vary around them). `(1 | group)` gives each group its own intercept; `(x | group)` also its own slope of `x`.
- The intraclass correlation is the share of variation between groups: how much groups differ from each other compared with how much members vary within them. The design effect shows how much less information grouped records carry.
- `lme4` prints no p-values: use confidence intervals, likelihood ratio tests with `anova()`, or the `lmerTest` package.
- Interactions with time show whether an effect grows or fades. `predict(..., re.form = NA)` gives predictions for the average case.
- A singular fit means the model is too complex for the data: simplify the random effects.
- `glmer()` fits mixed models for yes-or-no outcomes (`binomial`) and counts (`poisson`; `exp()` of the coefficients gives rate ratios).
- Mixed models use all available records, which helps when participants drop out.

10.8.2 - Key terms

Repeated measures, nested data, mixed-effects model, multilevel model, fixed effect, random effect, random intercept, random slope, intraclass correlation, design effect, likelihood ratio test, AIC, singular fit, generalised linear mixed model, Poisson model, count outcome, rate ratio.

10.9 - Exercises

The [playground](#) has these and more, with hints and solutions.

1. Fit a random intercept model for `sleep_hours` over time (`sleep_hours ~ time + (1 | student_id)`). Calculate the ICC, and describe whether sleep changes over the semesters.

2. Add a random slope of `time` to the model in Exercise 1, and compare the two models with `anova()`. Decide whether the random slope is worth keeping.
3. Fit `wellbeing ~ time * study_mode + (time | student_id)`, and describe whether part-time students' wellbeing changes at a different rate.
4. Fit a model with random intercepts for supervisors and students for `study_hours`, and calculate how much of its variation lies between supervisors.
5. Change the simulation at the start of the chapter so that supervisors differ more (`sd = 10` for `style`). Describe what happens to the false alarm rate of the ordinary regression, and explain why.
6. In your own words, explain to a fellow student why an ordinary regression of all the semester records would be wrong.

10.10 - Further reading

- *Data Analysis Using Regression and Multilevel/Hierarchical Models* [20] is a classic, readable introduction to multilevel models.
- “Fitting Linear Mixed-Effects Models Using lme4” [21] describes the `lme4` package and its formula syntax in detail.

11.13	Common misconceptions	251
11.14	Chapter review	252
11.15	Exercises	252
11.16	Further reading	253
12	Classification Models	255
12.1	Data and resampling	256
12.2	Decision trees	257
12.3	Random forests	259
12.4	k-nearest neighbours	260
12.5	Support vector machines	262
12.6	Comparing models fairly	263
12.7	The confusion matrix	264
12.8	Choosing the threshold	265
12.9	Imbalanced outcomes	270
12.10	Chapter review	272
12.11	Exercises	273
12.12	Further reading	274
13	Predictive Regression	275
13.1	Data for predicting final GPA	276
13.2	Measuring prediction error	277
13.3	Ordinary regression with many predictors	279
13.4	Regularised regression	280
13.5	Boosting	285
13.6	Comparing the models	286
13.7	The final test	287
13.8	Common misconceptions	290
13.9	Chapter review	290
13.10	Exercises	290
13.11	Further reading	291
14	Advanced Clustering	293
14.1	The profile data	294
14.2	What counts as a group	295
14.3	Gaussian mixture models	297
14.4	Clusters as dense regions	302
14.5	Evaluating a clustering	305
14.6	Common misconceptions	308
14.7	Chapter review	308
14.8	Exercises	309
14.9	Further reading	309
15	Neural Networks	311
15.1	A single neuron	311
15.2	From neurons to networks	312
15.3	A neural network for the dropout question	316
15.4	When neural networks are worth using	320
15.5	Deep learning	320
15.6	Common misconceptions	323
15.7	Chapter review	323
15.8	Exercises	324
15.9	Further reading	324
16	Time Series Forecasting	325
16.1	Time series data	326
16.2	The counselling service data	328
16.3	Decomposition	332
16.4	Forecasting	333
16.5	Common misconceptions	339
16.6	Chapter review	340
16.7	Exercises	340
16.8	Further reading	341

11. Introduction to Machine Learning in R

Some research questions are not about why something happens, but about what will happen. A hospital wants to know which patients are likely to be readmitted, a bank which loans are likely to fail, a university which students are likely to struggle. The goal is not to understand the causes, although that may help, but to make accurate predictions for new cases, early enough to act on them. Such questions need a different standard of evidence. For explanation, a model is judged by whether its coefficients are well estimated and make sense. For prediction, a model is judged by one thing only: how well it works on cases it has never seen.

Machine learning is the set of methods and practices for building predictive models and testing them fairly. This chapter introduces its core ideas: the difference between prediction and explanation, why a model must be tested on new data, how models overfit, and how cross-validation and tuning choose a model honestly. It uses the **tidymodels** framework, whose steps are best understood as the design of a fair test. In the story, Chapter 8 used logistic regression to explain who considers dropping out; Elaf’s supervisor now asks whether the university could identify, at the end of a student’s first semester, the students at risk, so that help could be offered in time (RQ9).

💡 By the end of this chapter you will be able to

- Explain supervised and unsupervised learning, and prediction versus explanation as different kinds of research question.
- Explain generalisation, overfitting, and the bias-variance trade-off, and show them by simulation.
- Split data into training and test sets, and explain why a model must be tested on data it has not seen.
- Prepare data for modelling with a recipe, and combine a recipe and a model in a workflow.
- Evaluate a classification model, and explain why accuracy alone can mislead.
- Use cross-validation to estimate performance honestly, tune a model’s hyperparameters, and evaluate the final model on the test set.

Note

Teaching & Practice: [Interactive Lecture Slides](#) · [Browser Playground Exercises](#)

11.1 - Machine learning and prediction

Machine learning covers methods that learn patterns from data in order to make predictions or find structure. In **supervised learning**, the data includes the outcome to be predicted, and the model learns the relationship between the predictors and that outcome. If the outcome is a category, such as “considering dropout: yes or no”, it is a **classification** problem; if it is a number, such as next semester’s GPA, it is a **regression** problem. Chapters 11 to 13 and Chapter 15 are about supervised learning. In **unsupervised learning**, there is no outcome, and the model looks for structure in the data itself. The clustering of Chapter 9 is unsupervised, and Chapter 14 returns to it.

Many machine learning methods are familiar statistics: logistic regression is both a statistical model and one of the most useful classifiers. What changes is the goal, and with it the way models are judged. Table 11.1 sets the two aims side by side.

Table 11.1: Explanation and prediction

	Explanation (Chapters 7 to 10)	Prediction (Chapters 11 to 16)
Question	<i>Why</i> does something happen?	<i>What</i> will happen for a new case?
Judged by	Sensible, well-estimated effects	Accuracy on new data
Typical models	Simple, interpretable	Anything that predicts well
Main danger	Confounding	Overfitting

A predictive model does not need to be causally correct to be useful: a model can predict dropout well from variables that do not cause it. Equally, a model that explains well may predict poorly, because the effects it estimates are small compared with everything it cannot see. Chapter 5 described prediction as a third aim of analysis, alongside description and explanation; this part of the book develops it.

11.2 - Generalisation and overfitting

The central idea of machine learning is **generalisation**: a model is useful only if what it learned from one set of data carries over to new data. Any dataset contains two things, the pattern that would appear again in new data and the noise that belongs to this sample only. A model that learns the pattern generalises; a model that also learns the noise

overfits, and looks better on its own data than it will ever be on new data.

A simulation makes the idea visible. The code below draws 15 points from a smooth curve with random noise added, and fits three models of increasing flexibility: a straight line, a gentle curve, and a very flexible curve (polynomials of degree 1, 3, and 12). It then draws new points from the same process, within the same range, to see how each model does on data it has not seen:

```
library(dplyr)
```

```
Attaching package: 'dplyr'
```

```
The following objects are masked from
'package:stats':
```

```
filter, lag
```

```
The following objects are masked from
'package:base':
```

```
intersect, setdiff, setequal, union
```

```
library(ggplot2)
```

```
true_curve <- function(x) sin(2 * x)
make_points <- function(n) {
  x <- runif(n, 0, 3)
  tibble(x = x, y = true_curve(x) + rnorm(n,
sd = 0.3))
}
```

```
set.seed(3)
train_points <- make_points(15)
new_points <- make_points(300) |>
  filter(x >= min(train_points$x), x <=
max(train_points$x)) # same range as training
```

```
grid <- tibble(x = seq(min(train_points$x),
max(train_points$x), length.out = 300))
curves <- bind_rows(lapply(c(1, 3, 12),
function(d) {
  fit <- lm(y ~ poly(x, d), data = train_points)
  tibble(grid, y = predict(fit, grid), model
= paste("Degree", d))
})))
curves$model <- factor(curves$model, levels =
paste("Degree", c(1, 3, 12)))
```

```
ggplot() +
  geom_point(data = new_points, aes(x, y),
colour = "grey75", size = 0.8) +
  geom_point(data = train_points, aes(x, y),
size = 2) +
  geom_line(data = curves, aes(x, y), colour
= "#b2182b", linewidth = 0.9) +
  facet_wrap(~ model) +
```

```
coord_cartesian(ylim = c(-2, 2)) +
theme_minimal(base_size = 11)
```




Figure 11.1: Three models fitted to the same 15 training points (black). The straight line is too simple, the degree-3 curve follows the pattern, and the degree-12 curve passes close to every training point but wanders far from the new points (grey).

The straight line is too simple to follow the wave: it **underfits**. The degree-3 curve captures the pattern. The degree-12 curve bends to pass close to every training point, and in doing so follows the noise, swinging far away from where new points actually lie. Measured on the training points, it would look like the best model; measured on new points, it is the worst. The prediction error of every degree from 1 to 12 shows the general pattern:

```
rmse <- function(observed, predicted)
sqrt(mean((observed - predicted)^2))

errors <- bind_rows(lapply(1:12, function(d) {
  fit <- lm(y ~ poly(x, d), data = train_points)
  tibble(degree = d,
         training = rmse(train_points$y,
                        predict(fit, train_points)),
         new_data = rmse(new_points$y,
                        predict(fit, new_points)))
}))
```

```
errors |>
  tidyr::pivot_longer(-degree, names_to =
    "data", values_to = "error") |>
  ggplot(aes(x = degree, y = error, colour =
    data)) +
  geom_line(linewidth = 1) +
  geom_point() +
  scale_x_continuous(breaks = 1:12) +
  scale_colour_manual(values = c(training =
    "grey50", new_data = "#b2182b"),
                     labels = c(training =
    "Training points", new_data = "New points")) +
  scale_y_log10() +
  labs(x = "Flexibility (polynomial degree)",
       y = "Prediction error (RMSE, log scale)",
       colour = NULL) +
  theme_minimal(base_size = 12)
```





Figure 11.2: Prediction error (RMSE) on the training points and on new points, for polynomials of degree 1 to 12. Training error keeps falling as the model becomes more flexible; error on new data falls, then rises.

The training error falls steadily as the model becomes more flexible: a more flexible model can always fit its own data more closely. The error on new data falls at first, reaches its minimum at degree 3, and then rises. This U shape is known as the **bias-variance trade-off**. A model that is too simple has high **bias**: it misses part of the real pattern, whatever data it is given. A model that is too flexible has high **variance**: it changes a great deal from one sample to another, because it follows each sample's noise. The best predictions come from a model in between, and the only way to find it is to measure performance on data the model did not learn from. The rest of the chapter builds that measurement into every step.

11.3 - Designing a fair test

The **tidymodels** framework is a collection of packages that together make a fair test of a predictive model easy to carry out. Each package handles one part of the design: **rsample** splits the data, **recipes** prepares it, **parsnip** specifies models, **workflows** combines the pieces, **tune** tunes models, and **yardstick** measures performance. Loading **tidymodels** loads them all:

```
library(tidymodels)
```

—	Attaching	packages
tidymodels 1.5.0 —		

✓ broom	1.0.13	✓ rsample	1.3.2
✓ dials	1.4.4	✓ tailor	0.1.0
✓ infer	1.1.0	✓ tidyr	1.3.2
✓ modeldata	1.6.0	✓ tune	2.1.0
✓ parsnip	1.6.0	✓ workflows	1.3.0
✓ purrr	1.2.2	✓ workflowsets	1.1.1
✓ recipes	1.4.0	✓ yardstick	1.4.0

—	Conflicts
---	-----------

```
tidymodels_conflicts() —  
* purrr::discard()   masks scales::discard()  
* dplyr::filter()    masks stats::filter()  
* dplyr::lag()        masks stats::lag()  
* .GlobalEnv::rmse() masks yardstick::rmse()  
* recipes::step()     masks stats::step()
```

```
tidymodels_prefer()
```

The function `tidymodels_prefer()` settles a few name clashes between packages in favour of `tidymodels`. Figure 11.3 shows how the pieces fit together: the test data is set aside before anything else happens, all choices are made with the training data alone, and the test data is used once, at the end. The rest of the chapter works through each step.




Figure 11.3: The machine learning workflow in tidymodels.

11.4 - Data for prediction

A predictive model can only use information that would be available when the prediction is made. The university wants to act at the end of the first semester, so the model may use the students' background, their questionnaire scores from the start of the year, and their first-semester records, but nothing from later. The questionnaire scores are calculated as before:

```
scores <- questionnaire |>
  mutate(
    stress_4 = 6 - stress_4,
    stress =
    rowMeans(pick(stress_1:stress_6), na.rm =
    TRUE),
    burnout =
    rowMeans(pick(burnout_1:burnout_6), na.rm =
    TRUE),
    support =
    rowMeans(pick(support_1:support_6), na.rm =
    TRUE),
    satisfaction = rowMeans(pick(satisfaction_1:satisfaction_4),
    na.rm = TRUE)
  ) |>
  select(student_id, stress, burnout, support,
  satisfaction)

dropout_data <- students |>
  left_join(scores, join_by(student_id)) |>
  left_join(semesters |> filter(semester == 1)
  |> select(-semester), join_by(student_id)) |>
  mutate(considering_dropout =
  factor(considering_dropout, levels = c("Yes",
  "No"))) |>
  select(-student_id, -supervisor_id, -
  workshop, -workshop_sessions)

dim(dropout_data)
```

```
[1] 600 21
```

Two details matter. The outcome must be a factor, and its first level is treated as the event to predict, so the levels are set to put "Yes" first. And the workshop variables are removed: the workshop took place after the first semester, so it would not be known when the prediction is made. Using information that would not be available at the time of prediction is a form of **data leakage**, and it makes a model look better than it can ever be in practice.

```
dropout_data |> count(considering_dropout) |>
  mutate(share = n / sum(n))
```

	considering_dropout	n	share
1	Yes	90	0.15
2	No	510	0.85

Only about 15% of students answer "Yes". Such an **imbalanced** outcome is common in real prediction problems, from rare diseases to fraud, and, as the evaluation below

shows, it catches out anyone who judges a model by accuracy alone.

11.5 - Splitting the data

The most important rule of machine learning follows from the simulation above: **never judge a model on the data it learned from**. The first step is therefore to split the data. The **training set** is used to build and tune the model. The **test set** is locked away and used only once, at the very end, to estimate how the final model will perform on new students. The function `initial_split()` does the split, and `strata` makes sure that both sets have the same share of “Yes” answers:

```
set.seed(2026)
dropout_split <- initial_split(dropout_data,
  prop = 0.75, strata = considering_dropout)
dropout_train <- training(dropout_split)
dropout_test  <- testing(dropout_split)

nrow(dropout_train)
```

```
[1] 449
```

```
nrow(dropout_test)
```

```
[1] 151
```

Three quarters of the students are used for training and a quarter are set aside for the final test.

11.6 - Preparing the data

Most models need the data prepared first: missing values filled in, categories turned into numbers, and numeric predictors put on a common scale. A **recipe** lists these steps:

```
dropout_recipe <- recipe(considering_dropout
  ~ ., data = dropout_train) |>
  step_impute_median(all_numeric_predictors())
  |>
  step_dummy(all_nominal_predictors()) |>
  step_normalize(all_numeric_predictors())
```

The first line says that `considering_dropout` is the outcome and every other column (.) is a predictor. The step `step_impute_median()` then replaces each missing numeric value with the median of that variable. The step `step_dummy()` turns each categorical predictor into **dummy variables**: 0/1 columns, one per category except a reference category, as `lm()` did automatically in Chapter 8. Finally, `step_normalize()` turns every numeric predictor into z-scores, so they are on the same scale; some models, such as k-nearest neighbours below, depend on distances and need this.

The recipe only describes the steps. To see what it produces, `prep()` estimates what it needs from the training

data (the medians, the means and standard deviations), and `bake()` applies it:

```
dropout_recipe |>
  prep() |>
  bake(new_data = NULL) |>
  glimpse()
```

```
Rows: 449
Columns: 25
 $ age                <dbl> 0.24508282,
-0.38372967, 0.87389531, 1.293103...
 $ financial_worry    <dbl> -0.7012103,
-0.7012103, -0.7012103, 1.9008013...
 $ stress             <dbl> -0.07328352,
0.78371678, 1.95645404, -0.29880...
 $ burnout            <dbl> -0.5660882,
0.1216109, 2.1847084, 1.0385432, ...
 $ support            <dbl> -1.48927219,
-1.70433127, 0.66131866, -0.8440...
 $ satisfaction       <dbl> -1.1166487,
-1.1166487, -0.4684583, -0.792553...
 $ gpa                <dbl> -0.20629899,
0.49821479, -1.55406449, -0.7576...
 $ sleep_hours        <dbl> 0.705824008,
-0.090694948, -1.185908512, -0.5...
 $ study_hours        <dbl> -0.09621851,
0.36256496, 0.89781234, -0.78439...
 $ exercise_days      <dbl> 1.5418905,
-0.1263236, -0.1263236, -0.1263236...
 $ caffeine_mg        <dbl> -0.41150631,
0.17158271, 1.00977318, -0.15640...
 $ supervisor_meetings <dbl> -1.50209363,
-0.42757418, -0.06940103, 2.0796...
 $ wellbeing          <dbl> 1.13636876,
-0.04432777, -1.22502430, -0.2973...
 $ considering_dropout <fct> No, No, No,
No, No, No, No, No, No, No, No, N...
 $ gender_Male        <dbl> -0.9489638,
-0.9489638, -0.9489638, 1.0514340...
 $ faculty_Health.Sciences <dbl> 1.6539530,
-0.6032655, 1.6539530, -0.6032655,...
 $ faculty_Humanities <dbl> -0.4072635,
-0.4072635, -0.4072635, 2.4499443...
 $ faculty_Natural.Sciences <dbl> -0.4365273,
-0.4365273, -0.4365273, -0.436527...
 $ faculty_Social.Sciences <dbl> -0.4861964,
-0.4861964, -0.4861964, -0.486196...
 $ programme_PhD      <dbl> -0.6445746,
1.5479556, -0.6445746, 1.5479556,...
 $ study_mode_Part.time <dbl> -0.6342138,
-0.6342138, 1.5732436, -0.6342138...
 $ employment_None    <dbl> 0.9966636,
-1.0011130, -1.0011130, -1.0011130...
 $ employment_Part.time.job <dbl> -0.7039606,
1.4173703, -0.7039606, 1.4173703,...
 $ has_children_Yes   <dbl> -0.5449999,
-0.5449999, -0.5449999, -0.544999...
 $ lives_away_Yes     <dbl> 1.1662426,
1.1662426, -0.8555448, 1.1662426, ...
```

In this code, `new_data = NULL` means “the training data the recipe was prepared on”. In the result, gender has

become `gender_Male`, `faculty` has become four dummy columns, and all the numbers are now z-scores.

! No peeking at the test data

The medians, means, and standard deviations are estimated from the **training** data only, and then applied unchanged to the test data. If they were calculated from all the data, information from the test set would leak into the model, and the test would no longer be a fair one. Recipes and workflows handle this automatically, which is one of the best reasons to use them.

11.7 - Specifying and fitting a model

The **parsnip** package specifies models in one consistent way, whatever package does the actual work. Here is logistic regression:

```
logistic_spec <- logistic_reg()
logistic_spec
```

```
Logistic Regression Model Specification
(classification)
```

```
Computational engine: glm
```

A **workflow** combines the recipe and the model, so they always travel together:

```
logistic_wf <- workflow() |>
  add_recipe(dropout_recipe) |>
  add_model(logistic_spec)

logistic_fit <- fit(logistic_wf, data =
  dropout_train)
```

The function `fit()` prepares the recipe on the training data and fits the model to the prepared data, in one step.

11.8 - Evaluating a classifier

The model is now judged on the test students. The function `augment()` adds the model's predictions to the test data: the predicted class (`.pred_class`) and the predicted probability of each class (`.pred_Yes`, `.pred_No`):

```
test_results <- augment(logistic_fit,
  new_data = dropout_test)
test_results |>
  select(considering_dropout, .pred_class, .pred_Yes)
|>
  head()
```

```
# A tibble: 6 × 3
  considering_dropout .pred_class .pred_Yes
  <fct>               <fct>       <dbl>
1 No                 No          0.216
```

2	No	No	0.0279
3	No	No	0.00744
4	Yes	No	0.424
5	No	No	0.0114
6	No	No	0.0298

The simplest measure is **accuracy**: the share of students classified correctly. The **yardstick** package calculates it:

```
test_results |> accuracy(truth = considering_dropout, estimate = .pred_class)
```

```
# A tibble: 1 × 3
  .metric .estimator .estimate
  <chr>   <chr>      <dbl>
1 accuracy binary      0.834
```

An accuracy of 83% sounds good. But a “model” that ignores every predictor and says “No” for every student would be right 85% of the time, because most students answer “No”, while finding not one of the students at risk. For an imbalanced outcome, accuracy is a misleading measure.

A better measure is the **ROC AUC** (the area under the ROC curve). It is the probability that, of one student who considered dropping out and one who did not, the model gives the first a higher predicted probability than the second. An AUC of 0.5 is no better than a coin toss, and 1.0 is perfect. Unlike accuracy, it is not fooled by an imbalanced outcome. It uses the predicted probabilities:

```
test_results |> roc_auc(truth = considering_dropout, .pred_Yes)
```

```
# A tibble: 1 × 3
  .metric .estimator .estimate
  <chr>   <chr>      <dbl>
1 roc_auc binary      0.840
```

An AUC of 0.84 means the model ranks an at-risk student above a not-at-risk student about 84% of the time: a useful model. Chapter 12 looks at classification measures in much more detail, including how to choose the threshold at which a student is flagged.

11.9 - Overfitting in practice

The simulation at the start of the chapter showed overfitting with a flexible curve. The same happens with real data. A dramatic example is **k-nearest neighbours** (k-NN) with $k = 1$. It predicts each student’s outcome by finding the single most similar student in the training data and copying their answer. On the training data, every student’s nearest neighbour is themselves, so it can never be wrong:

```
knn1_wf <- workflow() |>
  add_recipe(dropout_recipe) |>
  add_model(nearest_neighbor(neighbors = 1) |
  > set_mode("classification"))
knn1_fit <- fit(knn1_wf, data = dropout_train)
```



```
augment(knn1_fit, new_data = dropout_train) |
> roc_auc(considering_dropout, .pred_Yes)
```

```
# A tibble: 1 × 3
  .metric .estimator .estimate
  <chr>   <chr>       <dbl>
1 roc_auc binary         1
```

```
augment(knn1_fit, new_data = dropout_test) |
> roc_auc(considering_dropout, .pred_Yes)
```

```
# A tibble: 1 × 3
  .metric .estimator .estimate
  <chr>   <chr>       <dbl>
1 roc_auc binary     0.671
```

The model scores perfectly on the training data, and much more poorly on new students. The training score was an illusion: the model had memorised the training students, not learned anything that carries over. It is the degree-12 curve again, and the reason every model in this book is tested on data it has not seen.

11.10 - Cross-validation

The test set is used only once, at the end. During model building, however, models and settings often need to be compared, and neither the training data (that would reward overfitting) nor the test set (that would use it up) can be used to judge them. **Cross-validation** solves this by reusing the training data. The training data is split into, say, 10 equal parts, called **folds**. The model is fitted on 9 folds and its performance measured on the fold left out; this is repeated 10 times, leaving out each fold in turn, and the 10 performance measures are averaged. Every student is used for evaluation exactly once, always by a model that did not see them during fitting. Figure 11.4 shows the idea with 5 folds.




Figure 11.4: Five-fold cross-validation: in each round, the model is fitted on four folds (grey) and evaluated on the fifth (dark).

The function `vfold_cv()` creates the folds, and `fit_resamples()` fits and evaluates the workflow on each:

```
set.seed(2026)
dropout_folds <- vfold_cv(dropout_train, v =
  10, strata = considering_dropout)

logistic_cv <- fit_resamples(logistic_wf,
  resamples = dropout_folds,
  metrics =
  metric_set(roc_auc, accuracy))
collect_metrics(logistic_cv)
```

```
# A tibble: 2 × 6
  .metric .estimator mean n
std_err .config
  <chr>    <chr>    <dbl> <int> <dbl>
<chr>
1 accuracy binary    0.862    10 0.00806
pre0_mod0_post0
2 roc_auc binary    0.804    10 0.0303
pre0_mod0_post0
```

The cross-validated AUC, about 0.8, is an honest estimate of how the model will perform on new students, obtained without touching the test set. The standard error (`std_err`) shows how much the estimate varies across folds.

11.11 - Tuning a model

Many models have settings that are not learned from the data but must be chosen beforehand, such as k in k -nearest neighbours or the degree of the polynomials in the simulation. They are called **hyperparameters**, to distinguish them from the **parameters**, such as regression coefficients, that the model learns. The best value is found by **tuning**: trying several values and comparing them with cross-validation.

A **decision tree** (explained fully in Chapter 12) predicts by asking a series of yes-or-no questions, such as whether stress is above 3.5. Its depth, the number of questions in a row, is a hyperparameter that controls its flexibility. A shallow tree is too simple to capture the patterns and underfits; a very deep tree can fit the training data closely, noise included, and overfits. In the model specification, `tune()` marks the hyperparameter to be tuned:

```
tree_spec <- decision_tree(tree_depth =
  tune(), min_n = 10, cost_complexity = 0) |>
  set_mode("classification")

tree_wf <- workflow() |>
  add_recipe(dropout_recipe) |>
  add_model(tree_spec)
```

The function `tune_grid()` tries each value in a grid, using the same cross-validation folds:

```
set.seed(2026)
tree_tuning <- tune_grid(tree_wf, resamples =
  dropout_folds,
  grid = tibble(tree_depth
```

```

= 1:10),
                                     metrics
= metric_set(roc_auc))

autoplot(tree_tuning)
theme_minimal(base_size = 12)

```




Figure 11.5: Cross-validated ROC AUC of decision trees of different depths.

The pattern is the new-data curve of Figure 11.2 seen from the other side: a tree of depth 1 is no better than chance, the AUC rises as the tree grows, and beyond a moderate depth it levels off and dips slightly as the deeper trees start to overfit. The function `select_best()` picks the best depth, here 4:

```

best_depth <- select_best(tree_tuning, metric
= "roc_auc")
best_depth

```

```

# A tibble: 1 × 2
  tree_depth .config
    <int> <chr>
1         4 pre0_mod04_post0

```

11.12 - The final test

Only now, with the model chosen, is the test set used. The function `finalize_workflow()` plugs the best depth into the workflow, and `last_fit()` fits it on the whole training set and evaluates it once on the test set:

```

final_tree <- tree_wf |>
  finalize_workflow(best_depth) |>
    last_fit(dropout_split, metrics =
metric_set(roc_auc, accuracy))
collect_metrics(final_tree)

```

```

# A tibble: 2 × 4
  .metric .estimator .estimate .config
  <chr>    <chr>         <dbl> <chr>
1 accuracy binary         0.828 pre0_mod0_post0
2 roc_auc  binary         0.815 pre0_mod0_post0

```

The same is done for logistic regression, for comparison:

```
final_logistic <- last_fit(logistic_wf,
  dropout_split, metrics = metric_set(roc_auc,
  accuracy))
collect_metrics(final_logistic)
```

```
# A tibble: 2 × 4
  .metric .estimator .estimate .config
  <chr>   <chr>       <dbl> <chr>
1 accuracy binary      0.834 pre0_mod0_post0
2 roc_auc binary      0.840 pre0_mod0_post0
```

The tuned tree reaches an AUC of 0.81; plain logistic regression, 0.84. The more flexible model is not the better one. In this data, the risk of dropout rises steadily with stress and falls steadily with support, so a simple smooth model captures the pattern well, while a tree, which splits the data into boxes, captures it less well. Trying a simple model first, and making a complex one earn its place, is one of the most useful habits in machine learning.

Predictions about people

A model that flags students “at risk” affects real people. Before such a model is used, the consequences of its errors must be considered: a student flagged by mistake may be treated differently, and a student who is missed receives no help. The model should be checked separately for different groups, such as women and men or part-time and full-time students, because a model with a good AUC overall can still be unfair to a particular group. And its purpose matters: in this study, the right use is to offer support earlier, never to penalise.

In your field: business and finance

Banks use exactly these methods to predict which loan applicants will repay. The `credit_data` dataset from the **modeldata** package (installed with tidy-models) records real loan applications, with whether each loan turned out well or badly. The same workflow applies:

```
data(credit_data, package = "modeldata")

set.seed(1)
credit_split
<- initial_split(credit_data, strata =
  Status)

credit_wf <- workflow() |>
  add_recipe(
    recipe(Status ~ ., data =
  training(credit_split)) |>
    step_impute_median(all_numeric_predictors())
  |>
    step_impute_mode(all_nominal_predictors())
  |>
    step_dummy(all_nominal_predictors())
  ) |>
  add_model(logistic_reg())

last_fit(credit_wf, credit_split) |>
collect_metrics()
```

```
# A tibble: 3 × 4
  .metric .estimator .estimate .config
  <chr>    <chr>      <dbl> <chr>
1 accuracy binary      0.782 pre0_mod0_post0
2 roc_auc  binary      0.824 pre0_mod0_post0
3 brier_class binary     0.148 pre0_mod0_post0
```

The step `step_impute_mode()` fills missing categories with the most common one. Only the data and the outcome have changed; every step is the same. The output also shows the **Brier score**, `brier_class`, another measure of how good the predicted probabilities are: 0 is perfect, and lower is better.

11.13 - Common misconceptions

Predictive modelling has its own characteristic mistakes, most of them versions of testing a model on what it already knows.

- **“A model that fits the training data well will predict well.”** Training performance rewards overfitting; only performance on new data counts.
- **“A more complex model is a better model.”** Beyond a point, flexibility adds variance, not accuracy. Simple models often predict as well.

- **“High accuracy means a good classifier.”** For an imbalanced outcome, a model that predicts the common class every time can have high accuracy and no value.
- **“The test set can be used to choose between models.”** Every use of the test set for a decision makes it less of a test. Choices are made with cross-validation; the test set is used once.
- **“A good predictor is a cause.”** Prediction needs association, not causation; a predictive model says nothing about what would happen if a predictor were changed.

11.14 - Chapter review

11.14.1 - Summary

- Supervised learning predicts an outcome (classification for categories, regression for numbers); unsupervised learning finds structure. Prediction is a different research aim from explanation, judged by performance on new data.
- A model generalises when it learns the pattern and not the noise. Training error keeps falling as models become more flexible, while error on new data falls and then rises: the bias-variance trade-off.
- `tidymodels` designs a fair test: it splits data (`rsample`), prepares it (`recipes`), specifies models (`parsnip`), combines them (`workflows`), tunes them (`tune`), and measures performance (`yardstick`).
- Split the data into training and test sets, stratified by the outcome, and use the test set only once, at the end. Use only information available at the time of prediction.
- A recipe prepares data (imputation, dummy variables, normalisation) using the training data only, so no information leaks from the test set.
- Accuracy misleads for imbalanced outcomes; the ROC AUC is a better summary.
- Cross-validation estimates performance honestly without using the test set. Hyperparameters are chosen by tuning with cross-validation. Simple models often predict as well as complex ones; try them first.
- Predictions about people raise questions of fairness and use.

11.14.2 - Key terms

Machine learning, supervised learning, unsupervised learning, classification, regression, prediction, explanation, generalisation, overfitting, underfitting, bias-variance trade-off, imbalanced outcome, training set, test set, stratified split, recipe, imputation, dummy variable, normalisation, data leakage, model specification, workflow, accuracy, ROC AUC, k-nearest neighbours, cross-validation, fold, hyperparameter, parameter, tuning, decision tree.

11.15 - Exercises

The [playground](#) has these and more, with hints and solutions.

1. Remove the questionnaire scores (`stress`, `burnout`, `support`, `satisfaction`) from `dropout_data` and fit the logistic regression workflow again. Report how much the cross-validated AUC falls, and what that says about the questionnaire.
2. Split the data 80/20 instead of 75/25 with a different seed, and describe how much the test AUC of the logistic model changes, and why it might.
3. Tune k-nearest neighbours over `neighbors = c(5, 11, 21, 41)` with cross-validation. Identify the best value, and compare its AUC with logistic regression.
4. Add `step_zv(all_predictors())` to the recipe (it removes predictors with only one value). Look up what “zero variance” means, and explain when this step is useful.
5. In two or three sentences, explain to a university administrator why a model with 85% accuracy might still be useless for finding students at risk.
6. Repeat the polynomial simulation with 50 training points instead of 15. Describe how the error curve on new data changes, and explain why more data allows a more flexible model.

11.16 - Further reading

- *Tidy Modeling with R* [22], the book by the authors of `tidymodels`, is the definitive guide to the framework.
- *An Introduction to Statistical Learning* [18] explains the ideas behind overfitting, the bias-variance trade-off, cross-validation, and the methods of the next chapters.

12. Classification Models

A classification model is built to support decisions: which patients to screen further, which applications to review, which students to contact. Two questions follow. The first is technical: whether more flexible methods than logistic regression would rank the cases better. The second matters more, and is often forgotten: what happens when the model's predictions are turned into decisions. Every decision rule flags some cases wrongly and misses others, and the balance between those two errors is not a statistical matter alone. It depends on what each error costs, and on whom.

This chapter addresses both questions. It introduces four widely used families of classification models, decision trees, random forests, k-nearest neighbours, and support vector machines, each through the idea behind it and the way it separates the classes, and compares them fairly on the wellbeing data. It then looks beyond the AUC at how a model is used: the confusion matrix, precision and recall, the choice of the threshold at which a student is flagged, and what to do when one outcome is rare. In the story, Elaf shows her logistic regression to the student counselling service, which asks whether more powerful methods exist, and how many of the students it would contact are really at risk.

💡 By the end of this chapter you will be able to

- Explain how decision trees, random forests, k-nearest neighbours, and support vector machines classify.
- Fit each of them with `tidymodels`, and compare several models fairly with cross-validation.
- Read a confusion matrix, and calculate sensitivity (recall), specificity, precision, and the F1 score.
- Choose a classification threshold from the costs of the two kinds of error and the purpose of the model.
- Explain the options for an imbalanced outcome, and what resampling does and does not change.

Note

Teaching & Practice: [Interactive Lecture Slides](#) · [Browser Playground Exercises](#)

12.1 - Data and resampling

The data, split, recipe, and folds are exactly those of Chapter 11, so every model in this chapter is trained and tested on the same students:

```
library(tidymodels)
```

—	Attaching	packages
tidymodels 1.5.0	—	

✓ broom	1.0.13	✓ recipes	1.4.0
✓ dials	1.4.4	✓ rsample	1.3.2
✓ dplyr	1.2.1	✓ tailor	0.1.0
✓ ggplot2	4.0.3	✓ tidyr	1.3.2
✓ infer	1.1.0	✓ tune	2.1.0
✓ modeldata	1.6.0	✓ workflows	1.3.0
✓ parsnip	1.6.0	✓ workflowsets	1.1.1
✓ purrr	1.2.2	✓ yardstick	1.4.0

—	Conflicts
tidymodels_conflicts() —	
* purrr::discard()	masks scales::discard()
* dplyr::filter()	masks stats::filter()
* dplyr::lag()	masks stats::lag()
* recipes::step()	masks stats::step()

```
tidymodels_prefer()

scores <- questionnaire |>
  mutate(
    stress_4 = 6 - stress_4,
    stress =
  rowMeans(pick(stress_1:stress_6), na.rm =
TRUE),
    burnout =
  rowMeans(pick(burnout_1:burnout_6), na.rm =
TRUE),
    support =
  rowMeans(pick(support_1:support_6), na.rm =
TRUE),
    satisfaction = rowMeans(pick(satisfaction_1:satisfaction_4),
na.rm = TRUE)
  ) |>
  select(student_id, stress, burnout, support,
satisfaction)

dropout_data <- students |>
  left_join(scores, join_by(student_id)) |>
  left_join(semesters |> filter(semester == 1)
|> select(-semester), join_by(student_id)) |>
  mutate(considering_dropout =
factor(considering_dropout, levels = c("Yes",
"No"))) |>
  select(-student_id, -supervisor_id, -
workshop, -workshop_sessions)

set.seed(2026)
dropout_split <- initial_split(dropout_data,
```

```
prop = 0.75, strata = considering_dropout)
dropout_train <- training(dropout_split)
dropout_test  <- testing(dropout_split)

dropout_recipe <- recipe(considering_dropout
~ ., data = dropout_train) |>
  step_impute_median(all_numeric_predictors())
  |>
  step_dummy(all_nominal_predictors()) |>
  step_normalize(all_numeric_predictors())

set.seed(2026)
dropout_folds <- vfold_cv(dropout_train, v =
10, strata = considering_dropout)
```

12.2 - Decision trees

A **decision tree** classifies by asking a series of yes-or-no questions, like a flowchart. To see how a tree chooses its questions, take twelve imaginary students:

```
tiny <- tibble(
  stress = c(1.5, 2.0, 2.2, 2.8, 3.0, 3.1,
3.4, 3.6, 3.8, 4.0, 4.2, 4.5),
  support = c(4.5, 3.0, 4.0, 2.0, 4.2, 3.8,
2.1, 4.4, 1.8, 2.5, 3.9, 1.5),
  dropout = factor(c("No", "No", "No", "No",
"No", "No",
                    "Yes", "No", "Yes", "Yes",
"No", "Yes"))
)
```

Four of the twelve consider dropping out. The tree looks for the single question that best separates the “Yes” students from the “No” students. “Is support below 2.75?” puts seven students, all “No”, on one side, and five students, four “Yes” and one “No”, on the other. The one “No” student among the low-support group has the lowest stress, so a second question, “Is stress 3.1 or more?”, separates the rest perfectly:

```
library(rpart)
tiny_tree <- rpart(dropout ~ stress + support,
data = tiny, method = "class",
control = rpart.control(minsplit
= 2, cp = 0))
tiny_tree
```

```
n= 12

node), split, n, loss, yval, (yprob)
* denotes terminal node

1) root 12 4 No (0.6666667 0.3333333)
2) support>=2.75 7 0 No (1.0000000 0.0000000)
*
3) support< 2.75 5 1 Yes (0.2000000
0.8000000)
6) stress< 3.1 1 0 No (1.0000000 0.0000000)
*
```

```
7) stress>=3.1 4 0 Yes (0.0000000 1.0000000)
*
```

Each line is a **node**: the question that led to it, the number of students, the number misclassified, and the predicted class. The final nodes, marked *, are the **leaves**. To classify a new student, you follow the questions from the top until you reach a leaf. To choose each question, the tree tries every predictor and every possible cut-off and picks the one that makes the two groups most “pure”, usually measured by the **Gini impurity** (0 when a group contains only one class).

Decision trees have attractive properties: they are easy to explain, they need no dummy variables or normalisation (a question such as “is stress 3.1 or more?” works the same on any scale), and they capture interactions automatically (stress matters here only for low-support students). For the wellbeing data, a recipe that only fills in missing values is enough:

```
tree_recipe <- recipe(considering_dropout ~ .,
  data = dropout_train) |>
  step_impute_median(all_numeric_predictors())

tree_wf <- workflow() |>
  add_recipe(tree_recipe) |>
  add_model(decision_tree(tree_depth = 3,
    min_n = 10) |> set_mode("classification"))

tree_fit <- fit(tree_wf, data = dropout_train)
```

A depth of 3 keeps the tree small enough to read. Figure 12.1 draws it with `plot()` and `text()` from `rpart`; the `rpart.plot` package draws nicer trees if you install it.

```
tree_engine <- extract_fit_engine(tree_fit)
plot(tree_engine, uniform = TRUE, margin =
  0.08)
text(tree_engine, use.n = TRUE, cex = 0.75)
```




Figure 12.1: A decision tree of depth 3 for considering dropout. At each split, students who meet the condition go left.

The first question is about support, the strongest single predictor. The numbers under each leaf are the training

students in it (Yes/No). The tree is easy to read, but a single tree has a serious weakness: it is unstable. A slightly different sample of students can produce a completely different tree, and, as Chapter 11 showed, a deep tree overfits. The next method turns this weakness into a strength.

12.3 - Random forests

A **random forest** grows hundreds of trees, each on a slightly different version of the data, and lets them vote. Two sources of randomness make the trees different. Each tree is grown on a **bootstrap sample** of the training data (Chapter 7), students drawn at random with replacement. And at each split, the tree may only choose from a random subset of the predictors; the size of this subset, `mtry`, is a hyperparameter.

Each individual tree overfits in its own way, but because the trees differ, many of their errors cancel out when the trees vote. The predicted probability of “Yes” is the share of trees that vote “Yes”. Averaging many models in this way is called an **ensemble**; a random forest is an ensemble of trees [23].

The **ranger** package fits random forests quickly. `importance = "permutation"` asks it to also measure how much each predictor matters:

```
forest_spec <- rand_forest(trees = 500) |>
  set_engine("ranger", importance =
    "permutation") |>
  set_mode("classification")

forest_wf <- workflow() |>
  add_recipe(tree_recipe) |>
  add_model(forest_spec)

set.seed(2026)
forest_fit <- fit(forest_wf, data =
  dropout_train)
```

A forest of 500 trees cannot be drawn, so it is a “black box”. **Permutation importance** opens it a little: for each predictor in turn, the values are shuffled randomly, and the drop in the model’s accuracy is recorded. Shuffling an important predictor hurts the predictions; shuffling an unimportant one does not. Figure 12.2 shows the result.

```
importance <- extract_fit_engine(forest_fit)$variable.importance

tibble(predictor = names(importance),
  importance = importance) |>
  ggplot(aes(x = importance, y =
    reorder(predictor, importance))) +
  geom_col(fill = "#2f6793") +
  labs(x = "Permutation importance (drop in
    accuracy)", y = NULL) +
  theme_minimal(base_size = 12)
```




Figure 12.2: Permutation importance of each predictor in the random forest.

The forest relies most on support, stress, and burnout, and hardly at all on the students' background, which agrees with the logistic regression of Chapter 8. Importance says which predictors the model *uses*, not what *causes* dropout; everything Chapter 8 said about confounding still applies.

12.4 - k-nearest neighbours

k-nearest neighbours (k-NN), met briefly in Chapter 11, classifies a new student by finding the k most similar students in the training data and letting them vote. “Similar” means close together when the predictors are drawn as a map. Suppose a new student in the tiny example has a stress score of 3.5 and a support score of 3.0. The **distance** to each training student is the straight-line distance on the stress-support map:

```
tiny |>
  mutate(distance = sqrt((stress - 3.5)^2 +
    (support - 3.0)^2)) |>
  arrange(distance) |>
  head(3)
```

```
# A tibble: 3 × 4
  stress support dropout distance
  <dbl>   <dbl> <fct>      <dbl>
1     4       2.5 Yes       0.707
2     3.1     3.8 No        0.894
3     3.4     2.1 Yes       0.906
```

With $k = 3$, two of the three nearest students considered dropping out, so the predicted probability of “Yes” is 2/3. Figure 12.3 shows the picture.




Figure 12.3: k-nearest neighbours with $k = 3$. The new student (cross) is classified by the three closest students (circled).

Because k-NN is based on distances, the predictors must be on the same scale. Otherwise a variable measured in large numbers, such as caffeine in milligrams, would dominate the distance, and a variable measured on a 1-5 scale would hardly count. That is why the recipe normalises every predictor.

The number of neighbours, k , is a hyperparameter. With $k = 1$, the model copies the nearest student and overfits badly (Chapter 11); with a large k , it averages over many students and becomes smoother. Tuning shows which works best for the wellbeing data:

```
knn_wf <- workflow() |>
  add_recipe(dropout_recipe) |>
  add_model(nearest_neighbor(neighbors =
    tune()) |> set_mode("classification"))

set.seed(2026)
knn_tuning <- tune_grid(knn_wf, resamples =
  dropout_folds,
                        grid = tibble(neighbors
= c(1, 5, 11, 21, 41, 81)),
                        metrics
= metric_set(roc_auc))
collect_metrics(knn_tuning) |>
select(neighbors, mean, std_err)
```

```
# A tibble: 6 × 3
  neighbors mean std_err
  <dbl> <dbl> <dbl>
1         1 0.525 0.0200
2         5 0.623 0.0222
3        11 0.673 0.0256
4        21 0.708 0.0266
5        41 0.734 0.0286
6        81 0.746 0.0303
```

The AUC keeps rising up to the largest k tried. When the best value lies at the edge of the grid, the usual advice is to extend the grid, but here the result is informative in itself:

the more neighbours k-NN averages over, the smoother its predictions become, and the better it does. The well-being data has a smooth pattern (risk rises steadily with stress and falls steadily with support), which a smooth model such as logistic regression captures directly. k-NN is at its best when the pattern is irregular and there is a lot of data.

12.5 - Support vector machines

A **support vector machine** (SVM) separates the two classes with a boundary that is as far as possible from the students on either side. Imagine drawing a line between the “Yes” and “No” points on the stress-support map: of all the lines that separate them, the SVM picks the one with the widest empty “street” around it, the **margin**. Only the students closest to the boundary, the **support vectors**, determine where it goes. When the classes overlap, as they always do in real data, some students are allowed inside the margin or on the wrong side, at a cost set by a hyperparameter.

A straight boundary is a **linear SVM**. With a **kernel**, the SVM can draw curved boundaries: the popular **radial basis function** (RBF) kernel lets the boundary bend around groups of points. Both are available in tidymodels through the **kernlab** package:

```
svm_linear_spec <- svm_linear() |>
  set_engine("kernlab") |>
  set_mode("classification")

svm_rbf_spec <- svm_rbf() |>
  set_mode("classification")
```

Like k-NN, SVMs are based on distances, so they need normalised predictors. Figure 12.4 shows how differently the four families draw their boundaries, using just stress and support so the result can be drawn as a map.




Figure 12.4: Predicted probability of considering dropout from stress and support, according to four models fitted to the training data. Darker orange means higher risk; the points are the training students.

12.6 - Comparing models fairly

The fair way to compare the models is cross-validation on the same folds, and the **workflowsets** package (part of tidymodels) does it for several models at once. `workflow_set()` combines each recipe with each model, and `workflow_map()` runs cross-validation for every combination:

```
dropout_models <- workflow_set(
  preproc = list(tree_data = tree_recipe),
  models = list(
    tree = decision_tree(tree_depth = 4,
min_n = 10) |> set_mode("classification"),
    forest = rand_forest(trees
= 500) |> set_engine("ranger") |>
set_mode("classification")
  )
) |>
  bind_rows(workflow_set(
    preproc = list(normalised =
dropout_recipe),
    models = list(
      logistic = logistic_reg(),
      knn = nearest_neighbor(neighbors
= 41) |> set_mode("classification"),
      svm_linear = svm_linear_spec,
      svm_rbf = svm_rbf_spec
    )
  ))

dropout_comparison <-
workflow_map(dropout_models, "fit_resamples",
  resamples
= dropout_folds,
  metrics =
metric_set(roc_auc), seed = 2026)

rank_results(dropout_comparison, rank_metric
= "roc_auc") |>
  select(wflow_id, mean, std_err)
```

```
# A tibble: 6 × 3
  wflow_id      mean std_err
  <chr>      <dbl> <dbl>
1 normalised_logistic 0.804 0.0303
2 normalised_svm_linear 0.787 0.0313
3 tree_data_forest    0.781 0.0380
4 normalised_svm_rbf   0.762 0.0393
5 normalised_knn       0.734 0.0286
6 tree_data_tree      0.728 0.0250
```

The trees and forest use the simple recipe, and the others the normalised one. `bind_rows()` joins the two sets of workflows into one.

The winner, with a cross-validated AUC of 0.80, is plain logistic regression. The linear SVM (0.79), which also draws a straight boundary, comes close, followed by the random forest (0.78). The standard errors, between 0.02 and 0.04, show that the smaller differences at the top could be due to chance, but none of the flexible models beats the simple one. This is not a failure of the methods:

it says something about the data. The risk of considering dropout rises smoothly with stress and falls smoothly with support, with no sharp thresholds or complicated interactions for a tree or a forest to find. On data with such structure, and with larger samples, forests and SVMs often do win. The lesson is the one from Chapter 11: compare, and let the simpler model stand unless a complex one clearly does better.

Logistic regression is therefore kept. It predicts as well as any model here, and it can be explained to the counselling service in one sentence.

12.7 - The confusion matrix

The AUC measures how well a model *ranks* students. But the counselling service needs a *decision*: contact this student or not. A model makes a decision by comparing each predicted probability with a **threshold**; by default, a student is classified “Yes” when the probability of “Yes” is above 0.5. The **confusion matrix** counts how those decisions turn out.

Take a tiny example first. A screening questionnaire is given to 100 students, 10 of whom are truly at risk. It flags 12 students, 8 of whom are among the 10 at risk:

Table 12.1: The four outcomes of a classification

	Truly at risk	Not at risk	Total
Flagged	8 (true positives)	4 (false positives)	12
Not flagged	2 (false negatives)	86 (true negatives)	88
Total	10	90	100

Every classification ends in one of four cells. A **true positive** is a student at risk who is flagged; a **false negative** is a student at risk who is missed; a **false positive** is a student flagged by mistake; a **true negative** is a student correctly left alone.

The key measures come from these four counts, and each answers a different practical question. **Sensitivity**, also called **recall**, is the share of students at risk who are flagged, here $8/10 = 80\%$: it measures how many of the students who need help are found. **Specificity** is the share of students not at risk who are left alone, $86/90 = 96\%$. **Precision** is the share of flagged students who are truly at risk, $8/12 = 67\%$: it measures how often a contact is justified. The **F1 score** balances precision and recall in a single number (their harmonic mean), here 0.73, and it is high only when both are high. **Accuracy**, the share classified correctly, is $(8 + 86)/100 = 94\%$, but, as Chapter 11 showed, it is dominated by the large “No” group.

For the study’s logistic regression, `conf_mat()` produces the confusion matrix for the test students:

```
logistic_wf <- workflow() |>
  add_recipe(dropout_recipe) |>
```

```
add_model(logistic_reg())
logistic_fit <- fit(logistic_wf, data =
dropout_train)

test_results <- augment(logistic_fit,
new_data = dropout_test)
test_results |> conf_mat(truth =
considering_dropout, estimate = .pred_class)
```

	Truth	
Prediction	Yes	No
Yes	5	7
No	18	121

The yardstick package calculates the measures. The function `metric_set()` bundles several, and because “Yes” is the first level of the outcome, they treat “Yes” as the event of interest:

```
class_metrics <- metric_set(accuracy,
sensitivity, specificity, precision, f_meas)
test_results |> class_metrics(truth =
considering_dropout, estimate = .pred_class)
```

```
# A tibble: 5 × 3
  .metric .estimator .estimate
  <chr>    <chr>      <dbl>
1 accuracy binary      0.834
2 sensitivity binary     0.217
3 specificity binary     0.945
4 precision binary     0.417
5 f_meas   binary     0.286
```

This is sobering. Of the 23 test students who considered dropping out, the model flags only 5: a sensitivity of 22%. The high accuracy and high specificity come almost entirely from correctly leaving alone the large majority who were never at risk. The AUC showed the model ranks students well, so the problem is not the model but the threshold: with only 15% of students at risk, few students ever get a predicted probability above 0.5.

12.8 - Choosing the threshold

The threshold of 0.5 is a default, not a law. Lowering it flags more students: more of those at risk are found (higher recall), at the price of more false alarms (lower precision). Which balance is right depends on what the decision costs. Here, a flagged student is offered a conversation with a counsellor: cheap, and harmless if unnecessary. Missing a student who then leaves is costly. So a threshold well below 0.5 makes sense. If the decision were costly or stigmatising, a high threshold would be needed.

The threshold is a choice made while building the model, so it must be chosen without the test set. Cross-validation can provide predictions for every training student from a model that did not see them; `save_pred = TRUE` keeps them:

```

set.seed(2026)
logistic_cv <- fit_resamples(logistic_wf,
  resamples = dropout_folds,
                                control =
control_resamples(save_pred = TRUE))
cv_predictions
<- collect_predictions(logistic_cv)

```

The **ROC curve** shows every possible threshold at once. For each threshold, it plots the sensitivity against the false-positive rate ($1 - \text{specificity}$). A useless model follows the diagonal; a good one bends towards the top-left corner, and the AUC is the area under the curve:

```

roc_points <- cv_predictions |>
  roc_curve(truth =
considering_dropout, .pred_Yes)

marked <- roc_points |>
  slice(sapply(c(0.5, 0.2), \ (t)
which.min(abs(.threshold - t))))

ggplot(roc_points, aes(x = 1 - specificity, y
= sensitivity)) +
  geom_path(colour = "#2f6793", linewidth =
1) +
  geom_abline(linetype = "dashed", colour =
"grey60") +
  geom_point(data = marked, size = 3, colour
= "#e07b39") +
  geom_text(data = marked, aes(label =
paste("threshold", round(.threshold, 1))),
  hjust = -0.15, vjust = 1.2) +
  coord_equal() +
  labs(x = "False positive rate (1 -
specificity)", y = "Sensitivity (recall)") +
  theme_minimal(base_size = 12)

```




Figure 12.5: ROC curve of the logistic regression model, from cross-validated predictions on the training students. The points mark the thresholds 0.5 and 0.2.

A table makes the trade-off concrete. For several thresholds, it shows the share of students flagged, and the resulting recall, precision, and specificity:

```
threshold_table <- map(c(0.5, 0.4, 0.3, 0.2,
0.15, 0.1), \(t) {
  cv_predictions |>
    mutate(flag = factor(if_else(.pred_Yes >=
t, "Yes", "No"), levels = c("Yes", "No"))) |>
    summarise(
      threshold = t,
      flagged   = mean(flag == "Yes"),
      recall    =
sensitivity_vec(considering_dropout, flag),
      precision =
precision_vec(considering_dropout, flag),
      specificity =
specificity_vec(considering_dropout, flag)
    )
  }) |>
  list_rbind()

threshold_table
```

```
# A tibble: 6 × 5
  threshold flagged recall precision
specificity
  <dbl>   <dbl> <dbl>    <dbl>    <dbl>
1     0.5 0.0824 0.313  0.568  0.958
2     0.4 0.127 0.403  0.474  0.921
3     0.3 0.183 0.493  0.402  0.872
4     0.2 0.263 0.627  0.356  0.801
5    0.15 0.327 0.716  0.327  0.741
6     0.1 0.419 0.746  0.266  0.639
```

The function `map()` runs the same calculation for each threshold, and `list_rbind()` stacks the results into one table. The `_vec()` versions of the yardstick functions take two vectors instead of a data frame.

12.8.1 - Costs of the two errors

Moving the threshold trades one error for the other, so the choice depends on how much each error costs. The costs are not statistical quantities; they are judgements about consequences, and they should be made explicitly. Suppose the counselling service judges that missing a student who is at risk is ten times as costly as an unnecessary conversation. The total cost of each threshold can then be calculated from the cross-validated predictions:

```
cost_miss      <- 10 # a student at risk
                  who is not contacted
cost_false_alarm <- 1 # an unnecessary
                  conversation

cost_curve <- map(seq(0.02, 0.6, by = 0.02),
  \(t) {
    cv_predictions |>
      summarise(threshold = t,
                misses     = sum(.pred_Yes <
t & considering_dropout == "Yes"),
                false_alarms = sum(.pred_Yes >=
t & considering_dropout == "No"),
                flagged     = mean(.pred_Yes
>= t))
  }) |>
  list_rbind() |>
  mutate(total_cost = cost_miss * misses +
cost_false_alarm * false_alarms)

best_cost <- cost_curve |>
slice_min(total_cost, n = 1, with_ties =
FALSE)
best_cost
```

```
# A tibble: 1 × 5
  threshold misses false_alarms flagged
total_cost
  <dbl> <int>      <int> <dbl>    <dbl>
1    0.04     5      195  0.572    245
```

```
ggplot(cost_curve, aes(x = threshold, y =
total_cost)) +
  geom_line(linewidth = 1, colour = "#2f6793")
+
  geom_point(data = best_cost, size = 3, colour
= "#e07b39") +
  labs(x = "Threshold", y = "Total cost
(arbitrary units)") +
  theme_minimal(base_size = 12)
```




Figure 12.6: Total cost of the decisions at each threshold, when missing a student at risk costs ten times as much as an unnecessary conversation. The cost is lowest at a low threshold.

With these costs, the cheapest threshold is about 0.04, at which the model flags 57% of students. A simple rule from decision theory points in the same direction: when the predicted probabilities are accurate, the cost-minimising threshold is the cost of a false alarm divided by the sum of the two costs, here $1/(1 + 10) \approx 0.09$. The two do not match exactly. The rule assumes perfectly accurate probabilities, and the cost curve is estimated from only 67 students at risk, so its minimum is not precise: a threshold of 0.10, for example, costs 26% more. Both approaches agree on what matters for the decision: with these costs, the threshold belongs far below 0.5. Different costs give different thresholds. If a false alarm carried a real cost, such as a letter that stigmatised the student, the threshold would rise; if missing a student carried a greater one, it would fall further.

The costs, and therefore the threshold, are a research decision with an ethical side. They should be set with the people who will use the model, stated in the thesis, and, where the consequences are serious, examined for their effects on different groups of students.

Practical limits matter as well. The counselling service cannot talk to every student the cost analysis would flag, and after discussing the table with Elaf it settles on a threshold of 0.2: the model then flags about 26% of students, finds about 63% of those at risk, and about 36% of the students it flags are truly at risk. Only now is the chosen threshold applied to the test students:

```
test_results <- test_results |>
  mutate(flag = factor(if_else(.pred_Yes >=
    0.2, "Yes", "No"), levels = c("Yes", "No")))

test_results |> conf_mat(truth =
  considering_dropout, estimate = flag)
```

	Truth	
Prediction	Yes	No
Yes	13	22
No	10	106

```
test_results |> class_metrics(truth =
  considering_dropout, estimate = flag)
```

```
# A tibble: 5 × 3
  .metric .estimator .estimate
  <chr>    <chr>      <dbl>
1 accuracy binary      0.788
2 sensitivity binary      0.565
3 specificity binary      0.828
4 precision binary      0.371
5 f_meas  binary      0.448
```

On new students, the model now finds 13 of the 23 at risk instead of 5, while flagging 35 students in all. About one flagged student in 3 is truly at risk. The test numbers are a little lower than the cross-validated ones, as expected with only 23 at-risk students in the test set. Accuracy has fallen, and it does not matter: accuracy was never the goal.

! Report the threshold

Sensitivity, specificity, precision, and F1 all depend on the threshold, so a paper that reports them must say which threshold was used, and why. The AUC does not depend on the threshold, which is why it is the standard measure for comparing models, and the threshold-dependent measures are the ones for describing how a model will be used.

12.9 - Imbalanced outcomes

When one outcome is rare, as considering dropout is, models tend to predict the common one. Besides moving the threshold, there are two other common remedies. **Resampling** changes the training data so that the classes are balanced: **upsampling** repeats rare-class students, **downsampling** drops common-class students, and **SMOTE** creates new, artificial rare-class students between existing ones. The **themis** package adds all three to a recipe. **Class weights** instead make mistakes on the rare class count more when the model is fitted, and some model engines support them directly.

Upsampling with themis looks like this; `over_ratio = 1` repeats “Yes” students until there are as many as “No” students:

```
library(themis)

upsample_recipe <- dropout_recipe |>
  step_upsample(considering_dropout,
    over_ratio = 1)

upsample_wf <- workflow() |>
  add_recipe(upsample_recipe) |>
  add_model(logistic_reg())

set.seed(2026)
upsample_cv <- fit_resamples(upsample_wf,
```



```
resamples = dropout_folds,
                                metrics =
metric_set(roc_auc, sensitivity, precision))
collect_metrics(upsample_cv)      |>
select(.metric, mean)
```

```
# A tibble: 3 × 2
  .metric      mean
  <chr>      <dbl>
1 precision  0.314
2 roc_auc    0.789
3 sensitivity 0.671
```

Resampling steps are applied only to the training data; when the model predicts for new students, the step is skipped automatically, so the test set keeps its real balance.

With the usual 0.5 threshold, the upsampled model has a sensitivity of about 67%, far higher than before. But its AUC, 0.79, has not improved. Upsampling did not make the model better at telling students apart; it pushed all the predicted probabilities upwards, which has much the same effect as lowering the threshold. That is often how it works out, and for a model that produces probabilities, choosing the threshold directly is simpler and keeps the probabilities meaningful (after upsampling, a “probability” of 0.6 no longer means a 60% chance). Resampling is most useful for models that do not produce good probabilities, and when the rare class is very rare.

 Does the model work equally well for everyone?

Among the test students, the model with a threshold of 0.2 finds 7 of the 13 at-risk women and 6 of the 10 at-risk men, similar shares; but with only 23 at-risk students in the test set, any difference would be hard to detect. Before a model like this is used, its sensitivity and precision should be checked for each group that matters (gender, faculty, full- and part-time study) on as much data as possible, as Chapter 11 recommended. A threshold that works on average can still miss one group of students much more often than another.

In your field: ecology

Ecologists classify species from measurements. The `penguins` data in the `modeldata` package (installed with `tidymodels`) records the bill length, bill depth, flipper length, and body mass of 344 penguins of three species, measured on islands near Palmer Station in Antarctica [24]. A random forest classifies the species:

```
data(penguins, package = "modeldata")
penguins <- penguins |>
  drop_na(bill_length_mm, bill_depth_mm,
    flipper_length_mm, body_mass_g)

set.seed(1)
penguin_split <- initial_split(penguins,
  strata = species)

penguin_wf <- workflow() |>
  add_formula(species ~ bill_length_mm
+ bill_depth_mm + flipper_length_mm +
  body_mass_g) |>
  add_model(rand_forest(trees =
500) |> set_engine("ranger") |>
  set_mode("classification"))

penguin_fit <- last_fit(penguin_wf,
  penguin_split)
collect_predictions(penguin_fit) |>
conf_mat(truth = species, estimate
= .pred_class)
```

	Truth		
Prediction	Adelie	Chinstrap	Gentoo
Adelie	38	0	0
Chinstrap	0	17	0
Gentoo	0	0	31

With three classes, the confusion matrix has a row and a column for each species. Correct classifications lie on the diagonal, and any mistake would appear off it, showing exactly which species are confused; here, every penguin in the test set is classified correctly. The function `add_formula()` replaces a recipe when no preparation is needed. Unlike the wellbeing data, the classes here are well separated by the measurements, and flexible models work very well.

12.10 - Chapter review

12.10.1 - Summary

- A decision tree classifies by a series of yes-or-no questions. Trees are easy to explain and need no normalisation, but a single tree is unstable and overfits easily.
- A random forest averages hundreds of trees grown on bootstrap samples with random subsets of predictors. Permutation importance shows which predictors it relies on.

- k-nearest neighbours classifies by the vote of the k most similar training cases; it needs normalised predictors, and k must be tuned.
- A support vector machine separates the classes with the widest possible margin; kernels allow curved boundaries.
- Compare models with cross-validation on the same folds (`workflow_set()`). On the wellbeing data, with its smooth pattern, logistic regression predicts as well as any flexible model.
- The confusion matrix counts true and false positives and negatives. Sensitivity (recall) is the share of positives found; precision is the share of flagged cases that are positive; F1 balances the two.
- The threshold turns probabilities into decisions. Choose it from the costs of the two errors and the practical limits of its use, using cross-validated predictions, and report it.
- For imbalanced outcomes, moving the threshold, resampling, and class weights all shift the balance between recall and precision; resampling rarely improves the AUC.

12.10.2 - Key terms

Decision tree, node, leaf, Gini impurity, random forest, bootstrap sample, ensemble, `mtry`, permutation importance, k-nearest neighbours, distance, support vector machine, margin, support vector, kernel, radial basis function, workflow set, threshold, confusion matrix, true positive, false positive, true negative, false negative, sensitivity, recall, specificity, precision, F1 score, ROC curve, misclassification cost, upsampling, downsampling, SMOTE, class weights.

12.11 - Exercises

The [playground](#) has these and more, with hints and solutions.

1. Draw the decision tree for the wellbeing data with `tree_depth = 2`. Name the predictors it uses, and explain the tree in plain words, as you would to a counsellor.
2. Tune the random forest's `mtry` over `c(2, 5, 10)` with cross-validation, and compare the best value with the default.
3. Extend the k-NN grid to `neighbors = c(81, 161, 301)`, describe what happens to the AUC, and explain what happens to a k-NN model as k approaches the number of training students.
4. Using the tiny example's confusion matrix (Table 12.1), calculate the F1 score yourself from precision and recall ($F_1 = 2 \times \text{precision} \times \text{recall} / (\text{precision} + \text{recall})$).
5. Suppose the counselling service can only contact 10% of new students. Using `threshold_table`, recommend a threshold and report the recall and precision it would give.

6. Replace `step_upsample()` with `step_downsample()`, and compare the sensitivity and AUC with upsampling.
7. Repeat the cost analysis with a cost of 3 for a missed student instead of 10. Report the new cheapest threshold, and compare it with the decision-theory rule.

12.12 - Further reading

- *Tidy Modeling with R* [22] covers workflow sets, model comparison, and class imbalance in depth.
- *An Introduction to Statistical Learning* [18] explains trees, random forests, and support vector machines, with clear illustrations of how each draws its boundaries.
- Breiman's original paper on random forests [23] is still readable and explains why averaging many trees works.

13. Predictive Regression

Research data often contains many possible predictors for rather few cases. A questionnaire with dozens of items, several rounds of measurement, and a background survey can easily supply 50 predictors for a sample of 60 or 100 people. Ordinary regression estimates a coefficient for every predictor, and with that much freedom it can fit its own data almost perfectly while predicting new cases badly: the overfitting of Chapter 11, now in its most common form. The problem is made worse when predictors are correlated, as questionnaire items usually are.

This chapter introduces two families of methods that make many predictors usable. **Regularised regression** (ridge, lasso, and elastic net) keeps the familiar linear model but shrinks its coefficients towards zero, trading a little fit on the training data for much better predictions. **Boosting** builds a flexible model from many small decision trees, each correcting the errors of the ones before. The chapter also introduces the measures used to judge numeric predictions. The outcome is now a number, so this is a **regression** problem in the machine learning sense of Chapter 11. In the study, it answers the tenth research question (RQ10): the university's academic advisers meet every student at the end of the first year, and would like to know from year-one information which students are heading for a low final GPA, so that extra support can be planned.

💡 By the end of this chapter you will be able to

- Measure the accuracy of numeric predictions with RMSE, MAE, and R^2 , and compare them with a baseline.
- Explain why ordinary regression overfits when there are many predictors.
- Fit and tune ridge, lasso, and elastic net regression, and read which predictors the lasso keeps.
- Explain how boosting builds a model from many small trees, and tune an XGBoost model.
- Choose between models with cross-validation, and report the final model's accuracy on the test set.

Note

Teaching & Practice: [Interactive Lecture Slides](#) · [Browser Playground Exercises](#)

13.1 - Data for predicting final GPA

The advisers meet students at the end of year one, so the model may use only what is known by then: the students' background, their questionnaire answers, and their records from semesters 1 and 2. The outcome is GPA in semester 4.

This time, every questionnaire item is used as a separate predictor, rather than the four scale scores, to show how the methods cope with many related predictors. The semester records are in long format (one row per student per semester); for prediction, they are needed in wide format, one row per student, with a column for each measure in each semester. `pivot_wider()` from Chapter 3 does this:

```
library(tidymodels)
```

—	Attaching	packages
tidymodels 1.5.0 —		

✓ broom	1.0.13	✓ recipes	1.4.0
✓ dials	1.4.4	✓ rsample	1.3.2
✓ dplyr	1.2.1	✓ tailor	0.1.0
✓ ggplot2	4.0.3	✓ tidyr	1.3.2
✓ infer	1.1.0	✓ tune	2.1.0
✓ modeldata	1.6.0	✓ workflows	1.3.0
✓ parsnip	1.6.0	✓ workflowsets	1.1.1
✓ purrr	1.2.2	✓ yardstick	1.4.0

—	Conflicts
tidymodels_conflicts() —	
* purrr::discard()	masks scales::discard()
* dplyr::filter()	masks stats::filter()
* dplyr::lag()	masks stats::lag()
* recipes::step()	masks stats::step()

```
tidymodels_prefer()

year_one <- semesters |>
  filter(semester <= 2) |>
  pivot_wider(id_cols = student_id, names_from = semester,
              values_from = gpa:wellbeing,
              names_glue = "{.value}_s{semester}")

final_gpa <- semesters |>
  filter(semester == 4, !is.na(gpa)) |>
  select(student_id, final_gpa = gpa)

gpa_data <- students |>
  left_join(questionnaire,
            join_by(student_id)) |>
  left_join(year_one, join_by(student_id)) |>
  inner_join(final_gpa, join_by(student_id)) |>
  select(-student_id, -supervisor_id, -considering_dropout)
```

```
dim(gpa_data)
```

```
[1] 563  48
```

The argument `names_glue` builds the new column names, such as `gpa_s1` and `sleep_hours_s2`, and `inner_join()` keeps only students who have a final GPA. That is an important limitation: the 37 students who left the programme have no final GPA, so the model can only predict the final GPA of students who stay. (Chapter 6 showed that the leavers were not a random group, so the model should not be used to judge them.)

The split, folds, and recipe follow Chapter 11. With a numeric outcome, `strata` splits the outcome into quartiles and samples within each, so that both sets cover the full range of GPAs. The recipe adds `step_impute_mode()` to fill in any missing categories:

```
set.seed(2026)
gpa_split <- initial_split(gpa_data, prop =
  0.75, strata = final_gpa)
gpa_train <- training(gpa_split)
gpa_test  <- testing(gpa_split)

set.seed(2026)
gpa_folds <- vfold_cv(gpa_train, v = 10,
  strata = final_gpa)

gpa_recipe <- recipe(final_gpa ~ ., data =
  gpa_train) |>
  step_impute_median(all_numeric_predictors())
  |>
  step_impute_mode(all_nominal_predictors())
  |>
  step_dummy(all_nominal_predictors()) |>
  step_normalize(all_numeric_predictors())

gpa_recipe |> prep() |> bake(new_data = NULL)
|> ncol()
```

```
[1] 52
```

After the dummy variables are created, there are 51 predictors, plus the outcome.

13.2 - Measuring prediction error

The quality of a numeric prediction is judged by its errors. Five students whose final GPA a model has predicted show how:

```
tiny <- tibble(
  actual   = c(3.2, 2.8, 3.6, 3.0, 2.5),
  predicted = c(3.0, 2.9, 3.3, 3.1, 2.9)
)
tiny |> mutate(error = actual - predicted)
```

```
# A tibble: 5 × 3
  actual predicted error
```

	<dbl>	<dbl>	<dbl>
1	3.2	3	0.200
2	2.8	2.9	-0.100
3	3.6	3.3	0.300
4	3	3.1	-0.100
5	2.5	2.9	-0.4

Each **error** (or residual) is the actual value minus the prediction, and three measures summarise them. The **MAE** (mean absolute error) is the average size of the errors, ignoring their sign; here, $(0.2 + 0.1 + 0.3 + 0.1 + 0.4) / 5 = 0.22$ GPA points. The **RMSE** (root mean squared error) squares the errors, averages them, and takes the square root. Squaring gives large errors more weight, so the RMSE is always at least as large as the MAE, and much larger when there are a few big misses. R^2 , in yardstick, is the squared correlation between the predictions and the actual values: as in Chapter 8, the share of the variation in the outcome that the predictions capture, from 0 to 1. The yardstick package calculates all three:

```
regression_metrics <- metric_set(rmse, mae,
  rsq)
tiny |> regression_metrics(truth = actual,
  estimate = predicted)
```

```
# A tibble: 3 × 3
  .metric .estimator .estimate
  <chr>   <chr>       <dbl>
1 rmse   standard      0.249
2 mae    standard      0.220
3 rsq    standard      0.785
```

MAE and RMSE are in the units of the outcome, here GPA points, which makes them easy to explain: “the model’s predictions are typically off by about 0.2 GPA points”. Lower is better. R^2 has no units; higher is better.

A number on its own means little, so every model should be compared with a **baseline**: the simplest possible prediction. For a numeric outcome, the baseline predicts the same value, the training mean, for everyone, which is what `parsnip`’s `null_model()` does:

```
null_wf <- workflow(gpa_recipe,
  null_model(mode = "regression"))

set.seed(2026)
null_cv <- fit_resamples(null_wf, resamples =
  gpa_folds, metrics = metric_set(rmse, mae))
collect_metrics(null_cv)
```

```
# A tibble: 2 × 6
  .metric .estimator mean n
std_err .config
  <chr>   <chr>       <dbl> <int> <dbl>
<chr>
1 mae    standard    0.256    10 0.00738
pre0_mod0_post0
2 rmse    standard    0.323    10 0.0104
pre0_mod0_post0
```


Predicting the average GPA for everyone gives an RMSE of 0.32. Any useful model must do clearly better. (The R^2 of the baseline cannot be calculated, because its predictions do not vary, so it is left out here.)

13.3 - Ordinary regression with many predictors

The obvious model is the multiple regression of Chapter 8, using every predictor:

```
lm_wf <- workflow(gpa_recipe, linear_reg())

set.seed(2026)
lm_cv <- fit_resamples(lm_wf, resamples =
  gpa_folds, metrics = regression_metrics)
collect_metrics(lm_cv)
```

```
# A tibble: 3 × 6
  .metric .estimator   mean     n
std_err .config
  <chr>   <chr>       <dbl> <int> <dbl>
<chr>
1 mae     standard    0.172     10 0.00615
pre0_mod0_post0
2 rmse     standard    0.218     10 0.0105
pre0_mod0_post0
3 rsq      standard    0.564     10 0.0420
pre0_mod0_post0
```

A cross-validated RMSE of 0.218, much better than the baseline. With 422 training students and 51 predictors, ordinary regression works reasonably well. But watch what happens with fewer students. Many theses have samples of 50 or 60, so here is the same model fitted to 60 randomly chosen training students, and tested on the test set:

```
set.seed(3)
small_train <- gpa_train |> slice_sample(n =
  60)

lm_small <- fit(lm_wf, data = small_train)
augment(lm_small, new_data = small_train) |
  > regression_metrics(final_gpa, .pred)
```

```
# A tibble: 3 × 3
  .metric .estimator .estimate
  <chr>   <chr>       <dbl>
1 rmse     standard    0.115
2 mae      standard    0.0918
3 rsq      standard    0.876
```

```
augment(lm_small, new_data = gpa_test) |>
  regression_metrics(final_gpa, .pred)
```

```
# A tibble: 3 × 3
  .metric .estimator .estimate
  <chr>   <chr>       <dbl>
1 rmse     standard    0.526
```

2 mae	standard	0.430
3 rsq	standard	0.0456

On its own 60 students, the model looks excellent: an RMSE of 0.11. On new students, its RMSE is 0.53, *worse than predicting the average GPA for everyone*. With 51 coefficients estimated from 60 students, the model has enough freedom to fit the noise in its sample, and the coefficients that fit the noise do not carry over. In the extreme, with as many predictors as students, a regression can fit its data perfectly and predict nothing.

The problem is made worse by **multicollinearity**: many of the predictors are strongly correlated (six stress items, two semesters of GPA), and regression struggles to divide the credit between correlated predictors. Their coefficients become large, unstable, and sometimes of opposite signs, cancelling each other out on the training data but not on new data.

13.4 - Regularised regression

Regularisation tackles this directly: it fits the regression while penalising large coefficients. Ordinary regression chooses the coefficients that minimise the sum of squared errors. Regularised regression minimises

sum of squared errors + $\lambda \times \text{size of the coefficients}$

The **penalty** λ (lambda) sets how much large coefficients cost. With $\lambda = 0$, this is ordinary regression; as λ grows, the coefficients are pulled, or **shrunk**, towards zero. A little shrinkage costs almost nothing in fit to the training data, but makes the model much more stable on new data. Because the penalty depends on the size of the coefficients, the predictors must be on the same scale, which the recipe's `step_normalize()` ensures. There are three versions, which differ in how “size” is measured. **Ridge regression** uses the sum of the *squared* coefficients: it shrinks all coefficients towards zero and shares the credit among correlated predictors, but keeps every predictor in the model. The **lasso** (least absolute shrinkage and selection operator) uses the sum of the *absolute* coefficients [25]. It shrinks too, but it also sets some coefficients to exactly zero, removing those predictors, so it also *selects* predictors and leaves a simpler model. The **elastic net** mixes the two, with a **mixture** parameter that runs from 0 (pure ridge) to 1 (pure lasso).

In tidymodels, all three are `linear_reg()` with the **glmnet** engine, with `penalty` for λ and `mixture` for the mix. The 60 students are now refitted with a lasso:

```
lasso_small <- workflow(gpa_recipe,
  linear_reg(penalty
    = 0.03, mixture = 1) |> set_engine("glmnet"))
|>
  fit(data = small_train)

augment(lasso_small, new_data = gpa_test) |
  > regression_metrics(final_gpa, .pred)
```

```
# A tibble: 3 × 3
  .metric .estimator .estimate
  <chr>   <chr>         <dbl>
1 rmse    standard        0.204
2 mae     standard        0.168
3 rsq     standard        0.616
```

From the same 60 students, the lasso's RMSE on new students is 0.20, against 0.53 for ordinary regression. The penalty kept the model from chasing noise.

The difference lies in the coefficients. Figure 13.1 compares, predictor by predictor, the coefficients that ordinary regression and the lasso estimated from the same 60 students:

```
coefficient_comparison <- tidy(lm_small) |>
  select(term, ordinary = estimate) |>
  inner_join(tidy(lasso_small) |> select(term,
    lasso = estimate), join_by(term)) |>
  filter(term != "(Intercept)")
```

Attaching package: 'Matrix'

The following objects are masked from 'package:tidyr':

expand, pack, unpack

Loaded glmnet 4.1-10

```
ggplot(coefficient_comparison, aes(x =
  ordinary, y = lasso)) +
  geom_hline(yintercept = 0, colour = "grey70")
+
  geom_vline(xintercept = 0, colour = "grey70")
+
  geom_point(size = 2, alpha = 0.7, colour =
    "#2f6793") +
  labs(x = "Ordinary regression coefficient",
    y = "Lasso coefficient") +
  theme_minimal(base_size = 12)
```




Figure 13.1: Coefficients estimated from the same 60 students by ordinary regression (horizontal axis) and by the lasso (vertical axis), one point per predictor. Ordinary regression spreads large coefficients over many predictors; the lasso sets most of them to zero and keeps a few, smaller ones.

Ordinary regression gives every predictor a coefficient, many of them large and in both directions: with 60 students, it cannot tell real effects from the chance patterns of this sample, so it fits both. Some of its coefficients make no sense. Semester 1 GPA, one of the best predictors of final GPA, receives a negative coefficient (-0.03), and caffeine in semester 2 receives a larger one (0.17) than semester 2 GPA (0.16). The lasso gives both GPA measures positive coefficients (0.08 and 0.12) and almost nothing to caffeine. The lasso keeps only 12 of the 51 predictors and shrinks even those. This is the sense in which shrinkage trades fit for stability: coefficients pulled towards zero cannot chase the noise of one sample, so they carry over better to the next.

13.4.1 - Choosing the penalty

The penalty and mixture are hyperparameters, so they are tuned with cross-validation, now on the full training set. The grid tries 20 penalties, spread evenly on a logarithmic scale from 0.0001 to 1, for ridge, an even elastic net, and lasso:

```
glmnet_wf <- workflow(gpa_recipe,
                      linear_reg(penalty = tune(),
                                mixture = tune()) |> set_engine("glmnet"))

glmnet_grid <- expand_grid(penalty =
  10^seq(-4, 0, length.out = 20),
                        mixture = c(0,
                                    0.5, 1))

set.seed(2026)
glmnet_tuning <- tune_grid(glmnet_wf,
  resamples = gpa_folds, grid = glmnet_grid,
                        metrics
  = metric_set(rmse))
```

```
autoplot(glmnet_tuning)
theme_minimal(base_size = 12)
```




Figure 13.2: Cross-validated RMSE of ridge (mixture 0), elastic net (0.5), and lasso (1) regression for a range of penalties.

Figure 13.2 shows the typical pattern. With a tiny penalty, all three behave like ordinary regression. As the penalty grows, the RMSE improves, reaches a minimum, and then rises steeply once the penalty is so large that it shrinks away real effects too. At the far right, the lasso has shrunk every coefficient to zero and predicts the mean for everyone: its RMSE equals the baseline's.

```
show_best(glmnet_tuning, metric = "rmse", n = 3)
```

```
# A tibble: 3 × 8
  penalty mixture .metric .estimator mean
n std_err .config
  <dbl>   <dbl> <chr>   <chr>      <dbl>
<int>   <dbl> <chr>
1  0.0127     1 rmse    standard  0.204
10 0.00871 pre0_mod33_post0
2  0.0207     1 rmse    standard  0.205
10 0.00811 pre0_mod36_post0
3  0.0336    0.5 rmse    standard  0.205
10 0.00813 pre0_mod38_post0
```

The best combination is a mixture of 1 with a penalty of 0.013, but the top few are nearly identical. The function `select_best()` picks the winner, and `finalize_workflow()` plugs it in:

```
best_penalty <- select_best(glmnet_tuning,
metric = "rmse")
lasso_fit <- glmnet_wf |>
  finalize_workflow(best_penalty) |>
  fit(data = gpa_train)
```

13.4.2 - Predictors kept by the lasso

The function `tidy()` lists the coefficients, most of which are exactly zero:

```
lasso_coefs <- tidy(lasso_fit)
lasso_coefs |>
  filter(estimate != 0) |>
  arrange(desc(abs(estimate)))
```

```
# A tibble: 11 × 3
  term                estimate penalty
  <chr>              <dbl>   <dbl>
1 (Intercept)        3.10     0.0127
2 gpa_s2             0.145     0.0127
3 gpa_s1             0.106     0.0127
4 satisfaction_2      0.00986   0.0127
5 stress_3           -0.00890   0.0127
6 study_mode_Part.time 0.00812   0.0127
7 burnout_1          -0.00688   0.0127
8 burnout_4          -0.00622   0.0127
9 study_hours_s2      0.00557   0.0127
10 burnout_6          -0.00210   0.0127
11 support_3         -0.000169  0.0127
```

Of 51 predictors, the lasso keeps 10. Because the predictors are normalised, the coefficients are comparable: the change in predicted final GPA for a one-standard-deviation difference in the predictor. Year-one GPA dominates. A student whose semester 2 GPA is one standard deviation above average is predicted to finish about 0.15 GPA points higher; the other predictors add small adjustments.

Figure 13.3 shows how the lasso arrives at this. As the penalty decreases from left to right, predictors enter the model one by one, and the two GPA measures enter first.




Figure 13.3: Lasso coefficient paths. Each line is one predictor's coefficient as the penalty decreases from left to right. The dashed line marks the chosen penalty.

⚠ Selected is not the same as causal

It is tempting to read the lasso's choice as a list of what *matters* for grades. It is not. Sleep affects GPA in the wellbeing data (Chapter 8), yet the lasso drops the sleep variables, because their effect is already reflected in year-one GPA: once the model knows a student's past grades, knowing their sleep adds little to the *prediction*. Among correlated predictors, the lasso tends to keep one and drop the others, and which one it keeps can change from sample to sample. Use the lasso to predict, and use the methods of Chapters 8 and 10 to explain.

13.5 - Boosting

The second family builds a model in a completely different way. **Boosting** starts with a very simple prediction, the average, and then adds small decision trees one at a time, each fitted to the *errors* of the model so far. The first tree learns the biggest pattern in the errors; the next tree learns what the first missed; and so on. Each tree's contribution is scaled down by a **learning rate** (for example 0.1), so the model improves in small steps and no single tree dominates. Hundreds of trees, each weak on its own, add up to a strong model.

Figure 13.4 shows the idea with a single predictor, semester 2 GPA, and trees with a single split. After one tree, the prediction is a single step; after ten, a staircase; after a hundred, a curve that follows the data.




Figure 13.4: Boosting with one predictor. Predicted final GPA from semester 2 GPA after 1, 10, and 100 trees, each with a single split.

Boosting is closely related to the random forests of Chapter 12, which also combine many trees. The difference is that a forest grows its trees *independently* and averages them, while boosting grows them *in sequence*, each correcting the last. **XGBoost** (“extreme gradient boosting”) is a fast, popular implementation [26], and one of the most successful methods for tables of data. Its main hyperparameters are the number of trees, the learning rate, and the depth of each tree. A small learning rate needs more trees; deeper trees capture interactions between predictors. Here, 500 trees are used, and the learning rate and depth are tuned:

```
xgb_wf <- workflow(gpa_recipe,
  boost_tree(trees = 500,
  learn_rate = tune(), tree_depth = tune()) |>
  set_engine("xgboost") |>
  set_mode("regression"))

set.seed(2026)
xgb_tuning <- tune_grid(xgb_wf, resamples =
gpa_folds,
  grid = expand_grid(learn_rate
= c(0.01, 0.03, 0.1),
  tree_depth
= c(1, 2, 4)),
  metrics = metric_set(rmse))
show_best(xgb_tuning, metric = "rmse", n = 3)
```

```
# A tibble: 3 × 8
  tree_depth learn_rate .metric .estimator
mean      n std_err .config
<dbl>     <dbl> <dbl> <chr>    <chr>    <dbl>
<int>    <dbl> <chr>
1         1    0.03 rmse    standard 0.212
10 0.00968 pre0_mod2_post0
2         2    0.01 rmse    standard 0.213
10 0.00928 pre0_mod4_post0
3         1    0.01 rmse    standard 0.213
10 0.00834 pre0_mod1_post0
```

The best settings use trees of depth 1, and the best RMSE is 0.212. Trees of depth 1 use one predictor each, so a boosted model built from them adds up separate effects of each predictor, with no interactions. That the shallowest trees work best is another sign that the wellbeing data has no strong interactions for the trees to find.

13.6 - Comparing the models

All the models were evaluated with the same folds, so their cross-validated RMSEs can be compared directly:

```
bind_rows(
  collect_metrics(null_cv) |> mutate(model =
"Baseline (mean)",
  collect_metrics(lm_cv) |> mutate(model =
"Linear regression"),
  show_best(glmnet_tuning |>
filter_parameters(mixture == 0), metric =
"rmse", n = 1) |>
  mutate(model = "Ridge"),
  show_best(glmnet_tuning |>
filter_parameters(mixture == 0.5), metric =
"rmse", n = 1) |>
  mutate(model = "Elastic net"),
  show_best(glmnet_tuning |>
filter_parameters(mixture == 1), metric =
"rmse", n = 1) |>
  mutate(model = "Lasso"),
  show_best(xgb_tuning, metric = "rmse", n =
1) |> mutate(model = "XGBoost")
) |>
  filter(.metric == "rmse") |>
```



```
select(model, rmse = mean, std_err) |>
  arrange(rmse)
```

```
# A tibble: 6 × 3
  model      rmse std_err
  <chr>    <dbl>   <dbl>
1 Lasso      0.204 0.00871
2 Elastic net 0.205 0.00813
3 XGBoost     0.212 0.00968
4 Ridge      0.213 0.00938
5 Linear regression 0.218 0.0105
6 Baseline (mean) 0.323 0.0104
```

The function `filter_parameters()` keeps only the tuning results with a given mixture, so the best ridge, elastic net, and lasso can be picked out separately.

The regularised models come out on top, followed closely by XGBoost, and all of them improve on ordinary regression. The differences between the best models are smaller than their standard errors. The lasso is chosen: it predicts as well as any, and it uses only 10 predictors, which makes it the easiest to explain and to use.

13.7 - The final test

As always, the test set is used once, at the end. The function `last_fit()` fits the chosen model on the full training set and evaluates it on the test students:

```
final_lasso <- glmnet_wf |>
  finalize_workflow(best_penalty) |>
  last_fit(gpa_split, metrics =
    regression_metrics)
collect_metrics(final_lasso)
```

```
# A tibble: 3 × 4
  .metric .estimator .estimate .config
  <chr>   <chr>      <dbl> <chr>
1 rmse    standard     0.192 pre0_mod0_post0
2 mae     standard     0.161 pre0_mod0_post0
3 rsq     standard     0.648 pre0_mod0_post0
```

On new students, the lasso's predictions are off by 0.16 GPA points on average (MAE), with an RMSE of 0.19, and they capture 65% of the variation in final GPA. Figure 13.5 compares the predictions with the actual final GPAs.

```
collect_predictions(final_lasso) |>
  ggplot(aes(x = .pred, y = final_gpa)) +
  geom_abline(linetype = "dashed", colour =
    "grey50") +
  geom_point(alpha = 0.6, colour = "#2f6793") +
  coord_obs_pred() +
  labs(x = "Predicted final GPA", y = "Actual
  final GPA") +
  theme_minimal(base_size = 12)
```




Figure 13.5: Predicted and actual final GPA for the test students. Points on the dashed line are predicted perfectly.

The function `coord_obs_pred()` from `tune` gives both axes the same range, so the diagonal is at 45 degrees. The predictions are less spread out than the actual values: the model predicts the most extreme students closer to the average than they turn out to be. This is normal for any model that cannot predict perfectly, and it means that the model is most reliable for students in the middle of the range.

A last comparison shows how much all the year-one measures add to past grades. A plain regression on the two year-one GPAs alone gives:

```
gpa_only_wf <- workflow(recipe(final_gpa ~
  gpa_s1 + gpa_s2, data = gpa_train),
  linear_reg())
gpa_only_fit <- last_fit(gpa_only_wf,
  gpa_split, metrics = regression_metrics)
collect_metrics(gpa_only_fit)
```

```
# A tibble: 3 × 4
  .metric .estimator .estimate .config
  <chr>    <chr>         <dbl> <chr>
1 rmse    standard         0.193 pre0_mod0_post0
2 mae     standard         0.163 pre0_mod0_post0
3 rsq     standard         0.641 pre0_mod0_post0
```

Its test RMSE, 0.193, is practically the same as the lasso's 0.192. For the advisers, this is a useful, if humbling, result: past grades are by far the best predictor of future grades, and the questionnaire and semester records add little on top. That does not make sleep, stress, or support irrelevant to grades (Chapter 8 showed that they matter), but their influence is already visible in the year-one grades.

In your field: engineering

Engineers predict the strength of concrete from its recipe. The `concrete` data in the **modeldata** package records the compressive strength of 1030 concrete samples, with the amounts of cement, water, and six other ingredients, and the age of the sample in days [27]. Strength rises with age, but less and less, and the ingredients interact: exactly the kind of structure where boosting shines.

```
data(concrete, package = "modeldata")
set.seed(1)
concrete_split <- initial_split(concrete)

lm_concrete
<- workflow(compressive_strength ~ .,
linear_reg())
xgb_concrete
<- workflow(compressive_strength ~ .,
              boost_tree(trees
= 500, learn_rate = 0.05, tree_depth = 4)
|>
              set_engine("xgboost")
|>
set_mode("regression"))

lm_concrete_fit <- last_fit(lm_concrete,
concrete_split)
xgb_concrete_fit
<- last_fit(xgb_concrete, concrete_split)
collect_metrics(lm_concrete_fit)
```

```
# A tibble: 2 × 4
  .metric .estimator .estimate .config
  <chr>   <chr>       <dbl> <chr>
1 rmse      standard      10.6
pre0_mod0_post0
2 rsq      standard       0.603
pre0_mod0_post0
```

```
collect_metrics(xgb_concrete_fit)
```

```
# A tibble: 2 × 4
  .metric .estimator .estimate .config
  <chr>   <chr>       <dbl> <chr>
1 rmse      standard       4.54
pre0_mod0_post0
2 rsq      standard       0.926
pre0_mod0_post0
```

As the code shows, `workflow()` also accepts a formula in place of a recipe. Boosting cuts the RMSE from 10.6 to 4.5 megapascals. The lesson of this chapter and the last is the same from both sides: flexible models win when the data has curves and interactions, and simple models are just as good when it does not. Only a fair comparison can tell which case you are in.

13.8 - Common misconceptions

Predictive regression invites a few characteristic misreadings.

- **“More predictors always help.”** Each extra coefficient is another chance to fit noise; with few cases, adding predictors can make predictions worse.
- **“The predictors the lasso keeps are the ones that matter.”** The lasso keeps the predictors that help *prediction*, and among correlated predictors it keeps one more or less arbitrarily.
- **“A low RMSE means every prediction is accurate.”** The RMSE is a typical error; predictions for students at the extremes are pulled towards the average and can be far off.
- **“Boosting is always better than regression.”** Boosting wins when the data has curves and interactions; on smooth data, regularised regression predicts as well and is easier to explain.

13.9 - Chapter review

13.9.1 - Summary

- For numeric predictions, the MAE and RMSE measure the typical error in the units of the outcome (lower is better), and R^2 the share of variation captured. Always compare with a baseline that predicts the mean.
- With many predictors relative to the number of cases, and correlated predictors, ordinary regression overfits: excellent on the training data, poor on new data.
- Regularised regression penalises large coefficients. Ridge shrinks all coefficients; the lasso also sets some to zero, selecting predictors; the elastic net mixes the two. The penalty and mixture are tuned by cross-validation, and the predictors must be normalised.
- The predictors a lasso selects are useful for prediction, not evidence of causes.
- Boosting adds small trees one at a time, each fitted to the errors of the model so far. XGBoost’s main hyperparameters are the number of trees, the learning rate, and the tree depth.
- On the wellbeing data, regularised regression and boosting predict final GPA about equally well, and year-one GPA carries most of the information. On data with curves and interactions, boosting can be far better.

13.9.2 - Key terms

Regression (prediction), error, residual, MAE, RMSE, R^2 , baseline model, overfitting, multicollinearity, regularisation, penalty (lambda), shrinkage, ridge regression, lasso, elastic net, mixture, coefficient path, variable selection, boosting, weak learner, learning rate, XGBoost, tree depth.

13.10 - Exercises

The [playground](#) has these and more, with hints and solutions.

1. Calculate the MAE and RMSE of the tiny example by hand (or with basic R arithmetic), and check them against yardstick. Change one prediction so that it is off by 1.0 GPA point, and explain which measure changes more, and why.
2. Repeat the 60-student experiment with 150 students, and report the gap between ordinary regression and the lasso.
3. Fit a ridge regression (mixture 0) with the best ridge penalty, compare its coefficients with the lasso's, and count how many are exactly zero.
4. Replace the 22 questionnaire items with the four scale scores (as in Chapter 11), and describe how the cross-validated RMSE of the lasso changes.
5. Tune XGBoost with `trees = c(100, 500, 1000)` and a learning rate of 0.03, and describe what happens to the RMSE as trees are added.
6. Write two sentences for the academic advisers explaining what an RMSE of 0.2 means for an individual student's predicted GPA.

13.11 - Further reading

- *Tidy Modeling with R* [22] covers regularised regression and boosted trees in tidymodels, including more efficient tuning methods.
- *An Introduction to Statistical Learning* [18] explains ridge regression, the lasso, and boosting clearly, with the ideas behind why shrinkage helps.
- Tibshirani's paper introducing the lasso [25] is a classic of modern statistics.

14. Advanced Clustering

Every clustering method carries a definition of what a group is, usually without saying so. One definition says that a group is a set of cases close to a common centre. Another says that a group is a population with its own distribution, so that cases between groups can belong partly to each. A third says that a group is a dense region of data, separated from other groups by sparse regions, which leaves room for cases that belong to no group at all. The definition chosen decides what the method can find, and a clustering is only as meaningful as the definition behind it.

Chapter 9 used k-means, which takes the first definition. It divided the students of the wellbeing study into four lifestyle profiles, but left open questions that any examiner might ask: how sure one can be about which profile a student belongs to, whether four is really the right number, and whether some students fit no profile. This chapter addresses them with two methods that relax the assumptions of k-means. **Gaussian mixture models** give each student a *probability* of belonging to each profile, and choose the number of profiles with a statistical criterion. **DBSCAN** defines clusters as dense regions of data, and labels students in sparse regions as **noise**: people who fit no group. The chapter ends with the question behind all clustering: how to judge whether a clustering is any good.

💡 By the end of this chapter you will be able to

- Explain three definitions of a group, and the limitations of k-means: hard assignments, round clusters, and no room for outliers.
- Fit a Gaussian mixture model with `mclust`, choose the number of clusters with BIC, and interpret membership probabilities.
- Find dense clusters and noise points with DBSCAN, and choose its settings.
- Evaluate a clustering with the silhouette, agreement between methods (the adjusted Rand index), and interpretability.

Note

Teaching & Practice: [Interactive Lecture Slides](#) · [Browser Playground Exercises](#)

14.1 - The profile data

The data is exactly that of Chapter 9: each student's average sleep, study hours, caffeine, and exercise across the semesters, and their stress, support, and satisfaction scores, scaled to z-scores.

```
library(dplyr)
```

```
Attaching package: 'dplyr'
```

```
The following objects are masked from
'package:stats':
```

```
filter, lag
```

```
The following objects are masked from
'package:base':
```

```
intersect, setdiff, setequal, union
```

```
library(ggplot2)

q <- questionnaire
q$stress_4 <- 6 - q$stress_4

scores <- q |>
  mutate(
    stress =
    rowMeans(pick(stress_1:stress_6), na.rm =
    TRUE),
    support =
    rowMeans(pick(support_1:support_6), na.rm =
    TRUE),
    satisfaction = rowMeans(pick(satisfaction_1:satisfaction_4),
    na.rm = TRUE)
  ) |>
  select(student_id, stress, support,
  satisfaction)

profiles <- semesters |>
  summarise(across(c(sleep_hours,
  study_hours, caffeine_mg, exercise_days),
    ~ mean(.x, na.rm = TRUE)),
    .by = student_id) |>
  left_join(scores, join_by(student_id)) |>
  na.omit()

profile_data <- scale(profiles |> select(-
  student_id))

set.seed(123)
kmeans_clusters <- kmeans(profile_data,
  centers = 4, nstart = 25)
```

The last two lines repeat Chapter 9's k-means solution, for comparison.

14.2 - What counts as a group

k-means is simple and fast, but its definition of a group brings three strong assumptions. It assumes that **every case belongs to exactly one cluster**, with complete certainty, so a student halfway between two profiles is assigned to one of them as firmly as a student at the centre. It assumes that **clusters are round and similar in size**: because each case goes to the nearest centre, k-means draws straight boundaries halfway between centres, and elongated or unequal clusters are cut up wrongly. And it assumes that **every case belongs to some cluster**, with no way to say that a student fits nowhere. Real groups of people rarely satisfy these assumptions.

A simulation shows how much the definition matters. The code below creates data with a known structure: a small round group, a long thin group, and 15 scattered points that belong to neither. It then clusters the data in three ways, with k-means (groups around centres), a Gaussian mixture model (groups as distributions), and DBSCAN (groups as dense regions), the two methods this chapter introduces:

```
library(mclust)
```

```
Package 'mclust' version 6.1.2
Type 'citation("mclust")' for citing this R
package in publications.
```

```
Attaching package: 'mclust'
```

```
The following object is masked from
'package:dplyr':
```

```
count
```

```
library(dbSCAN)
```

```
Attaching package: 'dbSCAN'
```

```
The following object is masked from
'package:stats':
```

```
as.dendrogram
```

```
set.seed(8)
shapes <- bind_rows(
  tibble(x = rnorm(100, 0, 0.4), y = rnorm(100,
0, 0.4), truth = "Round group"),
  tibble(x = rnorm(100, 3, 2), y = rnorm(100,
1.4, 0.2), truth = "Long group"),
  tibble(x = runif(15, -2, 7), y = runif(15,
-2, 4), truth = "Scattered")
)
xy <- shapes |> select(x, y)
```

```

shape_clusters <- bind_rows(
  shapes |> mutate(method = "k-means", cluster
= factor(kmeans(xy, 2, nstart = 25)$cluster)),
  shapes |> mutate(method = "Mixture model",
cluster = factor(Mclust(xy, G = 2, verbose
= FALSE)$classification)),
  shapes |> mutate(method = "DBSCAN", cluster
= factor(dbscan(xy, eps = 0.5, minPts =
5)$cluster))
) |>
  mutate(method = factor(method, levels = c("k-
means", "Mixture model", "DBSCAN")),
        cluster = forcats::fct_recode(cluster,
noise = "0"))

ggplot(shape_clusters, aes(x, y, colour =
cluster)) +
  geom_point(size = 1.2) +
  facet_wrap(~ method) +
  scale_colour_manual(values = c("1" =
"#2f6793", "2" = "#e07b39", "3" = "#36a269",
"4" = "#8b68b5",
noise = "grey65")) +
  coord_equal() +
  theme_minimal(base_size = 11)

```




Figure 14.1: The same simulated data clustered by k-means, a Gaussian mixture model, and DBSCAN. The data contains a round group, a long thin group, and scattered points that belong to neither; each method's definition of a group decides what it finds.

The three methods see different groups in the same data. Their agreement with the two real groups can be measured with the adjusted Rand index, introduced at the end of the chapter, where 1 means perfect agreement and 0 means no more than chance. k-means scores 0.53: looking for round groups around two centres, it draws a straight boundary through the long group and gives its left end to the round group. The mixture model scores 0.98, because it allows the long group its elongated shape. DBSCAN scores 0.00, for a different reason: the end of the long group touches the round group, so the two form one continuous dense region, and by DBSCAN's definition that is one group. DBSCAN is, however, the only method that can say that points belong to no group: it labels 6 of the 15 scattered points as noise, while the other two methods must place every one of them in a group. No method is right in general. Each is right for data whose

groups match its definition, which is why the definition should be chosen deliberately, and stated.

14.3 - Gaussian mixture models

A **Gaussian mixture model** (GMM) assumes that the data comes from a mix of several groups, each with its own normal (Gaussian) distribution, and estimates each group's centre, spread, and size from the data.

14.3.1 - A one-variable example

The idea is easiest to see with one variable. R's `faithful` data records the waiting time, in minutes, between 272 eruptions of the Old Faithful geyser in Yellowstone National Park. The histogram in Figure 14.2 has two humps: short waits and long waits. The **mclust** package fits a mixture model with `Mclust()`:

```
geyser <- Mclust(faithful$waiting)
summary(geyser, parameters = TRUE)
```

```
-----
Gaussian finite mixture model fitted by EM
algorithm
-----
```

```
Mclust E (univariate, equal variance) model
with 2 components:
```

log-likelihood	n	df	BIC	ICL
-1034.002	272	4	-2090.427	-2099.576

```
Clustering table:
```

1	2
99	173

```
Mixing probabilities:
```

1	2
0.3609461	0.6390539

```
Means:
```

1	2
54.61675	80.09239

```
Variances:
```

1	2
34.44093	34.44093

The function tried mixtures of one to nine groups and chose 2. The output describes them: the **mixing probabilities** are the sizes of the groups (about 36% and 64% of eruptions), the **means** are their centres (about 55 and 80 minutes), and the **variances** their spreads. Figure 14.2 draws the two fitted normal curves over the data.




Figure 14.2: Waiting times between eruptions of Old Faithful, with the two normal distributions of the fitted mixture model.

The key difference from k-means appears between the groups. A waiting time of 67 minutes lies between them. Instead of forcing it into one, the model gives the probability that it came from each:

```
predict(geyser, newdata = c(50, 67, 85))$z |
> round(2)
```

```
      1      2
[1,] 1.00 0.00
[2,] 0.42 0.58
[3,] 0.00 1.00
```

A wait of 50 minutes almost certainly belongs to the short group, and 85 minutes to the long group, but 67 minutes is uncertain. These **membership probabilities** (also called soft assignments) are the main advantage of mixture models: they say not just which group a case is in, but how sure that is.

14.3.2 - Choosing the number of clusters with BIC

With more variables, each group is described by a centre and a **covariance matrix**, which sets its shape: round or elongated, and tilted in any direction. Groups can be allowed to differ in size (volume), shape, and orientation, or forced to be the same. `mclust` names each combination with three letters, such as `EEE` (all equal) or `VVV` (all variable), and fits them all for each number of groups.

To choose among them, it uses the **Bayesian information criterion** (BIC). BIC rewards a model for fitting the data well and penalises it for every parameter it needs, so a more complex model must earn its extra parameters. In `mclust`, *higher* BIC is better. For the profile data:

```
set.seed(123)
profile_gmm <- Mclust(profile_data, G = 1:8)
summary(profile_gmm)
```

```
-----
Gaussian finite mixture model fitted by EM
algorithm
```

```
Mclust EVE (ellipsoidal, equal volume and
orientation) model with 3 components:
```

```
log-likelihood  n df      BIC      ICL
-5128.189 599 63 -10659.28 -10752.69
```

```
Clustering table:
```

```
 1  2  3
231 174 194
```

```
library(factoextra)
```

```
Welcome! Want to learn more? See two
factoextra-related books at https://goo.gl/ve
3WBa
```

```
fviz_mclust(profile_gmm, what = "BIC")
```

```
Warning: `gather_()` was deprecated in tidyr
1.2.0.
```

```
i Please use `gather()` instead.
```

```
i The deprecated feature was likely used in
the factoextra package.
```

```
Please report the issue at <https://github.
com/kassambara/factoextra/issues>.
```

```
Warning: Using `size` aesthetic for lines was
deprecated in ggplot2 3.4.0.
```

```
i Please use `linewidth` instead.
```

```
i The deprecated feature was likely used in
the ggpubr package.
```

```
Please report the issue at <https://github.
com/kassambara/ggpubr/issues>.
```




Figure 14.3: BIC for Gaussian mixture models with one to eight clusters. Each line is one covariance structure; higher is better.

BIC chooses 3 clusters with the EVE structure (ellipsoidal clusters of equal size and orientation, but different shapes). The best four-cluster model is about 8.7 BIC points behind. A common rule of thumb reads a BIC difference above 6 as strong evidence and above 10 as

very strong, so the data favours three profiles, although four remain a reasonable alternative (exercise 1 explores them).

14.3.3 - Describing the profiles

As in Chapter 9, a cluster means something only once it is described:

```
profiles <- profiles |>
  mutate(gmm_cluster =
    profile_gmm$classification)

profiles |>
  summarise(students =
    n(), across(sleep_hours:satisfaction, ~
    round(mean(.x), 1)),
    .by = gmm_cluster) |>
  arrange(gmm_cluster)
```

	gmm_cluster	students	sleep_hours	study_hours	caffeine_mg	exercise_days	stress	support	satisfaction
1	1	231	6.7	17.9	151.7	2.2	3.3	2.8	2.7
2	2	174	7.2	25.0	126.2	3.7	2.6	3.6	3.9
3	3	194	5.5	41.7	284.1	1.3	3.6	3.3	3.1

The three profiles are clear. One is **balanced**: the most sleep and exercise, the least stress, and the highest support and satisfaction. One is **overloaded**: long study weeks, short sleep, a lot of caffeine, little exercise, and the most stress. The third, the largest, combines few study hours with low support and low satisfaction: students who seem **disengaged or isolated**. The cluster numbers are arbitrary labels. A cross-table compares the solution with k-means:

```
table(gmm = profile_gmm$classification,
      kmeans = kmeans_clusters$cluster)
```

	kmeans			
gmm	1	2	3	4
1	5	191	34	1
2	0	3	166	5
3	80	3	0	111

The mixture model's three profiles correspond closely to Chapter 9's four k-means clusters: the balanced and disengaged groups largely match, and the two k-means "overloaded" clusters, one of them with very high caffeine, are joined into one. The mixture model, with its more flexible cluster shapes, did not need to split the overloaded students in two.

14.3.4 - Certainty of assignment

The matrix `profile_gmm$z` holds each student's membership probabilities, one column per cluster, and `profile_gmm$uncertainty` is 1 minus the largest of them.

```
head(round(profile_gmm$z, 2))
```

```
  [,1] [,2] [,3]
1 0.63 0.37  0
2 0.99 0.00  0
3 0.14 0.86  0
4 0.00 0.00  1
5 0.10 0.90  0
6 1.00 0.00  0
```

```
certainty <- apply(profile_gmm$z, 1, max)
sum(certainty < 0.8)
```

```
[1] 79
```

The call `apply(..., 1, max)` takes the maximum of each row. Most students belong clearly to one profile, but 79 have less than an 80% probability for their most likely profile. Figure 14.4 shows where they are.

```
fviz_mclust(profile_gmm, what = "uncertainty")
```

```
Warning: The `<scale>` argument of `guides()`
cannot be `FALSE`. Use "none" instead as
of ggplot2 3.3.4.
i The deprecated feature was likely used in
the factoextra package.
Please report the issue at <https://github.
com/kassambara/factoextra/issues>.
```




Figure 14.4: The three mixture model clusters on the first two principal components. Larger points are students whose profile is more uncertain.

The uncertain students lie where the profiles meet. For a thesis, this is honest and useful: instead of claiming that every student has one profile, it can report the share

of students who clearly fit a profile, and treat the rest as mixtures. A later analysis could use the probabilities themselves, for example as weights, rather than the hard labels.

14.4 - Clusters as dense regions

DBSCAN (density-based spatial clustering of applications with noise) takes a different view. A cluster is a region where cases are packed closely together, separated from other clusters by sparser regions. Cases in sparse regions belong to no cluster; they are **noise**. DBSCAN does not need the number of clusters in advance, can find clusters of any shape, and can say “this case fits nowhere”.

It needs two settings: `eps` (epsilon), the radius of the neighbourhood around each case, and `minPts`, the number of cases a neighbourhood must contain for the case to be at the heart of a cluster. A case with at least `minPts` cases within `eps` is a **core point**. Core points within `eps` of each other are joined into the same cluster, and cases near a core point are added to its cluster as **border points**. Everything else is noise.

14.4.1 - A small example

Twelve points show how it works: two tight groups of five, and two isolated points.

```
tiny <- tibble(
  x = c(1.0, 1.2, 1.1, 0.9, 1.3, 4.0, 4.2,
        3.9, 4.1, 4.3, 2.5, 5.5),
  y = c(1.0, 1.1, 1.3, 1.2, 0.9, 3.0, 3.2,
        3.1, 2.8, 3.0, 4.5, 0.5)
)

tiny_db <- dbscan(tiny, eps = 0.5, minPts = 3)
tiny_db$cluster
```

```
[1] 1 1 1 1 1 2 2 2 2 2 0 0
```

The two tight groups become clusters 1 and 2, and the two isolated points are labelled 0: noise. k-means with two clusters would have had to put the isolated points into one of the groups.

14.4.2 - Choosing `eps`

The result depends heavily on `eps`. A common guide is the **k-nearest-neighbour distance plot**: for each case, the distance to its k -th nearest neighbour, sorted from smallest to largest. Most cases have close neighbours; the curve bends sharply upwards where the isolated cases begin, and that bend suggests a value for `eps`. With `minPts` set to 8, a common choice for data with seven variables (about the number of variables plus one or more), the plot uses $k = 7$:

```
kNNdistplot(profile_data, k = 7)
abline(h = 2, lty = "dashed")
```





Figure 14.5: Distance from each student to their seventh-nearest neighbour, sorted. The dashed line marks $\text{eps} = 2$.

The curve bends at a distance of about 2, so $\text{eps} = 2$ is used:

```
profile_db <- dbscan(profile_data, eps = 2,
  minPts = 8)
profile_db
```

```
DBSCAN clustering for 599 objects.
Parameters: eps = 2, minPts = 8
Using euclidean distances and borderpoints =
TRUE
The clustering contains 1 cluster(s) and 13
noise points.
```

```
0 1
13 586
```

```
Available fields: cluster, eps, minPts,
metric, borderPoints
```

DBSCAN finds 1 cluster and 13 noise points. At first this looks like a failure: no profiles at all. But it is an informative result. DBSCAN looks for dense regions separated by gaps, and the students form one continuous cloud, in which the profiles found by k-means and the mixture model overlap without gaps between them. (Smaller values of eps break the cloud into fragments and label hundreds of students as noise; try it in the exercises.) The profiles are real differences in *where* students sit in the cloud, not separate islands.

14.4.3 - Students who fit no profile

The noise points answer the last of the open questions: whether some students fit no profile. Their values are:

```
profiles |>
  filter(profile_db$cluster == 0) |>
  select(sleep_hours:satisfaction) |>
  round(1)
```

	sleep_hours	study_hours	caffeine_mg
exercise_days	stress	support	
1	3.6	60.5	772.5

0.5	3.2	4.0		
2		4.6	45.8	592.5
2.5	2.8	2.7		
3		9.7	2.2	26.2
5.2	1.7	2.7		
4		3.7	61.5	842.5
0.0	3.5	3.7		
5		9.4	5.5	23.8
6.2	2.7	3.7		
6		3.8	63.0	780.0
0.5	2.0	3.5		
7		3.7	63.2	742.5
0.0	1.8	4.3		
8		9.6	4.2	21.2
5.0	2.7	3.5		
9		9.8	5.2	18.8
5.8	2.6	2.2		
10		9.5	2.0	20.0
6.2	2.5	2.5		
11		4.0	65.5	735.0
0.5	3.2	3.3		
12		4.4	30.2	703.8
0.8	3.2	4.6		
13		3.8	35.8	641.2
1.0	3.0	1.8		
satisfaction				
1		4.2		
2		4.0		
3		4.2		
4		3.8		
5		3.7		
6		2.8		
7		4.7		
8		4.2		
9		3.5		
10		3.5		
11		3.2		
12		3.0		
13		2.8		

Two kinds of unusual student stand out. Among the noise points, 5 students study around 60 hours a week, sleep about 4 hours a night or less, drink over 700 mg of caffeine a day (about seven cups of coffee), and hardly exercise: more extreme than anyone in the overloaded profile. 5 others are the opposite: close to 10 hours of sleep, very few study hours, very little caffeine, and exercise almost every day. The remaining few are extreme versions of the overloaded profile. In Chapter 6, extreme values were checked one variable at a time; DBSCAN finds students who are unusual in their *combination* of values.

These students should not be deleted: they are real students, and the extreme workers are exactly the students a wellbeing service would want to know about. They are reported separately, as students who do not fit the profiles, and the profiles are checked for whether they change when these students are left out.

14.5 - Evaluating a clustering

In supervised learning, a model is judged against the true answers. In clustering there are usually no true answers, so evaluation rests on several kinds of evidence, none decisive on its own.

Internal measures judge how compact and well separated the clusters are, using the data alone. The **silhouette** of Chapter 9 is the most common, and BIC plays this role for mixture models. The `silhouette()` function from the `cluster` package calculates it for any clustering:

```
library(cluster)
profile_distances <- dist(profile_data)

mean(silhouette(kmeans_clusters$cluster,
profile_distances)[, "sil_width"])
```

```
[1] 0.2135627
```

```
mean(silhouette(profile_gmm$classification,
profile_distances)[, "sil_width"])
```

```
[1] 0.2272005
```

Both averages are low (the silhouette runs from -1 to 1), which confirms what DBSCAN suggested: the profiles overlap, and no clustering of this data will be crisp.

Agreement between methods asks whether different reasonable methods find similar groups. The **adjusted Rand index** (ARI) measures the agreement between two clusterings. It counts the pairs of cases that both clusterings put together or both put apart, and adjusts for the agreement expected by chance: 1 means identical groupings (whatever the labels), and 0 means no more agreement than chance. A tiny example shows that the labels themselves do not matter:

```
first <- c(1, 1, 1, 2, 2, 2)
second <- c("B", "B", "B", "A", "A", "A")
third <- c(1, 1, 2, 2, 3, 3)
adjustedRandIndex(first, second)
```

```
[1] 1
```

```
adjustedRandIndex(first, third)
```

```
[1] 0.2424242
```

The first two clusterings group the six people identically, so their ARI is 1, even though the labels differ. The third splits them differently, so its ARI is much lower. For the two methods applied to the profile data:

```
adjustedRandIndex(kmeans_clusters$cluster,
profile_gmm$classification)
```

```
[1] 0.6534227
```

An ARI of 0.65: substantial agreement, especially given that one solution has four clusters and the other three.

External validation compares clusters with known categories, also using the ARI. It is only possible when such categories exist, for example when clustering flowers of known species to test a method. For real profiles there is no answer key, which is exactly why they are being sought.

Stability asks whether the clusters survive small changes: a different random start, a bootstrap sample of the students, or leaving out the unusual students. Clusters that appear only under one setting should not be reported.

Interpretability and usefulness matter most in the end. The clusters should make sense in the light of theory, and they should differ on outcomes they were not built from, such as GPA or considering dropout. The second point can be checked directly:

```
profiles |>
  left_join(students |> select(student_id,
    considering_dropout), join_by(student_id)) |>
  summarise(students = n(),
    considering_dropout =
    round(mean(considering_dropout == "Yes"), 2),
    .by = gmm_cluster) |>
  arrange(gmm_cluster)
```

	gmm_cluster	students	considering_dropout
1	1	231	0.21
2	2	174	0.04
3	3	194	0.17

The profiles differ clearly in how many students consider dropping out, although dropout played no part in finding them. That is good evidence that they capture something real about students' situations.

! Report the choices, not just the clusters

A clustering depends on choices: which variables, how they are scaled, which method, how many clusters, and settings such as `eps`. Different reasonable choices can give different answers. In a thesis, report each choice and why it was made, show the evidence for the number of clusters (BIC, silhouette), say how clearly students fit (for example, the share with a membership probability above 0.8), and describe how robust the profiles are to other reasonable choices.

In your field: earth sciences

DBSCAN was designed for spatial data, where clusters can have any shape. R's `quakes` data records the location of 1,000 earthquakes near Fiji since 1964. Plotted on a map, they follow the long, curved lines of two ocean trenches: clusters that are neither round nor of equal size.

```
quake_db <- dbscan(quakes[, c("long",
"lat")], eps = 1.5, minPts = 10)
table(quake_db$cluster)
```

```
0  1  2
28 785 187
```

```
quakes |>
  mutate(cluster =
factor(quake_db$cluster)) |>
  ggplot(aes(x = long, y = lat, colour =
cluster, shape = cluster == "0")) +
  geom_point(alpha = 0.7) +
  scale_colour_manual(values = c("0" =
"grey60", "1" = "#2f6793", "2" =
"#e07b39")) +
  scale_shape_manual(values = c(`FALSE` =
16, `TRUE` = 4), guide = "none") +
  coord_quickmap() +
  labs(x = "Longitude", y = "Latitude",
colour = "Cluster") +
  theme_minimal(base_size = 12)
```




Figure 14.6: Earthquakes near Fiji, clustered by DBSCAN. Grey crosses are noise.

DBSCAN follows the two curved trenches and leaves the scattered earthquakes between and beyond them as noise. Here, unlike in the wellbeing data, there really are dense regions separated by empty space, which is exactly when DBSCAN works best. The function `coord_quickmap()` gives longitude and latitude the right proportions for a map.

14.6 - Common misconceptions

Advanced clustering methods answer some of the weaknesses of k-means, but not the underlying question of whether the groups are real.

- **“A more sophisticated method finds the true groups.”** Each method finds groups that match its own definition of a group; none can find groups that the data does not contain.
- **“Every case belongs to some group.”** Cases between groups are better described by membership probabilities, and cases far from all groups by the label “noise”.
- **“The best BIC proves the number of groups.”** BIC compares models; a model with more groups may be almost as good, and the choice should be reported with its evidence.
- **“DBSCAN failed because it found one cluster.”** Finding one continuous cloud is itself a finding: the groups found by other methods are regions of that cloud, not separate islands.

14.7 - Chapter review

14.7.1 - Summary

- Every clustering method has a definition of a group: around a centre (k-means), a distribution (mixture models), or a dense region (DBSCAN). The definition decides what the method can find.
- k-means assigns every case to exactly one round cluster, with no room for uncertainty or outliers.
- A Gaussian mixture model describes the data as a mix of normal distributions of different sizes and shapes. It gives each case a probability of belonging to each cluster, and BIC chooses the number of clusters and their shape (higher is better in mclust).
- On the wellbeing data, the mixture model finds three profiles (balanced, overloaded, and disengaged) that closely match Chapter 9’s k-means clusters, and shows which students fit no profile clearly.
- DBSCAN defines clusters as dense regions and labels cases in sparse regions as noise. It needs `eps` and `minPts`; the k-nearest-neighbour distance plot helps choose `eps`. It works best when clusters are separated by gaps.
- On the wellbeing data, DBSCAN finds one continuous cloud of students and flags a few extreme students who fit no profile.
- Evaluate clusterings with internal measures (silhouette, BIC), agreement between methods (adjusted Rand index), stability, and above all interpretability and relevance to outcomes. Report every choice.

14.7.2 - Key terms

Gaussian mixture model, mixture component, mixing probability, membership probability, soft assignment, covariance matrix, Bayesian information criterion (BIC), uncertainty, DBSCAN, density, `eps`, `minPts`, core point, border point, noise, k-nearest-neighbour distance plot,

internal validation, silhouette, adjusted Rand index, external validation, stability.

14.8 - Exercises

The [playground](#) has these and more, with hints and solutions.

1. Fit `Mclust(profile_data, G = 4)` and describe the four profiles. Identify which of the three-cluster profiles has been split, and compare the new solution with Chapter 9's k-means.
2. Count the students with a membership probability above 0.95 for their profile, and report the share of the sample.
3. Run DBSCAN on the profile data with `eps` of 1.2, 1.6, and 2.4, describe how the number of clusters and noise points change, and explain why a small `eps` labels so many students as noise.
4. Leave out the DBSCAN noise students, fit the mixture model again, and describe whether the profiles change.
5. Use `adjustedRandIndex()` to compare the three-cluster mixture model with hierarchical clustering (Ward's method, cut into three clusters) from Chapter 9.
6. Compare the three profiles on final GPA, and say whether the pattern is what their descriptions would lead you to expect.
7. In the simulation at the start of the chapter, move the long group upwards (a mean of 3 for `y`) so that a gap separates it from the round group. Describe how each method's result changes, and explain why DBSCAN now behaves differently.

14.9 - Further reading

- *An Introduction to Statistical Learning* [18] introduces k-means and hierarchical clustering; its chapter on unsupervised learning is a good companion to Chapters 9 and 14.
- The `mclust` paper [28] explains Gaussian mixture models, the covariance structures, and BIC, with many examples in R.
- The paper describing the `dbscan` package [29] explains DBSCAN, its settings, and its relatives, such as HDBSCAN and OPTICS.

15. Neural Networks

Neural networks are the method behind most of what is called artificial intelligence today: image recognition, speech recognition, translation, and the large language models of Chapter 18. Their reputation leads many researchers to assume that a neural network must predict better than a traditional model on any data. The assumption deserves to be tested rather than accepted, and testing it requires understanding what a neural network is. The answer is less mysterious than the reputation suggests: a neural network is built from units that each perform a logistic regression, and it learns by repeatedly adjusting its weights to reduce its errors.

This chapter builds a neural network from a single “neuron”, shows by hand what training does, fits a network to the dropout data with the same tidymodels workflow as Chapters 11 and 12, and compares it fairly with the simpler models. It ends with deep learning: what it is, when it helps, and where to learn more. In the story, a member of Elaf’s research group asks why she does not simply use AI, since a neural network would surely do better than logistic regression, and her supervisor suggests that she find out, carefully.

💡 By the end of this chapter you will be able to

- Explain how a neuron combines its inputs, and why a single neuron with a sigmoid activation is a logistic regression.
- Describe a neural network’s layers, weights, and activation functions, and how it is trained.
- Fit and tune a neural network with tidymodels (`mlp()`), using weight decay to prevent overfitting.
- Compare a neural network fairly with simpler models, and judge when one is worth using.
- Explain what deep learning is, and when it is (and is not) the right tool.

Note

Teaching & Practice: [Interactive Lecture Slides](#) · [Browser Playground Exercises](#)

15.1 - A single neuron

A **neuron** (or unit) in a neural network does three things. It multiplies each input by a **weight**, adds up the results together with a constant called the **bias**, and passes the total through an **activation function**, which turns it into the neuron’s output.

Take a neuron with two inputs, a student's stress and support scores, with weights 1.2 and -0.8 and a bias of -4 . For a student with stress 4 and support 2, the total is

$$-4 + 1.2 \times 4 - 0.8 \times 2 = -0.8. \quad (15.1)$$

The activation function here is the **sigmoid** (also called the logistic function), which squeezes any number into the range 0 to 1:

$$\text{sigmoid}(z) = \frac{1}{1 + e^{-z}}. \quad (15.2)$$

In R:

```
sigmoid <- function(z) 1 / (1 + exp(-z))

neuron <- function(stress, support) {
  sigmoid(-4 + 1.2 * stress - 0.8 * support)
}

neuron(stress = 4, support = 2)
```

```
[1] 0.3100255
```

```
neuron(stress = 2, support = 4)
```

```
[1] 0.008162571
```

The high-stress, low-support student gets an output of 0.31; the low-stress, high-support student, 0.008. If the output is read as the probability of considering dropout, this neuron is exactly a **logistic regression** (Chapter 8): the bias is the intercept, the weights are the coefficients, and the sigmoid turns the log odds into a probability. Everything a single neuron can do, Chapter 8 has already done.

15.2 - From neurons to networks

The power of neural networks comes from connecting many neurons in **layers**. In the most common design, the **multilayer perceptron**, the inputs feed into a **hidden layer** of neurons, each with its own weights; the outputs of the hidden neurons then feed into an **output layer**, which produces the prediction. Figure 15.1 shows a network with three inputs and three hidden neurons.




Figure 15.1: A neural network with one hidden layer. Every arrow carries a weight; each hidden neuron and the output neuron add up their weighted inputs and apply an activation function.

Each hidden neuron learns its own combination of the inputs: one might respond to high stress, another to the combination of low support and short sleep. The output neuron then combines these. Because each neuron applies a curved activation function, the network as a whole can represent curved relationships and interactions that a single logistic regression cannot. With enough hidden neurons, a network can approximate almost any relationship between inputs and output. That flexibility is its strength, and, as you will see, its danger.

The activation function matters. Figure 15.2 shows the two most common: the sigmoid, used in this chapter, and the **ReLU** (rectified linear unit), which simply replaces negative totals with zero and is the standard choice in deep networks.




Figure 15.2: Two activation functions: the sigmoid squeezes any input into the range 0 to 1; the ReLU sets negative inputs to 0 and passes positive inputs unchanged.

15.2.1 - How a network learns

A network starts with small random weights, so its first predictions are useless. **Training** adjusts the weights step by step to reduce the prediction error on the training data. At each step, an algorithm works out, for every weight, whether a small increase or decrease would reduce the error (the **gradient**), and moves all the weights a little in the helpful direction. This is **gradient descent**. One pass through the training data is an **epoch**, and training usually runs for hundreds of epochs.

The idea is clearest when one step is worked by hand. Take a single neuron, which is a logistic regression, and six students, with their stress measured from the group's average:

```
six <- data.frame(stress = c(2.0, 2.5, 3.0,
3.5, 4.0, 4.5) - 3.25,
                  dropout = c(0, 0, 1,
0, 1, 1))
```

The error to be reduced is the **log loss**, the measure of fit that logistic regression uses: it is small when the neuron gives high probabilities to the students who did consider dropping out and low probabilities to those who did not. Training starts with both weights at zero, so every student receives a probability of 0.5:

```
log_loss <- function(b0, b1) {
  p <- sigmoid(b0 + b1 * six$stress)
  -mean(six$dropout * log(p) + (1 -
six$dropout) * log(1 - p))
}

b0 <- 0
b1 <- 0
log_loss(b0, b1)
```

```
[1] 0.6931472
```

For this neuron, the gradient has a simple form: for each weight, the average of the prediction errors (predicted probability minus the actual outcome), each multiplied by the input that the weight belongs to (1 for the bias):

```
p <- sigmoid(b0 + b1 * six$stress)
gradient_b0 <- mean(p - six$dropout)
gradient_b1 <- mean((p - six$dropout) *
six$stress)
c(gradient_b0, gradient_b1)
```

```
[1] 0.0000000 -0.2916667
```

The gradient for the stress weight is negative, which means that increasing the weight would reduce the error: the students who considered dropping out have higher stress, and the neuron does not yet know it. One step of gradient descent moves each weight against its gradient, by an amount set by the **learning rate**:

```
learning_rate <- 1
b0 <- b0 - learning_rate * gradient_b0
```

```
b1 <- b1 - learning_rate * gradient_b1
c(b0 = b0, b1 = b1, loss = log_loss(b0, b1))
```

```
      b0      b1      loss
0.0000000 0.2916667 0.6157970
```

After one step, the stress weight has become positive and the loss has fallen from 0.693 to 0.616. Training is nothing more than this step, repeated. The loop below repeats it 2,000 times and records the loss:

```
b0 <- 0
b1 <- 0
loss_history <- numeric(2000)
for (step in 1:2000) {
  p <- sigmoid(b0 + b1 * six$stress)
  b0 <- b0 - learning_rate * mean(p -
    six$dropout)
  b1 <- b1 - learning_rate * mean((p -
    six$dropout) * six$stress)
  loss_history[step] <- log_loss(b0, b1)
}
rbind(gradient_descent = c(b0, b1),
      glm = coef(glm(dropout ~ stress, data =
        six, family = binomial)))
```

```
              (Intercept)  stress
gradient_descent 6.111907e-17 2.428055
glm              -1.790641e-16 2.428055
```

```
ggplot(data.frame(step = 1:2000, loss =
  loss_history), aes(x = step, y = loss)) +
  geom_line(colour = "#2f6793", linewidth =
    1) +
  scale_x_log10() +
  labs(x = "Training step", y = "Log loss") +
  theme_minimal(base_size = 12)
```




Figure 15.3: Log loss of a single neuron during 2,000 steps of gradient descent (logarithmic horizontal axis). The loss falls quickly at first and then levels off as the weights approach their best values.

After 2,000 small steps, the weights found by gradient descent are the coefficients that R's `glm()` calculates for the same logistic regression. A neural network with thousands of weights is trained in exactly this way. The

gradients are harder to write down, and a method called backpropagation calculates them efficiently, but each step still moves every weight a little in the direction that reduces the error.

Two consequences follow. First, because training starts from random weights, two runs can give slightly different networks, so `set.seed()` matters. Second, a network with many weights can keep reducing its training error long after it has learned the real pattern, by memorising the noise. The usual defence is **weight decay**: a penalty on large weights, exactly like the ridge penalty of Chapter 13. In `tidymodels` it is called `penalty`.

15.3 - A neural network for the dropout question

The data, split, recipe, and folds are those of Chapters 11 and 12. Neural networks need normalised inputs, like the penalised regressions of Chapter 13, and the recipe already provides them.

```
library(tidymodels)
```

Attaching		packages
tidymodels 1.5.0 —		
✓ broom	1.0.13	✓ rsample 1.3.2
✓ dials	1.4.4	✓ tailor 0.1.0
✓ dplyr	1.2.1	✓ tidyr 1.3.2
✓ infer	1.1.0	✓ tune 2.1.0
✓ modeldata	1.6.0	✓ workflows 1.3.0
✓ parsnip	1.6.0	✓ workflowsets 1.1.1
✓ purrr	1.2.2	✓ yardstick 1.4.0
✓ recipes	1.4.0	

Conflicts	
tidymodels_conflicts() —	
* purrr::discard()	masks scales::discard()
* dplyr::filter()	masks stats::filter()
* dplyr::lag()	masks stats::lag()
* .GlobalEnv::step()	masks recipes::step(), stats::step()

```
tidymodels_prefer()

scores <- questionnaire |>
  mutate(
    stress_4 = 6 - stress_4,
    stress
  ) %>%
  rowMeans(pick(stress_1:stress_6), na.rm =
    TRUE),
    burnout
  ) %>%
  rowMeans(pick(burnout_1:burnout_6), na.rm =
    TRUE),
    support
  ) %>%
  rowMeans(pick(support_1:support_6), na.rm =
    TRUE),
```

```

    satisfaction = rowMeans(pick(satisfaction_1:satisfaction_4),
na.rm = TRUE)
  ) |>
  select(student_id, stress, burnout, support,
satisfaction)

dropout_data <- students |>
  left_join(scores, join_by(student_id)) |>
  left_join(semesters |> filter(semester == 1)
|> select(-semester), join_by(student_id)) |>
  mutate(considering_dropout =
factor(considering_dropout, levels = c("Yes",
"No"))) |>
  select(-student_id, -supervisor_id, -
workshop, -workshop_sessions)

set.seed(2026)
dropout_split <- initial_split(dropout_data,
prop = 0.75, strata = considering_dropout)
dropout_train <- training(dropout_split)
dropout_test <- testing(dropout_split)

dropout_recipe <- recipe(considering_dropout
~ ., data = dropout_train) |>
  step_impute_median(all_numeric_predictors())
|>
  step_dummy(all_nominal_predictors()) |>
  step_normalize(all_numeric_predictors())

set.seed(2026)
dropout_folds <- vfold_cv(dropout_train, v =
10, strata = considering_dropout)

```

In `tidymodels`, a multilayer perceptron is `mlp()`. Its main arguments are `hidden_units` (the number of neurons in the hidden layer), `penalty` (the weight decay), and `epochs`. The **nnet** engine, which comes with R, fits networks with one hidden layer; `MaxNWts` raises its limit on the number of weights.

15.3.1 - A network that memorises

A network with 20 hidden neurons and no weight decay comes first, as a warning:

```

big_net <- mlp(hidden_units = 20, penalty = 0,
epochs = 1000) |>
  set_engine("nnet", MaxNWts = 5000) |>
  set_mode("classification")

set.seed(2026)
big_fit <- fit(workflow(dropout_recipe,
big_net), data = dropout_train)

augment(big_fit, new_data = dropout_train) |
> roc_auc(considering_dropout, .pred_Yes)

```

```

# A tibble: 1 × 3
  .metric .estimator .estimate
  <chr>   <chr>         <dbl>
1 roc_auc binary         1

```

```
augment(big_fit, new_data = dropout_test) |
> roc_auc(considering_dropout, .pred_Yes)
```

```
# A tibble: 1 × 3
  .metric .estimator .estimate
  <chr>   <chr>      <dbl>
1 roc_auc binary      0.641
```

A perfect AUC on the training data, and 0.64 on new students: the same overfitting as the one-nearest-neighbour model of Chapter 11. The network has 521 weights and only 449 training students, so it has more than enough freedom to memorise every one of them.

15.3.2 - Tuning size and weight decay

The number of hidden neurons and the weight decay are hyperparameters, tuned with cross-validation as in Chapters 11 to 13. Weight decay is tried on a logarithmic scale from 0.01 to about 30:

```
net_spec <- mlp(hidden_units = tune(), penalty
= tune(), epochs = 500) |>
  set_engine("nnet", MaxNWts = 5000) |>
  set_mode("classification")

net_wf <- workflow(dropout_recipe, net_spec)

net_grid <- expand_grid(hidden_units = c(1, 3,
5, 10, 20),
                      penalty = 10^seq(-2,
1.5, by = 0.5))

set.seed(2026)
net_tuning <- tune_grid(net_wf, resamples =
dropout_folds, grid = net_grid,
                      metrics
= metric_set(roc_auc))
```

```
autoplot(net_tuning)
theme_minimal(base_size = 12)
```




Figure 15.4: Cross-validated ROC AUC of neural networks with different numbers of hidden neurons and amounts of weight decay.

Figure 15.4 tells a clear story. With little weight decay, the networks overfit, and the bigger networks overfit most. With enough weight decay, networks of every size do about equally well. The penalty, not the size, is what matters most here.

```
show_best(net_tuning, metric = "roc_auc", n
= 3)
```

```
# A tibble: 3 × 8
  hidden_units penalty .metric .estimator mean
n std_err .config
  <dbl> <dbl> <chr> <chr> <dbl>
<int> <dbl> <chr>
1      10 3.16 roc_auc binary 0.804
10 0.0300 pre0_mod30_post0
2       10 1 roc_auc binary 0.804
10 0.0298 pre0_mod29_post0
3       20 3.16 roc_auc binary 0.803
10 0.0299 pre0_mod38_post0
```

```
best_net <- select_best(net_tuning, metric =
"roc_auc")
```

15.3.3 - A fair comparison

The best network has a cross-validated AUC of 0.804. For comparison, logistic regression on the same folds:

```
set.seed(2026)
logistic_cv
<- fit_resamples(workflow(dropout_recipe,
logistic_reg()),
resamples =
dropout_folds, metrics = metric_set(roc_auc))
collect_metrics(logistic_cv)
```

```
# A tibble: 1 × 6
  .metric .estimator mean n
std_err .config
  <chr> <chr> <dbl> <int> <dbl>
<chr>
1 roc_auc binary 0.804 10 0.0303
pre0_mod0_post0
```

The two are practically identical: 0.804 for the network and 0.804 for logistic regression. The final test on the held-out students confirms it:

```
set.seed(2026)
final_net <- net_wf |>
  finalize_workflow(best_net) |>
  last_fit(dropout_split, metrics =
metric_set(roc_auc, accuracy))
collect_metrics(final_net)
```

```
# A tibble: 2 × 4
  .metric .estimator .estimate .config
  <chr> <chr> <dbl> <chr>
1 accuracy binary 0.841 pre0_mod0_post0
2 roc_auc binary 0.853 pre0_mod0_post0
```

The network reaches a test AUC of 0.85, against 0.84 for logistic regression on the same test students (Chapter 11). With only 151 test students, of whom 23 considered dropping out, a difference of this size is well within chance. The answer to the research group is therefore that *a neural network does not predict dropout better*. Chapters 12 and 13 found the same for random forests, support vector machines, and boosting. The wellbeing data has a smooth pattern, which logistic regression captures fully, leaving nothing extra for a more flexible model to find.

And the neural network has real costs. Even the tuned network, with 10 hidden neurons, has 261 weights, which cannot be interpreted: there is no equivalent of an odds ratio saying how much each predictor matters. Without strong weight decay, its results depend on the random starting weights. And it needed careful tuning to avoid overfitting. Logistic regression is kept, and the thesis reports that a neural network was tried and did not improve on it, which is itself a useful finding.

15.4 - When neural networks are worth using

Neural networks are not magic, but they are not overrated either; they are powerful in particular situations. They excel with **unstructured data**, such as images, sound, and text, where the raw inputs (pixels, sound samples, words) mean little on their own and the useful features must be learned from them; here they outperform every other method by a wide margin. They need **very large datasets**, with many thousands or millions of cases, to fit models with thousands of weights reliably. And they can capture **complex, non-linear relationships** that simpler models miss, as the “In your field” box below shows.

For typical research data, a table with a few hundred cases and a few dozen variables, well-tuned regression, random forests, and boosting usually do as well or better, and are easier to explain. That is the situation of most theses.

15.5 - Deep learning

Deep learning means neural networks with many hidden layers, often dozens or hundreds, and millions or billions of weights. Each layer learns features built on the features of the layer before: in an image network, the first layers detect edges, later layers shapes, and the last layers whole objects. Special architectures suit special data: **convolutional networks** for images, and **transformers** for text. The large language models behind ChatGPT and Claude (Chapter 18) are transformer networks with billions of weights, trained on enormous amounts of text.

Deep learning needs more data, more computing power (often a graphics card), and more technical setup than anything else in this book. In R, there are two main tools. The **keras3** package is an interface to the Keras library, which runs on Python’s TensorFlow, JAX, or PyTorch; it is the most widely documented route,

but it requires Python to be installed alongside R. The **torch** package runs the same engine as Python's PyTorch without needing Python, and the **brulee** package lets `tidymodels` fit deep networks through `torch`, with the same `mlp()` interface used in this chapter.

For most researchers, the practical way to use deep learning is not to train a network from scratch, but to use a **pretrained model**, already trained by someone else on huge datasets, for a task such as transcribing interviews, recognising objects in images, or classifying text. Chapter 18 does exactly this with a language model, to code the study's open-ended survey answers.

In your field: engineering

The concrete data of Chapter 13 has the kind of curved, interacting relationships where a neural network can shine: strength rises steeply over the first weeks and then levels off, and the effect of each ingredient depends on the others. For regression, `mlp()` is used with `set_mode("regression")`:

```
data(concrete, package = "modeldata")
set.seed(1)
concrete_split <- initial_split(concrete)

concrete_recipe
<- recipe(compressive_strength ~ ., data
= training(concrete_split)) |>
  step_log(age) |>
step_normalize(all_numeric_predictors())

concrete_net <- mlp(hidden_units = 10,
penalty = 0.1, epochs = 1000) |>
  set_engine("nnet", MaxNWts = 5000) |>
  set_mode("regression")

set.seed(1)
concrete_lm_fit
<- last_fit(workflow(concrete_recipe,
linear_reg()), concrete_split)
set.seed(1)
concrete_net_fit
<- last_fit(workflow(concrete_recipe,
concrete_net), concrete_split)
collect_metrics(concrete_lm_fit)
```

```
# A tibble: 2 × 4
  .metric .estimator .estimate .config
  <chr>    <chr>         <dbl> <chr>
1 rmse      standard           7.21
pre0_mod0_post0
2 rsq      standard           0.815
pre0_mod0_post0
```

```
collect_metrics(concrete_net_fit)
```

```
# A tibble: 2 × 4
  .metric .estimator .estimate .config
  <chr>    <chr>         <dbl> <chr>
1 rmse      standard           5.04
pre0_mod0_post0
2 rsq      standard           0.909
pre0_mod0_post0
```

The step `step_log(age)` takes the logarithm of the age, which already straightens its curve for the linear model. Even so, the network's prediction error is clearly smaller: an RMSE of 5.0 megapascals, against 7.2 for the linear model. Here, there is structure for its flexibility to capture.

15.6 - Common misconceptions

The reputation of neural networks produces several misconceptions that a thesis should avoid.

- **“A neural network is always more accurate.”** On typical tables of research data it often predicts no better than regression, as this chapter found.
- **“Neural networks learn in a mysterious way.”** Training is gradient descent: many small steps that each reduce the error, as the six-student example showed.
- **“A network that fits the training data perfectly has learned the pattern.”** A perfect training score is usually a sign of memorisation; only data the network has not seen can show what it learned.
- **“Neural networks and AI are the same thing.”** Neural networks are one family of methods; large language models are a particular, very large kind of network.

15.7 - Chapter review

15.7.1 - Summary

- A neuron multiplies its inputs by weights, adds a bias, and applies an activation function. A single neuron with a sigmoid activation is a logistic regression.
- A multilayer perceptron connects neurons in layers: inputs, one or more hidden layers, and an output layer. Hidden neurons with curved activations let the network represent curves and interactions.
- Networks are trained by gradient descent: starting from random weights, each step moves every weight a little in the direction that reduces the error. For a single neuron, gradient descent arrives at the logistic regression coefficients. With many weights, networks overfit easily; weight decay (a penalty on large weights) prevents it.
- In `tidymodels`, `mlp()` with the `nnet` engine fits a network with one hidden layer; tune `hidden_units` and `penalty` with cross-validation.
- On the dropout data, a tuned neural network predicts no better than logistic regression, and is harder to interpret. Neural networks shine on images, sound, and text, on very large datasets, and on strongly non-linear problems.
- Deep learning uses many layers; in R, it is available through `keras3` (with Python) and `torch`. Pretrained models are often the practical route for researchers.

15.7.2 - Key terms

Neural network, neuron (unit), weight, bias, activation function, sigmoid, ReLU, multilayer perceptron, input layer, hidden layer, output layer, training, log loss, gradient, gradient descent, learning rate, epoch, weight decay, deep learning, convolutional network, transformer, pre-trained model.

15.8 - Exercises

The [playground](#) has these and more, with hints and solutions.

1. Change the neuron's weights so that support matters twice as much as in the chapter, recalculate the outputs for the two students, and explain what a negative weight means.
2. Fit a network with `hidden_units = 1` and a weight decay of 1, compare its cross-validated AUC with logistic regression, and explain why they might be so similar.
3. Fit the chapter's best network three times with different seeds, then do the same with a weight decay of 0.01. Explain when the seed matters, and why.
4. Train the 20-neuron network without weight decay for 10, 100, and 1000 epochs, and describe how the training and test AUCs change.
5. In two or three sentences, explain to the research group why the thesis does not use a neural network for the dropout model.
6. Repeat the gradient-descent loop with a learning rate of 0.1 and of 10. Describe how the loss curve changes, and explain what goes wrong when the learning rate is too large.

15.9 - Further reading

- *An Introduction to Statistical Learning* [18] has a clear chapter on deep learning, including convolutional networks, with the ideas explained through examples.
- *Deep Learning with R* [30] is the practical guide to keras in R, by the creator of Keras and the authors of its R interface.
- *Deep Learning* [31] is the standard, more mathematical textbook, free to read online.
- *Tidy Modeling with R* [22] shows how neural networks fit into the tidymodels workflow alongside other models.

16. Time Series Forecasting

Much of the data that organisations keep is recorded over time: patients admitted each day, products sold each month, rainfall each week. Such a sequence of measurements taken at regular intervals is a **time series**, and it calls for a different way of thinking from the surveys of the previous chapters. The observations are not independent: this week's value is related to last week's, and the order of the observations carries information. The patterns are also of a particular kind. A series usually combines a slow **trend**, a **seasonal** pattern that repeats over a fixed period, and irregular **noise**, and understanding a series largely means separating these parts. Finally, the future of a series can never be known exactly, so a forecast is honestly a range of plausible values, not a single number.

This chapter develops these ideas and the methods built on them, using the **fable** family of packages. The example comes from outside the thesis. The doctors at the university counselling service have heard about Elaf's wellbeing study and ask her to analyse their records: five years of weekly visit counts, with the service overwhelmed before exams and quiet in the summer, and staffing planned by guesswork. They want to know how many visits to expect next year, week by week. This is an **internal analysis**, done for the service rather than for publication, and a common situation for anyone who knows some data analysis: being asked to analyse someone else's data.

💡 By the end of this chapter you will be able to

- Explain dependence over time, the parts of a time series (trend, seasonality, and noise), and why forecasts are ranges.
- Store a time series as a tsibble, and plot it.
- Recognise trend, seasonality, and unusual observations, using time plots, seasonal plots, and autocorrelation.
- Decompose a series into trend, seasonal, and remainder components with STL.
- Produce benchmark forecasts, and forecasts from exponential smoothing (ETS) and ARIMA models, with fable.
- Evaluate forecasts on a test period, and report forecasts with prediction intervals.

Note

Teaching & Practice: [Interactive Lecture Slides](#) · [Browser Playground Exercises](#)

16.1 - Time series data

Most of this book's methods assume that observations are independent: one student's answers tell you nothing about another's. In a time series the opposite is true. The number of visits this week is closely related to the number last week, and the order of the observations is essential. The methods of this chapter are built around that dependence.

16.1.1 - A small example

Two weeks of daily visits to a small clinic, starting on a Monday, show how a time series is stored:

```
library(dplyr)
```

```
Attaching package: 'dplyr'
```

```
The following objects are masked from  
'package:stats':
```

```
filter, lag
```

```
The following objects are masked from  
'package:base':
```

```
intersect, setdiff, setequal, union
```

```
library(ggplot2)  
library(tsibble)
```

```
Attaching package: 'tsibble'
```

```
The following objects are masked from  
'package:base':
```

```
intersect, setdiff, union
```

```
library(fable)
```

```
Loading required package: fabletools
```

```
library(feasts)
```

```
clinic <- tsibble(  
  day = 1:14,  
  visits = c(12, 10, 9, 11, 7, 3, 2, 14, 11,  
            10, 12, 8, 4, 2),
```



```

  index = day
)
clinic

```

```

# A tibble: 14 x 2 [1]
  day visits
  <int> <dbl>
1     1     12
2     2     10
3     3      9
4     4     11
5     5      7
6     6      3
7     7      2
8     8     14
9     9     11
10    10     10
11    11     12
12    12      8
13    13      4
14    14      2

```

A **tsibble** (time series tibble) is a data frame that knows which column is time, its **index**. Here the index is the day number; the output header shows that the observations are one unit apart ([1]). tsibbles check that each time appears only once, and fable's functions use the index to keep everything in order. There is a clear weekly pattern: busy at the start of the week, quiet at the weekend.

16.1.2 - The parts of a pattern

A time series is easiest to understand as a sum of parts. The code below builds an artificial weekly series for three years from three ingredients: a trend that rises slowly, a seasonal pattern that repeats every 52 weeks, and random noise. Figure 16.1 shows each part and their sum.

```

set.seed(4)
week <- 1:156
parts <- tibble(
  week = week,
  trend = 30 + 0.05 * week,
  seasonal = 12 * sin(2 * pi * week / 52),
  noise = rnorm(156, sd = 3)
) |>
  mutate(series = trend + seasonal + noise)

parts |>
  tidyr::pivot_longer(-week, names_to =
    "part", values_to = "value") |>
  mutate(part = factor(part, levels =
    c("trend", "seasonal", "noise", "series"))) |>
  ggplot(aes(x = week, y = value)) +
  geom_line(colour = "#2f6793") +
  facet_wrap(~ part, ncol = 1, scales =
    "free_y") +
  labs(x = "Week", y = NULL) +
  theme_minimal(base_size = 11)

```




Figure 16.1: An artificial weekly series built from three parts: a rising trend, a seasonal pattern that repeats every 52 weeks, and random noise. The bottom panel is their sum.

In a real series, only the bottom panel is observed; the parts must be recovered from it. The trend says where the series is heading, the seasonal pattern says what happens at each point in the year, and the noise is what no pattern explains. Forecasting extends the trend and the seasonal pattern into the future. The noise cannot be forecast, and it is the main reason every forecast comes with a range. The decomposition later in this chapter recovers these parts from the counselling data.

16.2 - The counselling service data

The counselling records are in `counselling_visits`: the date of the Monday of each week, and the number of visits that week.

```
counselling_visits |> head()
```

	week_start	visits
1	2020-09-07	32
2	2020-09-14	28
3	2020-09-21	24
4	2020-09-28	37
5	2020-10-05	32
6	2020-10-12	39

```
range(counselling_visits$week_start)
```

```
[1] "2020-09-07" "2025-08-25"
```

The data covers five academic years, from September 2020 to August 2025, 260 weeks in all. The function `yearweek()` turns each date into a week, which becomes the index:

```
visits <- counselling_visits |>
  mutate(week = yearweek(week_start)) |>
```

```
as_tsibble(index = week)
visits
```

```
# A tsibble: 260 x 3 [1W]
  week_start visits    week
  <date>      <int>   <week>
1 2020-09-07     32 2020 W37
2 2020-09-14     28 2020 W38
3 2020-09-21     24 2020 W39
4 2020-09-28     37 2020 W40
5 2020-10-05     32 2020 W41
6 2020-10-12     39 2020 W42
7 2020-10-19     31 2020 W43
8 2020-10-26     36 2020 W44
9 2020-11-02     37 2020 W45
10 2020-11-09     34 2020 W46
# i 250 more rows
```

The header now reads [1W]: one week between observations. The first step with any time series is to plot it, and `autoplot()` from the `fable` family draws a **time plot**:

```
autoplot(visits, visits) +
  labs(x = NULL, y = "Visits per week") +
  theme_minimal(base_size = 12)
```

```
Warning: `autoplot.tbl_ts()` was deprecated in
fabletools 0.6.0.
i Please use `ggtime::autoplot.tbl_ts()`
instead.
i Graphics functions have been moved to the
{ggtime} package. Please use
`library(ggtime)` instead.
```




Figure 16.2: Weekly visits to the university counselling service, September 2020 to August 2025.

Figure 16.2 shows the parts of a pattern described above, and one feature the doctors did not mention. There is a **seasonal pattern** that repeats every academic year: peaks before and during the two exam periods, dips during the breaks, and very few visits in the summer, when the service is reduced. There is a gentle upward **trend**, each year a little busier than the last, and there is **noise**, week-to-week variation that follows no pattern. Finally, there is one **unusual week**, in the autumn of 2023, far busier than the weeks around it.

16.2.1 - Seasonal patterns

A **seasonal plot** puts the years on top of each other, so the seasonal pattern and the differences between years are easier to see. The academic year starts in September, so the weeks are numbered from the start of each academic year:

```
visits |>
  mutate(academic_year = paste0(20 +
    (row_number() - 1) %/% 52, "/", 21 +
    (row_number() - 1) %/% 52),
    week_of_year = (row_number() - 1) %
    % 52 + 1) |>
  ggplot(aes(x = week_of_year, y = visits,
    colour = academic_year)) +
  geom_line() +
  scale_colour_viridis_d(end = 0.9) +
  labs(x = "Week of the academic year", y =
    "Visits per week", colour = "Academic year") +
  theme_minimal(base_size = 12)
```




Figure 16.3: Seasonal plot: weekly visits by week of the academic year, one line per year.

Every year has the same shape, and later years lie slightly above earlier ones. The period of the seasonal pattern is 52 weeks.

16.2.2 - Autocorrelation

Dependence over time can be seen directly by plotting each week's visits against the visits one week earlier, as in Figure 16.4.

```
visits |>
  mutate(last_week = lag(visits)) |>
  ggplot(aes(x = last_week, y = visits)) +
  geom_point(alpha = 0.6, colour = "#2f6793") +
  labs(x = "Visits in the previous week", y =
    "Visits this week") +
  theme_minimal(base_size = 12)
```

```
Warning: Removed 1 row containing missing
values or values outside the scale range
(`geom_point()`).
```




Figure 16.4: Weekly visits against the visits of the previous week. Busy weeks tend to follow busy weeks, and quiet weeks quiet weeks.

The points rise from left to right: knowing last week's visits says a good deal about this week's. The **autocorrelation** at lag k puts a number on this relationship: it is the correlation between the series and itself k steps earlier, between each week and the week before (lag 1), two weeks before (lag 2), and so on. The function `ACF()` calculates it:

```
visits_acf <- visits |> ACF(visits, lag_max = 60)
visits_acf |> slice(c(1, 2, 26, 52))
```

```
# A tibble: 4 x 2 [1W]
  lag    acf
<cf_lag> <dbl>
1      1W 0.690
2      2W 0.572
3     26W -0.244
4     52W 0.684
```

The autocorrelation is 0.69 at lag 1, as the scatter plot suggested, and 0.68 at lag 52: a week tends to be like the same week a year earlier. At lag 26 it is negative: half a year apart, a busy term week is often matched with a quiet summer week. Figure 16.5 shows all the lags. Strong autocorrelation at the seasonal lag is the signature of seasonality; if a series had no autocorrelation at all, there would be nothing to forecast beyond its average.

```
autoplot(visits_acf) +
  theme_minimal(base_size = 12)
```

```
Warning: `autoplot.tbl_cf()` was deprecated in
feasts 0.4.2.
i Please use `ggtime::autoplot.tbl_cf()`
instead.
i Graphics functions have been moved to the
{ggtime} package. Please use
`library(ggtime)` instead.
```




Figure 16.5: Autocorrelation of weekly visits for lags of 1 to 60 weeks. The dashed lines mark the range expected for a series with no autocorrelation.

16.3 - Decomposition

Decomposition recovers the parts of a pattern from an observed series: a smooth **trend**, a repeating **seasonal** component, and the **remainder**, what is left over, which corresponds to the noise of Figure 16.1. **STL** (seasonal and trend decomposition using loess) is a flexible, widely used method. `season(period = 52)` tells it the length of the seasonal pattern, and `robust = TRUE` stops unusual weeks from distorting the trend and season:

```
visits_stl <- visits |>
  model(STL(visits ~ season(period = 52),
    robust = TRUE)) |>
  components()
```

```
autoplot(visits_stl)
theme_minimal(base_size = 11)
```

Warning: `autoplot.dcmp\_ts()` was deprecated in fabletools 0.6.0.
i Please use `ggtime::autoplot.dcmp\_ts()` instead.
i Graphics functions have been moved to the {ggtime} package. Please use `library(ggtime)` instead.




Figure 16.6: STL decomposition of weekly visits into trend, seasonal pattern, and remainder.

In Figure 16.6 the series is the sum of the three components below it, just as in the artificial example. The trend rises from about 28 to about 38 visits a week, levelling off in the last year; the seasonal component repeats the academic year; and the remainder is mostly small noise. The largest remainder is the unusual week:

```
visits_stl |>
  as_tibble() |>
  slice_max(abs(remainder), n = 3) |>
  select(week, visits, trend, season_52,
         remainder)
```

```
# A tibble: 3 × 5
   week visits trend season_52 remainder
<week> <int> <dbl>   <dbl>   <dbl>
1 2023 W43    71  35.2     8.08    27.7
2 2024 W48    32  38.4    12.1   -18.5
3 2024 W20    51  37.1    31.3   -17.5
```

In the week of 23 October 2023, there were 71 visits, about 28 more than trend and season would predict. When asked, the doctors remember at once: it was the university’s mental health awareness week, with posters everywhere encouraging students to seek help. The remainder is where such events show up. It is a real week, not an error, so it is kept, and the robust decomposition makes sure it does not distort the forecasts.

Notice also that the seasonal swings grow slightly as the trend rises: the peaks get higher and the summer lows stay low. When the seasonal pattern grows in proportion to the level of the series, it is usual to model the **logarithm** of the series, on which the swings become constant. *fable* handles this automatically: write `log(visits)` in the model, and the forecasts are transformed back to visits.

16.4 - Forecasting

16.4.1 - Simple benchmarks

Every forecasting method should be compared with simple **benchmark** methods, just as every predictive model in Chapters 11 to 13 was compared with a baseline. The three most common are the **mean** method, which forecasts the average of all past observations; the **naïve** method, which forecasts the last observed value; and the **seasonal naïve** method, which forecasts the value from the same season in the last cycle (the same day last week, or the same week last year). For the clinic example, with its weekly cycle of 7 days, they give:

```
clinic_fc <- clinic |>
  model(
    mean = MEAN(visits),
    naïve = NAIVE(visits),
    snaive = SNAIVE(visits ~ lag(7))
  ) |>
  forecast(h = 7)

clinic_fc |>
```

```
as_tibble() |>
  select(.model, day, .mean) |>
  tidyr::pivot_wider(names_from = .model,
    values_from = .mean)
```

```
# A tibble: 7 × 4
  day mean naive snaive
<dbl> <dbl> <dbl> <dbl>
1    15  8.21     2     14
2    16  8.21     2     11
3    17  8.21     2     10
4    18  8.21     2     12
5    19  8.21     2      8
6    20  8.21     2      4
7    21  8.21     2      2
```

The function `model()` fits several models at once, each named; `forecast(h = 7)` forecasts 7 steps ahead; and `.mean` holds the point forecasts. The mean method forecasts 8.2 visits every day; the naive method repeats the last value, a quiet Sunday, for the whole week; and the seasonal naive method repeats last week's pattern, which for this series is clearly the most sensible.

16.4.2 - Two families of models

Two families of models go beyond the benchmarks. **Exponential smoothing (ETS)** forecasts with weighted averages of past observations, with weights that decrease the further back the observations are, so that recent weeks count most. ETS models can track a changing level, a trend, and a seasonal pattern, each estimated from the data; the name stands for the three components it models, error, trend, and season. **ARIMA** models describe how each observation depends on earlier observations and on earlier random shocks, and use the autocorrelation of the series to forecast it.

The functions `ETS()` and `ARIMA()` in `fable` choose the details of each model automatically, by a criterion like the BIC of Chapter 14. Both work best with short seasonal periods, such as 4 quarters, 7 days, or 12 months. A 52-week season is long, so two strategies are common for weekly data. The first is to **decompose first**: remove the seasonal pattern with `STL`, forecast the **seasonally adjusted** series (trend and remainder) with ETS, and add the seasonal pattern back, all of which `decomposition_model()` does in one step. The second uses **Fourier terms**: the seasonal pattern is described with a few smooth waves (sines and cosines) of different lengths, used as predictors in an ARIMA model; `fourier(period = 52, K = 6)` uses six pairs of waves.

16.4.3 - Training and test periods

As in machine learning, forecasts must be judged on data the model has not seen. For time series, the test data must come *after* the training data, because a forecast can only use the past. The models are trained on the first four academic years and tested on the fifth:


```
visits_train <- visits |> filter(week_start
< as.Date("2024-09-01"))
visits_test  <- visits |> filter(week_start
>= as.Date("2024-09-01"))

visits_fit <- visits_train |>
  model(
    mean      = MEAN(visits),
    naive     = NAIVE(visits),
    snaive    = SNAIVE(visits ~ lag(52)),
    stl_ets   = decomposition_model(
      STL(log(visits) ~ season(period
= 52), robust = TRUE),
      ETS(season_adjust ~ season("N"))
    ),
    fourier   = ARIMA(log(visits) ~
fourier(period = 52, K = 6) + PDQ(0, 0, 0))
  )
```

In the STL model, `ETS(season_adjust ~ season("N"))` forecasts the seasonally adjusted series with no seasonal component of its own (“N” for none), because the seasonal pattern is added back from the decomposition. In the Fourier model, `PDQ(0, 0, 0)` tells ARIMA not to model the season itself, because the Fourier terms do that.

Each model forecasts the 52 weeks of the test year, and `accuracy()` compares the forecasts with what actually happened, using the RMSE and MAE of Chapter 13:

```
visits_fc <- visits_fit |> forecast(h = 52)

visits_fc |>
  accuracy(visits) |>
  select(.model, RMSE, MAE) |>
  arrange(RMSE)
```

```
# A tibble: 5 × 3
  .model  RMSE  MAE
<chr>   <dbl> <dbl>
1 stl_ets  7.40  5.57
2 snaive  9.16  6.96
3 fourier 10.3  7.81
4 mean   18.7 16.3
5 naive  27.5 23.5
```

The STL and ETS model is best: its forecasts are off by 5.6 visits a week on average (MAE), against 7.0 for the seasonal naive benchmark, which simply repeats the previous year. The mean and naive methods, which ignore the seasonal pattern, are far worse. The Fourier model does worse than the seasonal naive benchmark: the counselling service’s pattern has sharp steps (the summer drop happens from one week to the next), and smooth waves struggle to follow them. Figure 16.7 compares the two best models with the actual visits.

```
visits_fc |>
  filter(.model %in% c("stl_ets", "snaive"))
|>
  autoplot(visits_test, level = NULL) +
  labs(x = NULL, y = "Visits per week", colour
```

```
= "Model") +
  theme_minimal(base_size = 12)
```

```
Warning: `autoplot.fbl_ts()` was deprecated in
fabletools 0.6.0.
i Please use `ggtime::autoplot.fbl_ts()`
instead.
i Graphics functions have been moved to the
{ggtime} package. Please use
`library(ggtime)` instead.
```




Figure 16.7: Forecasts for the test year (September 2024 to August 2025) from the STL and ETS model and the seasonal naive benchmark, with the actual visits in black.

The seasonal naive forecast copies last year's noise along with its pattern: it even repeats the awareness week of October 2023 as a spike in October 2024. The STL and ETS model smooths the noise away and adds the trend, so its forecasts are both smoother and closer.

16.4.4 - Forecasting next year

With the model chosen, it is refitted on all five years, so the forecasts use the most recent information, and the next 52 weeks are forecast:

```
final_fit <- visits |>
  model(stl_ets = decomposition_model(
    STL(log(visits) ~ season(period = 52),
    robust = TRUE),
    ETS(season_adjust ~ season("N"))
  ))

next_year <- final_fit |> forecast(h = 52)
```

```
next_year |>
  autoplot(visits |> filter(week_start >=
as.Date("2023-09-01"))) +
  labs(x = NULL, y = "Visits per week") +
  theme_minimal(base_size = 12)
```




Figure 16.8: Forecast of weekly visits for the academic year 2025/26, with 80% and 95% prediction intervals, after the last two years of data.

A forecast is never a single number. The shaded bands in Figure 16.8 are **prediction intervals**: the ranges within which the actual visits are expected to fall with 80% and 95% probability. The function `hilo()` shows them as numbers:

```
next_year |>
  hilo(95) |>
  select(week, .mean, `95%`) |>
  head(4)
```

```
# A tsibble: 4 x 3 [1W]
   week .mean      `95%`
   <week> <dbl>      <hilo>
1 2025 W36 39.0 [23.60596, 60.76414]95
2 2025 W37 35.2 [21.32399, 54.89014]95
3 2025 W38 47.5 [28.74758, 73.99921]95
4 2025 W39 48.8 [29.58119, 76.14503]95
```

In the first week of the new year, about 39 visits are expected, but anything from about 24 to 61 would not be surprising. The doctors also asked for the year as a whole. Adding up the weekly forecasts gives the expected total, but not its uncertainty, because weeks do not vary independently. The function `generate()` solves this by **simulation**: it produces 1,000 possible futures from the model, and the totals of those futures show the range of likely totals:

```
set.seed(2026)
futures <- final_fit |> generate(h = 52, times
= 1000)

yearly_totals <- futures |>
  as_tibble() |>
  summarise(total = sum(.sim), .by = .rep)

quantile(yearly_totals$total, c(0.025, 0.5,
0.975)) |> round()
```

```
2.5%   50%  97.5%
1952  2111  2286
```

The report to the service therefore says: about 2,111 visits are expected in 2025/26 (95% interval 1,952 to 2,286), compared with 1,954 in 2024/25; the busiest weeks will

again be just before and during the two exam periods, when they should plan the most staff; and an awareness campaign can raise demand sharply for a week, so extra capacity should be planned for any such event.

 Forecasts assume the future resembles the past

Every forecast in this chapter assumes that the patterns of the last five years continue. A change the data cannot know about, such as a new online booking system, a change in the exam calendar, or a crisis like the COVID-19 pandemic, can make any forecast wrong. Report forecasts with their intervals, say what they assume, and update them as new data arrives. (You may also meet **Prophet**, a forecasting tool from Meta that was popular for business data; it is no longer actively developed, and fable’s models are a well-supported alternative.)

In your field: environmental science

R's `co2` dataset records the monthly concentration of carbon dioxide in the atmosphere, measured at the Mauna Loa observatory in Hawaii from 1959 to 1997: a famous series with a strong upward trend and a yearly cycle, as plants absorb CO₂ in the northern summer. With a 12-month season, ETS and ARIMA can model the seasonal pattern directly. Training on the years to 1990 and testing on 1991 to 1997:

```
co2_ts <- as_tsibble(co2)

co2_fit <- co2_ts |>
  filter(index <= yearmonth("1990 Dec"))
|>
  model(
    snaive = SNAIVE(value),
    ets    = ETS(value),
    arima  = ARIMA(value)
  )
```

```
Warning in sqrt(diag(best$var.coef)):
NaNs produced
```

```
co2_accuracy <- co2_fit |>
  forecast(h = "7 years") |>
  accuracy(co2_ts) |>
  select(.model, RMSE, MAE)
co2_accuracy
```

```
# A tibble: 3 × 3
  .model RMSE MAE
  <chr>   <dbl> <dbl>
1 arima  1.34  1.24
2 ets    1.16  1.03
3 snaive 6.07  5.26
```

The function `as_tsibble()` converts R's older time series objects, and `h = "7 years"` gives the forecast horizon in words. Both ETS and ARIMA forecast seven years ahead with an average error (MAE) of only 1.0 and 1.2 parts per million, while the seasonal naive method, which ignores the trend, falls further behind every year.

16.5 - Common misconceptions

Time series have their own traps, most of them versions of forgetting that time has an order.

- **“A forecast is a prediction of the exact value.”** A forecast is a range; the point forecast is only its centre, and the noise in a series makes exact prediction impossible.
- **“More history always gives better forecasts.”** Only if the old patterns still hold; a change in circumstances can make older data misleading.

- **“The test data can be any random subset.”** In a time series, the test period must come after the training period, since a forecast can only use the past.
- **“An unusual week should be removed.”** It is part of the record; robust methods keep it from distorting the forecast, and it may be exactly what the service needs to plan for.

16.6 - Chapter review

16.6.1 - Summary

- A time series is a sequence of observations over time, in which order matters and neighbouring observations are related. It can be understood as a sum of parts: trend, seasonal pattern, and noise. A tsibble stores it with a time index.
- Time plots, seasonal plots, and the autocorrelation function reveal trend, seasonality, and unusual observations.
- STL decomposes a series into trend, seasonal, and remainder components; unusual events show up in the remainder. A logarithm makes growing seasonal swings constant.
- Compare every forecasting method with benchmarks: mean, naive, and seasonal naive.
- ETS forecasts with weighted averages that favour recent observations; ARIMA uses the autocorrelation of the series. For long seasons such as 52 weeks, decompose first or use Fourier terms.
- Evaluate forecasts on a test period that comes after the training period, and report forecasts with prediction intervals. Simulating futures with `generate()` gives intervals for totals.
- Forecasts assume that past patterns continue.

16.6.2 - Key terms

Time series, tsibble, index, time plot, trend, seasonality, seasonal period, seasonal plot, autocorrelation, lag, decomposition, STL, remainder, seasonally adjusted series, benchmark forecast, mean method, naive method, seasonal naive method, exponential smoothing (ETS), ARIMA, Fourier terms, forecast horizon, prediction interval, simulation.

16.7 - Exercises

The [playground](#) has these and more, with hints and solutions.

1. Predict which of the three benchmarks would be best for a series with a trend but no seasonality, then test your prediction on `clinic` after adding a steady increase of one visit a day.
2. Plot the seasonally adjusted series (`season_adjust` in `visits_stl`), and describe what it shows that the original series hides.
3. Refit the Fourier model with `K = 2` and `K = 12`, and explain how and why the test accuracy changes.

4. Train the models on the first three years and test on the fourth, and check whether the STL and ETS model is still the best.
5. Use `generate()` to estimate how many visits to expect in the four weeks before the first exam period of 2025/26, with a 95% interval.
6. Write a short paragraph for the counselling service explaining what a 95% prediction interval means.
7. In the artificial series of Figure 16.1, double the standard deviation of the noise. Describe how the sum changes, and explain what this would mean for the width of forecast intervals.

16.8 - Further reading

- *Forecasting: Principles and Practice* [32], free online, is the standard introduction to forecasting, written around the fable packages used in this chapter.

IV Reproducible Research and Applications

17	Reproducible Research	345
17.1	Reproducibility and credibility	346
17.2	Documents that contain their own analysis . . .	346
17.3	A results chapter in Quarto	348
17.4	Keeping the same package versions	352
17.5	Interactive dashboards with Shiny	352
17.6	Open science	355
17.7	Keeping a history with git (optional)	358
17.8	Common misconceptions	359
17.9	Chapter review	359
17.10	Exercises	360
17.11	Further reading	360
18	Using AI with R	361
18.1	The reliability of coding	362
18.2	How AI assistants work	362
18.3	AI as a coding assistant	363
18.4	Coding the open-ended answers	364
18.5	Using AI responsibly	372
18.6	Common misconceptions	374
18.7	Chapter review	374
18.8	Exercises	375
18.9	Further reading	376
19	Putting It All Together	377
19.1	The shape of a research project	378
19.2	From raw export to clean data	380
19.3	Describing the sample	383
19.4	Answering the research questions	384
19.5	One figure for the thesis	386
19.6	The chain of reasoning	387
19.7	The results chapter	390
19.8	Choosing a method	390
19.9	Challenges every researcher meets	392
19.10	Where to go next	393
19.11	Common misconceptions	395
19.12	Chapter review	396
19.13	Exercises	396
19.14	Further reading	397
19.14	References	399

17. Reproducible Research

A research finding is credible only if it can be checked. A reader who doubts a result should be able to see exactly how it was obtained, and a second researcher who repeats the analysis should arrive at the same numbers. When that is not possible, the reader has to take the result on trust, and science is built on not having to. The failures that make results uncheckable are rarely dramatic. A table is copied into a draft before the data is corrected, a decision made while exploring is forgotten, a package changes its behaviour between versions, and nobody can say afterwards which of several analyses produced the published number. Reproducibility is therefore a matter of research integrity, not of technical tidiness.

The problem is easy to meet in a thesis. Three weeks before the deadline, Elaf's supervisor spots that twelve students appear twice in an early version of the survey export. They were removed in Chapter 3, but some tables were copied into the thesis draft before that, and now every table, figure, and number copied by hand from R into Word must be checked and replaced, one at a time, with the risk of missing one. This chapter shows a way of working in which that problem cannot happen. In a **reproducible** workflow, the text, the code, and the results live together, and the whole report is rebuilt from the raw data with one command, so every number is always up to date. The chapter introduces **Quarto** for writing such documents, from a single page to a thesis chapter; **renv** for keeping the same package versions; **Shiny** for turning an analysis into an interactive dashboard; and the practices of **open science** that make research checkable by others. An optional section introduces **git** for keeping the history of a project.

💡 By the end of this chapter you will be able to

- Explain why reproducibility is a matter of research integrity, and why copying results by hand and flexible analyses are risky.
- Write a Quarto document that combines text, code, results, tables, figures, citations, and equations, and render it to HTML, Word, and PDF.
- Keep a project's package versions with **renv**, and report software versions.
- Build a small Shiny dashboard with inputs, outputs, and reactive code.
- Apply open science practices: sharing data and code, protecting participants, preregistration, and citing software.

Note

Teaching & Practice: [Interactive Lecture Slides](#) · [Browser Playground Exercises](#)

17.1 - Reproducibility and credibility

An analysis is **reproducible** if someone else, or you in a year's time, can take the same data and code and get exactly the same results. It is **replicable** if a new study, with new data, reaches the same conclusions. Reproducibility is the minimum standard: if the original results cannot even be recomputed, there is little point asking whether they replicate [33].

Most irreproducibility is not fraud but everyday friction: a table copied before the data was corrected, a spreadsheet cell edited by hand, a result produced by clicking through menus that nobody wrote down, a package that changed its default between versions. The previous chapters have already built good habits against these. Every step is **written as code**, in scripts, so it can be rerun (Chapter 1). Each analysis lives in an **RStudio Project** with relative paths, so it runs on any computer (Chapter 1). The **raw data is never edited by hand**, and every cleaning decision is made and recorded in code (Chapter 3). Random steps use `set.seed()`, so they give the same answer every time (Chapter 7).

This chapter adds the last links: putting the results into the report automatically, fixing the software versions, and sharing everything responsibly. Good practices do not need to be perfect to be useful; even a few of them make research far easier to check [34].

17.2 - Documents that contain their own analysis

Quarto is a publishing system for documents that mix text and code. A Quarto document is a plain text file with the extension `.qmd`. When it is **rendered**, the code is run, and its results (numbers, tables, figures) are placed into the finished document, which can be a web page, a Word document, a PDF, a slide show, or a whole book. This book is written in Quarto: every number, table, and figure in it was produced by the code shown next to it.

i R Markdown

Before Quarto, the same job was done by **R Markdown** (`.Rmd` files), and you will meet it in many older projects, templates, and journal guides. The two are very similar: the text, code chunks, and inline code of this chapter work almost unchanged in R Markdown, where documents are rendered with the Knit button or `rmarkdown::render()`. Quarto is its successor, from the same developers, and works with Python and other languages as well as R. For new work, use Quarto.

17.2.1 - A small example

A complete Quarto document can be very short:

```
---
title: "Three friends' sleep"
author: "Elaf"
format: html
---

My three friends slept `r mean(sleep)` hours
on average last night.

```${r}
sleep <- c(6.5, 7, 5.5)
mean(sleep)
```
```

It has the three parts of every Quarto document. The **YAML header**, between the two `---` lines, holds settings: the title, the author, and the output **format**. The **text** is written in **Markdown**, a simple way of marking formatting: **`**bold**`**, *`*italic*`*, `# Heading`, and `-` for a bullet point. The **code chunks**, between ````${r}` and `````, are run when the document is rendered, and their code and output appear in the document.

The text also contains **inline code**, `r mean(sleep)`: R code inside a sentence, written between backticks and starting with the letter `r` and a space. When the document is rendered, the inline code is replaced by its result, so the sentence reads “My three friends slept 6.333333 hours on average last night.” (Rounding it, with `round(mean(sleep), 1)`, would be better.) If a friend’s number changes, the sentence changes with it. Inline code is the cure for the problem at the start of the chapter: numbers in the text are never typed by hand.

In RStudio, create a Quarto document with *File > New File > Quarto Document*, and render it with the **Render** button. RStudio also offers a **visual editor**, which shows the formatting as in a word processor while writing the same `.qmd` file underneath.

17.2.2 - What happens when you render

Figure 17.1 shows the steps. The **knitr** package runs the R code and writes a Markdown file with the results in place;

Pandoc, a universal document converter, then turns the Markdown into the requested format.




Figure 17.1: How a Quarto document is rendered.

Because the code is run from the beginning every time, a rendered document is a test of reproducibility in itself: if the code depends on something that is not in the document or the project, such as an object created by hand in the Console, rendering fails.

17.3 - A results chapter in Quarto

The descriptive results chapter of the thesis, rewritten as a Quarto document, begins like this:

```
---
title: "Chapter 4: Results"
author: "Elaf"
format: docx
bibliography: references.bib
```

```

execute:
  echo: false
  ---

  ```{r}
 #| label: setup
 #| message: false
 library(dplyr)
 library(ggplot2)
 library(data2thesis)
  ```

  The sample included `r nrow(students)`
  graduate students from
  `r n_distinct(students$faculty)` faculties.
  Their average wellbeing in the
  first semester was `r round(mean(semester$wellbeing[semester$semester
  == 1], na.rm = TRUE), 1)`
  on the wellbeing index (0 to 100). @tbl-faculty
  shows wellbeing by faculty.

  ```{r}
 #| label: tbl-faculty
 #| tbl-cap: "Wellbeing in the first semester,
 by faculty."
 semester |>
 filter(semester == 1) |>
 left_join(students, join_by(student_id)) |>
 summarise(Students = n(),
 Mean = mean(wellbeing, na.rm
= TRUE),
 SD = sd(wellbeing, na.rm = TRUE),
 .by = faculty) |>
 arrange(faculty) |>
 knitr::kable(digits = 1, col.names =
c("Faculty", "Students", "Mean", "SD"))
  ```

  Wellbeing declined over the programme (@fig-
  wellbeing), as found in
  earlier studies [@author2020].

```

Several new features appear. In the header, `format: docx` produces a Word document, the format most supervisors want for comments, and `bibliography:` names a file of references. The setting `execute: echo: false` hides the code in the finished document, so that a thesis shows only the results; the code is still in the `.qmd` file for anyone who wants to check it.

In the body, **chunk options**, the lines starting with `#|`, control each chunk: `label` names it, `message: false` hides package messages, and `tbl-cap` or `fig-cap` gives a caption. A chunk labelled `tbl-...` or `fig-...` becomes a numbered table or figure, and `@tbl-faculty` in the text becomes a **cross-reference**, “Table 1”, with the number filled in automatically. The function `knitr::kable()` turns a data frame into a formatted table. Finally, `[@author2020]` is a **citation**: Quarto looks up the key in the bibliography file, formats the citation, and adds the reference list at the end.

When the document is rendered, the first paragraph becomes:

The sample included 600 graduate students from 5 faculties. Their average wellbeing in the first semester was 60.5 on the wellbeing index (0 to 100). Table 1 shows wellbeing by faculty.

and the table chunk produces Table 17.1:

Table 17.1: Wellbeing in the first semester, by faculty.

| Faculty | Students | Mean | SD |
|------------------|----------|------|------|
| Education | 148 | 61.9 | 13.2 |
| Health Sciences | 154 | 59.0 | 11.6 |
| Humanities | 95 | 61.7 | 11.7 |
| Natural Sciences | 87 | 61.3 | 10.5 |
| Social Sciences | 116 | 58.9 | 12.1 |

If the data changes, the document is rendered again, and every number, table, and figure is updated. The supervisor's problem from the start of the chapter now takes one click to fix.

17.3.1 - References and citations

The bibliography file uses the **BibTeX** format, which almost every reference manager (Zotero, Mendeley, End-Note) can export, and Google Scholar can provide for any paper. One entry looks like this:

```
@book{kuhn2022,
  author   = {Kuhn, Max and Silge, Julia},
  title    = {Tidy Modeling with R},
  publisher = {O'Reilly Media},
  year     = {2022}
}
```

The key, `kuhn2022`, is what goes after the `@` in the text. A `csl:` line in the YAML header chooses the citation style (APA, Vancouver, Harvard, or any of thousands of journal styles from the Citation Style Language collection), so switching styles never means retyping references. With Zotero, RStudio's visual editor can insert citations directly from your library.

17.3.2 - Equations

Equations are written in **LaTeX** notation between dollar signs: `\bar{x} = \frac{1}{n} \sum x_i` inside a sentence gives $\bar{x} = \frac{1}{n} \sum x_i$, and double dollar signs set an equation on its own line:

```
$$
\text{logit}(p) = \beta_0 + \beta_1 \times \text{stress} + \beta_2 \times \text{support}
$$
```

$$\text{logit}(p) = \beta_0 + \beta_1 \times \text{stress} + \beta_2 \times \text{support} \quad (17.1)$$

The same notation works in HTML, Word, and PDF output. The equations in this book are all written this way.

17.3.3 - Output formats

Changing one line of the YAML header changes the output:

Table 17.2: Common Quarto output formats

| Format | YAML | Notes |
|----------|-------------------------------|---|
| Web page | <code>format: html</code> | Interactive; good for sharing results on-line |
| Word | <code>format: docx</code> | For comments from supervisors; reference-doc: <code>template.docx</code> applies your university's styles |
| PDF | <code>format: pdf</code> | Needs a LaTeX installation: run <code>quarto install tinytex</code> once in the Terminal |
| Slides | <code>format: revealjs</code> | For presentations |

Several formats can be listed at once, and one render then produces all of them. Many journals and universities provide Quarto templates, installed as **extensions**, that format a document to their requirements. A whole thesis can be written as a **Quarto book** (`project: type: book`), with one `.qmd` file per chapter, exactly like this book.

17.3.4 - Parameterised reports

The dean of each faculty would like a one-page summary for their own faculty. Rather than five reports, one report with a **parameter** serves them all:

```
---
title: "Wellbeing report"
format: html
params:
  faculty: "Education"
---

```{r}
faculty_semesters <- semesters |>
 left_join(students, join_by(student_id)) |>
 filter(faculty == params$faculty)
```

```

`r
This report describes the `r
n_distinct(faculty_semesters$student_id)`
students
of the Faculty of `r params$faculty`.

```

The value `params$faculty` is used in the code like any other value. Rendering from the Terminal with a different value produces each faculty's version:

```

quarto render faculty-report.qmd -P
faculty:"Health Sciences"

```

## 17.4 - Keeping the same package versions

Packages change. A function's default may change between versions, or a function may be removed, so an analysis that runs today may give different results, or fail, in two years. **renv** records the exact version of every package a project uses, and can restore them later on any computer. It has three main commands, run in the Console:

```

renv::init() # once: give the project its
 # own package library
renv::snapshot() # after installing or
 # updating packages: record the versions
renv::restore() # on another computer, or
 # later: install the recorded versions

```

The first command, `renv::init()`, creates a file called `renv.lock`, which lists every package and its version, and gives the project its own private library of packages, so updating a package for one project does not affect the others. Share `renv.lock` with your code, and anyone can rebuild your exact set of packages.

Even without `renv`, **report the versions** of R and the key packages in your methods section. R can tell you:

```
R.version.string
```

```
[1] "R version 4.4.3 (2025-02-28 ucrt)"
```

```
packageVersion("lme4")
```

```
[1] '1.1.37'
```

The function `sessionInfo()` lists everything that is loaded, for a full record.

## 17.5 - Interactive dashboards with Shiny

A report answers the questions its author thought of. Sometimes readers want to ask their own: “what about my faculty?”, “what about sleep instead of wellbeing?”.

**Shiny** turns R code into an interactive web application, without any knowledge of web programming.

A Shiny app rests on three ideas. **Inputs** are the controls the reader uses, such as drop-down menus, buttons, and sliders. **Outputs** are what the app shows in response, such as plots and tables. Every app therefore has two parts: the **user interface** (`ui`), which describes what the user sees, the inputs and the places for the outputs, and the **server** function, which contains the R code that produces the outputs from the inputs. The third idea is what makes Shiny work: **reactivity**. Whenever an input changes, Shiny works out which outputs depend on it and reruns only their code. Figure 17.2 shows the connections in a wellbeing dashboard: choosing a different faculty updates both the plot and the table, while choosing a different measure updates them without refiltering the data.




Figure 17.2: Reactivity in the wellbeing dashboard. When an input changes, everything downstream of it is updated.

### 17.5.1 - A wellbeing dashboard

The complete app, saved as a file called `app.R`, is:

```

library(shiny)
library(bslib)
library(dplyr)
library(ggplot2)
library(data2thesis)

Semester records with each student's faculty
and study mode
wellbeing_data <- semesters |>
 left_join(students |> select(student_id,
 faculty, study_mode), join_by(student_id))

measures <- c("Wellbeing (0-100)" =
 "wellbeing",
 "Sleep (hours a night)" =
 "sleep_hours",
 "Study (hours a week)" =
 "study_hours")

ui <- page_sidebar(

```

```

 title = "Graduate student wellbeing",
 sidebar = sidebar(
 selectInput("faculty", "Faculty", choices
= sort(unique(wellbeing_data$faculty))),
 radioButtons("measure", "Measure", choices
= measures)
),
 card(plotOutput("trend")),
 card(tableOutput("summary"))
)

 server <- function(input, output, session) {
 selected <- reactive({
 wellbeing_data |> filter(faculty ==
input$faculty)
 })

 output$trend <- renderPlot({
 selected() |>
 summarise(mean =
mean(.data[[input$measure]], na.rm = TRUE),
 .by = c(semester, study_mode))
 |>
 ggplot(aes(x = semester, y = mean, colour
= study_mode)) +
 geom_line(linewidth = 1) +
 geom_point(size = 3) +
 labs(x = "Semester", y = names(measures)
[measures == input$measure],
 colour = "Study mode") +
 theme_minimal(base_size = 14)
 })

 output$summary <- renderTable({
 selected() |>
 summarise(students =
n_distinct(student_id),
 mean =
mean(.data[[input$measure]], na.rm = TRUE),
 .by = semester)
 })
 }

 shinyApp(ui, server)

```

The code follows the three ideas. The layout comes first: the function `page_sidebar()` from the **bslib** package lays out the page, with a title, a sidebar with the inputs, and the main area with two **cards**, one for the plot and one for the table. The inputs are created by `selectInput()`, a drop-down menu whose value is available in the server as `input$faculty`, and by `radioButtons()`, which creates the `input$measure` choice; the names in `measures` are shown to the user, and the values are the column names. The outputs are reserved by `plotOutput("trend")` and `tableOutput("summary")`, and the server fills them by assigning to `output$trend` and `output$summary` with `renderPlot()` and `renderTable()`.

Reactivity is handled by `reactive()`, which creates `selected()`, the data for the chosen faculty. It is recalculated only when `input$faculty` changes, and both outputs use it. Finally, the expres-

sion `.data[[input$measure]]` picks the column whose name is stored in `input$measure`, the usual way to use a column chosen by the user inside `dplyr` and `ggplot2`.

Click **Run App** in RStudio, and the dashboard opens in a window. Figure 17.3 shows the plot it draws for the Faculty of Education and the wellbeing measure.




Figure 17.3: The dashboard’s plot for the Faculty of Education, showing average wellbeing by semester for full-time and part-time students.

To share an app with people who do not use R, it must run on a server: **shinyapps.io** offers free hosting for small apps, and many universities run **Posit Connect**. **Shinylive** can even turn a simple app into a web page that runs entirely in the reader’s browser, with no server, using the same webR technology as this book’s playground. Like any published result, a dashboard must protect participants: this one shows only averages, never individual students.

## 17.6 - Open science

Reproducibility inside your own project is the first step. **Open science** goes further: making research checkable and reusable by others.

### 17.6.1 - Sharing data and code

Share the code that produced every result, and the data if you can. Repositories such as **OSF** (the Open Science Framework) and **Zenodo** store them permanently and give them a **DOI**, a permanent identifier that can be cited in the thesis. A good package includes the raw data (or instructions for obtaining it), the cleaning and analysis code, the Quarto source of the report, the `renv.lock` file, and a **codebook** describing every variable, like the specification of the wellbeing dataset. A **licence** tells others what they may do with it: CC BY for data and text, and the MIT licence for code, are common choices.

### 17.6.2 - Protecting participants

Research data about people can only be shared if the people cannot be identified. Removing names and student numbers is not enough: combinations of ordinary variables can identify someone. In the wellbeing data, a few

combinations of faculty, study mode, gender, and having children describe only a handful of students:

```
small_cells <- students |>
 count(faculty, study_mode, gender,
has_children) |>
 arrange(n)
head(small_cells, 4)
```

	faculty	study_mode	gender	has_children	n
1	Education	Part-time	Male	Yes	3
2	Health Sciences	Part-time	Male	Yes	3
3	Humanities	Part-time	Female	Yes	3
4	Education	Part-time	Female	Yes	4

Only 3 students are, for example, part-time male students with children in the Faculty of Education; in a small department, such a description may point to recognisable people. Before sharing, such **small cells** are protected, for example by grouping categories, removing variables that are not needed, or sharing only summary data. Ethics approval and the consent form set what may be shared: if the consent form told students that only anonymised data would be shared, that is all that may be shared. When the real data cannot be shared at all, a **synthetic** dataset with the same structure, like the one in this book, lets others run the code.

### 17.6.3 - Preregistration

Every analysis involves choices that are reasonable either way: whether to exclude outliers, whether to transform a skewed variable, which test to use, which control variables to include. Chapter 5 warned that when these choices are made after seeing the data, they can be steered, often unconsciously, towards a significant result. A simulation shows how much this matters. The code below creates data with no real difference between two groups, analyses it in five defensible ways, and records whether the first analysis, and whether *any* of the five, gives a p-value below 0.05:

```
forking_paths <- function() {
 group <- rep(c("A", "B"), each = 30)
 score <- rexp(60, rate = 1 / 10) #
 a skewed score, no real group difference
 covariate <- rnorm(60)
 z <- abs(as.numeric(scale(score)))
 p <- c(
 t_test = t.test(score ~
group)$p.value,
 no_outliers = t.test(score[z < 2] ~
group[z < 2])$p.value,
 log_scale = t.test(log(score) ~
group)$p.value,
 rank_test = wilcox.test(score ~
group)$p.value,
 with_covariate = summary(lm(score ~ group
+ covariate))$coefficients[2, 4]
)
 c(first_analysis = p[["t_test"]] < 0.05,
```

```
any_analysis = any(p < 0.05))
}

set.seed(17)
false_positives <- rowMeans(replicate(2000,
 forking_paths()))
false_positives
```

```
first_analysis any_analysis
 0.0545 0.1085
```

A single planned analysis gives a false positive in about 5% of the simulated studies, as intended. Choosing afterwards among five reasonable analyses raises the rate to about 11%, although there is nothing to find. No individual choice is wrong; the problem is choosing after seeing the result. This is sometimes called the **garden of forking paths**, and preregistration is the way out of it.

**Preregistration** means recording the hypotheses, the design, and the planned analysis before seeing the data, in a time-stamped public registry such as OSF Registries or AsPredicted. It separates **confirmatory** analyses, planned in advance, from **exploratory** ones, found along the way, and protects against the temptation, conscious or not, to try analyses until one gives a significant result (Chapter 7). Exploratory findings are still welcome, but they are reported as such. In a **registered report**, a journal reviews and accepts the plan before the data is collected, so publication does not depend on the results.

#### 17.6.4 - Citing software

R and its packages are the work of researchers who depend on being cited. The function `citation()` gives the recommended citation for R itself, and `citation("lme4")` for a package:

```
citation("lme4")
```

To cite lme4 in publications use:

Douglas Bates, Martin Maechler, Ben Bolker,  
Steve Walker (2015).  
Fitting Linear Mixed-Effects Models Using  
lme4. Journal of  
Statistical Software, 67(1), 1-48.  
doi:10.18637/jss.v067.i01.

A BibTeX entry for LaTeX users is

```
@Article{,
 title = {Fitting Linear Mixed-Effects
Models Using {lme4}},
 author = {Douglas Bates and Martin M{"a}
chler and Ben Bolker and Steve Walker},
 journal = {Journal of Statistical
Software},
 year = {2015},
 volume = {67},
 number = {1},
```

```
pages = {1--48},
doi = {10.18637/jss.v067.i01},
}
```

A methods section might say: “Analyses were carried out in R version 4.4.3 [R Core Team], with mixed models fitted using lme4 version 1.1.37 [Bates et al., 2015].” Chapter 18 adds the question of disclosing the use of AI tools.

#### In your field: psychology and medicine

Psychology learned the importance of these practices the hard way. In 2015, the Open Science Collaboration repeated 100 studies published in leading psychology journals. Of the original studies, 97% had reported significant results; of the replications, only 36% did, and the effects were on average about half as large [35]. Many causes were identified, among them small samples, flexible analyses, and the publication of only significant results. The response has been a wave of preregistration, registered reports, and data and code sharing, now expected by many journals. Medicine went through a similar change earlier: since 2005, leading medical journals have required clinical trials to be registered before the first patient is enrolled, and reporting guidelines such as CONSORT set out what a trial report must contain.

## 17.7 - Keeping a history with git (optional)

As a project grows, files multiply: `analysis.R`, `analysis_v2.R`, `analysis_final.R`, `analysis_final_really.R`. **Version control** replaces this with a single copy of each file and a complete history of every change. **git** is the standard version control system, and **GitHub** is a website for storing git projects online, sharing them, and working on them with others [36].

git is not required for anything in this book, and it has a learning curve, so this section is an introduction for when you are ready. A **repository** is a project folder whose history git keeps. A **commit** is a saved snapshot of the project, with a short message saying what changed (“Remove duplicate survey rows”), and any commit can be returned to. **Pushing** copies the commits to GitHub, which is also an off-site backup, and **pulling** brings down changes made elsewhere.

RStudio has a **Git pane** that does all of this with buttons: tick the changed files, click *Commit*, write a message, and click *Push*. The **usethis** package sets things up from the Console: `usethis::use_git()` turns the current project into a repository, and `usethis::use_github()` connects it to GitHub. Never commit private data: list data files in the project’s `.gitignore` file, and git will leave them out.



## 17.8 - Common misconceptions

Reproducibility is often mistaken for something narrower or more technical than it is.

- **“Reproducibility is only about sharing code.”** It covers every step from raw data to reported number, including the decisions made along the way and the software versions used.
- **“Copying a few numbers by hand is harmless.”** It is the most common way that a thesis falls out of step with its analysis.
- **“If each analysis choice is defensible, the result is sound.”** Choosing among defensible analyses after seeing the results inflates false positives; the choices must be made in advance or reported as exploratory.
- **“Removing names makes data anonymous.”** Combinations of ordinary variables can identify people; small cells must be protected before data is shared.

## 17.9 - Chapter review

### 17.9.1 - Summary

- Reproducibility is a matter of research integrity: a result that cannot be checked must be taken on trust. Reproducible research can be recomputed from the same data and code; copying results by hand is the most common way that documents fall out of step with the analysis.
- A Quarto document combines a YAML header, Markdown text, code chunks, and inline code. Rendering runs the code and places the results in the output: HTML, Word, PDF, slides, or a book.
- Chunk options control each chunk; labelled tables and figures are numbered and cross-referenced; citations come from a BibTeX file; equations use LaTeX notation; parameters produce several versions of one report. R Markdown works in much the same way.
- `renv` records and restores package versions; report the versions of R and key packages, and cite them.
- A Shiny app has a user interface of inputs and outputs, and a server that computes the outputs; reactivity updates only what depends on a changed input.
- Choosing among reasonable analyses after seeing the data inflates false positives (the garden of forking paths); preregistration prevents it.
- Open science means sharing data and code with a DOI and a licence, protecting participants from identification, preregistering confirmatory analyses, and citing software.
- git and GitHub keep the history of a project; they are optional but valuable as projects grow.

### 17.9.2 - Key terms

Reproducibility, replicability, Quarto, R Markdown, render, YAML header, Markdown, code chunk, inline code, chunk option, knitr, Pandoc, cross-reference, citation, BibTeX, citation style (CSL), LaTeX, output format, parameter, `renv`, lockfile, Shiny, user interface, server, input, output, reactivity, reactive expression, open science, DOI,

codebook, licence, anonymisation, small cells, synthetic data, garden of forking paths, preregistration, confirmatory analysis, exploratory analysis, registered report, version control, git, repository, commit, push, GitHub.

## 17.10 - Exercises

The [playground](#) has these and more, with hints and solutions.

1. Create the tiny Quarto document of this chapter, render it, then change one friend's sleep and render it again. Round the mean to one decimal place with inline code.
2. Turn one of your analyses from an earlier chapter into a Quarto document with a numbered table, a numbered figure, and a sentence with inline numbers. Render it to HTML and to Word.
3. Add two references to a `references.bib` file, cite them in the document, and switch the citation style with a `csl:` file.
4. Add a slider to the Shiny dashboard that chooses the semesters shown.
5. Look at the four smallest cells in the table of faculty, study mode, gender, and children. Decide which variable you would remove or group before sharing the data, and explain why.
6. Add a sixth analysis to `forking_paths()`, such as a t-test that excludes the five highest scores, and describe how the rate of false positives for any analysis changes.

## 17.11 - Further reading

- The Quarto guide at [quarto.org](https://quarto.org) covers every feature of this chapter, with examples.
- *Mastering Shiny* [37], free online, is the definitive introduction to Shiny.
- “Good enough practices in scientific computing” [34] is a short, practical list of habits for researchers who are not programmers.
- “Excuse me, do you have a moment to talk about version control?” [36] makes the case for git in research, in plain language.

## 18. Using AI with R

Much research data is text: answers to open questions, interview transcripts, clinical notes, documents. To analyse it quantitatively, researchers use **qualitative coding**: reading each text and assigning it to one of a set of themes defined in a **codebook**. Coding is an act of judgement, and judgements differ. A central methodological question is therefore how to know whether a coder, whether a person, a keyword list, or a computer program, codes reliably. The standard answer is to compare the coder with another, and to measure how far they agree beyond what chance alone would produce.

In the wellbeing study, the final survey asked one open question: “What has been your biggest challenge during your studies?” 535 students answered in their own words, and the eleventh research question (RQ11) asks what those challenges are. Elaf coded 200 answers by hand, which took her most of two weeks, and the question is whether an AI tool could code the rest reliably.

Artificial intelligence (AI) tools now appear everywhere in research: assistants that write and explain code, and language models that read, summarise, and classify text. This chapter looks at both uses. It explains what these tools are and why they make mistakes, shows how to use an AI assistant to write R code and how to check what it produces, and then calls a language model from R with the **ellmer** package to code the students’ answers, testing the results against the hand coding exactly as Chapter 12 tested a classifier. It ends with the rules for using AI responsibly in research.

💡 By the end of this chapter you will be able to

- Explain how the reliability of coding is measured, and why agreement must be corrected for chance.
- Explain in plain terms how large language models work, and why they make confident mistakes.
- Use an AI assistant to write, explain, and debug R code, and check its answers.
- Call a language model from R with **ellmer**, and get structured answers for many texts.
- Validate AI coding of text against human coding with accuracy and Cohen’s kappa.
- Use AI responsibly: protect data, keep work reproducible, and disclose its use.

Note

**Teaching & Practice:** [Interactive Lecture Slides](#) · [Browser Playground Exercises](#)

## 18.1 - The reliability of coding

Two coders who read the same answers will not always choose the same theme. Their raw agreement, the share of answers on which they agree, seems the natural measure of reliability, but it can mislead. When one theme is very common, two coders will often agree simply because both choose it most of the time, even if they are guessing. A small example shows how much this matters. Two coders read 20 answers, 16 of which one coder puts under Workload:

```

coder_a <- c(rep("Workload", 16),
rep("Family", 4))
coder_b <- c(rep("Workload", 14),
rep("Family", 2), # the first 16 answers
rep("Workload", 2), rep("Family",
2)) # the last 4 answers

observed <- mean(coder_a == coder_b)
chance <- sum(prop.table(table(coder_a)) *
prop.table(table(coder_b)))
c(observed = observed, chance = chance, kappa
= (observed - chance) / (1 - chance))

```

observed	chance	kappa
0.800	0.680	0.375

The coders agree on 80% of the answers, which sounds good. But each coder puts 80% of the answers under Workload, so two coders who assigned themes at random in those proportions would agree on 68% of them by chance alone. **Cohen's kappa** measures agreement beyond chance: the observed agreement minus the chance agreement, as a share of the most that could be gained over chance [38]. Here it is only 0.37. A kappa of 0 means no better than chance, and 1 means perfect agreement. A common rule of thumb calls 0.41 to 0.60 moderate, 0.61 to 0.80 substantial, and above 0.80 almost perfect agreement [39]. The function `kap()` from the `yardstick` package gives the same result:

```

kap(tibble(a = factor(coder_a), b =
factor(coder_b)), truth = a, estimate = b)

```

```

A tibble: 1 × 3
 .metric .estimator .estimate
 <chr> <chr> <dbl>
1 kap binary 0.375

```

Kappa is the standard measure of **inter-rater reliability** (Chapter 5), and it applies to any pair of coders. Later in this chapter, one of the coders is a keyword list and then a language model, and the other is Elaf.

## 18.2 - How AI assistants work

The AI assistants that researchers use, such as ChatGPT, Claude, Gemini, and Copilot, are built on **large language models** (LLMs). A language model is a neural network (Chapter 15), with billions of weights, trained on an enor-

mous amount of text: books, websites, and code. During training it learns to do one thing extremely well: predict the next word (strictly, the next **token**, a word or part of a word) given the text so far. Answering a question means generating a reply one token at a time, each time choosing a likely continuation. Further training on examples of helpful answers turns this into an assistant that follows instructions.

This explains both their strengths and their weaknesses. They are fluent in R, because a huge amount of R code and documentation was in their training text, so they can write, explain, and fix code well, especially for common tasks. But they produce what is *plausible*, not what is *verified*: a made-up function name that looks like a real one, a wrong argument, or an outdated way of doing something can be delivered with complete confidence. Such invented content is called a **hallucination**. Their knowledge also stops at the date their training data was collected, so packages that changed since then, as `tidymodels` and `ggplot2` have, may be described as they used to be. And they know nothing about your data unless you tell them, and cannot run your code unless the tool is designed to do so.

The practical rule follows directly: **use AI to draft, and check everything it produces**. The checking is your job, and you remain responsible for the result.

## 18.3 - AI as a coding assistant

### 18.3.1 - Asking good questions

An assistant can only be as specific as the question. A good request for R code states **the goal** in words, such as “For each faculty, I want the percentage of students who considered dropping out”. It describes **the data**: the names and types of the relevant columns, for which the output of `glimpse(students)` (Chapter 3) is ideal, although real, identifiable data should never be pasted. It names **the tools**, such as “Use `dplyr` with the native pipe `|>`”, so that the answer matches a familiar style. And for errors, it gives the exact error message and the code that produced it.

### 18.3.2 - Always check the answer

Suppose an assistant is asked for the percentage of students in each faculty who considered dropping out, and suggests:

```
students |>
 summarise(percent = 100 *
sum(considering_dropout == "Yes") /
nrow(students),
 .by = faculty)
```

```
 faculty percent
1 Health Sciences 4.833333
2 Education 2.666667
3 Social Sciences 3.500000
```

```
4 Humanities 2.166667
5 Natural Sciences 1.833333
```

The code runs without an error and the numbers look reasonable. But they add up to only 15, the percentage of *all* students who considered dropping out: `nrow(students)` counts every student, not the students in each faculty. The question asked for the share *within* each faculty:

```
students |>
 summarise(percent = 100 *
 mean(considering_dropout == "Yes"),
 .by = faculty)
```

```
 faculty percent
1 Health Sciences 18.83117
2 Education 10.81081
3 Social Sciences 18.10345
4 Humanities 13.68421
5 Natural Sciences 12.64368
```

The mean of a TRUE/FALSE condition is the proportion of TRUE values (Chapter 6), calculated here within each faculty. The first answer was not a hallucination; it was a plausible misunderstanding, the most dangerous kind of error, because nothing looks wrong. Four checks catch most such errors. Code can be **run on a tiny example** whose answer is known, as every chapter of this book does. A number can be **checked by another route**; here, `students |> count(faculty, considering_dropout)` gives the counts to check against. The **help page** of any unfamiliar function (`?summarise`) confirms that it exists and does what the assistant says. And the assistant can be **asked to explain** each line, to see whether the explanation matches what was wanted.

Assistants are also excellent teachers: “explain this code line by line”, “why does this give NA?”, and “what does this error mean?” are among their most useful questions.

### 18.3.3 - AI inside the editor

Assistants can also work inside RStudio and Positron. **GitHub Copilot**, which can be switched on in RStudio’s options, suggests code as you type, completing a line or a whole block from a comment such as `# plot wellbeing by semester`. **Positron Assistant** and similar tools add a chat that can see your open files and, with permission, your data’s structure. Such suggestions arrive quickly and look authoritative, so the same rule applies: read each suggestion before accepting it. These tools change fast; the ideas in this chapter apply to whichever one you use.

## 18.4 - Coding the open-ended answers

A few of the answers show what the coding must handle:

```
open_responses |>
 slice(c(5, 9, 15, 55)) |>
 pull(biggest_challenge)
```

```
[1] "most of my income goes on
tuition fees. Besides that, I work
alone all the time. Things are slowly
improving."
[2] "I moved to a city where I
know no one for my studies and I
feel lonely. And I never have enough
time."
[3] "Supervisor."
[4] "The direction of my thesis changed twice
after comments from my supervisor. Sometimes
I rewrite the same section again and again
without knowing if it is right."
```

They vary in length and style, often mention more than one challenge, and sometimes contain spelling mistakes, as real answers do. The codebook has seven themes, and Elaf coded a random sample of 200 answers into the theme that each student presents as their main challenge:

```
open_responses_coded |>
 count(theme, sort = TRUE)
```

```
 theme n
1 Supervision 68
2 Workload 45
3 Isolation 31
4 Health 20
5 Family 14
6 Finances 14
7 Other 8
```

The hand-coded sample is the benchmark for any automatic method, exactly like the test set of Chapter 11.

### 18.4.1 - Coding by keywords

The simplest automatic method is a **dictionary**: a list of keywords for each theme. Each answer is assigned to the theme whose keywords it contains most often, and to “Other” if it contains none. The function `str_count()` from `stringr` counts the matches of a **regular expression**, a text pattern in which `|` means “or” and `\\b` marks the start of a word:

```
keywords <- c(
 Supervision = "\\b(supervis|feedback|
guidance|meeting)",
 Workload = "\\b(time|deadline|workload|
too much|assignment|reading|busy)",
 Finances = "\\b(money|fee|scholarship|
rent|afford|salary|income|funding|stipend|
loan)",
 Family = "\\b(famil|child|kid|son\\b|
daughter|baby|parent|mother|father|husband|
wife)",
 Health = "\\b(sleep|tired|stress|
```

```

anxi|health|ill\\b|burnout|headache|coffee|
exhaust)",
 Isolation = "\\b(lonel|alone|friends|
outsider|miss (home|my)|far from|know no
one)",
 Other = "\\b(ethic|participant|
statistic|library|procedure|english|power
cut|laptop|internet|software)"
)

code_by_keywords <- function(answer) {
 hits <- sapply(keywords, \(pattern)
str_count(str_to_lower(answer), pattern))
 if (all(hits == 0)) "Other" else
names(keywords)[which.max(hits)]
}

code_by_keywords("I miss my family and I
cannot afford the fees.")

```

```
[1] "Finances"
```

The example answer matches one Family keyword (“family”), one Isolation phrase (“miss my”), and two Finances keywords (“afford” and “fees”), so it is coded as Finances, although a human reader might well code it differently. Applying the function to the hand-coded answers:

```

themes <- names(keywords)

keyword_check <- open_responses_coded |>
 mutate(keyword_theme =
sapply(biggest_challenge, code_by_keywords),
 theme = factor(theme, levels = themes),
 keyword_theme = factor(keyword_theme,
levels = themes))

keyword_check |> accuracy(truth = theme,
estimate = keyword_theme)

```

```

A tibble: 1 × 3
 .metric .estimator .estimate
<chr> <chr> <dbl>
1 accuracy multiclass 0.61

```

```
keyword_check |> kap(truth = theme, estimate
= keyword_theme)
```

```

A tibble: 1 × 3
 .metric .estimator .estimate
<chr> <chr> <dbl>
1 kap multiclass 0.512

```

The keywords agree with the hand coding for 61% of the answers, and their Cohen’s kappa, the agreement beyond chance explained at the start of the chapter, is 0.51: moderate agreement. They fail because meaning is not in single words: “my supervisor” in “I feel I am letting down both my family and my supervisor” is about family, and “I never have enough time” can be about workload or about



children. Reading for meaning is exactly what language models are good at.

### 18.4.2 - Calling a language model from R

The **ellmer** package connects R to language models of two kinds.

**Online models** from providers: `chat_anthropic()` for Claude, `chat_openai()` for ChatGPT's models, and `chat_google_gemini()` for Gemini. They are the most capable, but using them from code requires an **API key**, a secret password linked to your account, and each request costs a small amount. Store the key in your personal `.Renv` file (`usethis::edit_r_environ()` opens it), never in a script that others may see:

```
ANTHROPIC_API_KEY=your-key-here
```

**Local models**, which run on your own computer through **Ollama** ([ollama.com](https://ollama.com)), a free program that downloads and runs open models. They need no key and cost nothing, and the text never leaves your computer. Install Ollama, download a model once in the Terminal (for example `ollama pull gemma4:e4b`, a model of about 10 GB), and use `chat_ollama()`. A computer with a graphics card (GPU) makes them much faster.

Either way, a chat then works much like a chat website, from R:

```
library(ellmer)
chat <- chat_ollama(model = "gemma4:e4b")
local
chat <- chat_anthropic(model = "claude-sonnet-5") # online, needs a key
chat$chat("In one sentence, what does Cohen's kappa measure?")
```

The `model` argument names the exact model. Always set it explicitly and record it: models are updated and retired, and different models give different answers.

For the study, a local model was chosen, Google's Gemma 4 (the **e4b** version), for two reasons: the students' answers never leave the computer, and anyone can repeat the analysis for free.

### 18.4.3 - Structured answers for many texts

Coding needs two things that a chat on a website does not give: the same instructions for every answer, and a reply in a fixed format that R can use. The **system prompt** holds the instructions: the task and the codebook, with a definition of each theme. With an online model, a **type** describes the required answer: here, an object with one field, `theme`, which must be one of the seven themes:

```
codebook <- "
You are helping a researcher code open-ended
survey answers from graduate
students, who were asked: 'What has been your
```

```

biggest challenge during your
studies?' Assign each answer to exactly ONE
theme: the main challenge the
student describes. If several challenges are
mentioned, choose the one the
student presents first or as most important.

Themes:
- Supervision: the supervisor; feedback,
guidance, meetings, disagreements.
- Workload: too much work or too little time;
deadlines, coursework, reading.
- Finances: money, fees, scholarships, living
costs, paid work to pay for study.
- Family: children, partners, parents, caring
duties, studying at home.
- Health: physical or mental health, sleep,
stress, anxiety, burnout, illness.
- Isolation: loneliness, being far from home,
no colleagues or friends.
- Other: anything else, such as ethics
approval, statistics, academic English.
"

theme_type <- type_object(
 theme = type_enum(themes, "The single main
theme of the answer.")
)

chat <- chat_anthropic(system_prompt =
codebook, model = "claude-sonnet-5",
 params = params(temperature
= 0))

ai_codes <- parallel_chat_structured(
 chat,
 prompts
= as.list(open_responses$biggest_challenge),
 type = theme_type
)

```

The function `type_enum()` restricts the answer to the listed themes, so the model cannot invent a new one or reply with a paragraph. The setting `temperature = 0` asks the model to choose its most likely answer every time, rather than sampling more freely, which makes the results as repeatable as possible. `parallel_chat_structured()` sends one request per answer, several at a time, and returns a data frame with one row per answer and a `theme` column.

Smaller local models do not always follow such a format. When `parallel_chat_structured()` was tried with the local model, it answered correctly (“Finances”, “Workload”) but as plain words, not in the structured format, and `ellmer` could not read the replies. The solution is to add one line to the end of the codebook, “Reply with the name of the theme only”, to ask for plain text with `parallel_chat_text()`, and to let R check each reply against the list of themes:

```
chat <- chat_ollama(system_prompt = codebook,
model = "gemma4:e4b",
 params = params(temperature
= 0))

replies <- parallel_chat_text(chat,
as.list(open_responses$biggest_challenge))
ai_codes <- themes[match(tolower(gsub("[^A-
Za-z]", "", replies)), tolower(themes))]
```

The function `gsub()` removes anything that is not a letter, such as a full stop, and `match()` finds the reply in the list of themes, ignoring capital letters. A reply that is not one of the themes becomes `NA`, so it is counted rather than silently accepted. Checking a model's output in code like this is good practice with any model.

The full script, `data-raw/run_ai_coding.R` in the book's repository, also records the provider, the model, the date, the software versions, and the time taken, and codes the 200 hand-coded answers a second time to check that the model gives the same answers when asked again. It was run once, and the results were saved to a file, so that this chapter can analyse them without calling the model again.

#### 18.4.4 - Agreement with the hand coding

The script used the model `gemma4:e4b`, run locally through Ollama, on 24 September 2026. The 735 requests took about 3 minutes, at no cost, and every reply was one of the seven themes. The results are joined to the hand-coded answers:

```
ai <- read.csv(ai_file)

ai_check <- open_responses_coded |>
 left_join(ai, join_by(student_id)) |>
 mutate(across(c(theme, ai_theme,
ai_theme_rerun), \(x) factor(x, levels =
themes)))

ai_check |> accuracy(truth = theme, estimate
= ai_theme)
```

```
A tibble: 1 × 3
 .metric .estimator .estimate
 <chr> <chr> <dbl>
1 accuracy multiclass 0.9
```

```
ai_check |> kap(truth = theme, estimate =
ai_theme)
```

```
A tibble: 1 × 3
 .metric .estimator .estimate
 <chr> <chr> <dbl>
1 kap multiclass 0.874
```

The model agrees with the hand coding on 90% of the answers, with a kappa of 0.87 (almost perfect agreement),

against 61% and 0.51 for the keywords. The confusion matrix of Chapter 12 shows where the two disagree:

```
ai_check |> conf_mat(truth = theme, estimate
= ai_theme)
```

	Truth			
Prediction	Supervision	Workload	Finances	
Family Health	Isolation	Other		
Supervision		60	2	0
0	0	2	0	
Workload		2	39	0
0	1	0	0	
Finances		1	0	14
0	0	0	0	
Family		2	1	0
14	1	0	0	
Health		2	1	0
0	16	0	0	
Isolation		1	2	0
0	1	29	0	
Other		0	0	0
0	1	0	8	

The diagonal holds the answers on which the hand coding and the model agree. The disagreements are worth reading one by one, because they show whether the model is wrong, the codebook is unclear, or the answer is genuinely ambiguous:

```
ai_check |>
 filter(theme != ai_theme) |>
 select(Answer = biggest_challenge, `Hand
coding` = theme, Model = ai_theme) |>
 head(5) |>
 knitr::kable()
```

Answer	Hand coding	Model
Getting feedback on my draft took over a month, and by then I had lost momentum. Also, home is not a quiet place to study.	Supervision	Isolation
Combining coursework, my job at a school and my thesis leaves me no time to rest. And my mental health has suffered. Things are slowly improving.	Health	Workload
I stopped exercising and I can feel the difference. The doctor told me to slow down, but I do not see how. On top of that, the university paperwork is slow. Things are slowly improving.	Other	Health
Writing my literature review took far longer than I planned. And I do not have friends in the department. I am coping, but only just.	Isolation	Workload
Nobody in my group works on anything close to my topic. At the same time, my supervisor is hard to reach. Otherwise the programme is fine.	Supervision	Isolation

Answer	Human coding	Model coding
		vi-sion

Look especially at answers that mention two challenges. If the model and the human coder chose different ones as the main challenge, that is a question for the codebook, not only for the model: a clearer rule (“code the challenge mentioned first”) helps both human and machine coders.

Asked a second time, the model gave exactly the same theme for every answer: with a temperature of 0, this local model is repeatable. That is not guaranteed for every model or setting (online models, in particular, can be updated between runs), which is why the outputs are saved and analysed from the file.

### 18.4.5 - What the students said

With its agreement checked, the model’s coding is used for all 535 answers:

```
ai |>
 count(ai_theme) |>
 mutate(share = n / sum(n)) |>
 ggplot(aes(x = share, y = reorder(ai_theme,
share))) +
 geom_col(fill = "#2f6793") +
 scale_x_continuous(labels =
scales::label_percent()) +
 labs(x = "Share of answers", y = NULL) +
 theme_minimal(base_size = 12)
```




Figure 18.1: Main challenge described by the students, coded by a language model.

Because the themes are now data, they can be linked to the rest of the study. Students whose main challenge is supervision should report less supervisor support in the questionnaire:

```
q_support <- questionnaire |>
 mutate(support =
rowMeans(pick(support_1:support_6), na.rm =
TRUE)) |>
 select(student_id, support)
```

```
ai |>
 left_join(q_support, join_by(student_id)) |>
 summarise(students = n(), support =
 round(mean(support, na.rm = TRUE), 2),
 .by = ai_theme) |>
 arrange(support)
```

	ai_theme	students	support
1	Supervision	177	2.82
2	Isolation	61	3.30
3	Other	23	3.31
4	Health	56	3.33
5	Family	47	3.43
6	Finances	49	3.45
7	Workload	122	3.46

This is a check of **validity**: the text and the numbers tell the same story. It also answers RQ11 in a way neither source could alone: what students struggle with, in their own words, and how that relates to their situation.

## 18.5 - Using AI responsibly

### 18.5.1 - Check, and stay responsible

Everything an AI tool produces is a draft. Check code by running it on known cases, check facts and references against their sources (assistants are known to invent plausible references that do not exist), and check coding against human coding, as this chapter did. Whatever appears in your thesis is your responsibility, whoever or whatever drafted it.

### 18.5.2 - Protect your data

Text sent to an online AI service leaves your computer and is processed, and possibly stored, by the provider. Before research data is sent, the **ethics approval and consent forms** must be checked: participants who agreed to have their answers read by the research team did not necessarily agree to have them sent to a company. The provider's terms matter too; many offer research or enterprise agreements under which data is not stored or used for training, but free accounts often do not. **Identifying information** must be removed before anything is sent: names, places, and details that could identify a person (Chapter 17). For sensitive data, a **local model** that runs on your own computer, through Ollama and `chat_ollama()`, keeps the data on the machine; local models are smaller and usually less accurate, so they must be validated in the same way.

The answers in the wellbeing study are anonymous, but a local model was used anyway, so they never left the computer: the simplest way to stay within any consent form.

### 18.5.3 - Keep it reproducible

Language models can give different answers to the same question on different days, and models are updated and retired. To keep AI-assisted work reproducible, the provider, the exact model name, the date, and the settings (such as the temperature) should be recorded, and the prompts (the system prompt and codebook) saved with the code. The model's outputs should be saved too, as in this chapter, and the saved outputs analysed rather than the model called again. Coding a sample twice checks consistency.

### 18.5.4 - Disclose it

Journals and universities increasingly require authors to disclose how they used AI. The consensus of publishers and of the Committee on Publication Ethics (COPE) is that an AI tool cannot be an author, because it cannot take responsibility for the work, and that its use must be described. A methods section might say: "Answers were coded into seven themes by a large language model (Gemma 4, e4b version, run locally with Ollama and the ellmer R package, temperature 0) using the codebook in Appendix X. Agreement with the author's hand coding of a random sample of 200 answers was assessed with Cohen's kappa." For help with code or language, a sentence in the acknowledgements or methods is usually enough; check your university's and journal's policy.

### 18.5.5 - Be aware of bias

Language models learn from human text, and they can reproduce its biases: in how they describe groups of people, in which answers they find typical, and in how well they understand non-standard English or answers written by non-native speakers. Check whether the model's accuracy differs between groups, as Chapter 11 recommended for any model that makes decisions about people.

## In your field: medicine and public health

Clinical notes, discharge letters, and incident reports hold information that is written as free text but needed as data: diagnoses, medications, doses, and dates. Language models can extract it into a fixed structure. With `ellmer`, the type describes the fields to extract:

```
medication_type <- type_object(
 drug = type_string("Name of the
medication"),
 dose_mg = type_number("Dose in
milligrams"),
 times_per_day = type_integer("How many
times a day it is taken")
)

chat <- chat_anthropic(model = "claude-
sonnet-5")
chat$chat_structured(
 "Patient to continue metformin 500 mg
twice daily with meals.",
 type = medication_type
)
```

The result is a list with `drug`, `dose_mg`, and `times_per_day`, ready to become a row of a data frame. Studies in this area validate the extraction against records checked by clinicians, report agreement for each field, and use local models or secure institutional services, because patient data must not be sent to public AI services.

## 18.6 - Common misconceptions

AI tools invite both too much trust and too little, and a few misunderstandings are especially common.

- **“High agreement means reliable coding.”** Raw agreement can be high by chance when one theme is common; kappa shows the agreement beyond chance.
- **“The AI knows the answer.”** A language model produces plausible text, not verified facts; its coding must be validated against human coding.
- **“A model’s answers are always the same.”** They can change with the model version, the settings, and the date, which is why models, prompts, and outputs are recorded and saved.
- **“Using AI must be hidden, or is not allowed.”** Most journals and universities allow it with disclosure; what they do not allow is presenting AI output as unchecked fact, or listing an AI as an author.

## 18.7 - Chapter review

### 18.7.1 - Summary

- The reliability of coding is measured by agreement between coders, corrected for chance with Cohen’s



kappa; raw agreement can look high even when coders agree little beyond chance.

- Large language models generate text by predicting the next token; they are fluent in R but produce what is plausible, not what is verified. They hallucinate, they may be out of date, and they know nothing about your data unless told.
- Ask assistants specific questions (goal, data structure, tools, exact error), and check every answer: run it on a tiny example, verify numbers by another route, and read the help.
- A dictionary of keywords is a simple, transparent baseline for coding text, but it misses meaning.
- `ellmer` calls language models from R. Store the API key in `.Renviron`; set the model explicitly; use a system prompt with a codebook, a type to fix the format of the answer, and `parallel_chat_structured()` for many texts.
- Validate AI coding against human coding with accuracy, a confusion matrix, and Cohen's kappa, and examine the disagreements.
- Use AI responsibly: you are responsible for the result; protect participants' data (consent, anonymisation, local models); record models, prompts, and outputs; disclose AI use; and watch for bias.

### 18.7.2 - Key terms

Artificial intelligence, large language model, token, hallucination, training cut-off, prompt, system prompt, AI coding assistant, qualitative coding, codebook, interrater reliability, dictionary method, regular expression, Cohen's kappa, API, API key, structured output, temperature, local model, disclosure.

## 18.8 - Exercises

The [playground](#) has these and more, with hints and solutions.

1. Ask an AI assistant to write code that calculates the average sleep in each faculty and semester. Check the answer with a tiny example and by another route, and check whether it handled missing values.
2. Add three keywords to the dictionary that you think would fix some of its mistakes, and report whether the kappa improves. Discuss whether you are now fitting the dictionary to these answers (Chapter 11's overfitting).
3. Read ten answers where the keyword coding disagrees with the hand coding, and say whether you would agree with the hand coding in every case.
4. Code 20 of the hand-coded answers with a language model (local or online), using the codebook in this chapter, and count how many agree with the hand coding. Then change one theme definition and code them again.
5. Write the paragraph for your own methods section describing how you would use and validate AI coding in a study of your own.

6. In the two-coder example at the start of the chapter, change the second coder so that it puts every answer under Workload. Calculate the raw agreement and kappa, and explain the result.

## 18.9 - Further reading

- The `ellmer` documentation at [ellmer.tidyverse.org](https://ellmer.tidyverse.org) covers providers, structured data, and tools, with many examples.
- Gilardi and colleagues compared language models with human coders for text annotation [40], a study that prompted much of the current interest in using them for research.
- Cohen's original paper on kappa [38] and Landis and Koch's benchmarks [39] are the standard references for agreement between coders.

## 19. Putting It All Together

A thesis is an argument, and its results chapter is the evidence. Each claim in it rests on a chain of reasoning that runs from a research question, through a hypothesis, a design, measured variables, and an analysis, to a result, a conclusion, and the limits of that conclusion. A thesis convinces when every link in that chain can be seen and checked, and every number in the text can be traced back to the raw data. The methods of this book are the links; this chapter joins them.

In the story, it is the final semester. Elaf's analyses are spread over two years of scripts, written as she learned each method, and her supervisor's advice for the last stretch is simple: "Before you write the results chapter, rebuild the whole thing, from the raw export to the last table, as one project that runs from start to finish. Then you will know that every number in your thesis is right, and you can answer any examiner's question by pointing to the code." This chapter does exactly that. It sets up a complete research project, runs the analysis from the messy survey export to the tables, figures, and sentences of a results chapter, and traces the chain of reasoning behind three of the research questions. It then steps back: how to report results to the standards journals expect, how to choose a method for a new question, the problems that every researcher meets, and where to go next.

💡 By the end of this chapter you will be able to

- Organise a research project so that it runs from raw data to finished results.
- Combine cleaning, description, modelling, and visualisation in one reproducible pipeline.
- Trace the chain of reasoning from research question to conclusion and limitation.
- Report results in the style expected in a thesis, following published reporting standards, with numbers taken directly from the code.
- Choose an appropriate method for a research question, using the map of this book.
- Recognise common challenges in real research projects, and know where to learn more.

Note

**Teaching & Practice:** [Interactive Lecture Slides](#) · [Browser Playground Exercises](#)

## **19.1 - The shape of a research project**

Every quantitative project follows roughly the same path, and this book has followed it too (Figure 19.1).




Figure 19.1: The path of a research project, with the chapters of this book that cover each step.

The path is not a straight line in practice. Exploring the data sends you back to cleaning when you find a problem; a model's diagnostics send you back to exploring. What matters is that each step is written in code, so that going back and rerunning everything is easy.

### 19.1.1 - Organising the project

A project that others (and your future self) can follow has a predictable structure. The final thesis project looks like this:

```
wellbeing-thesis/
├─ wellbeing-thesis.Rproj the RStudio
Project (Chapter 1)
├─ README.txt what the
project is and how to run it
├─ data-raw/
│ └─ wellbeing_raw.xlsx the survey
export, never edited by hand
├─ data/ clean data,
created by the scripts
├─ R/
│ ├── 01-clean-data.R raw export ->
clean tables (Chapter 3)
│ └─ 02-analysis.R models
and figures
├─ output/ figures and
tables, created by the scripts
├─ results.qmd the results
chapter (Chapter 17)
├─ references.bib
└─ renv.lock package versions
(Chapter 17)
```

Four principles lie behind it. **Raw data is read-only:** everything in `data/` and `output/` can be deleted and recreated by running the scripts. **Scripts are numbered** in the order they run, and each does one job. **Paths are relative to the project**, using `here::here()` (Chapter 1), so the project runs on any computer. And **the README** says in a few lines what the project is and how to run it.

The downloadable project for this chapter has the same structure, with working scripts (its `data/` and `output/` folders are created when the scripts run, and it has no `renv.lock`, which you create with `renv::init()` for your own project). The rest of the chapter walks through what the scripts do.

## 19.2 - From raw export to clean data

The cleaning of Chapter 3, gathered into one script, runs in a few seconds. Here is its core, reading the raw export and producing the clean questionnaire and semester tables:

```
library(dplyr)
```

```
Attaching package: 'dplyr'
```

The following objects are masked from 'package:stats':

filter, lag

The following objects are masked from 'package:base':

intersect, setdiff, setequal, union

```
library(tidyr)
library(stringr)
library(readr)
library(readxl)

raw <- read_excel(data2thesis_example("wellbeing_raw.xlsx"))

item_names <- c(paste0("stress_", 1:6),
 paste0("burnout_", 1:6),
 paste0("support_", 1:6),
 paste0("satisfaction_", 1:4))

responses <- raw |>
 filter(!str_detect(str_to_upper(`Q1_Student
ID`), "^TEST")) |>
 select(-`Response ID`) |>
 distinct() |>
 rename(student_id = `Q1_Student ID`,
 workshop = `Workshop group`,
 considering_dropout = `Y1_Considered
leaving?`) |>
 rename_with(~ item_names, .cols =
Q11_1:Q11_22) |>
 mutate(across(all_of(item_names), ~
as.integer(na_if(.x, "99"))))

questionnaire_clean <- responses |>
 select(student_id, all_of(item_names))

semesters_clean <- responses |>
 select(student_id, matches("_S[1-4]$")) |>
 pivot_longer(-student_id, names_to =
c(".value", "semester"), names_sep = "_S") |>
 rename(gpa = GPA, sleep_hours = Sleep,
study_hours = Study,
 exercise_days = Exercise, caffeine_mg
= Caffeine,
 supervisor_meetings = Meetings,
wellbeing = Wellbeing) |>
 mutate(
 semester = as.integer(semester),
 sleep_hours =
parse_number(str_replace(sleep_hours, " ",
".")),
 across(c(gpa, study_hours, exercise_days,
caffeine_mg, supervisor_meetings, wellbeing),
 parse_number),
 sleep_hours = if_else(sleep_hours > 24,
NA, sleep_hours),
 study_hours = if_else(study_hours > 168,
NA, study_hours),
```

```
gpa = if_else(gpa > 4, NA, gpa)
) |>
filter(!if_all(gpa:wellbeing, is.na))
```

Each step is explained in Chapter 3; the full script in the downloadable project also cleans the background variables. The most important line of any cleaning script is the check at the end, here a comparison with the clean data:

```
same <- function(mine, theirs) {
 mine <- as.data.frame(mine)[, names(mine)]
 theirs <- as.data.frame(theirs)[,
names(mine)]
 isTRUE(all.equal(mine, theirs,
check.attributes = FALSE))
}
same(questionnaire_clean
|> arrange(student_id), questionnaire |>
arrange(student_id))
```

```
[1] TRUE
```

```
same(semesters_clean |> arrange(student_id,
semester), semesters |> arrange(student_id,
semester))
```

```
[1] TRUE
```

Both are TRUE. In a real project there is no package to compare with, so the checks are the ones from Chapters 3 and 6: counts of categories, ranges of values, numbers of missing values, and a look at a few rows. From here on, the chapter uses the clean tables, together with the scale scores:

```
scores <- questionnaire |>
mutate(
 stress_4 = 6 - stress_4,
 stress =
rowMeans(pick(stress_1:stress_6), na.rm =
TRUE),
 burnout =
rowMeans(pick(burnout_1:burnout_6), na.rm =
TRUE),
 support =
rowMeans(pick(support_1:support_6), na.rm =
TRUE)
) |>
select(student_id, stress, burnout, support)

study <- students |>
left_join(scores, join_by(student_id)) |>
left_join(semesters |> filter(semester ==
1), join_by(student_id))
```

The table `study` has one row per student, with background, questionnaire scores, and first-semester records, as in Chapter 8.



## 19.3 - Describing the sample

Every results chapter begins by describing the participants, often in a table known as “Table 1”. Here the sample is described by programme:

```
mean_sd <- function(x) sprintf("%.1f (%.1f)",
 mean(x, na.rm = TRUE), sd(x, na.rm = TRUE))
percent <- function(x) sprintf("%.0f%%", 100
 * mean(x, na.rm = TRUE))

describe <- function(d) {
 tibble(
 Characteristic = c("Students", "Age, mean
 (SD)", "Women", "Part-time",
 "Stress (1-5), mean
 (SD)", "Support (1-5), mean (SD)",
 "Wellbeing in semester
 1, mean (SD)", "Considered dropping out"),
 Value = c(nrow(d), mean_sd(d$age),
 percent(d$gender == "Female"),
 percent(d$study_mode == "Part-
 time"), mean_sd(d$stress),
 mean_sd(d$support),
 mean_sd(d$wellbeing),
 percent(d$considering_dropout
 == "Yes"))
)
}

describe(filter(study, programme ==
 "Master's")) |>
 rename(`Master's` = Value) |>
 left_join(describe(filter(study, programme
 == "PhD")) |> rename(PhD = Value),
 join_by(Characteristic)) |>
 left_join(describe(study) |> rename(All =
 Value), join_by(Characteristic)) |>
 knitr::kable(align = "lrrr")
```

Table 19.1: Characteristics of the sample, by programme.

Characteristic	Master's	PhD	All
Students	426	174	600
Age, mean (SD)	27.9 (3.3)	34.2 (5.1)	29.7 (4.8)
Women	52%	53%	52%
Part-time	26%	39%	30%
Stress (1-5), mean (SD)	3.2 (0.7)	3.2 (0.7)	3.2 (0.7)
Support (1-5), mean (SD)	3.2 (0.8)	3.2 (0.7)	3.2 (0.8)
Wellbeing in semester 1, mean (SD)	60.9 (12.1)	59.3 (11.9)	60.5 (12.0)
Considered dropping out	15%	14%	15%

Two small helper functions, `mean_sd()` and `percent()`, format the numbers the way theses report them, and `describe()` builds the column for any group of students, so the same code makes all three columns. Writing a small function whenever you would otherwise copy and paste code is one of the best habits to take from this book. (The **gtsummary** package produces such tables automatically, with many options, if you prefer.)

## 19.4 - Answering the research questions

Three of the research questions are answered below with the methods of Parts 2 and 3, each with a sentence written the way it will appear in the thesis. A small function formats p-values in the usual style:

```
format_p <- function(p) if (p < 0.001) "p
< .001" else paste("p =", sub("^0", "",
sprintf("%.3f", p)))
```

### 19.4.1 - The workshop and its persistence (RQ3)

Students were randomly invited to the wellbeing workshop after the first semester. The mixed-effects model of Chapter 10 uses all four semesters and every student:

```
library(lme4)
```

```
Loading required package: Matrix
```

```
Attaching package: 'Matrix'
```

```
The following objects are masked from
'package:tidyr':
```

```
expand, pack, unpack
```

```
panel <- semesters |>
 left_join(students, join_by(student_id)) |>
 mutate(time = semester - 1)

workshop_model <- lmer(wellbeing ~
 factor(semester) * workshop + (time |
 student_id),
 data = panel)

gaps <- expand.grid(semester = 1:4, workshop
= c("Invited", "Not invited")) |>
 mutate(time = semester - 1)
gaps$wellbeing <- predict(workshop_model,
 newdata = gaps, re.form = NA)
gaps <- gaps |>
 pivot_wider(id_cols = semester, names_from
= workshop, values_from = wellbeing) |>
```

```
mutate(gap = Invited - `Not invited`)
gaps
```

```
A tibble: 4 × 4
 semester Invited `Not invited` gap
 <int> <dbl> <dbl> <dbl>
1 1 60.6 60.3 0.380
2 2 65.5 60.2 5.25
3 3 63.1 59.0 4.13
4 4 60.9 58.2 2.64
```

Wellbeing was similar in the two groups before the workshop (difference 0.4 points). After the workshop, invited students' wellbeing was 5.2 points higher in semester 2, but the difference narrowed to 4.1 points in semester 3 and 2.6 in semester 4.

### 19.4.2 - Explaining GPA (RQ5)

The multiple regression of Chapter 8, with **broom**'s `tidy()` giving the coefficients and their confidence intervals as a data frame:

```
library(broom)

gpa_model <- lm(gpa ~ sleep_hours +
 study_hours + stress + support, data = study)

gpa_table <- tidy(gpa_model, conf.int = TRUE)
gpa_table |>
 mutate(across(where(is.numeric), \ (x)
 round(x, 3)))
```

```
A tibble: 5 × 7
 term estimate std.error statistic
p.value conf.low conf.high
 <chr> <dbl> <dbl> <dbl>
<dbl> <dbl> <dbl>
1 (Intercept) 2.14 0.145 14.7
0 1.85 2.42
2 sleep_hours 0.106 0.014 7.49
0 0.078 0.134
3 study_hours 0.007 0.001 6.27
0 0.005 0.009
4 stress -0.084 0.018 -4.68
0 -0.119 -0.049
5 support 0.111 0.017 6.63
0 0.078 0.144
```

Each additional hour of sleep was associated with a GPA 0.11 points higher (95% CI 0.08 to 0.13,  $p < .001$ ), holding study hours, stress, and support constant. Together, the four predictors explained 25% of the variation in first-semester GPA ( $n = 587$ ).

### 19.4.3 - Considering dropout (RQ9)

The logistic regression of Chapter 8, with odds ratios:

```
dropout_model <- glm(
 I(considering_dropout == "Yes") ~ stress
+ support + financial_worry + employment +
 study_mode,
 data = study, family = binomial
)

dropout_table <- tidy(dropout_model, conf.int
= TRUE, exponentiate = TRUE)
dropout_table |>
 mutate(across(where(is.numeric), \ (x)
round(x, 2)))
```

```
A tibble: 7 × 7
 term estimate std.error
statistic p.value conf.low conf.high
<chr> <dbl> <dbl> <dbl> <dbl>
1 (Intercept) -3.26 0.19 0.02 1.23
2 stress 5.35 2.29 3.6 0.24
3 support -5.44 0.22 0.33 0.2
4 financial_worry 3.73 1.25 1.57 0.12
5 employmentNone -1.35 0.18 0.55 0.44
6 employmentPart-time j... -1.19 0.23 0.58 0.45
7 study_modePart-time 0.75 0.45 1.32 0.37
 0.75 0.45 0.63 2.68
```

The wrapper `I()` lets a condition be used directly as the outcome, and `exponentiate = TRUE` turns the log-odds into odds ratios and their confidence intervals.

Each one-point increase in stress (on the 1 to 5 scale) was associated with 3.6 times the odds of considering dropping out (95% CI 2.3 to 5.9), and each one-point increase in supervisor support with 0.33 times the odds (95% CI 0.22 to 0.49).

Chapters 11 to 15 went further with this question, asking how well dropout can be *predicted*; the thesis reports that logistic regression predicted as well as any machine learning model (a test AUC of about 0.84 in Chapter 11), which is itself a finding worth stating.

## 19.5 - One figure for the thesis

A thesis figure often combines panels. The **patchwork** package joins `ggplot2` plots with `+` (side by side) and `/` (one above the other), and `plot_annotation()` labels the panels:

```
library(ggplot2)
library(patchwork)

panel_a <- gaps |>
```

```

 pivot_longer(c(Invited, `Not invited`),
names_to = "workshop", values_to =
"wellbeing") |>
 ggplot(aes(x = semester, y = wellbeing,
colour = workshop)) +
 geom_line(linewidth = 1) +
 geom_point(size = 2.5) +
 scale_colour_viridis_d(end = 0.8) +
 labs(x = "Semester", y = "Predicted wellbeing
(0-100)", colour = "Workshop")

panel_b <- ggplot(study, aes(x = sleep_hours,
y = gpa)) +
 geom_point(alpha = 0.3) +
 geom_smooth(method = "lm", formula = y ~ x,
colour = "#2f6793") +
 labs(x = "Sleep (hours a night)", y = "GPA
(semester 1)")

(panel_a + panel_b) +
 plot_annotation(tag_levels = "A") &
 theme_minimal(base_size = 12)

```

Warning: Removed 6 rows containing non-finite outside the scale range (``stat_smooth()``).

Warning: Removed 6 rows containing missing values or values outside the scale range (``geom_point()``).




Figure 19.2: (A) Predicted wellbeing by semester and workshop group, from the mixed-effects model. (B) First-semester GPA against hours of sleep, with a regression line.

The operator `&` applies the theme to both panels, and `ggsave()` saves the figure at the size and resolution a thesis or journal requires (Chapter 4):

```

ggsave(here::here("output", "figure-1.png"),
width = 18, height = 8, units = "cm", dpi
= 300)

```

## 19.6 - The chain of reasoning

A result on its own is not yet a conclusion. Between the two lie the design that produced the data, the way the variables were measured, and the limits of both.

Table 19.2 traces the whole chain for the three research questions answered above, from the hypotheses stated in Chapter 5 to the limitations that a thesis discussion must acknowledge. Every result in it is filled in by the code of this chapter.

Table 19.2: The chain of reasoning for three research questions, from question to limitation

Step	Workshop (RQ3)	GPA (RQ5)	Considering dropout (RQ9)
<b>Question</b>	Does the workshop improve wellbeing?	What explains students' GPA?	Who considers dropping out?
<b>Hypothesis</b> (Chapter 5)	Invited students have higher wellbeing in semester 2	More sleep goes with a higher GPA, allowing for study hours, stress, and support	Higher stress raises the odds of considering dropout
<b>Design</b>	Randomised invitation within a longitudinal study	Observational, first semester	Observational, baseline and end of year 1
<b>Variables</b>	Wellbeing (0 to 100), invitation, semester	GPA, sleep, study hours, stress, support	Considering dropout (yes/no), stress, support, financial worry, employment, study mode
<b>Analysis</b>	Mixed-effects model (Chapter 10)	Multiple regression (Chapter 8)	Logistic regression (Chapter 8)
<b>Result</b>	5.2 points higher in semester 2, narrowing to 2.6 by semester 4	0.11 GPA points per hour of sleep (95% CI 0.08 to 0.13)	Odds ratio 3.6 per point of stress (95% CI 2.3 to 5.9)
<b>Conclusion</b>	The workshop raised wellbeing, and the effect faded; because the invitation was random, the effect is causal	Sleep is associated with GPA, independently of the other predictors; the null hypothesis is rejected, but the design does not show causation	Stress is associated with higher odds of considering dropout; the null hypothesis is rejected
<b>Limitation</b>	One university; the invitation,	Sleep does not self-report, and	The outcome is con-

Reading the table by columns shows how different the three conclusions are, although all three rest on “significant” results. Only the workshop conclusion is causal, because only the workshop was assigned at random (Chapter 5). The GPA and dropout conclusions are associations, and the limitations say what could still explain them. A discussion chapter that keeps each result attached to its design, in this way, claims exactly what the evidence supports.

## 19.7 - The results chapter

The final step is the one Chapter 17 prepared: the tables, figures, and sentences above go into a Quarto document, `results.qmd`, in which every number is written with inline code. The downloadable project contains such a document, built from this chapter.

What a results chapter must contain is not left to taste. **Reporting standards** list the information that readers need in order to judge a study. For quantitative research in psychology and neighbouring fields, the American Psychological Association’s Journal Article Reporting Standards (JARS) set out what to report about the participants, the measures, the analysis, and the results, including effect sizes and confidence intervals [41]. Randomised trials in health research follow CONSORT, and observational studies follow STROBE [42]. Checking a draft against the relevant standard is one of the most useful things a student can do before submission, and many journals require it. When the data changes, or an examiner asks for a different model, the code is changed and rendered again, and the whole chapter is updated.

## 19.8 - Choosing a method

The choice of method for a new research question depends on the goal, on the kind of **outcome**, and on how the **observations** are related. Figure 19.3 summarises the methods of this book as a guide.






Figure 19.3: A guide to choosing a method, with the chapters that cover each. Table 19.3 shows how each of the study's research questions was answered.

Table 19.3: The study's research questions and the methods used to answer them

Research question	Method	Chapters
RQ1 What does graduate life look like?	Plots, descriptive statistics	4, 6
RQ2 Do students sleep less than 7 hours?	One-sample t-test, confidence interval	7
RQ3 Does the workshop improve wellbeing?	Two-sample t-test; mixed-effects model	7, 10
RQ4 Do faculties and study modes differ?	ANOVA, post-hoc tests	7, 8
RQ5 What explains GPA?	Multiple regression	8, 13
RQ6 Does the questionnaire measure what it should?	Factor analysis, Cronbach's alpha	9
RQ7 Are there student profiles?	Cluster analysis, mixture models	9, 14
RQ8 How do wellbeing and GPA change?	Mixed-effects models	10
RQ9 Who considers dropping out?	Logistic regression; classification	8, 11, 12, 15
RQ10 Can final GPA be predicted?	Regularised regression, boosting	13
RQ11 What challenges do students describe?	Text coding with a language model	18
RQ12 How many counselling visits next year?	Time series forecasting	16

## 19.9 - Challenges every researcher meets

Real projects rarely go as planned, and some problems are almost universal. Every real dataset is **messy** and needs cleaning, which takes longer than the analysis;

it is done in code, checked after every step, and never applied to the raw file itself (Chapter 3). **Missing data** raises the question of *why* it is missing before anything is done about it, because dropping incomplete cases can bias results when the missing data is not random, as with the students who left the programme (Chapter 6). **Small samples** give wide confidence intervals and low power (Chapter 7) and make complex models overfit (Chapter 13); the remedies are to report effect sizes with intervals, keep models simple, and plan the sample size before collecting data with a power analysis (Chapter 5).

Analysis brings its own problems. **Many tests** produce many false alarms, so the main analyses are decided in advance, multiple comparisons are corrected where needed, and exploratory results are reported as exploratory (Chapters 7, 8, and 17). **Assumptions** are checked with plots rather than with tests alone, with robust tests, non-parametric tests, transformations, and mixed models as the alternatives (Chapters 7, 8, and 10). **Correlation is not causation**: observational data rarely proves causes, so confounders must be thought through, as sleep was behind the caffeine effect, and randomised designs used where possible, as in the workshop study (Chapters 5 and 8). **Prediction and explanation** are different aims: a model that predicts well may explain little, and the reverse (Chapter 11). Finally, results must be **communicated** to supervisors, co-authors, and examiners who may use other software; a clear results chapter with tables, figures, and an appendix of code serves them all, and data can be exported with `write_csv()` or `haven::write_sav()` when they need it (Chapter 2).

## 19.10 - Where to go next

This book is a beginning. Depending on your field, the next methods to learn may be:

Table 19.4: Topics for further study

Topic	What it is for	Packages
Structural equation modelling	Confirmatory factor analysis, path models, latent variables	lavaan
Bayesian statistics	Models with prior knowledge and full uncertainty	brms, rstanarm
Survival analysis	Time until an event, such as dropping out	survival
Meta-analysis	Combining results from several studies	metafor
Text analysis	Words, sentiment, and topics in documents	tidytext, quanteda
Spatial data	Maps and geographic data	sf
Publication tables	Automatic, formatted tables of results	gtsummary, modelsummary
Interpreting models	Predictions and effects from any model	marginaleffects

Some of the best resources are free online: *R for Data Science* [2] for data skills, *Regression and Other Stories* [17] for regression, *An Introduction to Statistical Learning* [18] and *Tidy Modeling with R* [22] for machine learning, *Forecasting: Principles and Practice* [32] for time series, and *Mastering Shiny* [37] for dashboards. Appendix E lists more.

You do not have to learn alone. The R community is known for being welcoming to beginners: the Posit Community forum and Stack Overflow answer questions; **R-Ladies** and local R user groups run meetings around the world; **TidyTuesday** publishes a new dataset every week for people to practise on and share their plots; and *R Weekly* collects news and tutorials. Asking a clear question, with a small reproducible example, is itself a skill, and the same skill that makes an AI assistant useful (Chapter 18).

## In your field: nutrition science

The same path works for any dataset, in a few lines. R's `ToothGrowth` data comes from an experiment on the growth of teeth in 60 guinea pigs given vitamin C at three doses, as orange juice (OJ) or as a vitamin C supplement (VC). A complete mini-project, from data to a reported result:

```
tooth <- ToothGrowth |>
 mutate(dose = factor(dose))

tooth |>
 summarise(mean = mean(len), sd = sd(len),
 n = n(), .by = c(supp, dose))
```

	supp	dose	mean	sd	n
1	VC	0.5	7.98	2.746634	10
2	VC	1	16.77	2.515309	10
3	VC	2	26.14	4.797731	10
4	OJ	0.5	13.23	4.459709	10
5	OJ	1	22.70	3.910953	10
6	OJ	2	26.06	2.655058	10

```
tooth_model <- aov(len ~ supp * dose, data
 = tooth)
summary(tooth_model)
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
supp	1	205.4	205.4	15.572	0.000231 ***
dose	2	2426.4	1213.2	92.000	< 2e-16 ***
supp:dose	2	108.3	54.2	4.107	0.021860 *
Residuals	54	712.1	13.2		
---					
Signif. codes:	0	'***'	0.001	'**'	0.01
	'*'	0.05	'.'	0.1	' ' 1

The two-way ANOVA of Chapter 8 shows that tooth length depends on the dose, the form of vitamin C, and their combination: orange juice helps more than the supplement at low doses, but not at the highest. The same steps (prepare, describe, model, report) would take a few more lines with a plot and a Quarto report. Whatever your field, the path from data to thesis is the one this book has followed.

## 19.11 - Common misconceptions

A few misunderstandings are common at the stage of writing up.

- **“A significant result answers the research question.”** It answers one link of the chain; the design and the measures decide what the result means.

- **“The discussion should defend the results.”** It should also state their limits; examiners trust a thesis more when it says clearly what it cannot show.
- **“Formatting numbers by hand is quicker.”** It is quicker once, and wrong the next time the data or the model changes.
- **“Choosing a method means choosing the most advanced one.”** The right method is the simplest one that matches the goal, the outcome, and the structure of the data.

## 19.12 - Chapter review

### 19.12.1 - Summary

- A research project follows a path from question and design through import, cleaning, exploration, and modelling to reporting and sharing, often going back to earlier steps. Writing every step in code makes going back easy.
- Organise projects with read-only raw data, numbered scripts, relative paths, a README, and a Quarto results document; everything else can be recreated.
- Check the cleaned data at the end of the cleaning script, then describe the sample, answer each research question with the method suited to it, and report results with numbers taken directly from the code.
- Every result rests on a chain of reasoning, from question and hypothesis through design, variables, and analysis to result, conclusion, and limitation. Reporting standards such as the APA’s JARS, CONSORT, and STROBE list what a report must contain.
- Small helper functions (for formatting means, percentages, and p-values) avoid copying and pasting code; broom turns model results into data frames; patchwork combines plots.
- Choose methods by the goal (describe, explain, predict, find structure, forecast), the type of outcome, and how the observations are related.
- Messy and missing data, small samples, many tests, assumptions, and causation are challenges in every project; the chapters of this book give tools for each.

### 19.12.2 - Key terms

Research workflow, project structure, raw data, pipeline, sample description (“Table 1”), helper function, broom, patchwork, reporting sentence, chain of reasoning, reporting standard, method choice.

## 19.13 - Exercises

The [playground](#) has these and more, and the downloadable project contains the complete thesis project.

1. Download the Chapter 19 project, run `R/01-clean-data.R` and `R/02-analysis.R`, and render `results.qmd`. Then change one cleaning rule (for example, treat ages above 70 as impossible), render again, and note which numbers change.

2. Add a row to the sample table (Table 19.1) for the share of students with children.
3. Write a reporting sentence, with inline numbers, for the effect of support on GPA in the regression model.
4. Add a third panel to the thesis figure: wellbeing by faculty as a box plot.
5. Choose a research question from your own field. Using Figure 19.3, decide which method you would use, and which chapter of this book you would reread first.
6. Complete the chain of reasoning (Table 19.2) for RQ2, whether students sleep less than 7 hours, using the results of Chapter 7.

## 19.14 - Further reading

- *R for Data Science* [2] covers the whole workflow of this chapter in more depth, and is the natural next book for most readers.
- *Regression and Other Stories* [17] is an excellent guide to building, checking, and reporting regression models in real research.





## References



## Appendices



# A. Appendix

## A: Installation and Setup

This appendix explains how to install everything the book uses, and how to fix the problems that most often get in the way. If you only want to start, the short version is: install **R**, then **RStudio**, then run the package installation command in Section A.5. If you are not ready to install anything yet, every chapter's exercises also run in your web browser, in the [playground](#).

Installers and websites change their appearance from year to year, so this appendix describes each step in words rather than with screenshots. The choices that matter are the same in every version.

### A.1 - What you need

Table A.1: The software used in this book

Software	What it is	Required?
<b>R</b>	The programming language and the engine that runs your code	Yes
<b>RStudio Desktop</b>	The program you work in: editor, console, plots, and help in one window	Yes (or Positron)
<b>Positron</b>	A newer editor from the makers of RStudio, an alternative to it	Optional
<b>Quarto</b>	Turns documents with code into reports and books (Chapter 17)	Comes with RStudio
<b>R packages</b>	Add-ons for specific tasks, installed from within R	Yes, as needed
<b>Rtools</b> (Windows) or <b>Xcode command line tools</b> (macOS)	Compilers for building packages from source	Only if asked

All of it is free. R must be installed first, because RStudio and Positron are programs *for* R: they need R to run your code. Install R, then the editor.

## A.2 - Installing R

R is downloaded from **CRAN**, the Comprehensive R Archive Network, at [cran.r-project.org](https://cran.r-project.org). Always download the latest release. This book was built with R 4.4.3; any later version works.

### A.2.1 - Windows

1. On the CRAN page, choose *Download R for Windows*, then *base*, then the link to download the latest version (a file such as `R-4.x.y-win.exe`).
2. Run the downloaded file. If Windows asks whether to allow the app to make changes, choose *Yes*. If you have no administrator rights on a university computer, the installer offers to install for your user only, which works just as well.
3. Accept the default settings on every screen. The defaults install R in `C:\Program Files\R\` and register it so that RStudio finds it automatically.

### A.2.2 - macOS

1. On the CRAN page, choose *Download R for macOS*.
2. Choose the right package for your Mac's processor. Newer Macs (from late 2020) have **Apple silicon** (M1, M2, and later): use the file marked *arm64*. Older Macs have **Intel** processors: use the file marked *x86\_64*. The Apple menu, *About This Mac*, shows which you have.
3. Open the downloaded `.pkg` file and follow the installer, accepting the defaults.

### A.2.3 - Linux

R is available in the package manager of every major Linux distribution, but the version there is often out of date. CRAN's *Download R for Linux* page gives up-to-date instructions for Ubuntu, Debian, Fedora, and others. On Ubuntu, for example, the instructions add CRAN's repository and then install R with `sudo apt install r-base r-base-dev`. Follow the page for your distribution, because the exact commands change with each release.

### A.2.4 - Checking the installation

Open R (from the Start menu on Windows or the Applications folder on macOS) and type:

```
R.version.string
```

```
[1] "R version 4.4.3 (2025-02-28 ucrt)"
```

If it prints a version number, R is working. You will rarely open R on its own again; from now on you will work in RStudio.

## A.3 - Installing RStudio

1. Go to [posit.co/download/rstudio-desktop](https://posit.co/download/rstudio-desktop). The page detects your operating system; step 1 on the page (installing R) is already done.
2. Download **RStudio Desktop** (the free version) and install it like any other program, accepting the defaults. On macOS, drag RStudio into the Applications folder.
3. Open RStudio. It finds R automatically and shows the R version in the Console pane, in the lower left (Chapter 1 takes you on a tour of the panes).

RStudio includes **Quarto**, so the reports of Chapter 17 work without installing anything else. For PDF output, Quarto also needs LaTeX; install a small version once by typing `quarto install tinytex` in RStudio's **Terminal** tab.

### A.3.1 - Positron, an alternative

**Positron** ([positron.posit.co](https://positron.posit.co)) is a newer editor from Posit, the company behind RStudio, built for both R and Python. Everything in this book works in Positron as well. RStudio is used in the book's instructions because it is still the most widely used, and most tutorials and university courses assume it.

### A.3.2 - No installation at all

If you cannot install software, for example on a locked-down computer, two options need only a web browser:

- The book's [playground](#) runs R in your browser, with exercises for every chapter.
- **Posit Cloud** ([posit.cloud](https://posit.cloud)) offers RStudio in the browser, with a free plan that is limited but enough to follow most chapters.

## A.4 - Recommended RStudio settings

A few settings make RStudio work the way this book recommends. Open *Tools > Global Options*:

- **General > Basic**: untick *Restore .RData into workspace at startup*, and set *Save workspace to .RData on exit* to **Never**. Every session then starts clean, and your scripts, not a hidden saved workspace, hold your work. This is the single most important setting for reproducible work (Chapter 17).
- **Code > Editing**: tick *Use native pipe operator, |>*, so that the pipe shortcut (Ctrl+Shift+M, or Cmd+Shift+M on a Mac) inserts the pipe used in this book (Chapter 3).
- **Code > Display**: tick *Show margin* and *Highlight R function calls*, which make code easier to read.
- **Appearance**: choose a font size and theme that are comfortable for you; dark themes are popular for long sessions.

## A.5 - Installing the book's packages

Packages are installed once, with `install.packages()`, and loaded in every session with `library()` (Chapter 1). Each chapter lists the packages it uses at the start; Table A.2 collects them all.

Table A.2: Packages used in each chapter

Chapters	Packages
1-3	here, readxl, haven, writexl, dplyr, tidy, stringr, readr (or the whole tidyverse)
4, 6-7	ggplot2, psych
8	broom, car
9	psych, GPArotation, corrplot, factoextra
10	lme4, lmerTest
11-13	tidymodels, kknn, ranger, kernlab, themis, glmnet, xgboost
14	mclust, dbscan, factoextra, cluster
15	tidymodels (nnet comes with R)
16	tsibble, fable, feasts, urca
17	rmarkdown, knitr, shiny, bslib, renv, usethis
18	ellmer, stringr, yardstick
19	patchwork, broom, lme4

To install everything at once, copy this command into the Console. It downloads a few hundred megabytes and can take ten minutes or more, so it is best done on a good connection:

```
install.packages(c(
 "tidyverse", "here", "readxl", "haven",
 "writexl",
 "psych", "GPArotation", "corrplot",
 "factoextra", "broom", "car",
 "lme4", "lmerTest",
 "tidymodels", "kknn", "ranger", "kernlab",
 "themis", "glmnet", "xgboost",
 "mclust", "dbscan", "patchwork",
 "tsibble", "fable", "feasts", "urca",
 "rmarkdown", "shiny", "bslib", "renv",
 "usethis",
 "ellmer"
))
```



The **tidyverse** package installs `dplyr`, `tidyr`, `stringr`, `readr`, `ggplot2`, and several others in one go. The first time you install packages, R may ask you to choose a **CRAN mirror** (choose *0-Cloud*, which is fast everywhere) and whether to use a personal library (answer *Yes*).

### A.5.1 - The student wellbeing data: the `data2thesis` package

Elaf's data is in the **data2thesis** package, which is installed from this book's website rather than from CRAN:

```
install.packages("https://polla-fattah.
github.io/data2thesis_r/downloads/data2
thesis_1.1.0.tar.gz",
 repos = NULL, type = "source")
```

The package contains only data, so it installs on every system without compiling anything. Check that it works:

```
library(data2thesis)
nrow(students)
```

```
[1] 600
```

The same data is also available as ordinary files (CSV, Excel, and SPSS) on the book's website, for readers who prefer to import files, as Chapter 2 does.

### A.5.2 - The versions used in this book

The results in this book were produced with these versions. Newer versions usually give the same results, but if a number in your output differs slightly from the book, a different package version is a likely reason.

Software	Version
R	4.4.3
<code>dplyr</code>	1.2.1
<code>tidyr</code>	1.3.2
<code>ggplot2</code>	4.0.3
<code>lme4</code>	1.1.37
<code>tidymodels</code>	1.5.0
<code>fable</code>	0.5.0
<code>ellmer</code>	0.4.0
<code>data2thesis</code>	1.1.0

## A.6 - Common problems and solutions

### “There is no package called ...”

The package is not installed, or the name is misspelled (names are case-sensitive: `GPArotation`, not `gparotation`). Install it with `install.packages("name")`, then load it with `library(name)`.

**“Package ... is not available for this version of R”**

Either the name is misspelled, or your R is too old for the current version of the package. Update R (see below). A few packages are not on CRAN at all; their documentation says how to install them.

**“Do you want to install from sources the package which needs compilation?”**

A newer version exists as source code than as a ready-made (binary) package. Answer *No* to install the slightly older binary version, which is fine for this book. Answering *Yes* needs **Rtools** on Windows (from CRAN’s *Download R for Windows* page, choosing the Rtools version that matches your R) or the **Xcode command line tools** on macOS (run `xcode-select --install` in the macOS Terminal).

**“Lib ... is not writable” or permission errors**

R cannot write to the system’s package folder, which is common on university computers. When R offers to create a personal library, answer *Yes*.

**Problems with folder names on Windows**

If your Windows user name or project path contains non-English characters, or your files are in a synchronised folder such as OneDrive, installation and file reading can fail in confusing ways. Keeping R projects in a simple local folder, such as `C:\research\`, avoids most of these problems.

**Downloads fail at the university**

Some networks block or slow down downloads. Try the *0-Cloud* mirror (`chooseCRANmirror()`), another network, or ask your IT service whether a proxy must be set.

**RStudio cannot find R**

Install R before RStudio, or reinstall R with the default settings. On Windows, *Tools > Global Options > General* lets you choose the R version RStudio uses.

**Everything worked yesterday, and today nothing does**

Restart R (*Session > Restart R*), then run your script from the top. Many problems come from objects left over from earlier work, which a clean start removes (see the settings above).

## A.7 - Keeping R up to date

Update your packages every few months with `update.packages()`, or the *Update* button in RStudio’s Packages pane. Update R itself once or twice a year by installing the new version exactly as the first time; your scripts keep working, but packages must be reinstalled for the new version, which the command in Section A.5 does in one go. Tools such as **rig** ([github.com/r-lib/rig](https://github.com/r-lib/rig)) can install and switch between several R versions, useful if an old project needs an old version. For projects that must keep exactly the same package versions, use **renv** (Chapter 17).

Two things do not belong in any script. API keys for AI services (Chapter 18) go in your personal `.Renv` file, opened with `usethis::edit_r_environ()`; and passwords never go into code at all.

## B. Appendix B: Glossary

This appendix collects the key terms from every chapter's review, with a short definition of each and the chapters where it is introduced or used, followed by a table of the R functions used most often in the book. Definitions are written in plain language; the chapters give the full explanations and examples.

### B.1 - Key terms

#### A

**accuracy**

The share of cases a classification model classifies correctly. Misleading when one outcome is rare. (Chapter 11)

**activation function**

The function a neuron applies to its weighted total to produce its output, such as the sigmoid or the ReLU. (Chapter 15)

**adjusted Rand index**

A measure of agreement between two clusterings or classifications, corrected for chance: 1 means identical groupings, 0 means chance agreement. (Chapter 14)

**adjusted R-squared**

R-squared corrected for the number of predictors, so that adding useless predictors does not make a model look better. (Chapter 8)

**AI coding assistant**

An AI tool that writes, explains, or completes code, in a chat or inside the editor. Its suggestions must be checked. (Chapter 18)

**AIC**

Akaike information criterion: a measure for comparing models that balances fit against complexity; lower is better. (Chapter 10)

**alternative hypothesis**

The claim that there is an effect or a difference, tested against the null hypothesis. (Chapters 5, 7)

**annotation**

Text, arrows, or shapes added to a plot to point out something specific. (Chapter 4)

**anonymisation**

Removing or altering information so that the people in a dataset cannot be identified. (Chapter 17)

**ANOVA**

Analysis of variance: a test of whether the means of three or more groups differ, by comparing variation between groups with variation within them. (Chapter 8)

**API**

Application programming interface: a way for one program to send requests to another, such as R sending text to a language model. (Chapter 18)

**API key**

A secret code that identifies you to an online service's API. Keep it in `.Renvi ron`, never in a script. (Chapter 18)

**argument**

A value given to a function inside its brackets, such as `na.rm = TRUE`, that controls what the function does. (Chapter 1)

**ARIMA**

A family of time series models that forecast from the autocorrelation of the series and past random shocks. (Chapter 16)

**artificial intelligence**

Computer systems that perform tasks usually needing human intelligence, such as understanding language; in this book, mainly large language models. (Chapter 18)

**aspect ratio**

The ratio of a graph's width to its height. It changes how steep lines and slopes appear. (Chapter 4)

**assignment**

Storing a value in an object with `<-`, as in `x <- 5`. (Chapter 1)

**attrition**

The loss of participants from a study over time. (Chapters 5, 6)

**autocorrelation**

The correlation between a time series and itself a number of steps (lags) earlier. (Chapter 16)

**B****bar chart**

A plot of counts or values as bars, one per category. (Chapter 4)

**baseline model**

The simplest possible model, such as predicting the mean for everyone, used as a benchmark for real models. (Chapter 13)

**Bayesian information criterion (BIC)**

A criterion for comparing models that rewards fit and penalises complexity. In `mclust`, higher BIC is better. Also: BIC. (Chapter 14)

**benchmark forecast**

A simple forecast, such as the mean or the value a season earlier, that any serious method should beat. (Chapter 16)

**between-group variation**

How much group means differ from the overall mean; the numerator of the ANOVA F statistic. (Chapter 8)

**bias**

In a neural network, the constant added to a neuron's weighted total (like an intercept). More generally, a systematic error. (Chapters 5, 15)

**bias-variance trade-off**

The balance between a model too simple to capture the pattern (high bias) and one so flexible that it follows the noise of each sample (high variance). Error on new data is lowest in between. (Chapter 11)

**BibTeX**

A text format for bibliographic references, used by Quarto to format citations and reference lists. (Chapter 17)

**bimodal**

Having two peaks. (Chapter 6)

- bin**  
One of the intervals into which a histogram divides the values. (Chapter 4)
- boosting**  
Building a model by adding many small models (usually trees) one at a time, each correcting the errors of the model so far. (Chapter 13)
- bootstrap**  
Estimating uncertainty by repeatedly resampling the data with replacement and recalculating a statistic. (Chapter 7)
- bootstrap sample**  
A sample of the same size as the data, drawn from it at random with replacement. (Chapter 12)
- border point**  
In DBSCAN, a case near a core point that belongs to its cluster but is not itself a core point. (Chapter 14)
- box plot**  
A plot that shows the median, the quartiles, and unusual values of a variable. (Chapter 4)
- broom**  
A package that turns model results into tidy data frames, with `tidy()`, `glance()`, and `augment()`. (Chapter 19)
- C**
- case**  
One unit that is measured in a study, such as a student, a patient, or a school; one row of a data table. (Chapter 2)
- causal question**  
A research question about what leads to what. It needs a design that rules out other explanations, ideally an experiment. (Chapter 5)
- central limit theorem**  
The result that the sampling distribution of a mean is approximately normal for large enough samples, whatever the shape of the data. (Chapter 7)
- centring**  
Subtracting the mean from a variable, so that zero means “average”; often used before fitting interactions. (Chapter 8)
- chain of reasoning**  
The sequence of steps behind a research conclusion: question, hypothesis, design, variables, analysis, result, conclusion, and limitation. (Chapter 19)
- chi-square test**  
A test for categorical data: whether counts fit expected proportions (goodness of fit), or whether two categorical variables are related (independence). (Chapter 7)
- chunk option**  
A setting for a code chunk in a Quarto document, written as `#| option: value`, such as `echo: false`. (Chapter 17)
- citation**  
A reference to a source in the text, written in Quarto as `[@key]`. (Chapter 17)
- citation style (CSL)**  
A file in the Citation Style Language that sets how citations and references are formatted, such as APA. (Chapter 17)
- class weights**  
Making errors on a rare class count more when a model is fitted, one remedy for imbalanced outcomes. (Chapter 12)

**classification**

Predicting a category, such as whether a student will consider dropping out. (Chapter 11)

**cluster analysis**

Methods that group cases so that cases in the same group are similar to each other. (Chapter 9)

**cluster sampling**

Sampling whole groups, such as all the students of randomly chosen supervisors. (Chapter 5)

**code chunk**

A block of code in a Quarto document that is run when the document is rendered. (Chapter 17)

**codebook**

A description of every variable in a dataset or, in qualitative coding, of every theme and how to assign it. (Chapters 2, 17, 18)

**coefficient path**

A plot of how each coefficient of a regularised model changes as the penalty changes. (Chapter 13)

**Cohen's *d***

An effect size for the difference between two means, in standard deviations. (Chapter 7)

**Cohen's kappa**

A measure of agreement between two coders or classifiers, corrected for the agreement expected by chance. (Chapter 18)

**comment**

Text in a script after #, which R ignores; used to explain the code. (Chapter 1)

**commit**

In git, a saved snapshot of a project, with a message describing the change. (Chapter 17)

**confidence interval**

A range of plausible values for a population quantity, calculated from a sample, with a stated level of confidence such as 95%. (Chapter 7)

**confirmatory analysis**

An analysis planned before seeing the data, to test a specific hypothesis. (Chapters 5, 17)

**confirmatory factor analysis**

Factor analysis that tests a factor structure specified in advance. (Chapter 9)

**confounder**

A variable related to both the predictor and the outcome, which can create or hide an apparent effect. (Chapters 5, 6, 8)

**confusion matrix**

A table of predicted against actual classes, showing true and false positives and negatives. (Chapter 12)

**Console**

The RStudio pane where R commands are run and results appear. (Chapter 1)

**construct**

An idea that cannot be observed directly, such as stress or wellbeing, and must be measured indirectly. (Chapters 5, 9)

**construct validity**

Whether a score behaves as the construct should: related to similar measures, less related to different ones. (Chapter 5)

**content validity**

Whether the items of a measure cover the whole construct. (Chapter 5)

**convenience sampling**

Sampling whoever is easy to reach. Common, and the weakest basis for generalising. (Chapter 5)

**convolutional network**

A type of deep neural network designed for images. (Chapter 15)

**core point**

In DBSCAN, a case with at least `minPts` cases within distance `eps`. (Chapter 14)

**correlation coefficient**

A number from  $-1$  to  $1$  measuring the strength and direction of a straight-line relationship between two variables. (Chapter 6)

**correlation matrix**

A table of the correlations between every pair of variables. (Chapter 9)

**count outcome**

An outcome that counts events (0, 1, 2, ...), usually modelled with a Poisson model. (Chapter 10)

**covariance matrix**

A table of the variances and covariances of several variables; in a mixture model it sets the shape of a cluster. (Chapter 14)

**Cronbach's alpha**

A measure of the internal consistency (reliability) of a scale made of several items. (Chapter 9)

**cross-loading**

An item that loads substantially on more than one factor. (Chapter 9)

**cross-reference**

A reference in a Quarto document, such as `@fig-wellbeing`, that becomes a numbered link. (Chapter 17)

**cross-sectional design**

A design that measures every variable once, at one time. (Chapter 5)

**cross-validation**

Estimating how well a model predicts new data by repeatedly fitting it on part of the training data and testing it on the rest. (Chapter 11)

**CSV file**

A plain-text file of comma-separated values, the most common format for sharing data tables. (Chapter 2)

**D****data cleaning**

Preparing raw data for analysis: removing test and duplicate cases, making categories consistent, marking missing values, and correcting or removing impossible values, with every decision recorded. (Chapter 3)

**data frame**

R's table of data: columns are variables, rows are cases. (Chapters 1, 2)

**data leakage**

Information from the test data reaching the model during training, which makes its performance look better than it is. (Chapter 11)

**data matrix**

The arrangement of data as a table with one row per case and one column per variable. (Chapter 2)

**data structure**

The way data is organised in R: vector, factor, data frame, matrix, or list. (Chapter 2)

**data type**

The kind of value: numeric, integer, character, or logical. (Chapter 1)

**DBSCAN**

A clustering method that finds dense regions of cases and labels cases in sparse regions as noise. (Chapter 14)

**decision tree**

A model that predicts by a series of yes-or-no questions about the predictors. (Chapters 11, 12)

**decomposition**

Splitting a time series into trend, seasonal, and remainder components. (Chapter 16)

**deep learning**

Neural networks with many hidden layers, used for images, sound, and text. (Chapter 15)

**dendrogram**

The tree diagram produced by hierarchical clustering. (Chapter 9)

**density**

In clustering, how closely packed cases are in a region. (Chapter 14)

**density plot**

A smooth version of a histogram, showing the shape of a distribution. (Chapter 4)

**descriptive question**

A research question about what is: how much, how many, how often. (Chapter 5)

**descriptive statistics**

Numbers that summarise data, such as means, medians, and standard deviations. (Chapter 6)

**design effect**

How many times more grouped observations are needed to give the same information as independent ones:  $1 + (m - 1) \times ICC$ , where  $m$  is the group size. (Chapter 10)

**deviation**

The distance of a value from the mean,  $x_i - \bar{x}$ . The standard deviation summarises the deviations of all values. (Chapter 6)

**diagnostic plot**

A plot used to check a model's assumptions, such as residuals against fitted values. (Chapter 8)

**dictionary method**

Classifying texts by counting keywords from a list for each category. (Chapter 18)

**diminishing returns**

A relationship that levels off, so each extra unit of the predictor adds less than the one before. (Chapter 8)

**directional hypothesis**

A hypothesis that predicts the direction of an effect, such as "higher" or "lower". (Chapter 5)

**disclosure**

Stating in a publication how AI tools or other aids were used. (Chapter 18)

**distance**

How different two cases are, calculated from their values on several variables. (Chapters 9, 12)

**diverging palette**

Colours running in two directions away from a meaningful midpoint, such as zero, used for values that can be positive or negative. (Chapter 4)

**DOI**

Digital object identifier: a permanent identifier for a publication or dataset, such as [10.1126/science.1213847](https://doi.org/10.1126/science.1213847). (Chapter 17)

**downsampling**

Balancing classes by randomly removing cases of the common class from the training data. (Chapter 12)



- dpi** Dots per inch: the resolution of a saved image; 300 dpi is usual for print. (Chapter 4)
- dummy variable** A 0/1 variable representing one category of a categorical predictor. (Chapter 11)

## E

- effect size** A measure of how large an effect is, independent of sample size, such as Cohen's  $d$  or eta squared. (Chapter 7)
- eigenvalue** In PCA, the amount of variance captured by a component. (Chapter 9)
- elastic net** Regularised regression that mixes the ridge and lasso penalties. (Chapter 13)
- elbow method** Choosing the number of clusters where adding more stops reducing within-cluster distance much. (Chapter 9)
- ensemble** A model that combines the predictions of many models, such as a random forest. (Chapter 12)
- epoch** One pass through the training data when training a neural network. (Chapter 15)
- eps** In DBSCAN, the radius of the neighbourhood around each case. (Chapter 14)
- error message** R's message when something goes wrong, which usually says what and where. (Chapter 2)
- eta squared** An effect size for ANOVA: the share of the variation in the outcome explained by the groups. (Chapter 8)
- expected count** In a chi-square test, the count expected in a cell if the null hypothesis were true. (Chapter 7)
- explanation** Using a model to understand why something happens, rather than to predict new cases. (Chapter 11)
- explanatory graph** A carefully designed graph made for readers, to show a finding clearly and accurately. (Chapter 4)
- exploratory analysis** Exploring data to find patterns and generate questions, rather than to test planned hypotheses. (Chapters 5, 6, 17)
- exploratory graph** A quick graph made by the researcher to understand the data: to check distributions, find unusual values, and notice patterns. (Chapter 4)
- exponential smoothing (ETS)** Forecasting with weighted averages of past observations, with more weight on recent ones; ETS models the error, trend, and season. Also: ETS. (Chapter 16)
- external validation** Checking clusters or predictions against known categories or outcomes. (Chapter 14)
- external validity** Whether a study's results apply beyond it, to other people, places, and times. (Chapter 5)

**F****F statistic**

In ANOVA, the ratio of between-group to within-group variation. (Chapter 8)

**F1 score**

A single measure combining precision and recall (their harmonic mean). (Chapter 12)

**facet**

A small panel of a plot showing one subgroup; `facet_wrap()` makes one panel per group. (Chapter 4)

**factor**

In R, a categorical variable with a fixed set of levels. In factor analysis, a hidden (latent) variable behind several items. (Chapters 2, 9)

**factor analysis**

A method that explains the correlations among items by a smaller number of hidden factors. (Chapter 9)

**false negative**

A case that belongs to the positive class but is predicted negative. (Chapter 12)

**false positive**

A case predicted positive that belongs to the negative class. (Chapter 12)

**falsifiability**

The property of a claim that some possible result would contradict it. A hypothesis that fits every possible result tells us nothing. (Chapter 5)

**fixed effect**

In a mixed model, an effect assumed the same for everyone, such as the average change over time. (Chapter 10)

**fold**

One of the parts into which the data is split for cross-validation. (Chapter 11)

**forecast horizon**

How far ahead a forecast is made. (Chapter 16)

**Fourier terms**

Pairs of sine and cosine waves used as predictors to describe a seasonal pattern. (Chapter 16)

**function**

A named piece of code that takes arguments and returns a result, such as `mean()`. (Chapter 1)

**G****garden of forking paths**

The many reasonable choices in an analysis. Choosing among them after seeing the results inflates the chance of a false positive. (Chapter 17)

**Gaussian mixture model**

A model that describes data as a mix of normal distributions, giving each case a probability of belonging to each cluster. (Chapter 14)

**generalisation**

The ability of a model to perform well on new data, not only on the data it learned from. (Chapter 11)

**generalised linear mixed model**

A mixed-effects model for outcomes that are not normal, such as counts or yes/no outcomes. (Chapter 10)

**geom**

In `ggplot2`, the geometric shape that represents data, such as points, lines, or bars. (Chapter 4)

**Gini impurity**

A measure of how mixed the classes are in a group; decision trees choose splits that reduce it. (Chapter 12)

- git** The standard version control system, which records the history of a project. (Chapter 17)
- GitHub** A website for storing git projects online, sharing them, and working on them together. (Chapter 17)
- gradient** The direction and rate at which the error changes as each weight changes. (Chapter 15)
- gradient descent** Training a model by repeatedly moving the weights a small step in the direction that reduces the error. (Chapter 15)
- grammar of graphics** The idea behind ggplot2: a plot is built from data, aesthetic mappings, and geometric layers. (Chapter 4)
- graphical perception** How people read values from graphs. Positions and lengths are judged most accurately, then angles, areas, and colours. (Chapter 4)
- grouped summary** Summary statistics calculated separately for each group, for example with `summarise(.by = ...)`. (Chapter 3)

## H

- hallucination** Plausible but invented content produced by a language model, such as a function or reference that does not exist. (Chapter 18)
- heat map** A grid of coloured cells showing values, such as a correlation matrix. (Chapter 4)
- helper function** A small function written to avoid repeating the same code, such as one that formats means. (Chapter 19)
- hidden layer** A layer of neurons between the inputs and the output of a neural network. (Chapter 15)
- hierarchical clustering** Clustering that repeatedly merges the most similar groups, producing a dendrogram. (Chapter 9)
- histogram** A plot of the distribution of a numeric variable, as bars counting the values in each bin. (Chapter 4)
- hyperparameter** A setting of a model that is chosen before fitting rather than learned from the data, such as the number of neighbours in k-NN. (Chapter 11)
- hypothesis** A prediction, stated before the data is analysed, of what the data will show; precise enough to be wrong. (Chapter 5)

## I

- identifier** A variable that uniquely identifies each case, such as `student_id`. (Chapter 3)
- imbalanced outcome** An outcome in which one class is much rarer than the other. (Chapter 11)
- imputation** Filling in missing values with estimated ones, such as the median. (Chapter 11)

**independence**

The assumption that observations do not influence each other; in a chi-square test, the hypothesis that two variables are unrelated. (Chapter 7)

**index**

In a tsibble, the column that holds time. (Chapter 16)

**inferential analysis**

Using a sample to draw conclusions about a population. (Chapter 6)

**inline code**

R code inside a sentence of a Quarto document, replaced by its result when rendered. (Chapter 17)

**input**

In Shiny, a control such as a menu or slider whose value the user chooses. (Chapter 17)

**input layer**

The inputs (predictors) of a neural network. (Chapter 15)

**interaction**

When the effect of one predictor depends on the value of another. (Chapter 8)

**interaction plot**

A plot of group means that shows whether the effect of one factor depends on another. (Chapter 8)

**intercept**

The predicted value of the outcome when all predictors are zero. (Chapter 8)

**internal consistency**

The extent to which the items of a scale agree with each other, often measured with Cronbach's alpha. (Chapters 5, 9)

**internal validation**

Judging clusters by how compact and separated they are, using the data alone, as with the silhouette. (Chapter 14)

**internal validity**

Whether a study can rule out other explanations for its results. Highest in randomised experiments. (Chapter 5)

**interquartile range**

The range of the middle 50% of the values: the third quartile minus the first. (Chapter 6)

**inter-rater reliability**

The extent to which two people coding or rating the same material agree, often measured with Cohen's kappa. (Chapters 5, 18)

**interval**

A level of measurement with equal distances between values but no true zero, such as a wellbeing index. (Chapter 5)

**intra-class correlation**

The share of the total variation that lies between groups (such as students or supervisors). (Chapter 10)

**J****join**

Combining two tables by matching rows on a key, such as `student_id`. (Chapter 3)

**K****Kaiser rule**

Keeping components with an eigenvalue above 1; a rough guide that often keeps too many. (Chapter 9)

**kernel**

In a support vector machine, a function that allows curved boundaries between classes. (Chapter 12)

**k-means**

A clustering method that assigns each case to the nearest of  $k$  cluster centres and moves the centres to the means. (Chapter 9)

**k-nearest neighbours**

Predicting a case from the  $k$  most similar cases in the training data. (Chapters 11, 12)

**k-nearest-neighbour distance plot**

A sorted plot of each case's distance to its  $k$ -th nearest neighbour, used to choose `eps` for DBSCAN. (Chapter 14)

**knitr**

The R package that runs the code in Quarto and R Markdown documents. (Chapter 17)

**Kruskal-Wallis test**

A non-parametric alternative to one-way ANOVA. (Chapter 8)

**L****lag**

The number of time steps between an observation and an earlier one it is compared with. (Chapter 16)

**large language model**

A neural network trained on huge amounts of text to predict the next token, used in AI assistants. (Chapter 18)

**lasso**

Regularised regression with a penalty on the absolute size of the coefficients, which sets some of them to exactly zero. (Chapter 13)

**latent variable**

A variable that cannot be observed directly, such as stress, and is assumed to cause part of the answers to the items that measure it. (Chapter 9)

**LaTeX**

A typesetting system whose notation is used to write equations, as in `$$\bar{x}$$`. (Chapter 17)

**layer**

In ggplot2, one part of a plot added with `+`, such as a set of points or a trend line. In a neural network, a group of neurons. (Chapter 4)

**leaf**

A final node of a decision tree, which gives the prediction. (Chapter 12)

**learning rate**

In boosting and neural networks, how large a step each update takes. (Chapters 13, 15)

**least squares**

Choosing a regression line that minimises the sum of the squared residuals. (Chapter 8)

**level**

One of the categories of a factor. (Chapter 2)

**level of measurement**

What the values of a variable mean (nominal, ordinal, interval, or ratio), which decides the summaries and tests that make sense. (Chapter 5)

**Levene's test**

A test of whether groups have equal variances. (Chapter 8)

**licence**

A statement of what others may do with shared data or code, such as CC BY or MIT. (Chapter 17)

**likelihood ratio test**

A test comparing two nested models by how much better the larger one fits. (Chapter 10)

**line chart**

A plot of values connected by lines, often over time. (Chapter 4)

**linear regression**

A model of a numeric outcome as a straight-line function of one or more predictors. (Chapter 8)

**list**

An R object that can hold elements of different types and sizes, such as the results of a test. (Chapter 2)

**loading**

The correlation between an item and a factor or component. (Chapter 9)

**local model**

A language model that runs on your own computer, so data does not leave it. (Chapter 18)

**lockfile**

A file, such as `renv's` `renv.lock`, that records the exact package versions of a project. (Chapter 17)

**log loss**

The measure of error used to fit logistic regression and classification networks: small when high probabilities are given to the cases that did occur and low probabilities to those that did not. (Chapter 15)

**logical indexing**

Selecting elements with a TRUE/FALSE condition, as in `x[x > 5]`. (Chapter 2)

**logistic regression**

A regression model for a yes/no outcome, which predicts the probability of “yes”. (Chapter 8)

**log-odds**

The logarithm of the odds; the scale on which logistic regression coefficients are estimated. (Chapter 8)

**long format**

Data with one row per observation, such as one row per student per semester. (Chapter 3)

**longitudinal design**

A design that measures the same people repeatedly over time. (Chapter 5)

**M****machine learning**

Methods that learn patterns from data to make predictions or find structure, judged by performance on new data. (Chapter 11)

**MAE**

Mean absolute error: the average size of the prediction errors. (Chapter 13)

**main effect**

The effect of one factor, averaged over the levels of another. (Chapter 8)

**Mann-Whitney U test**

A non-parametric test comparing two independent groups. (Chapter 7)

**mapping**

In `ggplot2`, linking a variable to a visual property (an aesthetic), such as position or colour. Also: aesthetic mapping. (Chapter 4)

**margin**

In a support vector machine, the empty band between the classes and the boundary. (Chapter 12)

**Markdown**

A simple way of formatting text with symbols, such as `bold` and `# Heading`. (Chapter 17)

- matrix**  
A two-dimensional table of values of one type. (Chapter 2)
- mean**  
The average: the sum of the values divided by their number. (Chapter 6)
- mean method**  
A benchmark forecast that predicts the average of past observations. (Chapter 16)
- median**  
The middle value when the values are sorted. (Chapters 4, 6)
- mediator**  
A variable on the path between a predictor and an outcome, through which the predictor has its effect. (Chapter 5)
- membership probability**  
The probability that a case belongs to each cluster, as given by a mixture model. Also: soft assignment. (Chapter 14)
- method choice**  
Choosing an analysis method from the goal, the type of outcome, and how observations are related. (Chapter 19)
- minPts**  
In DBSCAN, the number of cases a neighbourhood must contain for a case to be a core point. (Chapter 14)
- misclassification cost**  
The cost assigned to each kind of error of a classifier, such as missing a case or raising a false alarm. The costs determine the best threshold. (Chapter 12)
- missing at random**  
Missingness that depends only on observed variables, not on the missing values themselves. (Chapter 6)
- missing code**  
A value such as 99 or -9 used in raw data to mark a missing answer. (Chapter 3)
- missing completely at random**  
Missingness unrelated to any variable, observed or not. (Chapter 6)
- missing not at random**  
Missingness that depends on the missing values themselves. (Chapter 6)
- missing value (NA)**  
R's marker for a value that is not available. Also: NA. (Chapter 1)
- mixed-effects model**  
A regression model with both fixed effects and random effects, for repeated or nested data. (Chapter 10)
- mixing probability**  
In a mixture model, the share of cases belonging to a component. (Chapter 14)
- mixture**  
In regularised regression, the mix of lasso and ridge penalties, from 0 (ridge) to 1 (lasso). (Chapter 13)
- mixture component**  
One of the normal distributions in a Gaussian mixture model. (Chapter 14)
- mode**  
The most common value. (Chapter 6)
- model specification**  
In tidymodels, the description of a model (type, engine, mode) before it is fitted. (Chapter 11)
- moderator**  
A variable that changes the strength or direction of the relationship between a predictor and an outcome. (Chapter 5)

**mtry**

In a random forest, the number of predictors each split may choose from. (Chapter 12)

**multicollinearity**

Strong correlation among predictors, which makes regression coefficients unstable. (Chapter 13)

**multilayer perceptron**

A neural network with one or more hidden layers of neurons. (Chapter 15)

**multilevel model**

Another name for a mixed-effects model, emphasising nested levels. (Chapter 10)

**multiple regression**

Linear regression with more than one predictor. (Chapter 8)

**multiple testing**

Running many tests in one study. Each test carries its own risk of a false alarm, so the chance of at least one grows with the number of tests. (Chapter 7)

**multivariate analysis**

Methods that analyse many variables at once, such as PCA and factor analysis. (Chapter 9)

**N****naive method**

A benchmark forecast that repeats the last observed value. (Chapter 16)

**nested data**

Data in which cases are grouped within units, such as students within supervisors. (Chapter 10)

**neural network**

A model made of connected neurons in layers, whose weights are learned from data. (Chapter 15)

**neuron (unit)**

The basic element of a neural network: it weights its inputs, adds a bias, and applies an activation function. Also: neuron, unit. (Chapter 15)

**node**

A point in a decision tree where a question is asked or a prediction is made. (Chapter 12)

**noise**

In DBSCAN, cases in sparse regions that belong to no cluster. (Chapter 14)

**nominal**

A level of measurement for categories with no order, such as faculty. (Chapter 5)

**non-directional hypothesis**

A hypothesis that predicts a difference or relationship without saying in which direction. (Chapter 5)

**non-parametric test**

A test that does not assume a particular distribution, often based on ranks. (Chapter 7)

**normal distribution**

The symmetric, bell-shaped distribution described by a mean and a standard deviation. (Chapter 6)

**normalisation**

Rescaling variables, usually to z-scores, so that they are on the same scale. (Chapter 11)

**null distribution**

The distribution of a statistic that would be expected if the null hypothesis were true, used to judge how surprising the observed result is. (Chapter 7)

**null hypothesis**

The claim of no effect or no difference, which a test tries to reject. (Chapters 5, 7)



## O

**object**

A named value stored in R, such as a number, a vector, or a data frame. (Chapter 1)

**oblique rotation**

A factor rotation that allows the factors to correlate. (Chapter 9)

**odds**

The probability of an event divided by the probability of it not happening. (Chapter 8)

**odds ratio**

The ratio of the odds in two groups; in logistic regression, the multiplicative change in odds per unit of a predictor. (Chapter 8)

**one-sample t-test**

A test of whether a mean differs from a given value. (Chapter 7)

**one-tailed test**

A test that looks for a difference in one direction only. (Chapter 7)

**open science**

Practices that make research transparent and reusable: sharing data and code, preregistration, and open access. (Chapter 17)

**operationalisation**

The decision about how a construct is measured: which questions, records, or instruments turn it into a variable. (Chapter 5)

**ordered factor**

A factor whose levels have an order, created with `factor(..., ordered = TRUE)`. (Chapter 5)

**ordinal**

A level of measurement for ordered categories whose distances are unknown, such as none, part-time, and full-time employment. (Chapter 5)

**outcome**

The variable a model explains or predicts. (Chapters 5, 8)

**outlier**

A value far from the others. (Chapter 6)

**output**

In Shiny, a plot, table, or text that the server produces for the user interface. (Chapter 17)

**output format**

The kind of document Quarto produces, such as HTML, Word, or PDF. (Chapter 17)

**output layer**

The final layer of a neural network, which produces the prediction. (Chapter 15)

**overfitting**

A model learning the noise in its training data, so it performs well there and poorly on new data. (Chapters 11, 13)

**overplotting**

Points drawn on top of each other in a plot, hiding how many there are. (Chapter 4)

## P

**package**

A collection of R functions, data, and documentation that adds features to R. (Chapter 1)

**paired t-test**

A test comparing two measurements on the same cases. (Chapter 7)

**Pandoc**

The document converter that Quarto uses to produce HTML, Word, PDF, and other formats. (Chapter 17)

**parallel analysis**

Choosing the number of factors by comparing eigenvalues with those from random data. (Chapter 9)

**parameter**

A quantity describing a population (Chapter 7); a value learned by a model, such as a coefficient (Chapter 11); or an input to a Quarto report (Chapter 17). (Chapters 7, 11, 17)

**patchwork**

A package that combines ggplot2 plots into one figure. (Chapter 19)

**penalty (lambda)**

In regularised regression or a neural network, the strength of the penalty on large coefficients or weights. Also: penalty, lambda. (Chapter 13)

**permutation importance**

A predictor's importance measured by how much shuffling its values reduces a model's accuracy. (Chapter 12)

**permutation test**

A test that builds the null distribution by shuffling group labels many times, and compares the observed result with the shuffled ones. (Chapter 7)

**pipe**

The operator `|>`, which passes the result on its left to the function on its right. (Chapter 3)

**pipeline**

The series of steps, in code, from raw data to results. (Chapter 19)

**Poisson model**

A regression model for counts. (Chapter 10)

**population**

The whole group a study wants to draw conclusions about. (Chapters 5, 7)

**post-hoc test**

A test run after ANOVA to find which groups differ, such as Tukey's test. (Chapter 8)

**power**

The probability that a study detects an effect of a given size, if the effect is real. A common target is 80%. (Chapters 5, 7)

**precision**

The share of cases predicted positive that are truly positive. (Chapter 12)

**predicted probability**

A model's estimated probability of an outcome for a case. (Chapter 8)

**prediction**

Using a model to estimate the outcome for new cases. (Chapter 11)

**prediction interval**

A range within which a future observation is expected to fall with a stated probability. (Chapter 16)

**predictive analysis**

Analysis aimed at predicting new cases rather than explaining or testing. (Chapter 6)

**predictor**

A variable used to explain or predict the outcome. (Chapters 5, 8)

**preregistration**

Publicly recording hypotheses and planned analyses before seeing the data. (Chapters 5, 17)

**pretrained model**

A model already trained by others on large datasets, used as it is or adapted. (Chapter 15)

**principal component**

A combination of variables that captures as much of their variation as possible. (Chapter 9)

**principal component analysis**

A method that summarises many correlated variables with a few principal components. (Chapter 9)

**project structure**

The organisation of a project's folders and files, such as `data-raw/`, `R/`, and `output/`. (Chapter 19)

**prompt**

The text given to a language model. (Chapter 18)

**push**

In git, copying commits to an online repository such as GitHub. (Chapter 17)

**p-value**

The probability of results at least as extreme as those observed, if the null hypothesis were true. (Chapter 7)

**Q****Q-Q plot**

A plot comparing a variable's values with those expected from a normal distribution. (Chapter 6)

**quadratic term**

A squared predictor in a regression model, which allows a curved relationship. (Chapter 8)

**qualitative coding**

Assigning themes or categories to texts such as open-ended answers. (Chapter 18)

**qualitative palette**

A set of clearly different colours of similar strength, used to distinguish categories. (Chapter 4)

**quartile**

The values that divide sorted data into four equal parts. (Chapter 6)

**Quarto**

A publishing system that turns documents with text and code into reports, books, websites, and slides. (Chapter 17)

**quasi-experiment**

A study that compares groups receiving different treatments that were not assigned by chance. (Chapter 5)

**R****R**

The programming language and environment for statistics used in this book. (Chapter 1)

**R Markdown**

The predecessor of Quarto: documents (`.Rmd`) that combine text and R code. (Chapter 17)

**radial basis function**

A common kernel for support vector machines, which allows curved, rounded boundaries. (Chapter 12)

**random effect**

In a mixed model, variation between groups (such as students) described by a distribution rather than a separate estimate for each group. (Chapter 10)

**random error**

The difference between an estimate and the truth caused by which cases happened to be selected. It shrinks as the sample grows. (Chapter 5)

**random forest**

An ensemble of decision trees, each grown on a bootstrap sample with random subsets of predictors. (Chapter 12)

**random intercept**

A random effect that lets each group have its own average level. (Chapter 10)

**random slope**

A random effect that lets each group have its own effect of a predictor, such as its own rate of change. (Chapter 10)

**randomised experiment**

A study in which the researcher assigns the treatment by chance, so that the groups differ only by chance at the start. (Chapter 5)

**range**

The difference between the largest and smallest values. (Chapter 6)

**rate ratio**

In a Poisson model, the multiplicative change in the expected count per unit of a predictor. (Chapter 10)

**ratio**

A level of measurement with equal distances and a true zero, so that “twice as much” is meaningful, such as hours of sleep. (Chapter 5)

**raw data**

Data as it was collected or exported, before any cleaning; never edited by hand. (Chapter 19)

**reactive expression**

In Shiny, a calculation, made with `reactive()`, that is updated automatically when its inputs change. (Chapter 17)

**reactivity**

Shiny’s system for updating outputs automatically when inputs change. (Chapter 17)

**recall**

The share of truly positive cases that a model predicts as positive. Also: sensitivity. (Chapter 12)

**recipe**

In `tidymodels`, a list of data preparation steps, such as imputation and normalisation. (Chapter 11)

**reference category**

The category of a categorical predictor that the others are compared with. (Chapter 8)

**registered report**

A publication format in which a journal accepts a study’s plan before the data is collected. (Chapter 17)

**regression**

Modelling or predicting a numeric outcome. In machine learning, regression means predicting a number rather than a category. Also: Regression (prediction). (Chapters 11, 13)

**regular expression**

A pattern for matching text, such as “`\\b(money|fee)`”. (Chapter 18)

**regularisation**

Penalising large coefficients to prevent overfitting, as in ridge and lasso regression. (Chapter 13)

**relational question**

A research question about which variables go together. It establishes association, not cause. (Chapter 5)

**reliability**

How consistently a scale measures what it measures. (Chapters 5, 9)

**ReLU**

Rectified linear unit: an activation function that sets negative values to zero. (Chapter 15)

**remainder**

In a time series decomposition, what is left after removing trend and season. (Chapter 16)

- render**  
Running the code in a Quarto document and producing the finished output. (Chapter 17)
- renv**  
A package that records and restores the package versions used by a project. (Chapter 17)
- repeated measures**  
Several measurements of the same cases, such as each student in four semesters. (Chapter 10)
- replicability**  
Whether a new study, with new data, reaches the same conclusions. (Chapter 17)
- reporting sentence**  
A sentence in a results section that states a finding with its numbers, such as an estimate, interval, and p-value. (Chapter 19)
- reporting standard**  
A published list of the information a research report must contain, such as JARS for psychology, CONSORT for randomised trials, and STROBE for observational studies. (Chapter 19)
- repository**  
A project folder whose history git keeps; also its online copy on GitHub. (Chapter 17)
- reproducibility**  
Whether the same data and code give exactly the same results. (Chapter 17)
- reproducible research**  
Research whose results can be recomputed from the shared data and code. (Chapter 1)
- research question**  
A focused question that says precisely what a study will find out, narrow enough that data can answer it. (Chapter 5)
- research workflow**  
The path of a research project from question and data to reported results. (Chapter 19)
- residual**  
The difference between an observed value and the value a model predicts. Also: error. (Chapters 8, 13)
- reversed item**  
A questionnaire item worded in the opposite direction to the others, which must be reversed before scoring. (Chapters 3, 9)
- ridge regression**  
Regularised regression with a penalty on the squared coefficients, which shrinks them all towards zero. (Chapter 13)
- RMSE**  
Root mean squared error: the square root of the average squared prediction error. (Chapter 13)
- ROC AUC**  
The area under the ROC curve: how often a model ranks a positive case above a negative one; 0.5 is chance, 1 is perfect. (Chapter 11)
- ROC curve**  
A plot of sensitivity against the false positive rate across all classification thresholds. (Chapter 12)
- rotation**  
In factor analysis, turning the factors to make the loadings easier to interpret. (Chapter 9)
- R-squared**  
The share of the variation in the outcome that a model explains, from 0 to 1. Also:  $R^2$ . (Chapters 8, 13)
- RStudio**  
The most widely used editor for R, used throughout this book. (Chapter 1)

**RStudio Project**

A folder that RStudio treats as the home of one piece of work, setting the working directory. (Chapter 1)

**S****sample**

The part of a population that is observed. (Chapters 5, 7)

**sample description (“Table 1”)**

A table describing the participants, usually the first table of a results chapter. Also: sample description. (Chapter 19)

**sampling distribution**

The distribution of a statistic over many possible samples. (Chapter 7)

**sampling frame**

A list of the members of a population, from which a sample can be drawn. (Chapter 5)

**scale**

In ggplot2, the control of how data values are turned into positions, colours, or sizes. (Chapter 4)

**scale score**

A score combining several questionnaire items, usually their mean. (Chapter 3)

**scaling**

Converting variables to a common scale, usually z-scores, before calculating distances. (Chapter 9)

**scatter plot**

A plot of two numeric variables as points. (Chapter 4)

**scree plot**

A plot of eigenvalues, used to choose the number of components or factors. (Chapter 9)

**script**

A file of R code that can be saved and rerun. (Chapter 1)

**seasonal naive method**

A benchmark forecast that repeats the value from the same season a cycle earlier. (Chapter 16)

**seasonal period**

The length of a repeating pattern, such as 52 weeks or 12 months. (Chapter 16)

**seasonal plot**

A plot that overlays the cycles of a time series to show its seasonal pattern. (Chapter 16)

**seasonality**

A pattern in a time series that repeats at a fixed period. (Chapter 16)

**seasonally adjusted series**

A time series with the seasonal component removed. (Chapter 16)

**self-report**

A measure based on what participants say about themselves, such as their estimated hours of sleep. (Chapter 5)

**self-selection**

When people choose their own group, such as volunteering for a workshop, so that the groups may differ from the start. (Chapter 5)

**sequential palette**

Colours running from light to dark, used for quantities that go in one direction. (Chapter 4)

**server**

In Shiny, the function that computes the outputs from the inputs. (Chapter 17)

**setting**

In `ggplot2`, giving a visual property a fixed value, such as `colour = "blue"`, rather than mapping it to a variable. (Chapter 4)

**Shiny**

An R package for building interactive web applications and dashboards. (Chapter 17)

**shrinkage**

Pulling coefficients towards zero, as regularisation does. (Chapter 13)

**sigmoid**

The S-shaped function that turns any number into a value between 0 and 1. (Chapter 15)

**significance level**

The threshold (usually 0.05) below which a p-value counts as statistically significant. (Chapter 7)

**silhouette**

A measure of how much closer each case is to its own cluster than to the next nearest, from -1 to 1. (Chapters 9, 14)

**simple random sampling**

Sampling in which every member of the population has the same chance of being chosen. (Chapter 5)

**simulation**

Using random numbers to generate many possible outcomes, for example many possible futures of a time series. (Chapter 16)

**singular fit**

A warning that a mixed model estimates some variation as zero, often because the model is too complex for the data. (Chapter 10)

**skewness**

Asymmetry of a distribution; a long tail to the right is positive skew. (Chapter 6)

**slope**

The change in the predicted outcome for a one-unit change in a predictor. (Chapter 8)

**small cells**

Combinations of variables that describe very few people, who could be identified. (Chapter 17)

**SMOTE**

A method that balances classes by creating artificial cases of the rare class between existing ones. (Chapter 12)

**specificity**

The share of truly negative cases that a model predicts as negative. (Chapter 12)

**stability**

Whether clusters or results stay the same under small changes to the data or the method. (Chapter 14)

**standard deviation**

A measure of spread: roughly the typical distance of values from the mean. (Chapter 6)

**standard error**

The standard deviation of a sampling distribution; the uncertainty of an estimate. (Chapter 7)

**statistic**

A number calculated from a sample, such as a mean or a t value. (Chapter 7)

**statistically significant**

A result with a p-value below the significance level. (Chapter 7)

**STL**

Seasonal and trend decomposition using loess: a flexible method for decomposing a time series. (Chapter 16)

**stratified sampling**

Sampling in which the population is divided into groups (strata) and a random sample is drawn from each. (Chapter 5)

**stratified split**

A split of the data that keeps the same share of each outcome class in each part. (Chapter 11)

**structured output**

A language model's answer in a fixed format, such as one category from a list, that code can use directly. (Chapter 18)

**study design**

The plan for who is measured, on what, when, and under which conditions. (Chapter 5)

**supervised learning**

Machine learning with a known outcome to predict. (Chapter 11)

**support vector**

A case close to the boundary of a support vector machine, which determines where the boundary lies. (Chapter 12)

**support vector machine**

A classifier that separates classes with the widest possible margin. (Chapter 12)

**symmetrical**

Having the same shape on both sides of the centre. (Chapter 6)

**synthetic data**

Artificial data with the structure of real data, used when the real data cannot be shared. (Chapter 17)

**system prompt**

Instructions given to a language model before the conversation, such as a task and a codebook. (Chapter 18)

**T****temperature**

A setting of a language model that controls how random its answers are; 0 gives the most likely answer. (Chapter 18)

**test set**

Data set aside to evaluate a final model once, at the end. (Chapter 11)

**test-retest reliability**

The extent to which a measure gives similar results for the same people on two occasions. (Chapter 5)

**theme**

In ggplot2, the non-data appearance of a plot, such as fonts and background. In qualitative coding, a category of meaning. (Chapter 4)

**threshold**

The probability above which a classifier predicts the positive class. (Chapter 12)

**tibble**

The tidyverse's version of a data frame, which prints more neatly. (Chapters 2, 3)

**tidy data**

Data with one variable per column, one observation per row, and one value per cell. (Chapter 3)

**tidyverse**

A collection of R packages, such as dplyr and ggplot2, that share one way of working. (Chapter 3)

**time plot**

A plot of a time series against time. (Chapter 16)

**time series**

Observations of one quantity at regular intervals over time. (Chapter 16)



**token**

A word or part of a word, the unit of text a language model reads and writes. (Chapter 18)

**training**

Adjusting a model's weights or parameters to fit the training data. (Chapter 15)

**training cut-off**

The date after which a language model's training data contains no information. (Chapter 18)

**training set**

The data used to build and tune a model. (Chapter 11)

**transformer**

The neural network architecture behind large language models. (Chapter 15)

**tree depth**

The number of questions in a row a decision tree may ask. (Chapter 13)

**trend**

The long-term direction of a time series. (Chapter 16)

**trend line**

A line added to a scatter plot to show the overall relationship. (Chapter 4)

**true negative**

A negative case correctly predicted negative. (Chapter 12)

**true positive**

A positive case correctly predicted positive. (Chapter 12)

**tsibble**

A data frame for time series, which knows which column is time. (Chapter 16)

**Tukey's test**

A post-hoc test comparing every pair of groups, with a correction for multiple comparisons. (Chapter 8)

**tuning**

Choosing hyperparameters by comparing their performance, usually with cross-validation. (Chapter 11)

**two-sample t-test**

A test comparing the means of two independent groups. (Chapter 7)

**two-tailed test**

A test that looks for a difference in either direction. (Chapter 7)

**two-way ANOVA**

ANOVA with two factors, including their interaction. (Chapter 8)

**type conversion**

Changing a value from one data type to another, such as text to numbers. (Chapter 2)

**Type I error**

Rejecting a true null hypothesis: a false alarm. (Chapter 7)

**Type II error**

Failing to reject a false null hypothesis: a missed effect. (Chapter 7)

**U****uncertainty**

In a mixture model, 1 minus a case's largest membership probability. (Chapter 14)

**underfitting**

A model too simple to capture the real pattern in the data. (Chapter 11)

**unimodal**

Having one peak. (Chapter 6)

**unit of analysis**

What one row of the data represents, and what the question is about: a student, a semester, a supervisor. (Chapters 2, 5)

**unsupervised learning**

Machine learning without an outcome, looking for structure such as clusters. (Chapter 11)

**upsampling**

Balancing classes by repeating cases of the rare class in the training data. (Chapter 12)

**user interface**

In Shiny, what the user sees: inputs and places for outputs. (Chapter 17)

**V****validity**

Whether a measure measures what it claims to, or whether a study's conclusions are justified. (Chapter 5)

**value label**

A label attached to a coded value, such as 1 = "Strongly disagree", as in SPSS files. (Chapter 2)

**variable**

One characteristic measured on every case, such as age or faculty; one column of a data table. (Chapter 2)

**variable label**

A longer description attached to a variable, as in SPSS files. (Chapter 2)

**variable selection**

Choosing which predictors to keep in a model; the lasso does it automatically. (Chapter 13)

**variance**

The average squared distance from the mean; the square of the standard deviation. (Chapter 6)

**vector**

R's basic structure: a sequence of values of the same type. (Chapters 1, 2)

**vector format**

An image format, such as PDF or SVG, that stays sharp at any size. (Chapter 4)

**verb**

A dplyr function that does one thing to a data frame, such as `filter()` or `mutate()`. (Chapter 3)

**version control**

Keeping a history of every change to a project's files. (Chapter 17)

**W****Ward's method**

A way of merging clusters in hierarchical clustering that keeps them as compact as possible. (Chapter 9)

**weak learner**

A simple model, such as a small tree, that is only slightly better than chance on its own; boosting combines many. (Chapter 13)

**weight**

In a neural network, the number by which an input is multiplied; learned during training. (Chapter 15)

**weight decay**

A penalty on large weights in a neural network, which prevents overfitting. (Chapter 15)

**Welch test**

The version of the two-sample t-test that does not assume equal variances; R's default. (Chapter 7)

**wide format**

Data with repeated measurements side by side in columns, one row per case. (Chapter 3)

**Wilcoxon signed-rank test**

A non-parametric test comparing two measurements on the same cases. (Chapter 7)

**within-group variation**

How much values vary around their own group mean; the denominator of the ANOVA F statistic. (Chapter 8)

**workflow**

In `tidymodels`, a recipe and a model specification combined. (Chapter 11)

**workflow set**

In `tidymodels`, several workflows fitted and compared together. (Chapter 12)

**working directory**

The folder R reads files from and saves files to by default. (Chapter 1)

**X****XGBoost**

Extreme gradient boosting: a fast, widely used boosting method. (Chapter 13)

**Y****YAML header**

The settings at the top of a Quarto document, between two `---` lines. (Chapter 17)

**Z****z-score**

A value expressed as the number of standard deviations from the mean. (Chapter 6)

**B.2 - R functions**

Table B.1 lists the functions used most often in the book, grouped by task, with the package they come from and the chapters that use them. `base` R functions are available without loading any package. For any function, `?name` in the Console opens its help page.



## C. Appendix C: R Packages

This appendix lists every R package used in the book: what it is for, the chapters that use it, and the version used when the book was built. Appendix A shows how to install them all with one command. A dash in the Version column marks a package that is recommended to readers but was not needed to build the book itself.

A package is a collection of functions, data, and documentation that someone has written and shared. R comes with a set of **base packages**, such as stats (`t.test()`, `lm()`) and utils (`read.csv()`), which are always available; everything else is installed once with `install.packages()` and loaded in each session with `library()` (Chapter 1).

## C.1 - The packages in this book

Table C.1: The R packages used in this book.

Area	Package	What it is for	Chapters	Version
Data	data2life	Lifeis Graduate Wellbeing Study: the case-study data of this book	1, 2, 3, 4, 5, and later	1.1.0
Data	model	Example datasets for modeling, such as credit applications, concrete strength, and penguins	11, 12, 13, 15	1.6.0
Importing and exporting	readr	Reading and writing CSV and other text files, and <code>parse_number()</code>	3, 19	2.1.5
Importing and exporting	readxl	Reading Excel files	2, 3, 19	1.4.5
Importing and exporting	haven	Reading and writing SPSS, Stata, and SAS files, with their labels	2, 19	2.5.5
Importing and exporting	writexl	Writing Excel files	2	1.5.4
Importing and exporting	here	File paths relative to the project folder	1, 2, 3, 19	1.0.1
Data handling	tidyverse	Installs and loads the core tidyverse packages in one go	3	-
Data handling	dplyr	Filtering, selecting, creating, summarising, and joining data	3, 4, 5, 6, 7, and later	1.2.1
Data handling	tidyr	Reshaping data between wide and long formats, and other tidying tasks	3, 4, 5, 6, 7, 19, and later	1.3.2

Two entries are collections rather than single packages. **tidyverse** installs and loads `dplyr`, `tidyr`, `stringr`, `readr`, `ggplot2`, `purrr`, and a few others; **tidymodels** does the same for the modelling packages of Chapters 11 to 15 (`rsample`, `recipes`, `parsnip`, `workflows`, `tune`, and `yardstick`, among others). Several model packages (`ranger`, `knn`, `kernlab`, `glmnet`, `xgboost`, `nnet`) are rarely loaded by name: `tidymodels` calls them as **engines**, as in `set_engine("ranger")`, but they must be installed.

## C.2 - Finding and choosing packages

CRAN, R's official package archive, holds more than 20,000 packages, and many more are shared on GitHub. A few ways to find the right one:

- **CRAN Task Views** ([cran.r-project.org/web/views](https://cran.r-project.org/web/views)) are curated lists of packages by topic, such as psychometrics, survival analysis, or time series, maintained by experts in each field.
- A package's **vignettes**, long-form tutorials that come with it, are the best introduction: `browseVignettes("dplyr")` lists them. Many packages also have a website with examples.
- Methods papers in journals such as the *Journal of Statistical Software* and *The R Journal* describe many packages in depth.

Before relying on a package for your thesis, a few signs show whether it is trustworthy: it is on CRAN (which checks that packages install and run); it was updated in the last year or two; it has documentation and examples; it is described in a published paper or widely used in your field; and its authors are known in the area. The packages in this book meet all or most of these.

## C.3 - Citing packages

Package authors are researchers too, and citing their work is how they receive credit. `citation()` gives the recommended reference for R itself, and `citation("package")` for a package:

```
citation("psych")
```

To cite package 'psych' in publications use:

```
William Revelle (2026). _psych: Procedures
for Psychological,
Psychometric, and Personality Research_.
Northwestern University,
Evanston, Illinois. R package version 2.6.4,
<https://CRAN.R-project.org/package=psych>.
```

A BibTeX entry for LaTeX users is

```
@Manual{,
 title = {psych: Procedures for
Psychological, Psychometric, and Personality
```



```
Research},
 author = {{William Revelle}},
 organization = {Northwestern University},
 address = {Evanston, Illinois},
 year = {2026},
 note = {R package version 2.6.4},
 url = {https://CRAN.R-project.org/package=
psych},
}
```

Cite the packages that did substantial work in your analysis, and give their version numbers (Chapter 17). `packageVersion("psych")` shows the version installed on your computer.



## D. Appendix D: Common Errors and How to Fix Them

Everyone who writes R code sees error messages every day, experts included. An error is not a sign that you are doing badly; it is R telling you, as precisely as it can, what it could not do. This appendix collects the messages that readers of this book are most likely to meet, what each one means, and how to fix it. The messages below are real: each example was run when the book was built, so the wording is what your R will show (it can differ slightly between versions).

### D.1 - Reading a message

R produces three kinds of messages:

- An **error** stops the code: nothing after it runs. It starts with `Error`.
- A **warning** lets the code finish but tells you something may be wrong. It starts with `Warning`. Never ignore a warning without understanding it: several below mean your results contain missing values or were calculated differently from what you intended.
- A **message** is information only, such as a package telling you it has loaded.

When a long message appears, read it from the end: the last lines usually say what went wrong, and lines starting with `i` point to where. Newer packages (`dplyr`, `ggplot2`, `tidymodels`) write especially helpful messages, often with a suggestion (`Did you mean ...?`).

### D.2 - Starting out

#### D.2.1 - object ‘...’ not found

```
mean(sleep_hours)
```

```
Error in h(simpleError(msg, call)): error
in evaluating the argument 'x' in selecting
a method for function 'mean': object
'sleep_hours' not found
```

R does not know the name. The usual causes: a typing mistake (names are case-sensitive: `Sleep_hours` is not `sleep_hours`); the object was never created, because the line that creates it was not run; or the name is a column inside a data frame, which must be reached through the data frame, as in `mean(semester$sleep_hours)` or inside a `dplyr` verb. If a document fails to render with this error although the code works in the Console, the

object was created in the Console but not in the document (Chapter 17).

### D.2.2 - could not find function “..”

```
Mean(c(6.5, 7, 5.5))
```

```
Error in Mean(c(6.5, 7, 5.5)): could not find
function "Mean"
```

Either the function name is misspelled (here, `mean` with a capital M), or it belongs to a package that is not loaded. Load the package with `library()`, or write `package::function()`. The pipe `%>%` from older code gives the same error until `dplyr` or `magrittr` is loaded; the native pipe `|>` needs no package.

### D.2.3 - there is no package called ‘..’

```
library(ggplot)
```

```
Error in library(ggplot): there is no package
called 'ggplot'
```

The package is not installed, or its name is misspelled (the package is `ggplot2`). Install it once with `install.packages("ggplot2")`, then load it with `library(ggplot2)`. Appendix A lists every package the book uses.

### D.2.4 - unexpected symbol, unexpected ‘), unexpected string constant

```
mean(c(6.5, 7) na.rm = TRUE)
```

```
Error in parse(text = input): <text>:1:16:
unexpected symbol
1: mean(c(6.5, 7) na.rm
 ^
```

A **syntax error**: R cannot read the line at all. The `^` marks where R got lost; the mistake is usually just before it. Here a comma is missing before `na.rm`. Other common causes are an extra or missing bracket (`unexpected ' )`), a missing comma between two pieces of text (`unexpected string constant`), and a missing quotation mark. RStudio highlights matching brackets and marks syntax errors with a red cross in the margin before you even run the code.

### D.2.5 - The Console shows + and nothing happens

R is waiting for the rest of an unfinished command, usually because a bracket or quotation mark is not closed. Press **Esc** to cancel, fix the line, and run it again.

### D.2.6 - argument “...” is missing, with no default

```
rnorm()
```

```
Error in rnorm(): argument "n" is missing,
with no default
```

A required argument was not given. The help page (`?rnorm`) lists the arguments; those without a default value (here `n`, the number of values) must be supplied.

### D.2.7 - A misspelled argument, with no error at all

```
mean(c(6.5, NA, 7), na_rm = TRUE)
```

```
[1] NA
```

Not every mistake produces a message. The argument is `na.rm`, not `na_rm`, and because `mean()` accepts extra arguments, the misspelled one is silently ignored: the missing value is not removed, and the answer is `NA`. When a result looks wrong, check the spelling of every argument against the help page.

## D.3 - Files and folders

### D.3.1 - cannot open file ...: No such file or directory

```
read.csv("studnets.csv")
```

```
Warning in file(file, "rt"): cannot open file
'studnets.csv': No such file or
directory
```

```
Error in file(file, "rt"): cannot open the
connection
```

R looked for the file in the **working directory** and did not find it. Check the spelling of the file name (here `studnets`), check that the file is in the project folder, and build paths with `here::here()` inside an RStudio Project (Chapter 1), so they work on any computer. `list.files()` shows the files R can see.

### D.3.2 - cannot change working directory

```
setwd("C:/Users/Elaf/Documents/thesis")
```

```
Error in setwd("C:/Users/Elaf/Documents/
thesis"): cannot change working directory
```

The folder does not exist on this computer, which is exactly why `setwd()` with a full path breaks as soon as

a script is shared or moved. Use an RStudio Project and relative paths instead (Chapter 1).

## D.4 - Data

### D.4.1 - non-numeric argument to binary operator

```
"7" + 1
```

```
Error in "7" + 1: non-numeric argument to
binary operator
```

Arithmetic on text. The quotation marks make "7" a piece of text, not a number. In real data, this happens when a numeric column was imported as text, often because of one stray value such as "7 hrs" or a missing-value code. Check with `str()` or `glimpse()`, and convert with `as.numeric()` or `readr::parse_number()` (Chapters 2 and 3).

### D.4.2 - argument is not numeric or logical: returning NA

```
mean(c("6.5", "7"))
```

```
Warning in mean.default(c("6.5", "7")):
argument is not numeric or logical:
returning NA
```

```
[1] NA
```

A **warning**, not an error, and the result is `NA`: the same problem as above, a numeric column stored as text. Convert the column first.

### D.4.3 - NAs introduced by coercion

```
as.numeric(c("6.5", "7 hrs", "six"))
```

```
Warning: NAs introduced by coercion
```

```
[1] 6.5 NA NA
```

Values that could not be converted to numbers became `NA`. Look at which ones ("7 hrs" and "six" here) before going on: `parse_number()` can rescue "7 hrs", but "six" needs a decision (Chapter 3). Losing data silently through this warning is one of the most common problems in real analyses.

### D.4.4 - The result is NA

```
mean(c(6.5, NA, 7))
```

```
[1] NA
```

Not an error: any calculation that includes a missing value gives `NA`, because the true answer is unknown. Add `na.rm = TRUE` to calculate from the available values, and report how many were missing (Chapter 6).

### D.4.5 - object ‘...’ not found, inside dplyr

```
students |> filter(faculty == "Education")
```

```
Error in `filter()`:
i In argument: `faculty == "Education"`.
Caused by error:
! object 'faculty' not found
```

A column name is misspelled. dplyr says in which argument the problem is (i In argument: ...) and which name it could not find. `names(students)` lists the correct names.

### D.4.6 - arguments imply differing number of rows; replacement has ... rows

```
data.frame(student = 1:3, sleep = c(6.5, 7))
```

```
Error in data.frame(student = 1:3, sleep =
c(6.5, 7)): arguments imply differing number
of rows: 3, 2
```

Every column of a data frame must have the same length. The data behind a new column has more or fewer values than the data frame has rows; check the lengths with `length()` and `nrow()`.

### D.4.7 - \$ operator is invalid for atomic vectors

```
x <- c(sleep = 6.5, study = 30)
x$sleep
```

```
Error in x$sleep: $ operator is invalid for
atomic vectors
```

`$` works on data frames and lists, but `x` is a vector. Use `x["sleep"]` for a vector, and check what kind of object you have with `class()` or `str()`.

### D.4.8 - subscript out of bounds

```
results <- list(6.5, 7)
results[[3]]
```

```
Error in results[[3]]: subscript out of bounds
```

You asked for an element that does not exist: the third element of a list with two. Check the length with `length()`, and the names with `names()`.

### D.4.9 - invalid factor level, NA generated

```
answer <- factor(c("Yes", "No"))
answer[1] <- "Maybe"
```

```
Warning in `[<-.factor`(`*tmp*`, 1, value =
"Maybe"): invalid factor level, NA
generated
```

```
answer
```

```
[1] <NA> No
Levels: No Yes
```

A factor accepts only its existing levels (Chapter 2); any other value becomes *NA*. Add the level first with `levels()`, or work with the column as text and convert it to a factor at the end.

### D.4.10 - Detected an unexpected many-to-many relationship

```
students_small <- data.frame(id = c(1, 1),
group = c("A", "B"))
scores_small <- data.frame(id = c(1, 1), score
= c(3.2, 4.1))
left_join(students_small, scores_small,
join_by(id))
```

```
Warning in left_join(students_small,
scores_small, join_by(id)): Detected an
unexpected many-to-many relationship between
`x` and `y`.
i Row 1 of `x` matches multiple rows in `y`.
i Row 1 of `y` matches multiple rows in `x`.
i If a many-to-many relationship is expected,
set `relationship =
"many-to-many"` to silence this warning.
```

	id	group	score
1	1	A	3.2
2	1	A	4.1
3	1	B	3.2
4	1	B	4.1

A **join** (Chapter 3) found identifiers that appear more than once in *both* tables, so every copy was matched with every other, multiplying rows. Usually one of the tables should have had one row per identifier: check for duplicates with `count(id) |> filter(n > 1)`, and remove them or join by more columns.

### D.4.11 - Can't combine ... and ...

```
tidyr::pivot_longer(students, cols = c(age,
gender))
```



```
Error in `tidyr::pivot_longer()`:
! Can't combine `age` <integer> and `gender`
<character>.
```

`pivot_longer()` puts the chosen columns into one column, so they must be of the same type: here a number and a text column. Reshape only columns of the same type, or convert them first.

## D.5 - Plots

### D.5.1 - Cannot use + with a single argument

```
plot <- ggplot(students, aes(x = age))
+ geom_histogram()
```

```
Error:
! Cannot use `+` with a single argument.
i Did you accidentally put `+` on a new line?
```

In `ggplot2`, the `+` must come at the *end* of a line, not the start of the next one. Otherwise R thinks the first line is complete, and the second line starts with a stray `+`. The message even asks: “Did you accidentally put `+` on a new line?”

### D.5.2 - mapping must be created by `aes()` ... Did you use `%>%` or `|>` instead of `+`?

```
ggplot(students, aes(x = age)) |>
geom_histogram()
```

```
Error in `geom_histogram()`:
! `mapping` must be created by `aes()`.
* You've supplied a <ggplot2::ggplot> object.
i Did you use `%>%` or `|>` instead of `+`?
```

Layers of a `ggplot` are added with `+`, not with the pipe. The pipe passes data *into* `ggplot()`; after that, it is `+` all the way.

### D.5.3 - object ‘...’ not found, in a plot

```
ggplot(students) + geom_point(x = age, y =
financial_worry)
```

```
Error: object 'age' not found
```

Columns must be mapped inside `aes()`: `geom_point(aes(x = age, y = financial_worry))`. Outside `aes()`, R looks for objects called `age` and `financial_worry` and does not find them (Chapter 4 on mapping versus setting).

### D.5.4 - `stat_count()` must only have an x or y aesthetic

```
ggplot(students, aes(x = faculty, y = age))
+ geom_bar()
```

```
Error in `geom_bar()` :
! Problem while computing stat.
i Error occurred in the 1st layer.
Caused by error in `setup_params()` :
! `stat_count()` must only have an x or y
aesthetic.
```

`geom_bar()` counts the rows in each category, so it takes only x. To plot a value you have calculated, such as a mean per faculty, use `geom_col()`.

## D.6 - Statistical tests and models

### D.6.1 - grouping factor must have exactly 2 levels

```
t.test(age ~ faculty, data = students)
```

```
Error in t.test.formula(age ~ faculty, data =
students): grouping factor must have exactly
2 levels
```

A two-sample t-test compares exactly two groups, but `faculty` has five. Use ANOVA for more than two groups (Chapter 8), or filter the data to the two groups you want to compare.

### D.6.2 - not enough 'x' observations

```
t.test(c(6.5))
```

```
Error in t.test.default(c(6.5)): not enough
'x' observations
```

The test needs more data than it was given, often because filtering or missing values left too few cases. Check the number of cases with `nrow()` or `sum(!is.na(x))`.

### D.6.3 - contrasts can be applied only to factors with 2 or more levels

```
lm(age ~ gender, data = filter(students,
gender == "Female"))
```

```
Error in `contrasts<-`(`*tmp*`, value =
contr.funs[1 + isOF[nn]]): contrasts can be
applied only to factors with 2 or more levels
```

A categorical predictor has only one value in the data used, here because the data was filtered to women only, so there is nothing to compare. Remove the predictor, or check the filtering.

### D.6.4 - Chi-squared approximation may be incorrect

```
chisq.test(matrix(c(3, 1, 2, 4), nrow = 2))
```

```
Warning in chisq.test(matrix(c(3, 1, 2, 4),
nrow = 2)): Chi-squared
approximation may be incorrect
```

```
Pearson's Chi-squared test with Yates'
continuity correction
```

```
data: matrix(c(3, 1, 2, 4), nrow = 2)
X-squared = 0.41667, df = 1, p-value = 0.5186
```

Some expected counts are below 5, so the chi-square test's p-value may be inaccurate (Chapter 7). Use Fisher's exact test (`fisher.test()`), or combine small categories.

### D.6.5 - glm.fit: fitted probabilities numerically 0 or 1 occurred; algorithm did not converge

```
glm(passed ~ hours, family = binomial,
 data = data.frame(hours = 1:10, passed =
c(0, 0, 0, 0, 0, 1, 1, 1, 1, 1)))
```

```
Warning: glm.fit: algorithm did not converge
```

```
Warning: glm.fit: fitted probabilities
numerically 0 or 1 occurred
```

In a logistic regression, a predictor separates the outcomes perfectly (here, everyone above 5 hours passed), so the coefficients become enormous and meaningless, and the fitting algorithm may also report that it did not converge. With real data, this usually means a very small group or a predictor that is almost the outcome itself. Check the cross-table of the outcome and the predictor, and consider simplifying the model.

### D.6.6 - boundary (singular) fit: see help('isSingular')

```
lmer(wellbeing ~ semester + (semester |
supervisor_id),
 data = left_join(semesters, students,
join_by(student_id)))
```

```
boundary (singular) fit: see
help('isSingular')
```

A mixed-effects model estimated some random-effect variance as zero, or a correlation as exactly  $\pm 1$ : the random effects part is too complex for the data (Chapter 10). Simplify it, for example by removing the random slope: `(1 | supervisor_id)`.

### D.6.7 - Observations deleted due to missingness

Not an error, but a line in the output of `summary()` for `lm()` and `glm()`: rows with a missing value in any variable of the model are left out. Check how many with `nobs(model)`, report the number used, and think about whether the missing rows differ from the others (Chapters 6 and 8).

## D.7 - Machine learning

### D.7.1 - For a classification model, the outcome should be a

```
library(tidymodels)
```

—	Attaching	packages
tidymodels 1.5.0 —		
✓ broom	1.0.13	✓ rsample 1.3.2
✓ dials	1.4.4	✓ tailor 0.1.0
✓ infer	1.1.0	✓ tidyr 1.3.2
✓ modeldata	1.6.0	✓ tune 2.1.0
✓ parsnip	1.6.0	✓ workflows 1.3.0
✓ purrr	1.2.2	✓ workflowsets 1.1.1
✓ recipes	1.4.0	✓ yardstick 1.4.0

—	Conflicts
tidymodels_conflicts() —	
* purrr::discard()	masks scales::discard()
* tidyr::expand()	masks Matrix::expand()
* dplyr::filter()	masks stats::filter()
* dplyr::lag()	masks stats::lag()
* tidyr::pack()	masks Matrix::pack()
* recipes::step()	masks stats::step()
* tidyr::unpack()	masks Matrix::unpack()
* recipes::update()	masks Matrix::update(), stats::update()

```
fit(logistic_reg(), considering_dropout ~
age, data = students)
```

```
Error in `check_outcome()`:
! For a classification model, the outcome
should be a <factor>, not a
character vector.
```

tidymodels needs a categorical outcome to be a factor. Convert it first, and set the event of interest as the first level (Chapter 11): `mutate(considering_dropout = factor(considering_dropout, levels = c("Yes", "No")))`.

### D.7.2 - Can't select columns that don't exist ... `.pred_Yes`

```
tibble(truth = factor(c("Yes", "No")),
 probability = c(0.2, 0.8)) |>
 roc_auc(truth, .pred_Yes)
```

```
Error in `roc_auc()`:
! Can't select columns that don't exist.
* Column `.pred_Yes` doesn't exist.
```

A column name does not exist in the data given to a yardstick function. The prediction columns created by `augment()` are named after the outcome's levels (`.pred_Yes`, `.pred_No`), so check the names with `names()`, and check that the predictions were added to the data.

## D.8 - When you are stuck

1. **Read the message**, from the end, and look at the line it points to.
2. **Restart R and run the script from the top** (*Session > Restart R*). Many mysterious errors come from objects left over from earlier work.
3. **Check your data** with `str()`, `glimpse()`, or `summary()`: wrong types and unexpected missing values cause most errors in real analyses.
4. **Make the problem small**. Reproduce it with a few rows of data or a tiny example, like those at the start of every chapter. Often the cause becomes obvious; if not, a small example is what others need to help you. The `reprex` package formats such an example for sharing.
5. **Search for the message** in quotation marks, or ask an AI assistant, including the code, the full message, and a description of your data (Chapter 18). Check the answer before trusting it.
6. **Ask a person**: a colleague, your supervisor, the Posit Community forum, or Stack Overflow. A clear question with a small example usually gets a quick answer.



## E. Appendix E: Further Learning Resources

This book is a starting point. This appendix suggests where to go next: books, courses, and communities, almost all of them free. The books cited in the chapters are listed first, by topic; the rest are chosen because they suit researchers who are not programmers.

### E.1 - Books

#### E.1.1 - Getting started and data skills

- *R for Data Science* [2], free at [r4ds.hadley.nz](https://r4ds.hadley.nz): the standard introduction to importing, cleaning, transforming, and visualising data with the tidyverse. The natural next book after Part 1 of this one.
- *R for Researchers: An Introduction* by Tyson Barrett, free at [tysonbarrett.com/Rstats](https://tysonbarrett.com/Rstats): a short, friendly introduction written for researchers in the social and health sciences, and a good companion to Chapters 1 to 8.
- *ggplot2: Elegant Graphics for Data Analysis* [6], free at [ggplot2-book.org](https://ggplot2-book.org): everything about ggplot2, for when Chapter 4 leaves you wanting more.
- *Happy Git and GitHub for the useR* by Jenny Bryan, free at [happygitwithr.com](https://happygitwithr.com): a gentle, practical guide to git and GitHub with RStudio (Chapter 17).

#### E.1.2 - Statistics

- *Introduction to Modern Statistics* [13], free at [open-intro-ims.netlify.app](https://open-intro-ims.netlify.app): a modern introductory statistics textbook built around simulation and real data, like Chapters 6 and 7.
- *Learning Statistics with R* by Danielle Navarro, free at [learningstatisticswithr.com](https://learningstatisticswithr.com): a thorough, readable introduction to statistics for psychology and the social sciences, using R.
- *Statistical Inference via Data Science (ModernDive)* by Chester Ismay and Albert Kim, free at [moderndive.com](https://moderndive.com): regression and inference with the tidyverse.
- *Regression and Other Stories* [17]: an excellent guide to building, checking, and interpreting regression models in real research, with a free PDF from the authors.
- *Statistical Power Analysis for the Behavioral Sciences* [15]: the classic reference on effect sizes and power.

#### E.1.3 - Mixed models, machine learning, and forecasting

- *Data Analysis Using Regression and Multilevel/Hierarchical Models* [20]: multilevel models from the ground up (Chapter 10).

- *An Introduction to Statistical Learning* [18], free at [statlearning.com](http://statlearning.com): the clearest introduction to the ideas of machine learning (Chapters 11 to 15).
- *Tidy Modeling with R* [22], free at [tmwr.org](http://tmwr.org): the tidy-models framework in depth.
- *Forecasting: Principles and Practice* [32], free at [otexts.com/fpp3](http://otexts.com/fpp3): the standard introduction to forecasting with fable (Chapter 16).
- *Deep Learning with R* [30]: for readers who want to go further with neural networks (Chapter 15).

### E.1.4 - Reporting and dashboards

- The Quarto guide at [quarto.org](http://quarto.org): documentation and examples for every kind of document (Chapter 17).
- *Mastering Shiny* [37], free at [mastering-shiny.org](http://mastering-shiny.org): the definitive guide to Shiny apps.

The **Big Book of R** ([bigbookofr.com](http://bigbookofr.com)) indexes hundreds of free R books by topic, from psychology and ecology to economics and text analysis: a good way to find a book for your own field.

## E.2 - Courses and practice

- **Posit Cheatsheets** ([posit.co/resources/cheatsheets](http://posit.co/resources/cheatsheets)): one- or two-page summaries of dplyr, ggplot2, tidyr, Quarto, Shiny, and more, worth keeping next to you while you work. Several have been translated into other languages by volunteers.
- **The Carpentries** ([carpentries.org](http://carpentries.org)): free, hands-on lessons for researchers, including *R for Social Scientists* and *R for Ecologists*, and two-day workshops held at universities around the world.
- **swirl** ([swirlstats.com](http://swirlstats.com)): interactive lessons that run inside R itself: `install.packages("swirl")`, then `library(swirl)` and `swirl()`.
- **TidyTuesday** ([github.com/rfordatascience/tidytuesday](https://github.com/rfordatascience/tidytuesday)): a new real dataset every week to practise on, with thousands of shared solutions to learn from.
- This book's [playground](#): exercises for every chapter, in the browser and as downloadable projects.

## E.3 - Communities

You do not have to learn alone, and the R community is known for welcoming beginners.

- **Posit Community** ([forum.posit.co](http://forum.posit.co)): a friendly forum for questions about R, RStudio, the tidyverse, Quarto, and Shiny.
- **Stack Overflow** ([stackoverflow.com/questions/tagged/r](http://stackoverflow.com/questions/tagged/r)): the largest archive of answered R questions; search it before asking, and include a small reproducible example when you ask (Appendix D).
- **R-Ladies** ([rladies.org](http://rladies.org)): a worldwide organisation promoting gender diversity in the R community, with local chapters that run meetings and workshops open to learners.
- **R user groups** meet in many cities and universities; the R Consortium lists them at [r-consortium.org](http://r-consortium.org).



- **R Weekly** ([rweekly.org](http://rweekly.org)): a weekly digest of R news, tutorials, and new packages.
- **useR!**, the annual international R conference, posts its talks online.

Whatever you learn from, the most effective way to learn R is to use it on your own data, a little every day, one question at a time, as Elaf did.



# Bibliography

- [1] R. Ihaka and R. Gentleman, “R: A Language for Data Analysis and Graphics,” *Journal of Computational and Graphical Statistics*, vol. 5, no. 3, pp. 299–314, 1996, doi: [10.1080/10618600.1996.10474713](https://doi.org/10.1080/10618600.1996.10474713).
- [2] H. Wickham, M. Çetinkaya-Rundel, and G. Grolemund, *R for Data Science: Import, Tidy, Transform, Visualize, and Model Data*, 2nd ed. O'Reilly Media, 2023. [Online]. Available: <https://r4ds.hadley.nz/>
- [3] H. Wickham, “Tidy Data,” *Journal of Statistical Software*, vol. 59, no. 10, pp. 1–23, 2014, doi: [10.18637/jss.v059.i10](https://doi.org/10.18637/jss.v059.i10).
- [4] F. J. Anscombe, “Graphs in Statistical Analysis,” *The American Statistician*, vol. 27, no. 1, pp. 17–21, 1973, doi: [10.1080/00031305.1973.10478966](https://doi.org/10.1080/00031305.1973.10478966).
- [5] W. S. Cleveland and R. McGill, “Graphical Perception: Theory, Experimentation, and Application to the Development of Graphical Methods,” *Journal of the American Statistical Association*, vol. 79, no. 387, pp. 531–554, 1984, doi: [10.1080/01621459.1984.10478080](https://doi.org/10.1080/01621459.1984.10478080).
- [6] H. Wickham, *ggplot2: Elegant Graphics for Data Analysis*, 2nd ed. New York: Springer, 2016. doi: [10.1007/978-3-319-24277-4](https://doi.org/10.1007/978-3-319-24277-4).
- [7] K. Healy, *Data Visualization: A Practical Introduction*. Princeton, NJ: Princeton University Press, 2018. [Online]. Available: <https://socviz.co/>
- [8] K. R. Popper, *The Logic of Scientific Discovery*. London: Hutchinson, 1959.
- [9] S. S. Stevens, “On the Theory of Scales of Measurement,” *Science*, vol. 103, no. 2684, pp. 677–680, 1946, doi: [10.1126/science.103.2684.677](https://doi.org/10.1126/science.103.2684.677).
- [10] L. J. Cronbach and P. E. Meehl, “Construct Validity in Psychological Tests,” *Psychological Bulletin*, vol. 52, no. 4, pp. 281–302, 1955, doi: [10.1037/h0040957](https://doi.org/10.1037/h0040957).
- [11] B. A. Nosek, C. R. Ebersole, A. C. DeHaven, and D. T. Mellor, “The Preregistration Revolution,” *Proceedings of the National Academy of Sciences*, vol. 115, no. 11, pp. 2600–2606, 2018, doi: [10.1073/pnas.1708274114](https://doi.org/10.1073/pnas.1708274114).
- [12] E. W. Crampton, “The Growth of the Odontoblasts of the Incisor Tooth as a Criterion of the Vitamin C Intake of the Guinea Pig,” *The Journal of Nutrition*, vol. 33, no. 5, pp. 491–504, 1947, doi: [10.1093/jn/33.5.491](https://doi.org/10.1093/jn/33.5.491).
- [13] M. Çetinkaya-Rundel and J. Hardin, *Introduction to Modern Statistics*, 2nd ed. OpenIntro, 2024. [Online]. Available: <https://openintro-ims.netlify.app/>
- [14] W. R. Shadish, T. D. Cook, and D. T. Campbell, *Experimental and Quasi-Experimental Designs for Generalized Causal Inference*. Boston: Houghton Mifflin, 2002.

- [15] J. Cohen, *Statistical Power Analysis for the Behavioral Sciences*, 2nd ed. Hillsdale, NJ: Lawrence Erlbaum Associates, 1988.
- [16] Student, “The Probable Error of a Mean,” *Biometrika*, vol. 6, no. 1, pp. 1–25, 1908, doi: [10.1093/biomet/6.1.1](https://doi.org/10.1093/biomet/6.1.1).
- [17] A. Gelman, J. Hill, and A. Vehtari, *Regression and Other Stories*. Cambridge: Cambridge University Press, 2020. doi: [10.1017/9781139161879](https://doi.org/10.1017/9781139161879).
- [18] G. James, D. Witten, T. Hastie, and R. Tibshirani, *An Introduction to Statistical Learning: With Applications in R*, 2nd ed. New York: Springer, 2021. doi: [10.1007/978-1-0716-1418-1](https://doi.org/10.1007/978-1-0716-1418-1).
- [19] G. Belenky *et al.*, “Patterns of Performance Degradation and Restoration During Sleep Restriction and Subsequent Recovery: A Sleep Dose-Response Study,” *Journal of Sleep Research*, vol. 12, no. 1, pp. 1–12, 2003, doi: [10.1046/j.1365-2869.2003.00337.x](https://doi.org/10.1046/j.1365-2869.2003.00337.x).
- [20] A. Gelman and J. Hill, *Data Analysis Using Regression and Multilevel/Hierarchical Models*. Cambridge: Cambridge University Press, 2007. doi: [10.1017/CBO9780511790942](https://doi.org/10.1017/CBO9780511790942).
- [21] D. Bates, M. Mächler, B. Bolker, and S. Walker, “Fitting Linear Mixed-Effects Models Using lme4,” *Journal of Statistical Software*, vol. 67, no. 1, pp. 1–48, 2015, doi: [10.18637/jss.v067.i01](https://doi.org/10.18637/jss.v067.i01).
- [22] M. Kuhn and J. Silge, *Tidy Modeling with R: A Framework for Modeling in the Tidyverse*. O'Reilly Media, 2022. [Online]. Available: <https://www.tmwr.org/>
- [23] L. Breiman, “Random Forests,” *Machine Learning*, vol. 45, no. 1, pp. 5–32, 2001, doi: [10.1023/A:1010933404324](https://doi.org/10.1023/A:1010933404324).
- [24] A. M. Horst, A. P. Hill, and K. B. Gorman, “Palmer Archipelago Penguins Data in the palmerpenguins R Package: An Alternative to Anderson’s Irises,” *The R Journal*, vol. 14, no. 1, pp. 244–254, 2022, doi: [10.32614/RJ-2022-020](https://doi.org/10.32614/RJ-2022-020).
- [25] R. Tibshirani, “Regression Shrinkage and Selection via the Lasso,” *Journal of the Royal Statistical Society: Series B (Methodological)*, vol. 58, no. 1, pp. 267–288, 1996, doi: [10.1111/j.2517-6161.1996.tb02080.x](https://doi.org/10.1111/j.2517-6161.1996.tb02080.x).
- [26] T. Chen and C. Guestrin, “XGBoost: A Scalable Tree Boosting System,” in *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 2016, pp. 785–794. doi: [10.1145/2939672.2939785](https://doi.org/10.1145/2939672.2939785).
- [27] I.-C. Yeh, “Modeling of Strength of High-Performance Concrete Using Artificial Neural Networks,” *Cement and Concrete Research*, vol. 28, no. 12, pp. 1797–1808, 1998, doi: [10.1016/S0008-8846\(98\)00165-3](https://doi.org/10.1016/S0008-8846(98)00165-3).
- [28] L. Scrucca, M. Fop, T. B. Murphy, and A. E. Raftery, “mclust 5: Clustering, Classification and Density Estimation Using Gaussian Finite Mixture Models,” *The R Journal*, vol. 8, no. 1, pp. 289–317, 2016, doi: [10.32614/RJ-2016-021](https://doi.org/10.32614/RJ-2016-021).
- [29] M. Hahsler, M. Piekenbrock, and D. Doran, “dbscan: Fast Density-Based Clustering with R,” *Journal of Statistical Software*, vol. 91, no. 1, pp. 1–30, 2019, doi: [10.18637/jss.v091.i01](https://doi.org/10.18637/jss.v091.i01).

- [30] F. Chollet, T. Kalinowski, and J. J. Allaire, *Deep Learning with R*, 2nd ed. Manning, 2022.
- [31] I. Goodfellow, Y. Bengio, and A. Courville, *Deep Learning*. MIT Press, 2016. [Online]. Available: <https://www.deeplearningbook.org/>
- [32] R. J. Hyndman and G. Athanasopoulos, *Forecasting: Principles and Practice*, 3rd ed. Melbourne, Australia: OTexts, 2021. [Online]. Available: <https://otexts.com/fpp3/>
- [33] R. D. Peng, “Reproducible Research in Computational Science,” *Science*, vol. 334, no. 6060, pp. 1226–1227, 2011, doi: [10.1126/science.1213847](https://doi.org/10.1126/science.1213847).
- [34] G. Wilson, J. Bryan, K. Cranston, J. Kitzes, L. Nederbragt, and T. K. Teal, “Good Enough Practices in Scientific Computing,” *PLOS Computational Biology*, vol. 13, no. 6, p. e1005510, 2017, doi: [10.1371/journal.pcbi.1005510](https://doi.org/10.1371/journal.pcbi.1005510).
- [35] Open Science Collaboration, “Estimating the Reproducibility of Psychological Science,” *Science*, vol. 349, no. 6251, p. aac4716, 2015, doi: [10.1126/science.aac4716](https://doi.org/10.1126/science.aac4716).
- [36] J. Bryan, “Excuse Me, Do You Have a Moment to Talk About Version Control?,” *The American Statistician*, vol. 72, no. 1, pp. 20–27, 2018, doi: [10.1080/00031305.2017.1399928](https://doi.org/10.1080/00031305.2017.1399928).
- [37] H. Wickham, *Mastering Shiny*. O'Reilly Media, 2021. [Online]. Available: <https://mastering-shiny.org/>
- [38] J. Cohen, “A Coefficient of Agreement for Nominal Scales,” *Educational and Psychological Measurement*, vol. 20, no. 1, pp. 37–46, 1960, doi: [10.1177/001316446002000104](https://doi.org/10.1177/001316446002000104).
- [39] J. R. Landis and G. G. Koch, “The Measurement of Observer Agreement for Categorical Data,” *Biometrics*, vol. 33, no. 1, pp. 159–174, 1977, doi: [10.2307/2529310](https://doi.org/10.2307/2529310).
- [40] F. Gilardi, M. Alizadeh, and M. Kubli, “ChatGPT Outperforms Crowd Workers for Text-Annotation Tasks,” *Proceedings of the National Academy of Sciences*, vol. 120, no. 30, p. e2305016120, 2023, doi: [10.1073/pnas.2305016120](https://doi.org/10.1073/pnas.2305016120).
- [41] M. Appelbaum, H. Cooper, R. B. Kline, E. Mayo-Wilson, A. M. Nezu, and S. M. Rao, “Journal Article Reporting Standards for Quantitative Research in Psychology: The APA Publications and Communications Board Task Force Report,” *American Psychologist*, vol. 73, no. 1, pp. 3–25, 2018, doi: [10.1037/amp0000191](https://doi.org/10.1037/amp0000191).
- [42] E. von Elm, D. G. Altman, M. Egger, S. J. Pocock, P. C. Gøtzsche, and J. P. Vandenbroucke, “The Strengthening the Reporting of Observational Studies in Epidemiology (STROBE) Statement: Guidelines for Reporting Observational Studies,” *The Lancet*, vol. 370, no. 9596, pp. 1453–1457, 2007, doi: [10.1016/S0140-6736\(07\)61602-X](https://doi.org/10.1016/S0140-6736(07)61602-X).



# From Data To Thesis Using R

## Turn messy data into a defended, publication-ready thesis.

For many postgraduate and academic researchers, collecting data is only the beginning of the real challenge. The hardest hurdle is turning spreadsheets, surveys, and lab observations into robust statistical models and defensible written arguments. Commercial point-and-click tools leave your analysis opaque, while standard programming textbooks bury you in abstract computer science theory.

**From Data to Thesis** bridges that divide. Written specifically for graduate students, doctoral researchers, and academic professionals, this book provides a complete, modern roadmap for empirical research using R. Rather than abstract mathematics, every concept is taught through a single, realistic running case study following 600 students across four semesters.

### Every chapter guides you through five clear steps:

1. A Concrete Research Question.
2. A Minimal Working Example .
3. The Case-Study Analysis.
4. In Your Field.
5. Thesis Reporting & Interpretation .

From exploratory data analysis and hypothesis testing to mixed-effects models, machine learning, time-series forecasting, and ethical AI assistance, this book gives you the tools, code, and confidence to defend your research with scientific rigor.

### Inside you will discover:

- ❖ How to clean, reshape, and visualize complex data using modern **tidyverse** workflows.
- ❖ How to choose the right statistical test—and how to diagnose and fix common assumption violations.
- ❖ Practical machine learning and predictive modeling without the computational overwhelm.
- ❖ Reproducible research workflows using **Quarto** that generate thesis-ready tables and figures.
- ❖ Interactive browser playgrounds to practice every method without installing a single package.

*Whether you are writing your first Master's dissertation, completing your doctoral defense, or preparing an empirical journal article, this book is your desk companion from raw data to submitted thesis.*



by Polla Fattah  
and Collection of LLMs